

Essays in Applied Econometrics and Labor Economics
(with updates of Chapters 1-2 for Zellner Thesis Award, 2022)

Leonard Goff

Submitted in partial fulfillment of the
requirements for the degree of
Doctor of Philosophy
under the Executive Committee
of the Graduate School of Arts and Sciences

COLUMBIA UNIVERSITY

2021

Abstract

Essays in Applied Econometrics and Labor Economics

Leonard Goff

Recent decades have seen great advances in the methods we use to understand cause and effect in the world of work. Building on that tradition, this dissertation explores two broad topics in econometrics as tools to address specific questions in labor economics. The main econometric contributions are to extend identification results for research designs based on bunching (Chapter 1) and those that make use of instrumental variables (Chapters 2 and 3). The empirical questions that compel them are described below.

Chapter 1 examines the effect of overtime regulation on hours of work in the United States, extending a recently popularized technique that uses bunching observed at kinks in agents' choice sets for identification. In the U.S., most workers are required to be paid one-and-a-half times their typical rate of pay for any hours in excess of forty within a week. While prominent and long-standing, this policy has not been meaningfully reformed since it was first established at the federal level in 1938. As a result, few studies have been able to leverage causal research designs to assess its labor market impacts. I use bunching in the distribution of weekly hours at forty—where the policy introduces a convex “kink” in firms’ costs—to estimate this effect. To do so, I develop a framework in which bunching at a choice-set kink is informative about causal effects under substantially weaker assumptions than those maintained in existing work. This allows the effect of the overtime policy to be partially identified without making parametric assumptions about firms’ objective functions, or about the distribution of hours they would set in the absence of the policy. Using an administrative dataset of weekly hours derived from payroll records, I find that the bounds are informative and that covered hourly workers in the U.S. work an average of at least half an hour less as a result, in affected weeks.

Chapter 2 turns to a still-more popular strategy in applied microeconomics: the instrumental variables research design. I propose a new method for estimating causal effects

when a researcher has more than one such instrument, and apply it to reassess the labor market returns to college education. The method is motivated by the following issue. When treatment effects are heterogeneous, it is known that instruments can be used to identify local average treatment effects under an assumption known as “monotonicity”. However, when a researcher wishes to use multiple instruments together, this assumption can become quite restrictive, and empirical conclusions may be misleading if it is violated. I propose an alternative assumption that I call “vector monotonicity”, which is quite natural in typical settings with multiple instruments. I show that vector monotonicity leads to identification of a useful class of treatment effect parameters, but the two-stage-least-squares estimator popular in applied work does not consistently estimate them. I propose an alternative estimator, and apply it to the classic question of the returns to schooling. I find that the approach based upon vector monotonicity reveals new patterns of heterogeneity in the earnings effect of college education.

Chapter 3, with coauthors Ashna Arora and Jonas Hjort, considers the effects of a worker’s first job on outcomes later in their career. This is typically a difficult question to answer empirically, as workers entering the labor force are not randomly assigned to employers. We make use of a unique opportunity to study this question in the context of medical residencies in Norway. For decades, medical school graduates in Norway were matched to residencies based on a random serial dictatorship mechanism, in which doctors could choose—in an order determined by lottery—among available positions in the country. We develop an econometric framework in which the random choice set a doctor is presented with provides a collection of instruments for their choice of residency hospital, and hence first job as a doctor. Because we only observe choices and not a doctor’s full preferences, this requires new methods—related to those of Chapter 2. We find persistent effects of a doctor’s first job on earnings, specializations, and mid-career moves. We use the estimates to assess the replacement of the serial-dictatorship by a decentralized labor market in 2013, which we find led to a small increase in resident welfare.

Table of Contents

Chapter 1: Treatment Effects in Bunching Designs: The Impact of the Federal Over-time Rule on Hours	1
Chapter 2: A Vector Monotonicity Assumption for Multiple Instruments	121
Chapter 3: The Career Impact of First Jobs: Evidence and Labor Market Design Lessons from Randomized Choice Sets	196

**Chapter 1: Treatment Effects in Bunching Designs: The Impact of the
Federal Overtime Rule on Hours**

Treatment Effects in Bunching Designs: The Impact of the Federal Overtime Rule on Hours

Leonard Goff*

This version: May 12, 2022
For current version [click here](#)

Abstract

The 1938 Fair Labor Standards Act mandates overtime premium pay for most U.S. workers, but limited variation in the rule has made assessing its impacts on the labor market difficult. With data from individual paychecks, I use the extent to which firms bunch workers at the overtime hours threshold of 40 to estimate the rule's effect on work hours. Generalizing previous methods, I show that bunching at a choice-set kink partially identifies an average causal response to the policy switch at the kink, under nonparametric assumptions about preferences and heterogeneity. The bounds indicate a small elasticity of demand for hours.

*Department of Economics, University of Georgia. I thank my PhD co-advisors Simon Lee and Suresh Naidu, as well as Michael Best, Daniel Hamermesh, and Bernard Salanié for their gracious advice and support throughout this project. I have also benefited from discussions with Christopher Ackerman, Doug Almond, Joshua Angrist, Jushan Bai, Iain Bamford, Marc Bellemare, Sandra Black, Carol Caetano, Brant Callaway, Ivan Canay, Gregory Cox, Junlong Feng, Bhargav Gopal, Jonas Hjort, Wojciech Kopczuk, Bentley MacLeod, Matthew Masten, Serena Ng, José Luis Montiel Olea, Dilip Ravindran, Ian Schmutte, Miguel Urquiola, and seminar participants at Columbia, Duke, Lehigh, Microsoft, the University of Calgary, the University of Georgia, the University of North Carolina Greensboro, the US Census Bureau, and Washington State University, as well as audiences at the Canadian Economics Association and Southern Economics Association meetings, and the European Winter Meeting of the Econometric Society. I thank my main data provider as well as Ted To and the Bureau of Labor Statistics. Online Appendices available [here](#).

1 Introduction

Many countries require premium pay for long work hours, in an effort to limit excessive work schedules and encourage hours to be spread over more workers. In the U.S., such regulation comes through the “time-and-a-half” rule of the Fair Labor Standards Act (FLSA): firms must pay a worker one and a half times their normal hourly wage for any hours worked in excess of 40 within a single week. Although many salaried workers are exempt, the time-and-a-half rule applies to a majority of the U.S. workforce, including nearly all of its over 80 million hourly workers. Workers in many industries average multiple overtime hours per week, making overtime the largest form of supplemental pay in the U.S. (Hart, 2004; Bishow, 2009).

Nevertheless, only a small literature has studied the effects of the FLSA overtime rule on the labor market. This stands in marked contrast to the large body of work on the minimum wage, which was also introduced at the federal level by the FLSA in 1938. A key reason for this gap is that the overtime rule has varied little since then: the policy has remained as time-and-a-half after 40 hours in a week, for now more than 80 years. Reforms to overtime policy have been rare and have focused on eligibility, leaving the central parameters of the rule unaffected. This lack of variation has afforded few opportunities to leverage research designs that exploit policy changes to identify causal effects.¹

This paper assesses the effect of the FLSA overtime rule on hours of work, taking a new approach that makes use of variation *within* the rule itself. The policy introduces a sharp discontinuity in the marginal cost of a worker-hour—a convex “kink” in firms’ costs—which provides firms with an incentive to set workers’ hours exactly at 40. Optimizing behavior by firms predicts that the resulting mass of workers working 40 hours in a given week will be larger or smaller depending on how responsive firms are to the wage increase imposed by the time-and-a-half rule. Combining this observation with assumptions about the shape of the distribution of hours that would be chosen absent the FLSA, I use the bunching mass to identify the effect of the overtime rule on hours.

To do so, I develop a generalization of the “bunching design” identification strategy, which has previously used bunching at kinks in income tax liability to identify the elasticity of labor supply (Saez 2010; Chetty et al. 2011).² I give new identification results under weakened assumptions that may be suitable to a variety of empirical contexts, showing that the bunching design can be useful for program-evaluation questions such as assessing the effect of the FLSA.

¹A few studies that have used difference-in-differences approaches to estimating effects of U.S. overtime policy on hours: Hamermesh and Trejo (2000) consider the expansion of a daily overtime rule in California to men in 1980, while Johnson (2003) use a supreme court decision on the eligibility of public-sector workers in 1985. Costa (2000) studies the initial phase-in of the FLSA in the years following 1938. See footnotes 30 and 31 for a comparison of my results to these papers. Quach (2021) looks at very recent reforms to eligibility criteria for exemption from the FLSA, estimating effects of the expansion on employment and the incomes of salaried workers, but not on hours of work.

²The same basic model has since been applied in a range of settings beyond income taxation. This paper considers only the bunching design for kinks, and not a related method for bunching at *notches* (e.g. Kleven and Waseem 2013).

In income tax settings, the promise of the bunching design is to overcome endogeneity in the marginal tax rates that apply to different individuals, while requiring only the cross-sectional distribution of income near a threshold between tax brackets. Analogously, my starting point in the overtime setting is to construct the distribution of weekly work hours. Administrative hours data at the weekly level has previously been unavailable, and studies of overtime in the U.S. have typically relied on self-reported integer hours from surveys such as the Current Population Survey. I instead obtain detailed data via individual paycheck records from a large payroll processing company. Among workers paid weekly, these paychecks report the exact number of hours that the worker was paid for in a given week, allowing me to construct the distribution of hours-of-pay without rounding or other sources of measurement error.

With these novel data in hand, the goal is to translate features of the observed hours distribution into estimates of the overtime rule's causal effect, under credible assumptions about how weekly working hours are determined. This requires moving beyond the standard bunching-design model popularized in public-finance applications, in which decision-makers have parametric "isoelastic" preferences and strong restrictions are placed on heterogeneity. In the overtime setting, I show that bunching is informative about firms as the decision-maker, choosing the hours of each of their workers in a given week. With this in mind, the identifying assumptions of the bunching design can be separated into two parts: i) assumptions about how individual agents (firms) would make choices given counterfactual choice sets—a *choice model*, and ii) assumptions about the distribution of heterogeneity in choices across observational units (paychecks).

As a first methodological contribution, I show that the class of choice models under which the bunching-design can be used is considerably more general than the benchmark isoelastic model and its variants. In particular, I find that the method does not rest upon the researcher positing any explicit functional form for decision-makers' (firms') utility; rather, the main prediction about choice driving identification comes from *convexity* of preferences (e.g. weekly profits). In my formulation, agents in the bunching design can furthermore have multiple underlying margins of choice, which might be unobserved to the researcher and vary by observational unit.³ These findings establish an important robustness property for the bunching design: it rests on a prediction about choice behavior that remains broadly valid even when the isoelastic utility model typically employed in the bunching design is misspecified.

This generality is accomplished by recasting the bunching design in the language of potential outcomes, defining the parameter of interest in terms of a pair of counterfactual *choices* rather than as a preference parameter from a parametric choice model. In the overtime setting these potential outcomes correspond to: a) the number of hours the firm would choose for the worker this week if

³This property of my choice model also generalizes Blomquist et al. (2021), who discuss a bunching-design setup with nonparametric utility but with a scalar choice variable.

the worker’s normal wage rate applied to all of this week’s hours; and b) the number that the firm would choose if the worker’s overtime rate applied to all hours this week. I show that choice from a kinked choice set can be fully characterized by this pair of counterfactuals: agents either choose one of them or they choose the location of the kink. Bunching at the kink then directly identifies a feature of the joint distribution of the potential outcomes, allowing one to make statements about treatment effects purged of selection bias.⁴

While generalizing the choice model underlying the bunching design, I also propose a new approach to weakening assumptions about heterogeneity required by the method. Blomquist et al. (2021) emphasize that identification from bunching rests on assumptions regarding the distribution of heterogeneity that cannot be directly verified in the data. In my formulation, such assumptions take the form of extrapolating the marginal distributions of each of the two potential outcomes, which are both observed in a censored manner. To perform this extrapolation I impose a natural nonparametric shape constraint—*bi-log-concavity*—on the distribution of each potential outcome. Bi-log-concavity nests many previously proposed distributional assumptions for bunching analyses, and is in-part testable. The restriction affords partial identification of a conditional average treatment effect among units located at the kink, a parameter I call the “buncher ATE”. In the overtime context, the buncher ATE represents an average reduced-form wage elasticity of hours demand, which I then use to assess the overall average effect of the FLSA.

I also show that the data in the bunching design are informative about counterfactual policies that change the location or “sharpness” of a kink. To do so, I extend a characterization of bunching from Blomquist et al. (2015), and show that when combined with a general continuity relationship (Kasy, 2022) the result yields bounds on the derivative of bunching and mean hours with respect to policy parameters. I use this to evaluate proposed reforms to the FLSA: for example lowering the overtime threshold below 40 hours, or increasing the premium pay factor from 1.5 to 2.⁵

My results supplement other partial identification approaches recently proposed for the bunching design. Notably, the bounds I derive for the buncher ATE are substantially narrowed by making extrapolation assumptions separately for each of *two* counterfactuals. By contrast, existing approaches operate by constraining the distribution of a single scalar heterogeneity parameter, a simplification afforded by the isoelastic choice model. In the context of that model, Bertanha et al. (2020) and Blomquist et al. (2021) obtain bounds on the elasticity when the researcher is willing to put an explicit limit on how sharply the density of heterogeneous choices can rise or fall. My approach based on bi-log-concavity avoids the need to choose any such tuning parameters, and is applicable in the general choice model.

The empirical setting of overtime pay involves confronting two challenges that are not typical of

⁴This echoes Kline and Tartari’s 2016 approach to studying labor supply, but in reverse. They use observed marginal distributions of counterfactual choices to identify features of their joint distribution, assuming optimizing behavior.

⁵For example, the bill HR4728 introduced in 2021 would establish a 32 hour rule for overtime pay.

existing bunching-design analyses. Firstly, 40 hours is not an “arbitrary” point and bunching there could arise in part from factors other than it being the location of the kink. I use two strategies to estimate the amount of bunching that would exist at 40 absent the FLSA, and deliver clean estimates of the rule’s effect. My preferred approach exploits the fact that when a worker makes use of paid-time-off hours these do not count towards that week’s overtime threshold, shifting the location of the kink week-to-week in a plausibly idiosyncratic way. A second feature of the overtime setting is that work hours may not always be set unilaterally by one party: in principle either the firm or the worker could have control over a given worker’s schedule. I provide evidence that week-to-week variation in hours tends to be driven by firms, but show that even when bargaining weight between workers and firms is arbitrary and heterogeneous, bunching at 40 hours is informative about labor demand rather than supply.

Empirically, I find that the FLSA overtime rule does in fact reduce hours of work among hourly workers, despite the theoretical possibility that offsetting wage adjustments might eliminate any such effect (Trejo, 1991). My preferred estimate suggests that about one quarter of the bunching observed at 40 among hourly workers is due to the FLSA, and those working at least 40 hours work, on average, about 30 minutes less in a week than they would absent the time-and-a-half rule. Across specifications, I obtain estimates of the local wage elasticity of weekly hours demand near 40 hours in the range -0.04 to -0.19 , indicating that firms are fairly resistant to changing hours to avoid overtime payments. A back-of-the-envelope calculation using these effects suggests that FLSA regulation creates about 700,000 jobs (relative to an estimated 100 million non-exempt workers), despite a reduction in total hours.

The structure of the paper is as follows. Section 2 lays out a motivating conceptual framework for work hours that relates my approach to existing literature on overtime. Section 3 introduces the payroll data I use in the empirical analysis. In Section 4 I develop the generalized bunching-design approach, with Appendix A generalizing some of the supporting formal results. Section 5 applies these results to estimate effect of the FLSA overtime rule on hours worked, as well as the effects of proposed reforms to the FLSA. Section 6 discusses the empirical findings from the standpoint of policy objectives, and 7 concludes.

2 Conceptual framework

This section outlines a framework useful for reasoning about the determination of weekly hours, which motivates the identification strategy of Section 4. Readers primarily interested in the bunching design may wish to skip directly to that section.

The conceptual framework is centered around two observations from the data detailed in Section 3: weekly hours vary considerably between pay periods for an individual hourly worker, and a given

worker's hourly wage tends to change infrequently. I propose to view this as a two stage-process. In a first step, workers are hired with an hourly wage set along with an "anticipated" number of hours they will work per week. Then, with that hourly wage fixed in the short-run, final scheduling of hours is controlled by the firm and varies by week given shocks to the firm's demand for labor. Given the FLSA overtime rule and a worker's fixed wage, their employer thus faces a kinked cost schedule when choosing hours in a given week, as pictured in Figure 1.

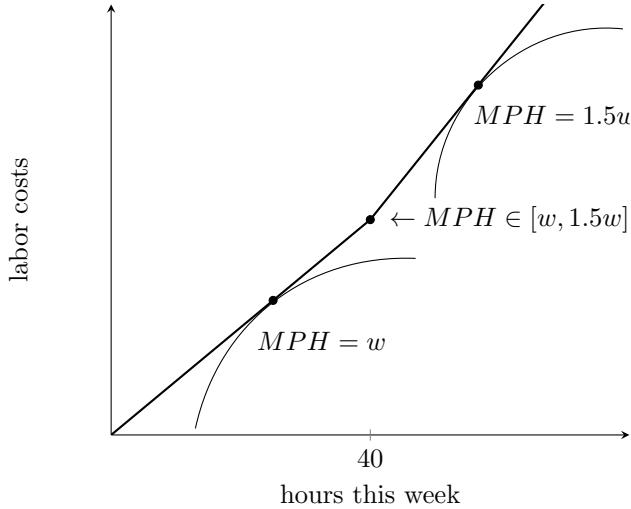


FIGURE 1: With a given worker's straight-time wage fixed at w , labor costs as a function of hours have a convex kink at 40 hours, given the overtime rule. Simple models of week-by-week hours choice (see Section 4.2) yield bunching when for a mass of workers the marginal product of an hour at 40 is between w and $1.5w$.

Wages and anticipated hours set at hiring

We begin with the hiring stage, which pins down the worker's wage. Throughout the analysis, I focus on workers paid on an hourly basis. Following the literature I refer to the hourly rate of pay that applies to the first 40 of a worker's hours as their *straight-time wage* or simply *straight wage*. This section provides a benchmark model to endogenize such straight wages, which yields predictions about how these wages may themselves be affected by the overtime rule. However, the basic bunching design strategy of Section 4 only requires that *some* straight-time wage is agreed upon and fixed in the short-run for each worker, as can be observed directly in the data.

Suppose that firms hire by posting an earnings-hours pair (z, h) , where z is total weekly compensation offered to each worker, and h is the advertised number of hours they will be each expected to work per week. The firm faces a labor supply function $N(z, h)$ determined by workers' preferences over the labor-leisure tradeoff,⁶ and makes a choice of (z^*, h^*) given this labor supply

⁶This labor supply function can be viewed as an equilibrium object that reflects both worker preferences and the

function and their production technology. For simplicity, workers are here taken to be homogeneous in production, paid hourly, and all covered by the overtime rule.⁷

While labor supply has above been viewed as a function over *total* compensation z and hours, there is always a unique straight wage associated with a particular (z, h) pair, such that h hours of work yields earnings of z , given the FLSA overtime rule:

$$w_s(z, h) := \frac{z}{h + 0.5 \cdot \mathbb{1}(h > 40)(h - 40)} \quad (1)$$

To distinguish between the two main views on the likely effects of overtime policy, let us suppose that a workers' straight-time wage is set according to Eq. (1), given values z^* and h^* that the firm and worker agree upon at the time of hiring. Trejo (1991) calls these two views the *fixed-job* and the *fixed-wage* models of overtime.

The *fixed-job* view observes that for a generic smooth labor supply function $N(z, h)$ (and smooth revenue production function with respect to hours), the optimal job package (z^*, h^*) for the firm to post will be *the same* as the optimal one absent the FLSA, as the hourly wage rate simply adjusts to fully neutralize the overtime premium.⁸ Suppose that workers then in fact work exactly h^* hours in all weeks (abstracting away from any reasons for the firm to deviate from h^* in a given week). Then the FLSA would have no effect on earnings, hours or employment, provided that $w_s(z^*, h^*)$ is above any applicable minimum wage (Trejo, 1991).

On the *fixed-wage* view, the firm instead faces an exogenous straight-time wage when determining (z^*, h^*) . Versions of this idea are considered in Brechling (1965), Rosen (1968), Ehrenberg (1971), Hamermesh (1993), Hart (2004) and Cahuc and Zylberberg (2014). This can be captured by a discontinuous labor supply function $N(z, h)$ that exhibits perfect competition on the quantity $w_s(z, h)$. I show in Appendix F.1 that in this case h^* and z^* are pinned down by the concavity of production with respect to hours and the scale of fixed costs (e.g. training for each worker) that do not depend on hours. The fixed-wage job makes the clear prediction that the FLSA will cause a reduction in hours, and bunching at 40.⁹

Existing work has investigated whether the fixed-job or fixed-wage model better accords with the observed joint distribution of hourly wages and hours (Trejo, 1991; Barkume, 2010). These competitive environment for labor. In Appendix F.2, I endogenize this function in a simple extension of the imperfectly competitive Burdett and Mortensen (1998) search model.

⁷I use the term "covered" to indicate workers whose employer is covered by the FLSA and who are not exempt from the overtime rule.

⁸In Appendix F.1 I give a closed-form expression for (z^*, h^*) when both labor supply and production are iso-elastic: hours and earnings are each increasing in the elasticity of labor supply with respect to earnings, and decreasing in the magnitude of the elasticity of labor supply with respect to pay.

⁹A fixed-wage model tends to predict an overall positive effect on employment given plausible assumptions on substitution between labor and capital (Cahuc and Zylberberg, 2014), though the total number of labor-hours will decrease (Hamermesh, 1993).

papers find that wages do tend to be lower among jobs that have overtime pay provisions and more overtime hours, but by a magnitude smaller than would be predicted by the fixed-jobs model. These estimates could however be driven by selection, e.g. of lower-skilled workers into covered jobs with longer hours. In Appendix C.2, I construct an empirical test of Equation (1) that is instead based on assuming that the conditional distribution of pay across paychecks is smooth across 40 hours. I find that roughly one quarter of paychecks around 40 hours reflect the wage/hours relationship predicted by the fixed-job model.

This finding is consistent with a model in which hours remain flexible week-to-week, while straight-wages remain fairly static over time after being initially set according to Equation (1).¹⁰ In common with the fixed wage model, this two-stage framework allows for the possibility that the overtime rule affects hours, and predicts bunching at 40. However, this is driven by short-run rigidity in straight-wages, rather than by perfect competition.

Dynamic adjustment to hours by week

Table 1 from the next section shows that workers' hours in my sample indeed vary considerably week-to-week. I assume that this week-level variation reflects choices made by their employers. There are many reasons to expect such variation: for instance shocks to product demand or productivity change the number of weekly hours that would be optimal that week from the firm's perspective. If demand for the firm's products is seasonal or volatile, it may not be worthwhile to hire additional workers only to reduce employment later. Similarly, variation in productivity across workers may only become apparent to supervisors after workers' straight wages have been set, and it may be profitable to increase the hours of the most productive ones.

Throughout Section 4, I maintain a strong version of the assumption that a firm—rather than a worker—chooses the final hours that I observe on each paycheck. In this benchmark model workers' preferences *do* matter in the determination of each worker's straight wage at hiring, but final scheduling rights week-to-week belong to the firm.¹¹ This simplification eases notation and emphasizes the intuition behind my identification strategy. Appendix D presents a generalization in which some fraction of workers choose their hours, along with intermediate cases in which the firm and worker bargain over hours each week. The results there show that if some workers have control of their weekly hours, the bunching-design strategy will only be informative about effects of the FLSA among workers whose final hours are chosen by the firm.

Available survey evidence suggests that this latter group is the dominant one: a relatively small

¹⁰This dovetails other recent evidence of uniformity and discretion in wage-setting, e.g. nominal wage rigidity (Grigsby et al., 2021), wage standardization (Hjort et al., 2020) and bunching at round numbers (Dube et al., 2020).

¹¹This can be rationalized on the basis of workers generally having less bargaining power: if the worker and firm fail to agree on a worker's hours, the worker's outside option may be unemployment while the firm's outside option is having one less worker (Stole and Zwiebel, 1996).

share of workers report that they choose their own schedules. For example, the 2017–2018 Job Flexibilities and Work Schedules Supplement of the American Time Use Survey asks workers whether they have some input into their schedule, or whether their firm decides it. Only 17% report that they have some input. In a survey of firms, only 10% report that most of their employees have control over which shifts they work (Matos et al., 2017).

3 Data and descriptive patterns

The main dataset I use comes from a large payroll processing company. They provided anonymized paychecks for the employees of 10,000 randomly sampled employers, for all pay periods in the years 2016 and 2017. At the paycheck level, I observe the check date, straight wage, and amount of pay and hours corresponding to itemized pay types, including normal (straight-time) pay, overtime pay, sick leave, holiday pay, and paid time off. The data also include state and industry for each employer and for employees: age, tenure, gender, state of residence, pay frequency and their salary if one is stored in the system.

3.1 Sample description

I construct a final sample for analysis based on two desiderata: a) the ability to observe hours within a single week; and b) a focus on workers who are non-exempt from the FLSA overtime rule. For the purposes of a), I drop paychecks from workers who are not paid on a weekly basis (roughly half of the workers in the sample). To achieve b) I keep paychecks only from hourly workers, since nearly all workers who are paid hourly are subject to the overtime rule. I also drop any workers who have no variation in hours or never receive overtime pay during the study period. The final sample includes 630,217 paychecks for 12,488 workers across 566 firms. See Appendix C.1 for further details of the sample construction.

Table 1 shows how the sample compares to survey data that is constructed to be representative of the U.S. labor force. Column (1) reports means from the final sample used in estimation, while (2) reports means before sampling restrictions. Column (3) reports means from the Current Population Survey (CPS) for the same years 2016–2017, among individuals reporting hourly employment. The “gets overtime” variable for the CPS sample indicates that the worker usually receives overtime, tips, or commissions. Column (4) reports means for 2016–2017 from the National Compensation Survey (NCS), a representative establishment-level dataset accessed on a restricted basis from the Bureau of Labor Statistics. The NCS reports typical overtime worked at the quarterly level for each job in an establishment (drawn from firm administrative data when possible).¹²

¹²The hourly wage variable for the CPS may mix straight-time and overtime rates, and is only present in outgoing

	(1) Estimation sample	(2) Initial sample	(3) CPS	(4) NCS
Tenure (years)	3.21	2.81	6.34	.
Age (years)	37.15	35.89	39.58	.
Female	0.23	0.46	0.50	.
Weekly hours	38.92	27.28	36.31	35.70
Gets overtime	1.00	0.37	0.17	0.52
Straight-time wage	16.16	22.17	18.09	23.31
Weekly overtime hours	3.56	0.94	.	1.04
Number of workers in sample	12488	149459	63404	228773

TABLE 1: Comparison of the sample with representative surveys. Columns 1 and 2 average across periods within worker from the administrative payroll sample, and then present means across workers. Column 2 presents means of worker-level data from the Current Population Survey and Column 3 averages representative job-level data from the National Compensation Survey.

The sample I use is somewhat more male, earns lower straight-time wages, and works more overtime than a typical hourly worker in the U.S. Column (2) in Table 1 reveals that my sampling restrictions can explain why the estimation sample tilts male and has higher overtime hours than the workforce as a whole. In particular, conditioning on workers that are paid on a weekly basis over-samples industries that tend to have more men, and tend to pay somewhat lower wages. Appendix C compares the industry and regional distributions of the estimation sample to the CPS.

3.2 Hours and wages in the sample

I turn now to the main variables to be used in the analysis. Figure 2 reports the distribution of hours of work in the final sample of paychecks. The graphs indicate a large mass of individuals who were paid for exactly 40 hours that week, amounting to about 11.6% of the sample.¹³ Appendix Figure 10 shows that overtime pay is present in nearly all weekly paychecks that report more than 40 hours, in line with the presumption that workers in the final sample are not FLSA-exempt.

Table 2 documents that while the hours paid in 70% of all pay checks in the final estimation sample differ from those of the last paycheck by at least one hour, just 4% of all paychecks record a different straight-time wage than the previous paycheck for the same worker. Among the roughly 22,500 wage change events, the average change is about a 45 cent raise per hour. When hours change the magnitude is about 7 hours on average, and roughly symmetric around zero (see Appendix Figure 12 for the distribution of hours changes).

rotation groups. The tenure variable comes from the 2018 Job Tenure Supplement. The NCS does not distinguish between hourly and salaried workers, reporting an average hourly rate that includes salaried workers, who tend to be paid more. This likely explains the higher value than the CPS and payroll samples.

¹³The second largest mass occurs at 32 hours, and is explained by paid-time-off, holiday, and sick pay hours as discussed in Section 5.

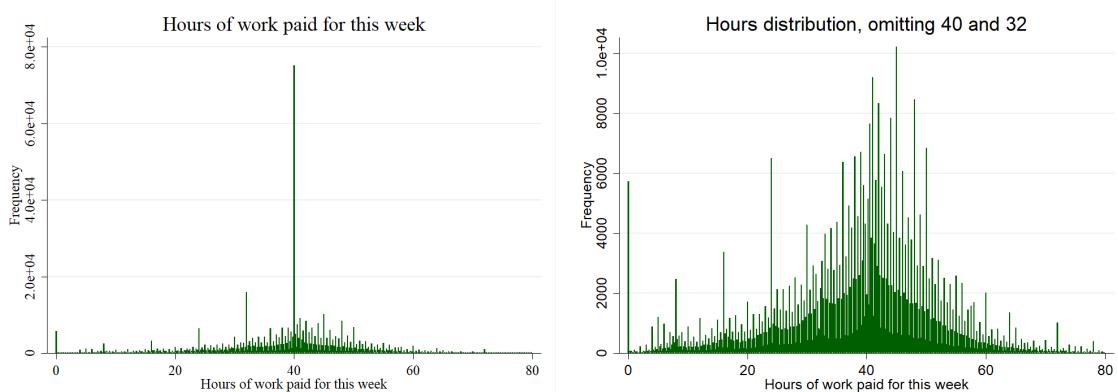


FIGURE 2: Empirical densities of hours worked pooling all paychecks in final estimation sample. Sample is restricted to hourly workers receiving overtime pay at some point (to ensure nearly all are non-exempt from FLSA, see text), and workers having hours variation. The right panel omits the points 40 and 32 to improve visibility elsewhere. Bins have a width of 1/8 of an hour, based on the observed granularity of hours (see Appendix Figure 12 for details).

	Mean	Std. dev.	N
Indicator for hours changed from last period	0.84	0.37	630,217
Indicator for hours changed by at least 1 hour	0.70	0.46	630,217
Indicator for wage changed from last period	0.04	0.19	630,217
Indicator for wage changed, if hours changed	0.04	0.19	529,791
Absolute value of hours difference, if hours changed	6.83	8.23	529,791
Difference in wage, if wage changed	0.45	26.46	22,501

TABLE 2: Changes in hours or straight wages between a worker's consecutive paychecks.

Appendix C reports some further details on the variation in hours and wages present in the data. Appendix Table 2 regresses hours, overtime hours, and an indicator for bunching on worker observables, and shows that after controlling for worker and date fixed effects bunching and overtime hours are both predicted by recent hiring at the firm, lending further evidence for the assumption that shocks to labor demand drive variation in hours. Appendix Table 3 shows that overall, about 63% of variation in total hours can be explained by worker and employer-by-date fixed effects. Appendix Table 1 documents heterogeneity in the prevalence of overtime pay across industry classifications. Appendix Figure 5 studies the joint distribution of wages and hours and reproduces Bick et al.'s (2022) finding that wages increase with hours until just beyond 40, then decline some.

4 Empirical strategy: a generalized kink bunching design

Let us now turn to the firm choosing the hours of a given worker in a particular week, with that worker's wage fixed and costs a kinked function of hours as depicted in Figure 1. This section shows

that under weak assumptions, firms facing such kinks will make choices that can be completely characterized by choices they *would* make under two counterfactual linear cost schedules that differ with respect to wage. I relate the observable bunching at 40 hours to a treatment effect defined from these two counterfactuals, which I then use to estimate the impact of the FLSA on hours.

The identification results in this section hold in a much more general setting in which a decision-maker faces a choice set with a possibly multivariate kink and has “nearly” convex preferences. I present the general version of this model in Appendix A. Throughout the present section I refer to a worker i in week t as a *unit*: an observation of h_{it} for unit it is thus the hours recorded on a single paycheck. Probability statements are to be understood with respect to such paycheck-level units.

4.1 A general choice model

Let us start from the conceptual framework introduced in Section 2. In choosing the hours h_{it} of worker i in week t , worker i ’s employer faces a kinked cost schedule, given the worker’s straight-time wage this week w_{it} . If the firm chooses less than 40 hours, it will pay $w = w_{it}$ for each hour, and if the firm chooses $h > 40$ it will pay $40w$ for the first 40 hours and $1.5w(h - 40)$ for the remaining hours, giving the convex shape to Figure 1. We can write the kinked pay schedule for unit it , as a function of hours this week h , as

$$B_{kit}(h) = w_{it}h + .5w_{it}\mathbb{1}(h > 40)(h - 40) = \max\{B_{0it}(h), B_{1it}(h)\}$$

where $B_{0it}(h) = w_{it}h$ and $B_{1it}(h) = 1.5w_{it}h - 20w_{it}$. The kinked pay schedule $B_{kit}(h)$ is equal to $B_{0it}(h)$ for values $h \leq 40$ and $B_{kit}(h)$ is equal to $B_{1it}(h)$ for values $h \geq 40$. The functions B_0 and B_1 recover the two segments in Figure 1 when restricted to these domains respectively (see Appendix Figure 1). The following definition is generalized in Appendix A:

Definition (potential outcomes). Let h_{0it} denote the hours of work that of unit it would be paid for if instead of $B_{kit}(h)$, the pay schedule for week t ’s hours were $B_{0it}(h)$. Similarly, let h_{1it} denote the hours of pay that would occur for unit it if the pay schedule were $B_{1it}(h)$.

The potential outcomes h_0 and h_1 thus imagine what would happen if instead of the kinked piecewise pay schedule $B_k(h)$, one of $B_0(h)$ or $B_1(h)$ applied globally for all values of h .

Let h_{it} denote the actual hours for which unit it is paid. Our first assumption is that actual hours and potential outcomes reflect choices made by the firm:

Assumption CHOICE. Each of h_{0it} , h_{1it} and h_{it} reflect choices the firm would make under the pay schedules $B_{0it}(h)$, $B_{1it}(h)$, and $B_{kit}(h)$ respectively.

CHOICE reflects the assumption that hours are perfectly manipulable by firms. Note that if firm preferences over a unit’s hours are quasi-linear with respect to costs (e.g. if they maximize weekly

profits), the term $-20w_{it}$ appearing in B_{1it} plays no role in firm choices. As such, I will often refer to h_{1it} as choice made under linear pay at the overtime rate $1.5w_{it}$, keeping in mind that the exact definition for B_1 given above is necessary for the interpretation if preferences are not quasi-linear.

Our second assumption is that each unit's firm optimizes some vector \mathbf{x} of choice variables that pin down that unit's hours. As a leading case, we may think of hours of work as a single component of firms' choice vector \mathbf{x} (Appendix A.3 gives some examples of this). Firm preferences are taken to be convex in \mathbf{x} and the unit's wage costs z :

Assumption CONVEX. *Firm choices for unit it maximize some $\pi_{it}(z, \mathbf{x})$, where π_{it} is strictly quasiconcave in (z, \mathbf{x}) and decreasing in z . Hours are a continuous function of \mathbf{x} for each unit.*

For the sake of brevity, I here state a version of CONVEX that is a bit stronger than necessary for the identification results below. In particular, Appendix A relaxes CONVEX to allow for “double-peaked” preferences with one peak located exactly at the kink (this is relevant if firms have a special preference for a 40 hour work week). The appendix also shows that bunching still has some identifying power under no assumptions about convexity of preferences. The assumption that firms rather than workers choose hours enters in the claim that π is decreasing (rather than increasing) in z , but Appendix D relaxes this to allow some workers to set their hours.

Observables in the bunching design

The starting point for our analysis of identification in the bunching design is the following mapping between actual hours h_{it} and the counterfactual hours choices h_{0it} and h_{1it} . Appendix Lemma 1 shows that Assumptions CHOICE AND CONVEX imply that:

$$h_{it} = \begin{cases} h_{0it} & \text{if } h_{0it} < 40 \\ 40 & \text{if } h_{1it} \leq 40 \leq h_{0it} \\ h_{1it} & \text{if } h_{1it} > 40 \end{cases} \quad (2)$$

That is, a worker will work h_{0it} hours when the counterfactual choice h_{0it} is less than 40, and h_{1it} hours when h_{1it} is greater than 40. They will be found at the corner solution of 40 if and only if the two counterfactual outcomes fall on either side, “straddling” the kink.¹⁴ Figure 3 depicts the implications of Eq. (2) for what is therefore observable by the researcher in the bunching design: censored distributions of both h_0 and of h_1 , and a point-mass of size $\mathcal{B} = P(h_{1it} \leq 40 \leq h_{0it})$ at the kink.

¹⁴“Straddling” can only occur in one direction, with $h_{1it} \leq k \leq h_{0it}$. The other direction: $h_{0it} \leq k \leq h_{1it}$ with at least one inequality strict, is ruled out by the weak axiom of revealed preference (see Appendix A).

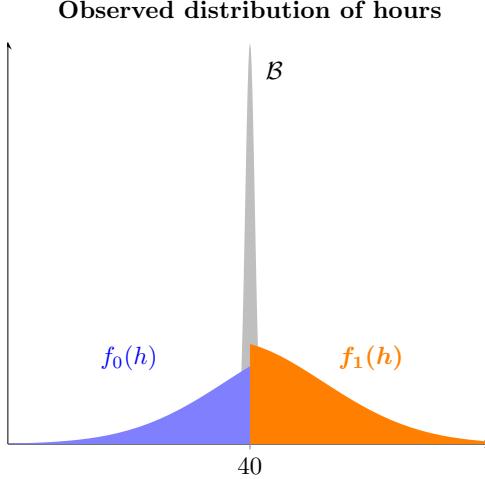


FIGURE 3: Observables in the bunching design, given Equation (2). To the left of the kink at 40, the researcher observes the density $f_0(h)$ of the counterfactual h_{0it} , up to values $h = 40$. To the right of the kink, the researcher observes the density $f_1(h)$ of h_{1it} for values $h > 40$. At the kink, one observes a point-mass of size $\mathcal{B} := P(h_{it} = 40) = P(h_{1it} \leq 40 \leq h_{0it})$.

Equation (2) represents a key departure from previous approaches to the bunching design, which characterize bunching in terms of the counterfactual h_0 only.¹⁵ I show below that such is a simplification afforded by the benchmark isoelastic utility model, but in a generic choice model, both h_0 and h_1 are necessary to pin down actual choices h_{it} . Appendix A shows that Eq. (2) also holds in settings with possibly non piecewise-linear kinked choice sets of the form: $z \geq \max\{B_0(\mathbf{x}), B_1(\mathbf{x})\}$ where B_0 and B_1 are weakly convex in the full vector \mathbf{x} , and z any “cost” decision-makers dislike.

Intuition for Equation (2) in the overtime setting

As an intuitive illustration of Equation (2), suppose that firms balance the cost $B_{kit}(h)$ against the value of h hours of the worker’s labor, in order to maximize that week’s profits. Then Eq. (2) can be written:

$$h_{it} = \begin{cases} MPH_{it}^{-1}(w_{it}) & \text{if } MPH_{it}(40) < w_{it} \\ 40 & \text{if } MPH_{it}(40) \in [w_{it}, 1.5w_{it}] \\ MPH_{it}^{-1}(1.5w_{it}) & \text{if } MPH_{it}(40) > 1.5w_{it} \end{cases} \quad (3)$$

where denotes $MPH_{it}(h)$ is the marginal product of an hour of labor for unit it , as a function of that unit’s hours h . Assuming that production is strictly concave, the function $MPH_{it}(h)$ will be strictly decreasing in h , and we have that $h_{0it} = MPH_{it}^{-1}(w_{it})$ and $h_{1it} = MPH_{it}^{-1}(1.5w_{it})$.

¹⁵Blomquist et al. (2015) also derive an expression for \mathcal{B} in terms of agents’ choices given all intermediate slopes between those occurring on either side of the kink. I discuss this and offer a generalization in Appendix Lemma 2.

Figure 1 depicts Eq. (3) visually. Consider for example a worker with a straight-wage of \$10 an hour. If there exists a value $h < 40$ such that the worker's MPH is equal to \$10, then the firm will choose this point of tangency. This happens if and only if the marginal product of an hour at 40 hours this week is less than \$10. If instead, the marginal product of an hour is still greater than \$15 at $h = 40$, the firm will choose the value $h > 40$ such that MPH equals \$15. The third possibility is that the MPH at $h = 40$ is *between* the straight and overtime rates \$10 and \$15. In this case, the firm will choose the corner solution $h = 40$, contributing to bunching at the kink.

While Eq. (3) provides a natural nonparametric characterization of when the firm will ask a worker to work overtime (when the ratio of productivity to wages is high), it is still more restrictive than necessary for the purposes of the bunching design. Appendix A.3 provides some examples that use the full generality of Assumption CONVEX, in which firms simultaneously consider *multiple* margins of choice aside from a given unit's hours. For example, the firm may attempt to mitigate the added cost of overtime by reducing bonuses when a worker works many overtime hours. Eq. (2) remains valid even when such additional margins of choice are unmodeled and unobserved by the econometrician, varying possibly by unit.

Note that if production depends jointly on the hours of all workers within a firm, we may expect the function $MPH_{it}(h)$ in Eq. (3) to depend on the hours of worker i 's colleagues in week t . In this case the quantities h_{0it} and h_{1it} hold the hours of i 's colleagues fixed at their *realized* values: they contemplate *ceteris paribus* counterfactuals in which the pay schedule for a single unit it is varied, and nothing else. In the baseline isoelastic model that we consider in the next section, such interdependencies between workers' hours are ruled out by assuming that production is linearly separable across units. Section 4.4 considers how in general, interdependencies affect the interpretation of our treatment effects, while Appendix E discusses the impact of nonseparable production functions in more detail.

4.2 The benchmark isoelastic model

The canonical approach from the bunching-design literature (Saez, 2010; Chetty et al., 2011; Kleven, 2016), strengthens Assumption CONVEX to suppose that $\mathbf{x} = h$ and decision-makers' utility follows an isoelastic functional form, with preferences identical between units up to a scalar heterogeneity parameter. This corresponds to a model in which firm profits from unit it are:

$$\pi_{it}(z, h) = a_{it} \cdot \frac{h^{1+\frac{1}{\epsilon}}}{1 + \frac{1}{\epsilon}} - z \quad (4)$$

where $\epsilon < 0$ is common across units, and z represents wage costs for worker i in week t . Eq. (4) is analogous to the isoelastic, quasilinear labor *supply* model used in the context of tax kinks.

Under a linear pay schedule $z = wh$, the profit maximizing number of hours is $(w/a_{it})^\epsilon$, so ϵ yields the elasticity of hours demand with respect to a linear wage. Let $\eta_{it} = a_{it}/w_{it}$ denote the ratio of a unit's current productivity factor a_{it} to their straight wage. In the isoelastic model

$$h_{0it} = MPH_{it}^{-1}(w_{it}) = \eta_{it}^{-\epsilon} \quad \text{and} \quad h_{1it} = MPH_{it}^{-1}(1.5w_{it}) = 1.5^\epsilon \cdot \eta_{it}^{-\epsilon},$$

and by Eq. (3) actual hours h_{it} are ranked across units in order of η_{it} . If η_{it} is continuously distributed with support overlapping the interval $[40^{-1/\epsilon}, 1.5 \cdot 40^{-1/\epsilon}]$, then the observed distribution of h_{it} will feature a point mass at 40—“bunching”—and a density elsewhere. In the isoelastic model whether a worker works overtime in a given week is determined by the scalar η_{it} : a worker with a wage w_{it} fixed throughout the year may for example work overtime only in periods when a_{it} is relatively high due to seasonally elevated productivity.

Identification in the isoelastic model

In the context of the isoelastic model, a natural starting place for evaluating the FLSA is to estimate the parameter ϵ . Ignoring for the moment any effects of the policy on straight-wages, the effect of the time-and-a-half rule on unit it 's hours will simply be the difference $h_{it} - h_{0it}$, what we might call the *effect of the kink*. It follows from the above that the effect of the kink is $h_{it} \cdot (1 - 1.5^{-\epsilon})$ for any unit such that $h_{it} > 40$. Provided the value of ϵ , we could thus evaluate the effect of the time-and-a-half rule for any paycheck recording overtime using that unit's observed hours.

The classic bunching-design method pioneered by Saez (2010) identifies ϵ by relating it to the observable bunching probability:

$$\mathcal{B} := P(h_{it} = 40) = \int_{40}^{1.5^{|\epsilon|} \cdot 40} f_0(h) \cdot dh \tag{5}$$

where f_0 is the density of h_0 . If the function f_0 were known, the value of ϵ could be pinned down by simply solving Eq. (5) for $|\epsilon|$. However, f_0 is not globally identified from the data: from Figure 3 we can see that f_0 is only identified to the left of the kink, while the density of h_1 is identified to the right of the kink. Since $h_{1it} = 1.5^\epsilon \cdot h_{0it}$, it is convenient in the isoelastic model to analyze observables after applying a log transformation to hours: the quantity $\delta = \ln h_{0it} - \ln h_{1it} = |\epsilon| \cdot \ln 1.5$ is homogeneous across all units it , and the density of $\ln h_{1it}$ is thus a simple leftward shift of the density of $\ln h_{0it}$, by δ , as shown in Figure 4.

Standard approaches in the bunching design make parametric assumptions that interpolate f_0 through the missing region of Figure 4 to point-identify ϵ .¹⁶ The approach of Saez (2010) assumes

¹⁶Bertanha et al. (2020) note that given a full parametric distribution for f_0 , the entire model could be estimated by maximum likelihood. This approach would enforce (5) automatically while enjoying the efficiency properties of MLE.

for example that the density of h_0 is linear through the missing region $[40, 40 \cdot e^\delta]$ of Figure 4. The popular method of Chetty et al. (2011) instead fits a global polynomial, using the distribution of hours outside the missing region to impute the density of h_0 within it. Neither approach is particularly suitable in the overtime context. The linear method of Saez (2010) implies monotonicity of the density in the missing region, which is unlikely to hold given that 40 appears to be near the mode of the h_0 latent hours distribution. Meanwhile, the method of Chetty et al. (2011) ignores the “shift” by δ in the right panel of Figure 4. Both of these approaches rely on parametric assumptions, and sufficient conditions for each are outlined in Appendix G.2.

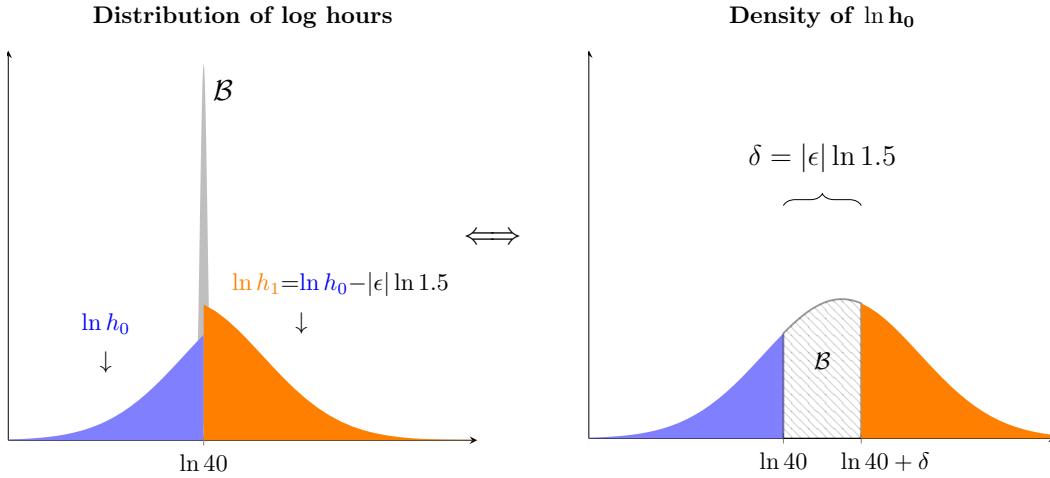


FIGURE 4: The left panel depicts the distribution of observed log hours $\ln h_{it}$ in the isoelastic model, while the right panel depicts the underlying full density of $\ln h_{0it}$. Specializing from the general setting of Figure 3, we have in the isoelastic model that $f_1(h) = f_0(h + |\epsilon| \cdot \ln 1.5)$. Thus, the full density of f_0 is related to the observed distribution by “sliding” the observed distribution for $h > 40$ to the right by the unknown distance $\delta = |\epsilon| \ln 1.5$, leaving a missing region in which f_0 is unobserved. The total area in the missing region from $\ln 40$ to $\ln 40 + \delta$ must equal the observed bunching mass \mathcal{B} .

If in the other extreme, the researcher is unwilling to assume anything about the density of h_0 in the missing region of Figure 4, then the data are compatible with any finite $\epsilon < 0$, a point emphasized by Blomquist et al. (2021) and Bertanha et al. (2020). In particular, given (5), an arbitrarily small $|\epsilon|$ could be rationalized by a density that spikes sufficiently high just to the right of 40, while an arbitrarily large $|\epsilon|$ can be reconciled with the data by supposing that the density of h_0 drops quickly to some very small level throughout the missing region. I find a middle ground by imposing a nonparametric shape constraint on h_0 : *bi-log-concavity* (BLC), leading to partial identification. A detailed discussion of BLC is given in Section 4.3.

Limitations of the isoelastic model

Compared with the isoelastic model, the general choice model from Section 4.1 allows a wide range of underlying choice models that might drive a firm's hours response to the FLSA. This robustness over structural models is important in the overtime context. As reported in Appendix C.5, assuming the isoelastic model and that h_0 and h_1 are BLC (which nests a linear density as a special case) suggests that $\epsilon \in [-.179, -.168]$.¹⁷ Such values are implausible when interpreted through the lens of Equation (4): $\epsilon = -.2$ for example would imply an hours production function of $f(h) = -\frac{1}{4}h^{-4}$ (up to an affine transformation), which features an unrealistic degree of concavity. Allowing a more general production function $f(h)$ (separable between units) is also not much help, as the standard bunching design approach then estimates an averaged local inverse elasticity of $f(h)$ (see Appendix C.5).

In short, the observed bunching at 40 hours is too small to be reconciled with a model in which ϵ parameterizes the concavity of weekly production with respect to hours. This motivates a model like the one presented in Section 4.1, in which we can interpret the estimand of the bunching design as a *reduced-form* elasticity of the demand for hours. As described through some examples in Appendix A.3, this elasticity may reflect adjustment by firms along additional margins that can attenuate the hours response, and thus reduce the magnitude of bunching.

4.3 Identifying treatment effects in the general choice model

In this section I turn to identification in the general choice model of Section 4.1. Without a single preference parameter like ϵ that characterizes responsiveness to incentives for all units, we face the question of what quantity might be identifiable from the data without the restrictive isoelastic model, but still help us to evaluate the effect of the FLSA on hours.

Let us refer to the difference $\Delta_{it} := h_{0it} - h_{1it}$ between h_0 and h_1 as unit i 's *treatment effect*. Recall that h_0 and h_1 are interpreted as potential outcomes, indicating what *would* have happened had the firm faced either of two counterfactual pay schedules instead of the kink. Δ_{it} thus represents the causal effect of a one-period 50% increase in worker i 's wage on their hours in week t : the difference between the hours that unit's firm would choose if the worker were paid at their straight-time rate versus at their overtime rate for all hours in that week (assuming quasi-linearity of firm preferences). As such we would expect that $\Delta_{it} \geq 0$. A unit's treatment effect can be contrasted with the "effect of the kink" quantity $h_{it} - h_{0it}$ introduced before, but importantly the two are related: by Eq. (2) the effect of the kink is $-\Delta_{it}$ for all units working overtime.

¹⁷The width of these bounds is about 4 times smaller than if BLC is assumed for h_0 only. These estimates attribute all of the bunching observed at 40 to the FLSA: attributing just a portion of the bunching at 40 to the FLSA (as I do in Section 5.1) would only further reduce the magnitude of ϵ . Industry-specific bounds on ϵ range from -0.26 to -0.06 .

In the isoelastic model $\Delta_{it} = h_{0it} \cdot (1 - 1.5^\epsilon)$, representing a special case in which treatment effects are homogenous across units after a log transformation of the outcome: $\ln h_{0it} - \ln h_{1it} = |\epsilon| \cdot \ln 1.5$. In general we can expect Δ_{it} to vary much more flexibly across units, and a reasonable parameter of interest becomes a summary statistic of Δ_{it} of some kind. In particular, Eq. (2) suggests that bunching is informative about the distribution of Δ_{it} among units “near” the kink. To see this, let $k = 40$ denote the location of the kink, and write the bunching probability as:

$$\mathcal{B} = P(h_{1it} \leq k \leq h_{0it}) = P(h_{0it} \in [k, k + \Delta_{it}]) = P(h_{1it} \in [k - \Delta_{it}, k]), \quad (6)$$

i.e. units bunch when their h_0 potential outcome lies to the right of the kink, but within that unit’s individual treatment effect of it. Note that by Eq. (2) we can also write bunching in terms of the marginal distributions of h_0 and h_1 : $\mathcal{B} = F_1(k) - F_0(k)$, provided that each potential outcome is continuously distributed and with F_0 and F_1 their cumulative distribution functions.

Parameter of interest: the buncher ATE

I focus my identification analysis on the average treatment effect among units who locate at exactly 40 hours, a parameter I call the “buncher ATE”. In the overtime setting some care is needed in defining this parameter, allowing for the possibility that a mass of units would still work exactly 40 hours, even absent the FLSA. Let us indicate such “counterfactual bunchers” by an (unobserved) binary variable $K_{it}^* = 1$, and define the buncher ATE to be:

$$\Delta_k^* = \mathbb{E}[\Delta_{it}|h_{it} = k, K_{it}^* = 0],$$

That is, Δ_k^* is the average value of Δ_{it} among bunchers who bunch in response to the FLSA kink, and would not locate at 40 hours otherwise. In evaluating the FLSA, I suppose that all counterfactual bunchers have a zero treatment effect, such that $h_{0it} = h_{1it} = k$. Since $\Delta_{it} = 0$ for these units by assumption, we can move back and forth between Δ_k^* and $\mathbb{E}[\Delta_{it}|h_{it} = k]$, provided the counterfactual bunching mass $p := P(K_{it}^* = 1)$ is known. In this section, I treat p as given, and present a strategy estimate it empirically in Section 5.1.

To simplify the discussion, suppose for the moment that there are no counterfactual bunchers, so that $\Delta_k^* = \mathbb{E}[\Delta_{it}|h_{it} = k]$. Our goal is to invert (6) in some way to learn about the buncher ATE from the observable bunching probability \mathcal{B} . In Figure 4, we’ve seen the intuition for this exercise in the context of the isoelastic model, in which there is only a scalar dimension of heterogeneity and $h_{1it} = h_{0it} \cdot 1.5^\epsilon$. The key implication of the isoelastic model that aids in identification is *rank invariance* between h_0 and h_1 . Rank invariance (Chernozhukov and Hansen 2005) says that $F_0(h_{0it}) = F_1(h_{1it})$ for all units, i.e. increasing each unit’s wage by 50% does not change any

unit's rank in the hours distribution (for example, a worker at the median of the h_0 distribution also has a median value of h_1). Rank invariance is satisfied by models in which there is perfect positive co-dependence between the potential outcomes (left panel of Figure 5).

Rank invariance is useful because it allows us to translate statements about Δ_{it} into statements about the *marginal* distributions of h_{0it} and h_{1it} . In particular, under rank invariance the buncher ATE is equal to the quantile treatment effect $Q_0(u) - Q_1(u)$ averaged across all u between $F_0(k)$ and $F_1(k) = F_0(k) + \mathcal{B}$, where Q_d is the quantile function of h_{dit} , i.e.:

$$\Delta_k^* = \frac{1}{\mathcal{B}} \int_{F_0(k)}^{F_1(k)} [Q_0(u) - Q_1(u)] du, \quad (7)$$

so long as $F_0(y)$ and $F_1(y)$ are continuous and strictly increasing. I focus on partial identification of the buncher ATE, for which it is sufficient to place point-wise bounds on the quantile functions $Q_0(u)$ and $Q_1(u)$ throughout the range $u \in [F_0(k), F_1(k)]$ as depicted in Figure 6.

While rank invariance already relaxes the isoelastic model used thus far in the literature, a still weaker assumption proves sufficient for Eq. (7) to hold:

Assumption RANK. *There exist fixed values Δ_0^* and Δ_1^* such that $h_{0it} \in [k, k + \Delta_{it}]$ iff $h_{0it} \in [k, k + \Delta_0^*]$, and $h_{1it} \in [k - \Delta_{it}, k]$ iff $h_{1it} \in [k - \Delta_1^*, k]$.*

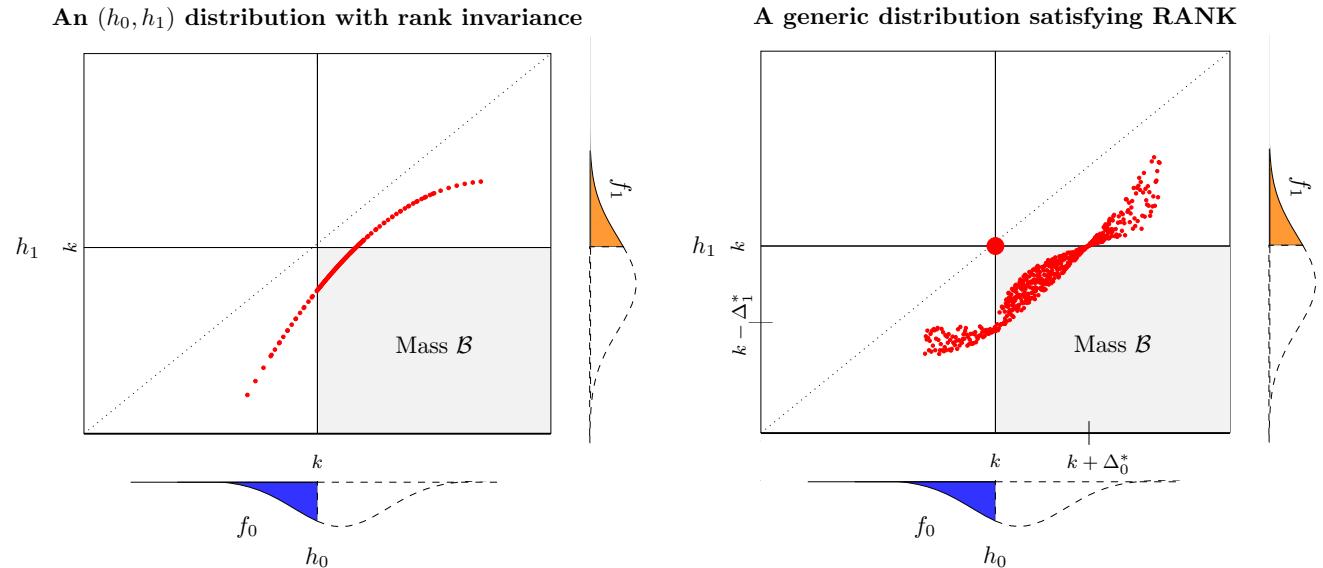


FIGURE 5: The joint distribution of (h_{0it}, h_{1it}) (in red), comparing an example satisfying rank invariance (left) to a case satisfying Assumption RANK (right). RANK allows the support of the joint distribution to “fan-out” from perfect co-dependence of h_0 and h_1 , except when either outcome is equal to k . The large dot in the right panel indicates a possible mass p of counterfactual bunchers. The observable data identifies the shaded portions of each outcome’s marginal distribution (depicted along the bottom and right edges), as well as the total mass \mathcal{B} in the (shaded) south-east quadrant.

Unlike (strict) rank invariance, Assumption RANK allows ranks to be reshuffled by treatment among bunchers and among the group of units that locate on each side of the kink.¹⁸ For example, suppose that a 50% increase in the wage of worker i would result in their hours being reduced from $h_{0it} = 50$ to $h_{1it} = 45$. If another worker j 's hours are instead reduced from $h_{0jt} = 48$ to $h_{1jt} = 46$ under a 50% wage increase, workers i and j will switch ranks, without violating RANK. Note that RANK is also compatible with the existence of counterfactual bunchers $p > 0$.

The right panel of Figure 5 shows an example of a distribution satisfying RANK, which requires the support of (h_0, h_1) to narrow to a point as it crosses $h_0 = k$ or $h_1 = k$. If this is not perfectly satisfied, Appendix A.5 demonstrates how the RHS of Equation (7) will then yield a lower bound on the true buncher ATE (and can still be interpreted as an averaged quantile treatment effect). Appendix Figure 15 generalizes RANK to case in which some workers choose their hours, resulting in mass also appearing in the north-west quadrant of Figure 5.

4.3.1 Bounds on the buncher ATE via bi-log-concavity

Given Eq. (7), I obtain bounds on the buncher ATE by assuming that both h_0 and h_1 have *bi-log-concave* distributions. Bi-log-concavity is a nonparametric shape constraint that generalizes log-concavity, a property of many familiar parametric distributions:

Definition (BLC). A distribution function F is *bi-log-concave (BLC)* if both $\ln F$ and $\ln(1 - F)$ are concave functions.

If F is BLC then it admits a strictly positive density f that is itself differentiable with locally bounded derivative: $\frac{-f(h)^2}{1-F(h)} \leq f'(h) \leq \frac{f(h)^2}{F(h)}$ (Dümbgen et al., 2017). Intuitively, this rules out cases in which the density of h_0 or h_1 ever spikes or falls *too* quickly on the interior of its support, leading to non-identification of the type discussed in Section 4.2.¹⁹

The family of BLC distributions includes parametric distributions assumed by previous bunching design studies, such as those with uniform or linear densities (Saez, 2010), or those with polynomial densities as in Chetty et al. 2011 (provided they have real roots). All globally log-concave distributions are BLC, though BLC distributions do not need to be unimodal (as log-concave distributions do). Importantly, the BLC property is partially testable in the bunching design, since $F_0(y)$ is identified for all $h < k$ and $F_1(h)$ is identified for all $h > k$. Appendix Figure 9 shows that the observable portions of F_0 and F_1 indeed satisfy BLC. Identification requires us to believe that BLC *also* holds in the unobserved portions of F_0 and F_1 .

¹⁸When $p = 0$ Assumption RANK is equivalent to an instance of the *rank-similarity* assumption of Chernozhukov and Hansen (2005), in which the conditioning variable is which of the three cases of Equation (2) hold for the unit. Specifically, for both $d = 0$ and $d = 1$: $U_d|(h < k) \sim \text{Unif}[0, F_0(k)]$, $U_d|(h = k) \sim \text{Unif}[F_0(k), F_1(k)]$, and $U_d|(h > k) \sim \text{Unif}[F_1(k), 1]$.

¹⁹Bertanha et al. (2020) propose bounds in the isoelastic model by specifying a Lipschitz constant on the density of $\ln \eta_{it}$. This yields global rather than local bounds on f' , based on a tuning parameter value that must be chosen.

We are now ready to state the main identification result. Its logic is summarized by Figure 6: given the general choice model, RANK converts identification of the buncher ATE into a pair of extrapolation problems, each of which are approached by assuming bi-log-concavity of the corresponding marginal potential outcome distribution. Let $F(h) := P(h_{it} \leq h)$ be the CDF of observed hours.

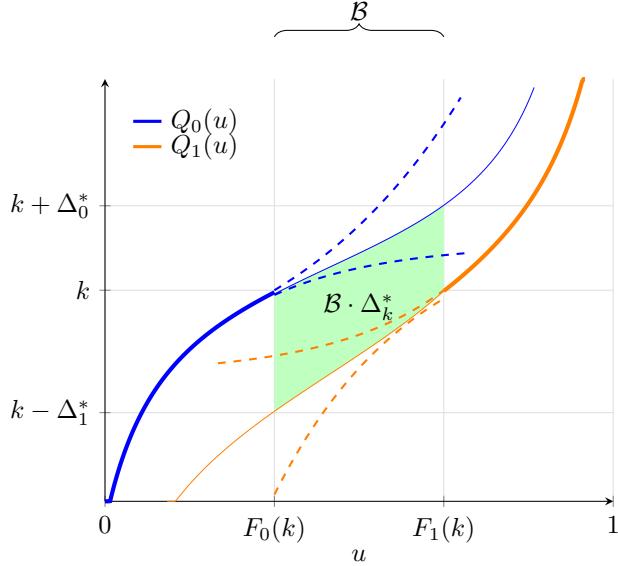


FIGURE 6: Extrapolating the quantile functions for h_0 and h_1 (blue and orange, respectively) to place bounds on the buncher ATE (case depicted has no counterfactual bunchers). The observed portions of each quantile function are depicted by thick curves, while the unobserved portions are indicated by thinner curves. The dashed curves represent upper and lower bounds for this unobserved portion coming from an assumption such as bi-log-concavity. The buncher ATE is equal to the area shaded in green, divided by the bunching probability \mathcal{B} .²⁰ The quantities Δ_0^* and Δ_1^* are defined in Assumption RANK below.

Theorem 1 (bi-log-concavity bounds on the buncher ATE). Assume CHOICE, CONVEX, RANK and that h_{0it} and h_{1it} have bi-log-concave distributions conditional on $K_{it}^* = 0$. Then:

1. $F(h)$, $F_0(h)$ and $F_1(h)$ are continuously differentiable for $h \neq k$. $F_0(k) = \lim_{h \uparrow k} F(h) + p$, $F_1(k) = F(k)$, $f_0(k) = \lim_{h \uparrow k} f(h)$ and $f_1(k) = \lim_{h \downarrow k} f(h)$, where if $p > 0$ we define the density of h_{dit} at $y = k$ to be $f_d(k) = \lim_{h \rightarrow k} f_d(h)$, for each $d \in \{0, 1\}$.
2. The buncher ATE Δ_k^* lies in the interval $[\Delta_k^L, \Delta_k^U]$, where:

$$\Delta_k^L := g(F_0(k) - p, f_0(k), \mathcal{B} - p) + g(1 - F_1(k), f_1(k), \mathcal{B} - p)$$

and

$$\Delta_k^U := -g(1 - F_0(k), f_0(k), p - \mathcal{B}) - g(F_1(k) - p, f_1(k), p - \mathcal{B})$$

with $g(a, b, x) = \frac{a}{bx} (a + x) \ln(1 + \frac{x}{a}) - \frac{a}{b}$. The bounds Δ_k^L and Δ_k^U are sharp.

Proof. See Appendix B. □

Combining Items 1 and 2 of Theorem 1, it follows that the sharp bounds Δ_k^L and Δ_k^U on the buncher ATE are identified, given the CDF $F(h)$ of hours and p .²¹ Inspection of the expressions appearing in Theorem 1 reveals that the bounds become wider the larger the net bunching probability $\mathcal{B} - p$. If $f_0(k) \approx f_1(k)$ and $p \approx 0$, the bounds will tend to be narrower when $F_0(k)$ is closer to $(1 - \mathcal{B})/2$, i.e. the kink is close to the median of the latent hours distribution. This helps explain why the estimated bounds in Section 5 turn out to be quite informative.

Comparison to existing results. The existing bunching design literature does contain a few results circling the common intuition that when responsiveness to incentives varies by observational unit, bunching is informative about a local average responsiveness. For instance, Saez (2010) and Kleven (2016) consider a “small-kink” approximation that $\mathbb{E}[\Delta_{it}|h_{0it} = k] \approx \mathcal{B}/f_0(k)$. The result requires f_0 to be constant throughout the region $[k, k + \Delta_{it}]$ conditional on each value of Δ_{it} , an assumption that is hard to justify except in the limit that the distribution of Δ_{it} concentrates around zero (Appendix Proposition 8 and Lemma SMALL make the above claims precise). A kink that produces only tiny responses is unlikely to provide a good approximation in a context like overtime, in which treatment corresponds to a 50% increase in the hourly cost of labor. Nevertheless, even in a “small-kink” setting, Theorem 1 offers a refinement to the small-kink approximation: a second-order expansion of $\ln(1 + \frac{x}{a})$ shows that when \mathcal{B} is small, the bounds Δ_k^L and Δ_k^U converge around $\Delta_k^* \approx \frac{\mathcal{B}-p}{2f_0(k)} + \frac{\mathcal{B}-p}{2f_1(k)}$.

A second existing result comes from Blomquist et al. (2015), who show that bunching identifies a certain weighted average of compensated elasticities in a nonparametric labor supply model, if the density of choices at an income tax kink is assumed to be linear across counterfactual tax rates. But as these authors point out, such a parametric assumption would be difficult to motivate.²²

4.4 Estimating policy relevant parameters

The buncher ATE yields the answer to a particular causal question, among a well-defined subgroup of the population. Namely: how would hours among workers bunched at 40 hours by the overtime

²⁰It is worth noting that BLC of h_1 and h_0 implies bounds on the treatment effect $Q_1(u) - Q_0(u)$ at *any* quantile u . But these bounds widen quickly as one moves away from the kink. When $f_0(k) \approx f_1(k)$, the narrowest bounds for a single rank u are obtained for a “median” buncher roughly halfway between $F_0(k)$ and $F_1(k)$. However, averaging over a larger group is more useful for meaningful ex-post evaluation of the FLSA (Sec. 4.4), and reduces the sensitivity to departures from RANK (see Figure 3). In the other extreme, one could drop RANK entirely and bound $\mathbb{E}[h_{0it} - h_{it}]$ directly via BLC of h_0 alone, but the bounds are *very* wide. The buncher ATE balances this tradeoff.

²¹Since the bounds depend only on the density around k and the total amount mass to its left/right, point masses elsewhere in the distributions of h_0 and h_1 do not effect on the bounds provided that they are well-separated from k .

²²In particular, the data identifies the density at the kink for two particular tax rates only, so cannot provide evidence of such linearity. Theorem 1 instead requires assumptions only about the two counterfactuals that are in fact observed.

rule be affected by a counterfactual change from linear pay at their straight-time wage to linear pay at their overtime rate? This section discusses how we may then use this quantity to both evaluate the overall ex-post effect of the FLSA on hours, as well as forecast the impacts of proposed changes to the FLSA. This requires some additional assumptions, which I continue to approach from a partial identification perspective.

4.4.1 From the buncher ATE to the ex-post hours effect of the FLSA

To consider the overall ex-post hours effect of the FLSA among covered workers, I proceed in two steps. I first relate the buncher ATE to the overall average effect of introducing the overtime kink, holding fixed the distributions of counterfactual hours h_{0it} and h_{1it} . Then, I allow straight-time wages to be affected by the FLSA, using the buncher ATE again to bound the additional effect of these wage changes on hours.

To motivate this strategy, let us first define the parameter of interest to be the difference in average weekly hours among hourly workers with and without the FLSA: $\theta := \mathbb{E}[h_{it}] - \mathbb{E}^*[h_{it}^*]$, where h_{it}^* indicates the hours unit it would work absent the FLSA, and the second expectation \mathbb{E}^* is over units corresponding to workers that would exist in the no-FLSA counterfactual and be covered were it introduced.²³ Defining θ in this way allows us to remain agnostic as to whether the FLSA changes employment, and hence the population of workers it applies to. However, I assume that the hours among any workers who enter or exit employment due to the FLSA are not systematically different from those who would exist without it, so that we may rewrite θ as $\theta = \mathbb{E}[h_{it} - h_{it}^*]$, averaging over individual-level causal effects in the population that does exist given the FLSA.

Next, decompose θ as:

$$\begin{aligned} \theta &= \mathbb{E}[h_{it}(w_{it}, \mathbf{h}_{-i,t}) - h_{0it}(w_{it}^*, \mathbf{h}_{-i,t}^*)] = \mathbb{E}\underbrace{[h_{it}(w_{it}, \mathbf{h}_{-i,t}) - h_{0it}(w_{it}, \mathbf{h}_{-i,t})]}_{\text{"effect of the kink"}} \\ &\quad + \mathbb{E}\underbrace{[h_{0it}(w_{it}, \mathbf{h}_{-i,t}) - h_{0it}(w_{it}^*, \mathbf{h}_{-i,t}^*)]}_{\text{"wage effects"}} + \mathbb{E}\underbrace{[h_{0it}(w_{it}^*, \mathbf{h}_{-i,t}^*) - h_{0it}(w_{it}^*, \mathbf{h}_{-i,t}^*)]}_{\text{"interdependencies"}}, \end{aligned} \quad (8)$$

where the notation makes explicit the dependence of h and h_0 on the worker's straight-time wage w_{it} , and possibly the hours \mathbf{h}_{-i} of other workers in their firm this week. In the notation of the last section: $h_{it} = h_{it}(w_{it}, \mathbf{h}_{-i,t})$, $h_{0it} = h_{0it}(w_{it}, \mathbf{h}_{-i,t})$ and $h_{1it} = h_{1it}(w_{it}, \mathbf{h}_{-i,t})$. I have used that $h_{it}^* = h_{0it}(w_{it}^*, \mathbf{h}_{-i,t}^*)$, since pay is linear in hours in the no-FLSA counterfactual.

The first term in Equation (8) reflects the "effect of the kink" quantity $h_{it} - h_{0it}$ examined in Section 4.2, and I view it as the first-order object of interest. The second term reflects that straight-time wages w_{it} may differ from those that workers would face without the FLSA, denoted by w_{it}^* .

²³Note that h_{it}^* in this section differs from the "anticipated" hours quantity h^* in Section 2.

The third term is zero when firms’ choice of hours for their workers decomposes into separate optimization problems for each unit, as in the benchmark model from Section 4.2. More generally, it will capture any interdependencies in hours across units, for instance due to different workers’ hours being not linearly separable in production. In Appendix E I provide evidence that such effects do not play a large role in θ , and I thus treat this term as zero when estimating θ .²⁴

Turning first to the “effect of the kink” term, note that with straight-wages and the hours of other units fixed, the kink only has such direct effects on those units working at least $k = 40$ hours:

$$h_{it} - h_{0it} = \begin{cases} 0 & \text{if } h_{it} < k \\ k - h_{0it} & \text{if } h_{it} = k \\ -\Delta_{it} & \text{if } h_{it} > k \end{cases} \quad (9)$$

and thus $\mathbb{E}[h_{it} - h_{0it}] = \mathcal{B} \cdot \mathbb{E}[k - h_{0it}|h_{it} = k] - P(h_{it} > k)\mathbb{E}[\Delta_{it}|h_{it} > k]$. To identify this quantity we must extrapolate from the buncher ATE to obtain an estimate of $\mathbb{E}[\Delta_{it}|h_{it} > k]$, the average effect for units who work overtime. To do this, I assume that the Δ_{it} of units working more than 40 hours are at least as large on average as those who work exactly 40, but that the reduced-form *elasticity* of their response is no greater than that of the bunchers. The logic is as follows: assuming a constant percentage change between h_0 and h_1 over units would imply responses that grow in proportion to h_1 , eventually becoming implausibly large. On the other hand, it would be an underestimate to assume high-hours workers, say at 60 hours, have the same effect in levels $h_0 - h_1$ as those closer to 40. Finally, I use bi-log-concavity of h_0 to put bounds on the average effect of the kink among bunchers $\mathcal{B} \cdot \mathbb{E}[k - h_{0it}|h_{it} = k]$. Details are provided in Appendix H.9.

The “wage effects” term in Equation (8) arises because the straight-time wages observed in the data may reflect some adjustment to the FLSA, as we would expect on the basis of the conceptual framework in Section 2. While the “effect of the kink” term is expected to be negative, this second term will be positive if the FLSA causes a reduction in the straight-time wages set at hiring. However, both terms ultimately depend on the same thing: responsiveness of hours to the cost of an hour of work. I can thus use the buncher ATE to compute an approximate upper bound on wage effects by assuming that all straight-time wages are adjusted according to Equation (1) and that the hours response is iso-elastic in wages, with anticipated hours approximated by h_{it} . Appendix H.9 provides a visual depiction of the logic. A lower bound on the “wage effects” term, on the other hand, is zero. In practice, the estimated size of the wage effect $\mathbb{E}[h_{0it} - h_{0it}^*]$ is appreciable but still

²⁴In particular, I fail to find evidence of contemporaneous hours substitution in response to colleague sick pay, in an event study design. Another piece of evidence comes from obtaining similar “effect of the kink” estimates across small, medium and large firms, which suggests that a firm’s capacity to reallocate hours between existing workers does not tend to drive their hours response to the FLSA. See Appendix E. If the third term of Eq. (8) is not zero, my strategy still estimates the average of a unit-level labor demand elasticity in which the hours of a worker’s colleagues are fixed.

small relative to $\mathbb{E}[h_{it} - h_{0it}]$ (cf. Appendix Table 11).

4.4.2 Forecasting the effects of policy changes

Apart from ex-post evaluation of the overtime rule, policymakers may also be interested in predicting what would happen if the parameters of overtime regulation were modified. Reforms that have been discussed in the U.S. include decreasing “standard hours” k at which overtime pay begins from 40 hours to 35 hours,²⁵ or increasing the overtime premium from time-and-a-half to “double-time” (Brown and Hamermesh, 2019). This section builds upon Sections 4.1 and 4.3 to show that the bunching-design model is also informative about the impact of such reforms on hours.

Let us begin by considering changes to standard hours k , for now holding the distributions of h_0 and h_1 fixed across the policy change. Inspection of Equation (2) reveals that as the kink is moved upwards, say from $k = 40$ hours to $k' = 44$ hours, some workers who were previously bunching at k now work h_{0it} hours: namely those for whom $h_{0it} \in [k, k']$. By the same token, some individuals with values of $h_{1it} \in [k, k']$ now bunch at k' . Some individuals who were bunching at k now bunch at k' —namely those for whom $h_{1it} \leq k$ and $h_{0it} \geq k'$. In the case of a reduction in overtime hours, say to $k' = 35$, this logic is reversed. Figure 8 depicts both cases, assuming that the mass of counterfactual bunchers p remains at $k = 40$ after the shift.²⁶

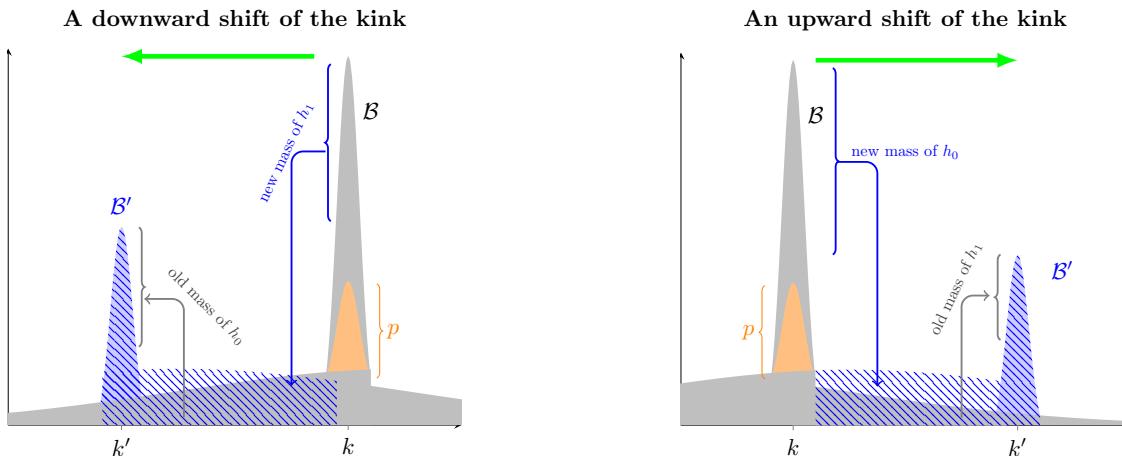


FIGURE 7: The left panel depicts a shift of the kink point downwards from k to k' , while right panel depicts a shift of the kink point upwards. See text for details.

Quantitatively assessing a change to double-time pay requires us to move beyond the two counterfactual choices h_{0it} and h_{1it} : hours that would be worked under straight-wages or under time-and-a-half pay. Let $h_{it}(\rho)$ be the hours that it would work if their employer faced a linear pay

²⁵Some countries have indeed changed standard hours in recent decades; see Brown and Hamermesh (2019).

²⁶It is conceivable that some or all counterfactual bunchers locate at 40 because it is the FLSA threshold, while still being non-responsive to the incentives introduced there by the kink. In this case, we might imagine that they would all coordinate on k' after the change. The effects here could then be seen as short-run effects before that occurs.

schedule at rate $\rho \cdot w_{it}$ (with w_{it} and hours of other units fixed at their realized levels). In this notation, $h_{0it} = h_{it}(1)$ and $h_{0it} = h_{it}(1.5)$. Now consider a new overtime policy in which a premium pay factor of ρ_1 is due from employers for hours in excess of k , e.g. $\rho_1 = 2$ for a “double-time” policy. Let $h_{it}^{[k,\rho_1]}$ denote realized hours for unit it under this overtime policy as a function of k and ρ_1 , and let $\mathcal{B}^{[k,\rho_1]} := P(h_{it}^{[k,\rho_1]} = k)$ be the observable bunching that would occur.

Theorem 2 obtains expressions for the effects of small changes to k or ρ_1 on hours. I continue to assume that counterfactual bunchers $K_{it}^* = 1$ stay at $k^* := 40$, regardless of ρ and k . Let $p(k) = p \cdot \mathbb{1}(k = k^*)$ denote the possible mass of counterfactual bunchers as a function of k .

Theorem 2 (marginal comparative statics in the bunching design). *Under Assumptions CHOICE, CONVEX, SEPARABLE and SMOOTH:*

1. $\partial_k \left\{ \mathcal{B}^{[k,\rho_1]} - p(k) \right\} = f_1(k) - f_0(k)$
2. $\partial_k \mathbb{E}[h_{it}^{[k,\rho_1]}] = \mathcal{B}^{[k,\rho_1]} - p(k)$
3. $\partial_{\rho_1} \mathcal{B}^{[k,\rho_1]} = -k f_{\rho_1}(k) \mathbb{E} \left[\frac{dh_{it}(\rho_1)}{d\rho} \middle| h_{it}(\rho_1) = k \right]$
4. $\partial_{\rho_1} \mathbb{E}[h_{it}^{[k,\rho_1]}] = - \int_k^\infty f_{\rho_1}(h) \mathbb{E} \left[\frac{dh_{it}(\rho_1)}{d\rho} \middle| h_{it}(\rho_1) = h \right] dh$

Proof. See Appendix A. □

The final two assumptions above are given in Appendix A: SEPARABLE requires firm preferences to be quasi-linear in costs, while SMOOTH is a set of regularity conditions which imply that $h_{it}(\rho)$ admits a density $f_\rho(h)$ for all ρ . Theorem 2 also uses a slightly stronger version of Assumption CHOICE that applies to all ρ rather than just ρ_0 and ρ_1 . The proof of Theorem 2 builds on results from Blomquist et al. (2015) and Kasy (2022)—see Appendix A for details.

Beginning from the actual FLSA policy of $k = 40 = k^*, \rho_1 = 1.5$, the RHS of Items 1 and 2 are in fact point identified from the data, provided that p is known. Item 1 says that if the location k of the kink is changed marginally, the kink-induced bunching probability will change according to the difference between the densities of h_{1i} and h_{0i} at k^* , which are in turn equal to the left and right limits of the observed density $f(h)$ at the kink. This result is intuitive: given continuity of each potential outcome’s density, a small increase in k will result in a mass proportional to $f_1(k)$ being “swept in” to the mass point at the kink, while a mass proportional to $f_0(k)$ is left behind. Item 2 aggregates this change in bunching with the changes to non-bunchers’ hours as k is increased: the combined effect turns out to be to simply transport the mass of inframarginal bunchers to the new value of k .²⁷ Making use of Theorem 2 for a discrete policy change like reducing standard hours

²⁷Intuitively, “marginal” bunchers who would choose exactly k under one of the two cost functions B_0 or B_1 cease to “bunch” as k increases, but in the limit of a small change they also do not change their realized h . Moore (2021)

to 35 requires integrating across the actual range of hypothesized policy variation. We lose point identification, but I use bi-log-concavity of the marginal distributions of h_0 and h_1 to retain bounds.

Now consider the effect of moving from time-and-a-half to double time on average hours worked, in light of Item 4. This scenario, similar to the effect of the kink term in Eq. (8), requires making assumptions about the response of individuals who may locate far above the kink, and for whom the buncher ATE is less directly informative. Integrating Item 4 over ρ we obtain an expression for the average effect of this reform in terms of local average elasticities of response:

$$\mathbb{E}[h_{it}^{[k,\rho_1]} - h_{it}^{[k,\bar{\rho}_1]}] = \int_{\rho_1}^{\bar{\rho}_1} d \ln \rho \int_k^\infty f_\rho(h) \cdot h \cdot \mathbb{E} \left[\frac{d \ln h_{it}(\rho)}{d \ln \rho} \middle| h_{it}(\rho) = h \right] dh$$

Recall that in the isoelastic model the elasticity quantity $\frac{d \ln h_{it}(\rho)}{d \ln \rho} = \frac{dh_{it}(\rho)}{d\rho} \frac{\rho}{h_{it}(\rho)}$ is constant across ρ and across units, and it is partially identified under BLC. Just as a constant proportional response is likely to overstate responsiveness at large values of hours, it is likely to *understate* responsiveness to larger values of ρ . This yields a lower bound on the effect of moving to double-time. For an upper bound on the magnitude of the effect, I assume rather that in levels $\mathbb{E}[h_{it}(\rho_1) - h_{it}(\bar{\rho}_1)|h_{1it} > k]$ is at least as large as $\mathbb{E}[h_{0it} - h_{1it}|h_{1it} > k]$, and that the increase in bunching from a change of ρ_1 to $\bar{\rho}_1$ is as large as the increase from ρ_0 to ρ_1 . Additional details are provided in Appendix H.9.

5 Implementation and Results

This section implements the empirical strategy described in Section 4 with the sample of administrative payroll data described in Section 3.

5.1 Identifying counterfactual bunching at 40 hours

To deliver final estimates of the effect of the FLSA overtime rule on hours, it is necessary to first return to an issue raised in the introduction and allowed for in Section 4: that there are other reasons to expect bunching at 40 hours, in addition to being the location of the FLSA kink. For one, 40 may reflect a kind of *status-quo* choice, being chosen even when it is not exactly profit maximizing for the firm. This effect could be amplified by firms synchronizing the schedules of different workers, requiring some common number of hours per week to coordinate around. Finally, if any salaried workers were not successfully removed from the sample, hours for such workers might be recorded as 40 even as actual hours worked vary.

gives a closely-related result, derived independently of this work. In the context of a tax kink with x a scalar and $p(k) = 0$, the result of Moore (2021) generalizes Item 2 of Theorem 2, showing that bunching is a sufficient statistic for the effect of a marginal change in k on tax revenue.

In terms of the empirical strategy from Section A.2, all of these alternative explanations manifest in the same way: a point mass p at 40 in the distribution of hours that would occur even if pay did not feature a kink at 40. In the notation introduced in Section 4.3, these “counterfactual bunchers” are demarcated by $K_{it}^* = 1$. Let us refer to the $K_{it}^* = 0$ individuals who also locate at the kink as “active bunchers”. The mass of active bunchers is $\mathcal{B} - p$. Theorem 1 shows that we can still partially identify the buncher ATE in the presence of counterfactual bunchers, so long as we know what portion of the total bunchers are active versus counterfactual.

I leverage two strategies to provide plausible estimates for the mass of counterfactual bunchers p . My preferred estimate makes use of the fact that when an employee is paid for hours that are not actually worked—including sick time, paid time off (PTO) and holidays—these hours do not contribute to the 40 hour overtime threshold of the FLSA that week. For example, if a worker applies PTO to miss a six hour shift, then they are not required to be paid overtime until they reach 46 total paid hours in that week. Thus while the kink remains at 40 hours *worked*, non-work hours like PTO shift the location of the kink in hours of *pay*.

The identifying assumption that I rely on is that individuals who still work 40 hours a week, even when they have non-work hours (and are hence paid for more than 40), are all active bunchers: they would not be located at forty hours in the counterfactuals h_{0it} and h_{1it} . This reflects the idea that additional explanations for bunching at 40 hours operate at the level of hours paid, rather than hours worked. Letting n_{it} indicate non-work hours of pay for paycheck it , I make two assumptions:

1. $P(h_{it} = 40|n_{it} > 0) = P(h_{it} = 40 \text{ and } K_{it}^* = 0|n_{it} > 0)$
2. $P(h_{it} = 40 \text{ and } K_{it}^* = 0|n_{it} > 0) = P(h_{it} = 40 \text{ and } K_{it}^* = 0|n_{it} = 0)$

The first item reflects the above logic, and allows me to identify the mass of active bunchers in the $n_{it} > 0$ conditional distribution of hours. The second item says that this conditional mass is representative of the unconditional mass of active bunchers. To increase the plausibility of this assumption, I focus on η as paid time off because it is generally planned in advance, yet has somewhat idiosyncratic timing.²⁸

Together, the two assumptions above imply that $p = P(K_{it}^* = 1 \text{ and } h_{it} = 40)$ is identified as $\mathcal{B} - P(h_{it} = 40|\eta_{it} > 0)$. Figure 8 shows the conditional distribution of hours paid for work when the paycheck contains a positive number of PTO hours ($n_{it} > 0$). The figure reveals that when moving from the unconditional (left panel) to positive-PTO conditional (right panel) distribution, most of the point mass at 40 hours moves away, largely concentrating now at 32 hours (corresponding to the PTO covering eight hours). Of the total bunching of $\mathcal{B} \approx 11.6\%$ in the unconditional distribution, I estimate that only $P(h_{it} = 40|n_{it} > 0) \approx 2.7\%$ are active bunchers,

²⁸By contrast, sick pay is often unanticipated so the firm may not be able to re-optimize total hours within the week in which a worker calls in sick. Holiday pay is known in advance, but holidays are unlikely to be representative in terms of other factors important for hours determination (e.g. product demand).

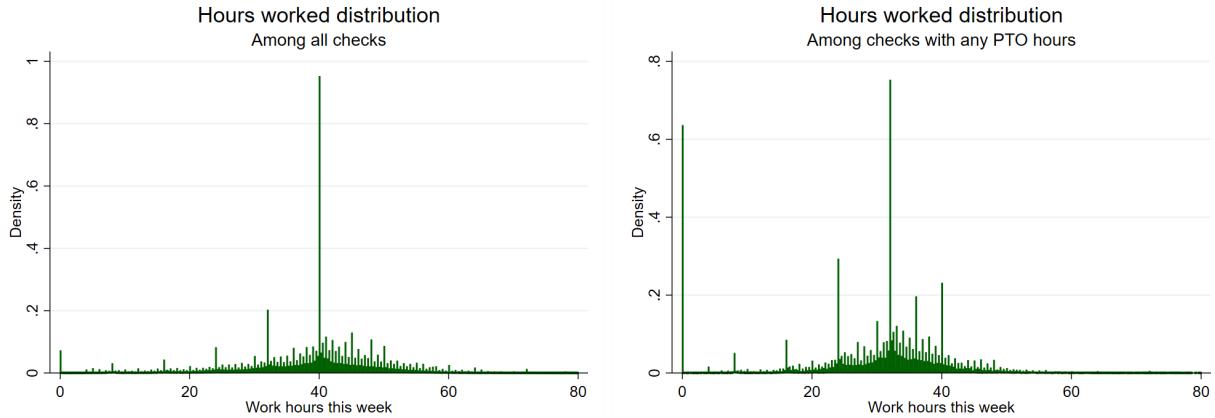


FIGURE 8: The right panel shows a histogram of hours worked when paid time off hours are positive ($\eta_{it} > 0$). The left panel shows the unconditional distribution. While $\mathcal{B} \approx 11.6\%$, $P(h_{it} = 40 | n_{it} > 0) \approx 2.7\%$.

leaving $p \approx 8.9\%$. Thus roughly three quarters of the individuals at 40 hours are counterfactual rather than active bunchers.

As a secondary strategy, I estimate an upper bound for p by using the assumption that the potential outcomes of counterfactual bunchers are relatively “sticky” over time. If the hours of counterfactual bunchers are at 40 for behavioral or administrative reasons, it is reasonable to assume that these external considerations are fairly static, preventing latent hours h_{0it} from changing much between adjacent weeks. In particular, assume that in a given week t nearly all of the counterfactual bunchers are also “non-changers” of hours from week $t - 1$. Then:

$$p = P(h_{0it} = 40) \approx P(h_{0it} = h_{0it-1} = 40) \leq P(h_{it} = h_{i,t-1} = 40),$$

where the inequality follows from $(h_{0it} = 40) \implies (h_{it} = 40)$ by Lemma 1. The probability $P(h_{it} = h_{i,t-1} = 40)$ can be directly estimated from the data, yielding $p \leq 6\%$.

5.2 Estimation and inference

Given Theorem 1 and a value of p , computing bounds on the buncher ATE requires estimates of the right and left limits of the CDF and density of hours at the kink. I use the local polynomial density estimator of Cattaneo, Jansson and Ma (2020) (CJM), which is well-suited to estimating a CDF and its derivatives at boundary points. A local-linear CJM estimator of the left limit of the CDF and density at k , for instance, is:

$$(\hat{F}_-(k), \hat{f}_-(k)) = \operatorname{argmin}_{(b_1, b_2)} \sum_{it: h_{it} < k} (F_n(h_{it}) - b_1 - b_2 h_{it})^2 \cdot K\left(\frac{h_{it} - k}{\alpha}\right) \quad (10)$$

where $F_n(y) = \frac{1}{n} \sum_{it} \mathbb{1}(h_{it} \leq y)$ is the empirical CDF of a sample of size n , $K(\cdot)$ is a kernel function, and α is a bandwidth. The right limits $F_+(k)$ and $f_+(k)$ are estimated analogously using observations for which $h_{it} > k$. I use a triangular kernel, and choose h as follows: first, I use CJM's mean-squared error minimizing bandwidth selector to produce a bandwidth choice using the data on either side of $k = 40$ (for the left and right limits, respectively). I then average the two bandwidths, and use this as the bandwidth in the final calculation of both the right and left limits. In the full sample, the bandwidth chosen by this procedure is about 1.7 hours, and is somewhat larger for subsamples that condition on a single industry.

To construct confidence intervals for parameters that are partially identified (e.g. the buncher ATE), I use adaptive critical values proposed by Imbens and Manski (2004) and Stoye (2009) that are valid for the underlying parameter. To easily incorporate sampling uncertainty in all of $\hat{F}_-(k)$, $\hat{f}_-(k)$, $\hat{F}_+(k)$, $\hat{f}_+(k)$ and \hat{p} , I estimate variances by a cluster nonparametric bootstrap that resamples at the firm level. This allows arbitrary autocorrelation in hours across pay periods for a single worker, and between workers within a firm. All standard errors use 500 bootstrap samples.

5.3 Results of the bunching estimator: the buncher ATE

Table 3 reports treatment effect estimates based on Theorem 1, when p is either assumed to be zero or is estimated by one of the two methods described in Section 5.1. The first row reports the corresponding estimate of the net bunching probability $\mathcal{B} - p$, while the second row reports the bounds on the buncher ATE $\mathbb{E}[h_{0it} - h_{1it}|h_{it} = k, K_{it}^* = 0]$. Within a fixed estimate of p , the bounds on the buncher ATE based on bi-log-concavity are quite informative: the upper and lower bounds are close to each other and precisely estimated.

	$p=0$	p from non-changers	p from PTO
Net bunching:	0.116 [0.112, 0.120]	0.057 [0.055, 0.058]	0.027 [0.024, 0.030]
Buncher ATE	[2.614, 3.054] [2.493, 3.205]	[1.324, 1.435] [1.264, 1.501]	[0.640, 0.666] [0.574, 0.736]
Num observations	630217	630217	630217
Num clusters	566	566	566

TABLE 3: Estimates of net bunching $\mathcal{B} - p$ and the buncher ATE: $\Delta_k^* = \mathbb{E}[h_{0it} - h_{1it}|h_{it} = k, K_{it}^* = 0]$, across various strategies to estimate counterfactual bunching $p = P(K_{it}^* = 1)$. Unit of analysis is a paycheck, and 95% bootstrap confidence intervals (in gray) are clustered by firm.

The PTO-based estimate of p provides the most conservative treatment effect estimate, attributing roughly one quarter of the observed bunching to active rather than counterfactual bunchers. Nevertheless, this estimate still yields a highly statistically significant buncher ATE of about 2/3 of an hour, or 40 minutes. This estimate has the following interpretation: consider the group of workers that are in fact working 40 hours in a given pay period and are not counterfactual bunchers. Firms would ask this group to work on average about 40 minutes more that week if they were paid their straight-time wage for all hours, compared with a counterfactual in which they are paid their overtime rate for all hours. If we instead attribute all of the observed bunching mass to active bunchers ($p = 0$), then this buncher ATE parameter is estimated to be at least 2.6 hours. In Appendix C I report estimates based on alternative shape constraints and assumptions about effect heterogeneity (with similar results), as well as separate estimates of the buncher ATE for each of the largest industries in the sample.

5.4 Estimates of policy effects

I now use estimates of the buncher ATE and the results of Section 4.4 to estimate the overall causal effect of the FLSA overtime rule, and simulate changes based on modifying standard hours or the premium pay factor. Table 4 first reports an estimate of the buncher ATE expressed as a reduced-form hours demand elasticity,²⁹ which I use as an input in these calculations. The next two rows report bounds on $\mathbb{E}[h_{it} - h_{it}^*]$ and $\mathbb{E}[h_{it} - h_{it}^* | h_{1it} \geq 40, K_{it}^* = 0]$, respectively. The second row is the overall ex-post effect of the FLSA on hours, averaged over workers and pay periods, and the third row conditions on paychecks reporting at least 40 hours (omitting counterfactual bunchers). The final row reports an estimate of the effect of moving to double-time pay. I provide details of the calculations in Appendix H.9.

Taking the PTO-based estimate of p as yielding a lower bound on treatment effects, the estimates suggest that workers work at least about 1/4 of an hour less in any given week than they would absent overtime regulation: about one third the magnitude of the buncher ATE in levels. When I focus on those workers that are directly affected in a given week, the figure is about twice as high: roughly 30 minutes. Since my data has been restricted to hourly workers paid on a weekly basis, these estimates should be interpreted as holding for that population only. While one might assume that similar effects hold for hourly workers paid at other intervals (e.g. bi-weekly), speaking to the hours effects of the FLSA on salary workers is beyond the scope of this study.

Table 4 also suggests that a move to double-time pay would introduce a further reduction in hours comparable to the existing overall ex-post effect, but with substantially wider bounds. These

²⁹ This is $\hat{\Delta}_k^*/(40 \ln(1.5))$ where $\hat{\Delta}_k$ is the estimate of the buncher ATE presented in Table 3. This is numerically equivalent to the elasticity implied by the buncher ATE in logs $\mathbb{E}[\ln h_{0it} - \ln h_{1it} | h_{it} = k, K_{it}^* = 0] / (\ln 1.5)$ estimated under assumption that $\ln h_0$ and $\ln h_1$ are BLC.

	$p=0$	p from non-changers	p from PTO
Buncher ATE as elasticity	[-0.188,-0.161] [-0.198,-0.154]	[-0.088,-0.082] [-0.093,-0.078]	[-0.041,-0.039] [-0.045,-0.035]
Average effect of FLSA on hours	[-1.466, -1.026] [-1.535, -0.977]	[-0.727, -0.486] [-0.762, -0.463]	[-0.347, -0.227] [-0.384, -0.203]
Avg. effect among directly affected	[-2.620, -1.833] [-2.733, -1.750]	[-1.453, -0.972] [-1.518, -0.929]	[-0.738, -0.483] [-0.812, -0.434]
Double-time, average effect on hours	[-2.604, -0.569] [-2.707, -0.547]	[-1.239, -0.314] [-1.285, -0.300]	[-0.580, -0.159] [-0.638, -0.143]

TABLE 4: Estimates of the buncher ATE expressed as an elasticity, the average ex-post effect of the FLSA $\mathbb{E}[h_{it} - h_{it}^*]$,²⁹ the effect among directly affected units $\mathbb{E}[h_{it} - h_{it}^* | h_{it} \geq k, K_{it}^* = 0]$ and predicted effects of a change to double-time. 95% bootstrap confidence intervals in gray, clustered by firm.

estimates include the effects of possible adjustments to straight-time wages, which tend to attenuate the impact of the policy change. Appendix Table 11 replicates Table 4 neglecting these wage adjustments, which might be viewed as a short-run response to the FLSA before wages have time to adjust.

Figure 9 breaks down estimates of the ex-post effect of the overtime rule by major industries, revealing considerable heterogeneity between them. The estimates suggest that Real Estate & Rental and Leasing as well as Wholesale Trade see the highest average reduction in hours. The least-affected industries are Health Care and Social Assistance and Professional Scientific and Technical, with the average worker working just about 6 minutes less per week due to the overtime rule. Appendix Figure 8 compares the hours distribution for Real Estate & Rental and Leasing with the distribution for Professional Scientific and Technical, showing that the difference in their effects is explained both by a larger value of $\mathcal{B} - p$ and a lower density of hours close to the kink for Real Estate & Rental and Leasing. Appendix C reports estimates broken down by gender, finding that the FLSA has considerably higher effects on the hours of men compared with women.

Appendix Figure 14 looks at the effect of changing the threshold for overtime hours k from 40 to alternative values k' . The left panel reports estimates of the identified bounds on $\mathcal{B}^{[k', \rho_1]}$ as well as point-wise 95% confidence intervals (gray) across values of k' between 35 and 45, for each of the three approaches to estimating p . In all cases, the upper bound on bunching approaches zero as k' is moved farther from 40. This is sensible if the h_0 and h_1 distributions are roughly unimodal with modes around 40: straddling of potential outcomes becomes less and less likely as one moves away from where most of the mass is. Appendix Figure 13 shows these bounds as k' ranges all the way

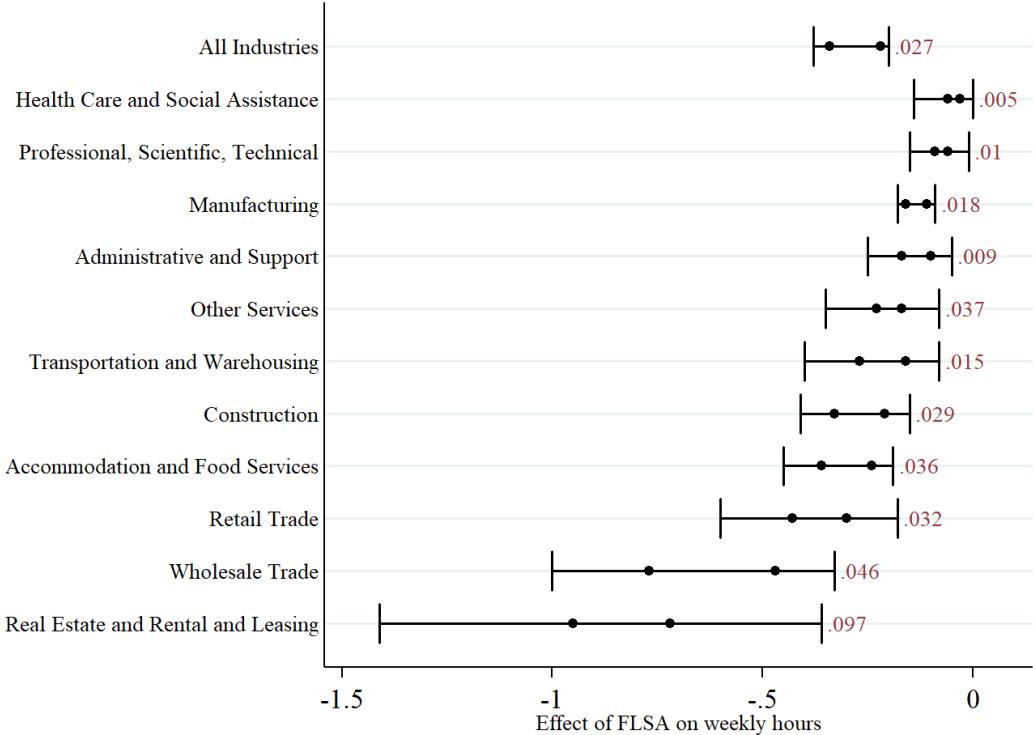


FIGURE 9: 95% confidence intervals for the effect of the FLSA on hours by industry, using PTO-based estimates of p for each. Dots are point estimates of the upper and lower bounds. The number to the right of each range is the point estimate of the net bunching $\mathcal{B} - p$ for that industry.

from 0 to 80, for the $p = 0$ case. These estimates do not account for adjustment to straight-time wages, so should be viewed as quantifying short-run responses.

When p is estimated using PTO or non-changers between periods, we see that the upper bound of the identified set for $\mathcal{B}^{[k', \rho_1]}$ in fact reaches zero quite quickly in k' . Moving standard hours to 35 is thus predicted to completely eliminate bunching due to the overtime kink in the short run, before any adjustment to latent hours (e.g. through changes to straight-time wages). The right panel of Appendix Figure 14 shows estimates for the average effect on hours of changing standard hours, inclusive of wage effects (see Appendix H.9 for details). Increases to standard hours cause an increase in hours per worker, as overtime policy becomes less stringent, and reductions to standard hours reduce hours.³⁰ The size of these effects is not precisely estimated for changes larger than a couple of hours, however the range of statistically significant effects depends on p . Even for the preferred estimate of p from PTO, increasing the overtime threshold as high as 43 hours is estimated to increase average working hours by an amount distinguishable from zero.

³⁰ The magnitudes are consistent with estimates by Costa (2000), that hours fell by 0.2-0.4 on average during the phased introduction of the FLSA in which standard hours declined by 2 hours in 1939 and 1940.

6 Implications of the estimates for overtime policy

The estimates from the preceding section suggest that FLSA regulation indeed has real effects on hours worked, in line with labor demand theory when wages do not fully adjust to absorb the added cost of overtime hours. When averaged over affected workers and across pay periods, I find that hourly workers in my sample work at least 30 minutes less per week than they would without the overtime rule. This lower bound is broadly comparable to the few causal estimates that exist in the literature, including Hamermesh and Trejo (2000) who assess the effects of expanding California's daily overtime rule to cover men in 1980, and Brown and Hamermesh (2019) who use the erosion of the salary threshold for exemption of white-collar jobs in real terms over the last several decades.³¹ By contrast, my estimates use an identification strategy that does not require focusing on the sub-population affected by a natural experiment, and are based on recent and administrative data.

My estimates speak to the substitutability of hours of labor between workers. The primary justifications for overtime regulation have been to reduce excessive workweeks, while encouraging hours to be distributed over more workers (Ehrenberg and Schumann, 1982). How well this (and related policies that have received recent interest like work-sharing programs) plays out in practice hinges on how easily an hour of work can be moved from one worker to another or across time, from the perspective of the firm. The results of this paper find hours demand to be relatively inelastic: hours cannot be easily so reallocated between workers or weeks. This suggests that ongoing efforts to expand coverage of the FLSA overtime rule may have limited scope to dramatically affect the hours of U.S. workers.

Nevertheless, the overall impact of the FLSA overtime rule on workers is still notable. The data suggest that at least about 3% and as many as about 12% of workers' hours are adjusted to the threshold introduced by the policy, indicating that it may have distortionary impacts for a significant portion of the labor force. The policy may also have important effects on unemployment. While a full assessment of the employment effects of the FLSA overtime rule is beyond the scope of this paper, my estimates of the hours effect can be used to build a back-of-the-envelope calculation. Following Hamermesh (1993), I assume a value for the rate at which firms substitute labor for capital to obtain a "best-guess" estimate that the FLSA overtime rule creates about 700,000 jobs (see Appendix C.6 for details). To get an overall upper bound on the size of employment effects, one can instead attribute all of the bunching at 40 to the FLSA and assume that the total number of worker-hours is not reduced by the FLSA. By this estimate the FLSA increases employment by at most 3 million jobs, or roughly 3% among covered workers. A reasonable range of param-

³¹ Hamermesh and Trejo (2000) and Brown and Hamermesh (2019) report estimates of -0.5 and -0.18 for the elasticity of overtime hours with respect to the overtime rate. My preferred estimate of -0.04 for the buncher ATE as an elasticity is the elasticity of *total* hours, including the first 40. An elasticity of overtime hours can be computed from this using the ratio of mean hours to mean overtime hours in the sample, resulting in an estimate of roughly -0.45 .

eter values in this simple calculation rules out that the FLSA overtime rule has negative overall employment effects on hourly workers.

7 Conclusion

This paper has provided a new interpretation of the popular bunching-design method in the language of treatment effects, showing that the basic identifying power of the method is robust to a variety of specific underlying choice models. Across such modeling choices, the parameter of interest remains a reduced-form local average treatment effect between two appropriately-defined counterfactual choices, which can be partially identified by making nonparametric assumptions about the counterfactuals' distributions. This provides conditions under which the bunching design can be useful to answer program evaluation questions in a variety of contexts, particularly beyond those in which the researcher is prepared to posit a parametric model of decision-makers' preferences.

By leveraging these insights with a new payroll dataset recording exact weekly hours paid at the individual level, I estimate that U.S. workers subject to the Fair Labor Standard Act work shorter hours due to its overtime provision, which may lead to positive employment effects. Given the large amount of within-worker variation in hours observed in the data, the modest size of the FLSA effects estimated in this paper suggest that firms do face significant incentives to maintain longer working hours, countervailing against the ones introduced by policies intended to reduce them.

References

- BARKUME, A. (2010). "The Structure of Labor Costs with Overtime Work in U.S. Jobs". *Industrial and Labor Relations Review* 64 (1).
- BERTANHA, M., MCCALLUM, A. H. and SEEGERT, N. (2020). "Better Bunching , Nicer Notching". *SSRN Working Paper*.
- BICK, A., BLANDIN, A. and ROGERSON, R. (2022). "Hours and Wages". *The Quarterly Journal of Economics* (forthcoming).
- BISHOW, J. L. (2009). "A Look at Supplemental Pay: Overtime Pay, Bonuses, and Shift Differentials". *Monthly Labor Review*. Publisher: Bureau of Labor Statistics, U.S. Department of Labor.
- BLOMQUIST, S., KUMAR, A., LIANG, C.-Y. and NEWHEY, W. (2021). "On Bunching and Identification of the Taxable Income Elasticity". *Journal of Political Economy* 129 (8).

- BLOMQUIST, S., KUMAR, A., LIANG, C.-Y. and NEWHEY, W. K. (2015). "Individual heterogeneity, nonlinear budget sets and taxable income". *The Institute for Fiscal Studies Working Paper* CWP21/15.
- BRECHLING, F. P. R. (1965). "The Relationship Between Output and Employment in British Manufacturing Industries". *The Review of Economic Studies* 32 (3), p. 187.
- BROWN, C. and HAMERMESH, D. S. (2019). "Wages and Hours Laws: what do we know? what can be done?" *The Russell Sage Foundation Journal of the Social Sciences* 5 (5), pp. 68–87.
- BURDETT, K. and MORTENSEN, D. T. (1998). "Wage Differentials, Employer Size, and Unemployment". *International Economic Review* 39 (2), p. 257.
- BUREAU OF LABOR STATISTICS (2020). *National Compensation Survey (Restricted-Use Microdata)*. www.bls.gov/ncs/ncs-data-requests.htm.
- CAHUC, P. and ZYLBERBERG, A. (2014). *Labor economics*. 2nd. Cambridge, Mass.: MIT Press.
- CATTANEO, M. D., JANSSON, M. and MA, X. (2020). "Simple Local Polynomial Density Estimators". *Journal of the American Statistical Association* 115 (531), pp. 1449–1455.
- CHERNOZHUKOV, V. and HANSEN, C. (2005). "An IV Model of Quantile Treatment Effects". *Econometrica* 73 (1), pp. 245–261.
- CHETTY, R., FRIEDMAN, J. N., OLSEN, T. and PISTAFERRI, L. (2011). "Adjustment Costs, Firm Responses, and Micro vs. Macro Labor Supply Elasticities: Evidence from Danish Tax Records." *Quarterly Journal of Economics* 126 (2), pp. 749–804.
- COSTA, D. L. (2000). "Hours of Work and the Fair Labor Standards Act: A Study of Retail and Wholesale Trade, 1938–1950". *Industrial and Labor Relations Review*, p. 17.
- DÜMBGEN, L., KOLESNYK, P. and WILKE, R. A. (2017). "Bi-log-concave distribution functions". *Journal of Statistical Planning and Inference* 184, pp. 1–17.
- DUBE, A., MANNING, A. and NAIDU, S. (2020). "Monopsony, Misoptimization, and Round Number Bunching in the Wage Distribution". *NBER Working Paper* w24991.
- EHRENBERG, R. and SCHUMANN, P. (1982). *Longer hours or more jobs? : an investigation of amending hours legislation to create employment*. New York State School of Industrial and Labor Relations, Cornell University.

- EHRENBERG, R. G. (1971). "The Impact of the Overtime Premium on Employment and Hours in U . S . Industry". *Economic Inquiry* 9 (2).
- GRIGSBY, J., HURST, E. and YILDIRMAZ, A. (2021). "Aggregate Nominal Wage Adjustments: New Evidence from Administrative Payroll Data". 11 (2), pp. 428–71.
- HAMERMESH, D. S. (1993). *Labor demand*. Princeton, NJ: Princeton Univ. Press.
- HAMERMESH, D. S. and TREJO, S. J. (2000). "The Demand for Hours of Labor : Direct Evidence from California". *The Review of Economics and Statistics* 82 (1), pp. 38–47.
- HART, R. A. (2004). *The economics of overtime working*. Cambridge, UK: Cambridge University Press.
- HJORT, J., LI, X. and SARSONS, H. (2020). "Across-Country Wage Compression in Multinationals". *NBER Working Paper* w26788.
- IMBENS, G. W. and MANSKI, C. F. (2004). "Confidence Intervals for Partially Identified Parameters". *Econometrica* 72, p. 14.
- JOHNSON, J. (2003). "The Impact of Federal Overtime Legislation on Public Sector Labor Markets". *Journal of Labor Economics* 21 (1), pp. 43–69.
- KASY, M. (2022). "Who wins, who loses? Identification of the welfare impact of changing wages". *Journal of Econometrics* 226 (1), pp. 1–26.
- KLEVEN, H. J. (2016). "Bunching". *Annual Review of Economics* 8, pp. 435–464.
- KLEVEN, H. J. and WASEEM, M (2013). "Using notches to uncover optimization frictions and structural elasticities: Theory and evidence from Pakistan". *The Quarterly Journal of Economics* 128 (2), pp. 669–723.
- KLINE, P. and TARTARI, M. (2016). "Bounding the Labor Supply Responses to a Randomized Welfare Experiment: A Revealed Preference Approach". *American Economic Review* 106 (4), pp. 972–1014.
- MATOS, K., GALINSKY, E. and BOND, J. (2017). "National Study of Employers". *Society for Human Resource Management*, p. 79.
- MOORE, D. T. (2021). "Evaluating Tax Reforms without Elasticities: What Bunching Can Identify". *Mimeo*, p. 61.

- QUACH, S. (2021). “The Labor Market Effects of Expanding Overtime Coverage”. *MPRA Paper* 100613.
- ROSEN, S. (1968). “Short-Run Employment Variation on Class-I Railroads in the U.S., 1947-1963”. *Econometrica* 36 (3), p. 511.
- SAEZ, E. (2010). “Do Taxpayers Bunch at Kink Points?” *American Economic Journal: Economic Policy* 2 (3), pp. 180–212.
- STOLE, L. A. and ZWIEBEL, J. (1996). “Intra-Firm Bargaining under Non-Binding Contracts”. *The Review of Economic Studies* 63 (3), pp. 375–410.
- STOYE, J. (2009). “More on Confidence Intervals for Partially Identified Parameters”. *Econometrica* 77 (4), pp. 1299–1315.
- TREJO, B. S. J. (1991). “The Effects of Overtime Pay Regulation on Worker Compensation”. *American Economic Review* 81 (4), pp. 719–740.

Online Appendices for “Treatment Effects in Bunching Designs: The Hours Impact of the Federal Overtime Rule”

Leonard Goff

Last updated: March 8, 2022

Contents

A Identification in a generalized bunching design	2
A.1 The policy environment	3
A.2 Potential outcomes as counterfactual choices	5
A.3 Examples from the general choice model in the overtime setting	7
A.4 Observables in the kink bunching design	8
A.5 Treatment effects in the bunching design	10
A.6 Policy changes in the bunching-design	12
B Main proofs	14
B.1 Proof of Lemma 1	14
B.2 Proof of Theorem 1	16
B.3 Proof of Theorem 2	18
C Additional empirical information and results	19
C.1 Sample restrictions	19
C.2 A test of the Trejo (1991) model of straight-time wage adjustment	20
C.3 Further characteristics of the sample	21
C.4 Additional treatment effect estimates and figures	24
C.5 Estimates from the iso-elastic model	32
C.6 Results of the employment effect calculation	33
D Incorporating workers that set their own hours	34
E Interdependencies among hours within the firm	38

F Modeling the determination of wages and “typical” hours	42
F.1 A simple model with exogenous labor supply	42
F.2 Endogenizing labor supply in an equilibrium search model	44
G Additional identification results for the bunching design	53
G.1 A generalized notion of homogeneous treatment effects	53
G.2 Parametric approaches with homogeneous treatment effects	54
G.3 Nonparametric approaches with homogeneous treatment effects	55
G.4 Alternative identification strategies with heterogeneous treatment effects . .	56
G.5 Two bunching design settings from the literature	59
H Additional proofs	61
H.1 Proof of Propositions 1 and 2	61
H.2 Proof of Lemma 2	61
H.3 Proof of Lemma SMALL	62
H.4 Proof of Appendix D Proposition 3	64
H.5 Proof of Appendix D Theorem 1*	67
H.6 Proof of Appendix F.1 Proposition 4	69
H.7 Proof of Appendix Proposition 5	69
H.8 Proof of Appendix Proposition 6	70
H.9 Details of calculations for policy estimates	70
H.10 Details of calculating wage correction terms	76

A Identification in a generalized bunching design

This section presents some generalizations of the bunching-design model used in the main text. While the FLSA will provide a running example throughout, I largely abstract from the overtime context to emphasize the general applicability of the results.

To facilitate comparison with the existing literature on bunching at kinks – which has mostly considered cross-sectional data – I throughout this section suppress time indices and use the single index i to refer to each unit of observation (a paycheck in the overtime setting). Further, the “running variable” of the bunching design is typically denoted by Y rather than h , and so the random variable Y_i will play the role of h_{it} from the main text. This is done to emphasize the link to the treatment effects literature, while allowing a distinction that is in some cases important (e.g. models in which hours of pay for work differ from actual hours of work).

A.1 The policy environment

Here we abstract from the conventional piece-wise linear kink setting that appears in tax examples as well as the main body of this paper. Consider a population of observational units indexed by i . For each i , a decision-maker $d(i)$ chooses a point (z, \mathbf{x}) in some space $\mathcal{X} \subseteq \mathbb{R}^{m+1}$ where z is a scalar and \mathbf{x} a vector of m components, subject to a constraint of the form:

$$z \geq \max\{B_{0i}(\mathbf{x}), B_{1i}(\mathbf{x})\} \quad (1)$$

The functions $B_{0i}(\mathbf{x})$ and $B_{1i}(\mathbf{x})$ are taken to be continuous and weakly convex functions of the vector \mathbf{x} , and assume that there exist continuous scalar functions $y_i(\mathbf{x})$ and a scalar k such that:

$$B_{0i}(\mathbf{x}) > B_{1i}(\mathbf{x}) \text{ whenever } y_i(\mathbf{x}) < k \quad \text{and} \quad B_{0i}(\mathbf{x}) < B_{1i}(\mathbf{x}) \text{ whenever } y_i(\mathbf{x}) > k$$

The value k is taken to be common to all units i , and is assumed to be known by the researcher.¹ In the overtime setting, $y_i(\mathbf{x})$ represents the hours of work for which a worker is paid in a given week, $k = 40$, and $B_{0i}(\mathbf{x}) = w_i y_i(\mathbf{x})$ and $B_{1i}(\mathbf{x}) = 1.5w_i y_i(\mathbf{x}) - 20w_i$. In most applications of the bunching design, the decision-maker $d(i)$ is simply i themselves, for example a worker choosing their labor supply subject to a tax kink. In the overtime application however i is a worker-week pair, and $d(i)$ is the worker's firm.

Let X_i be i 's realized outcome of \mathbf{x} , and $Y_i = y_i(X_i)$. I assume that Y_i is observed by the econometrician, but not that X_i is. In the overtime setting this means that the econometrician observes hours for which workers are paid, but not necessarily all choices made by firms that pin down those hours (for example, how many hours to allow the worker to stay “on the clock” during paid breaks—see Section A.3).

In general, the functions B_{0i} , B_{1i} will represent a schedule of some kind of “cost” as a function of the choice vector \mathbf{x} , with two regimes of costs that are separated by the condition $y_i(\mathbf{x}) = k$, characterizing the locus of points at which the two cost functions cross. Let $B_{ki}(\mathbf{x}) := \max\{B_{0i}(\mathbf{x}), B_{1i}(\mathbf{x})\}$ denote the actual constraint function that applies to z . A budget constraint like Eq. $z \geq B_{ki}(\mathbf{x})$ is typically “kinked” because while the function $B_{ki}(\mathbf{x})$ is continuous, it will generally be non-differentiable at the \mathbf{x} for which $y_i(\mathbf{x}) = k$.² While the functions B_0 , B_1 and y can all depend on i , I will often suppress this dependency for clarity of notation.

Discussion of the general model:

¹This comes at little cost of generality since with heterogeneous k_i this could be subsumed as a constant into the function $y_i(\mathbf{x})$, so long as the k_i are observed by the researcher.

²In particular, the subgradient of $\max\{B_{0i}(\mathbf{x}), B_{1i}(\mathbf{x})\}$ will depend on whether one approaches from the $y_i(\mathbf{x}) > k$ or the $y_i(\mathbf{x}) < k$ side. With a scalar x and linear B_0 and B_1 , the derivative of $B_{ki}(x)$ discontinuously rises at x for which $y_i(\mathbf{x}) = k$.

In the most common cases from the literature, no distinction is made between the “running variable” y of the kink and any underlying choice variables \mathbf{x} . This corresponds to a setting in which \mathbf{x} is a scalar and $y_i(x) = x$. For example, the seminal bunching design papers Saez (2010) and Chetty et al. (2011) considered progressive taxation with z being tax liability (or credits), $y = x$ corresponding to taxable income, and B_0 and B_1 linear tax functions on either side of a threshold y between two adjacent tax/benefit brackets. Similarly, in the overtime context, the functions B_0 and B_1 are linear and only depend on hours $y_i(\mathbf{x})$, as depicted in Figure 1. Appendix G discusses a tax setting in the literature in which the functions B_0 and B_1 are linear but depend directly on a vector \mathbf{x} of two components.³ This represents a non-standard bunching-design setting, but fits naturally within the framework of this section.

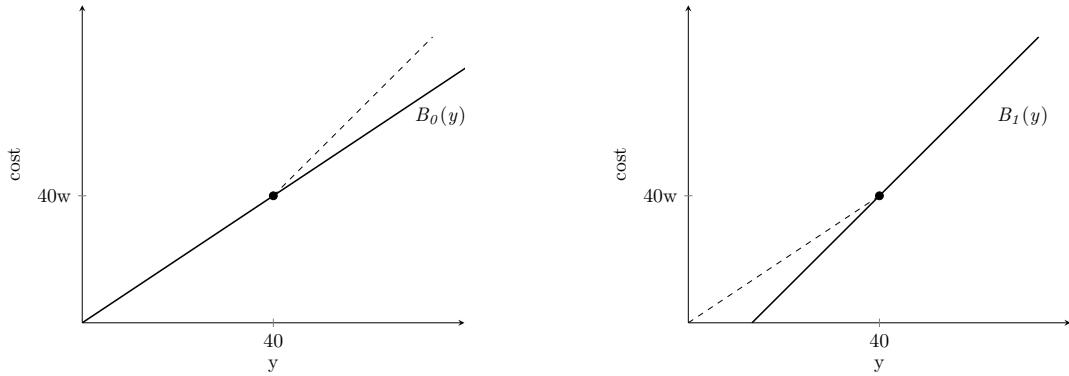


Figure 1: Definition of counterfactual cost functions B_0 and B_1 that firms could have faced, absent the overtime kink. Regardless of what choice variables are in \mathbf{x} , these functions only depend on $y_i(\mathbf{x})$ and are thus depicted as a function of y . Dashed lines show the rest of actual kinked-cost function in comparison to the counterfactual as a solid line. Note that we use the notation y here to indicate hours, rather than the h used in the main text.

Even when the functions B_0 and B_1 only depend on \mathbf{x} through $y_i(\mathbf{x})$, as in standard settings, the bunching design is compatible with models in which multiple margins of choice respond to the incentives provided by the kink. As discussed in the overtime context, the econometrician may be agnostic as to even what the full set of components of \mathbf{x} are, with $B_{0i}(\cdot)$, $B_{1i}(\cdot)$, and $y_i(\cdot)$ depending only on various subsets of the \mathbf{x} that are possibly heterogeneous by i (this is allowable because y need only be continuous in \mathbf{x} , and the cost functions only need to be continuous and *weakly* convex in \mathbf{x} , both of which are compatible with zero dependence on some of its components). Appendix G.5 gives an example in which the over-

³Best et al. (2015) study firms in Pakistan that pay either a tax on output or a tax on profit, whichever is higher. The two tax schedules cross when the ratio of profits to output crosses a certain threshold that is pinned down by the two respective tax rates. In this case, the variable y depends both on production and on reported costs, leading to two margins of response to the kink: one from choosing the scale of production and the other from choosing whether and how much to misreport costs. In this setting a distinction between y and \mathbf{x} cannot be avoided. The authors use features of the function $y_i(\mathbf{x})$ to argue that the bunching reveals changes mostly to reported costs rather than to output (see Appendix G.5 for details).

time kink gives firms an incentive to reduce bonuses, which appear in firm costs but not in the kink the variable y .

In general, the bunching design allows us to conduct causal inference on $Y_i = y_i(X_i)$, but not directly on the underlying choice variables X_i . For example in the overtime setting with possible evasion (see Sec. A.3), bunching at 40 hours will be informative about the effect of a move from B_0 to B_1 on reported hours worked y . However, it will not disentangle whether the effect on hours actually worked is attenuated by, for example, an increase in hours worked off-the-clock. The empirical setting of Best et al. (2015) provides another environment in which this point is relevant (see Appendix G.5).

A.2 Potential outcomes as counterfactual choices

Here I restate slightly more general versions of assumptions CONVEX and CHOICE from Section 4, in the present notation. As in Section 4, let us define a pair of potential outcomes as what would occur if the decision-maker faced either of the functions B_0 or B_1 globally, without the kink.

Definition (potential outcomes). *Let Y_{0i} be the value of $y_i(\mathbf{x})$ that would occur for unit i if $d(i)$ faced the constraint $z \geq B_0(\mathbf{x})$, and let Y_{1i} be the value that would occur under the constraint $z \geq B_1(\mathbf{x})$.*

I again make explicit the assumption that these potential outcomes reflect choices made by the decision-maker. For any function B let Y_{Bi} be the outcome that would occur under the choice constraint $z \geq B(\mathbf{x})$, with Y_{0i} and Y_{1i} shorthands for $Y_{B_{0i}i}$ and $Y_{B_{0i}i}$, respectively. In this notation, the actual outcome Y_i observed by the econometrician is equal to $Y_{B_{ki}i}$.

Assumption CHOICE (perfect manipulation of y). *For any function $B(\mathbf{x})$, $Y_{Bi} = y_i(\mathbf{x}_{Bi})$, where $(z_{Bi}, \mathbf{x}_{Bi})$ is the choice that $d(i)$ would make under the constraint $z \geq B(\mathbf{x})$.*

Assumption CHOICE rules out for example optimization error, which could limit the decision-maker's ability to exactly manipulate values of \mathbf{x} and hence y . It also takes for granted that counterfactual choices are unique, and rules out some kinds of extensive margin effects in which a decision-maker would not choose any value of Y at all under B_1 or B_0 . Note that CHOICE here is slightly stronger than the version given in the main text in that it applies to all functions B , not just B_0 , B_1 and B_k (this is useful for Theorem 2).

The central behavioral assumption that allows us to reason about the counterfactuals Y_0 and Y_1 is that decision-makers have convex preferences over (c, \mathbf{x}) and dislike costs z :

Assumption CONVEX (strictly convex preferences except at kink, decreasing in z). *For each i and any function $B(\mathbf{x})$, choice is $(z_{Bi}, \mathbf{x}_{Bi}) = \text{argmax}_{z, \mathbf{x}} \{u_i(z, \mathbf{x}) : z \geq B(\mathbf{x})\}$*

where $u_i(z, \mathbf{x})$ is weakly decreasing in z and satisfies

$$u_i(\theta z + (1 - \theta)z^*, \theta \mathbf{x} + (1 - \theta)\mathbf{x}^*) > \min\{u_i(z, \mathbf{x}), u_i(z^*, \mathbf{x}^*)\}$$

for any $\theta \in (0, 1)$ and points $(z, \mathbf{x}), (z^*, \mathbf{x}^*)$ such that $y_i(\mathbf{x}) \neq k$ and $y_i(\mathbf{x}^*) \neq k$.

Note: The function $u_i(\cdot)$ represents preferences over choice variables for unit i , but the preferences are those of the decision maker $d(i)$. I avoid more explicit notation like $u_{d(i),i}(\cdot)$ for brevity. In the overtime setting with firms choosing hours, $u_i(z, \mathbf{x})$ corresponds to the firm's profit function π as a function of the hours of a particular worker this week, and costs this week z for that worker.

Note: The second part of Assumption CONVEX is implied by strict quasi-concavity of the function (z, \mathbf{x}) , corresponding to strictly convex preferences. However it also allows for decision-makers preferences to have "two peaks", provided that one of the peaks is located exactly at the kink. This is useful in cases in which the kink is located at a point that has particular value to decision-makers, such as firms setting weekly hours. For example, suppose that firms choose hours only $\mathbf{x} = h$, and have preferences of the form:

$$u_i(z, h) = af(h) + \phi \cdot \mathbb{1}(h = 40) - z \quad (2)$$

where $f(h)$ is strictly concave. This allows firms to have a behavioral "bias" towards 40 hours, or to extract extra profits when $h = 40$ exactly. Figure 2 depicts an example of such preferences, given an arbitrary linear budget function $B(h)$. Note that if a mass of firms were to have preferences of this form, then it would be natural to expect bunching in the distributions of h_{0it} and h_{1it} , which I allow in Section 5.

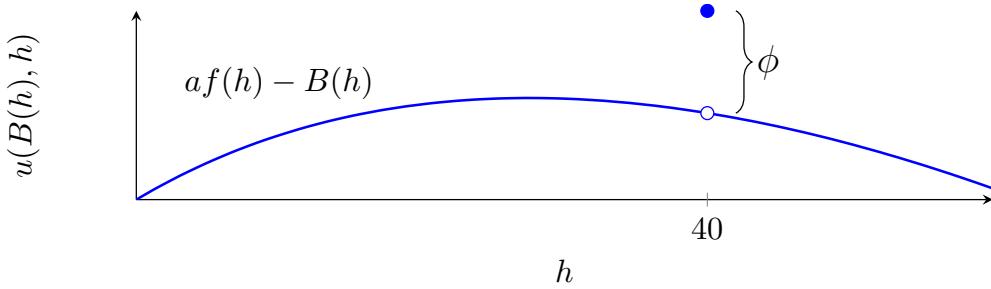


Figure 2: An example of preferences that satisfy CONVEX but are not strictly convex, cf. Eq. (2).

Note: Some departures from CONVEX are allowable without compromising its main implication for the bunching-design, which is given in Lemma 1 below. If B_0 and B_1 are linear in \mathbf{x} and the constraint $z \geq B_k(\mathbf{x})$ can be assumed to bind (hold as an equality), then the assumption that u_i is decreasing in z from CONVEX can be dropped (see Assumption

CONVEX* in Appendix D). If by contrast B_0 and B_1 were strictly (rather than weakly) convex, strict convexity of preferences could be replaced with weakly convex preferences along with an assumption that u_i are strictly decreasing in z (see Eq. (5) in the Proof of Lemma 1).

Note: The notation of Assumption CONVEX does not make explicit any dependence of the functions $u_i(\cdot)$ on the choices made for other observational units $i' \neq i$. When the functions $u_i(\cdot)$ are indeed invariant over such counterfactual choices, we have a version of the no-interference condition of the stable unit treatment values assumption (SUTVA). Maintaining SUTVA is not necessary to define treatment effects in the bunching design, provided that the variables y and z can be coherently defined at the individual unit i level (see Appendix E for details). Nevertheless, the interpretation of the treatment effects identified by the bunching design is most straightforward when SUTVA does hold. This assumption is standard in the bunching design.⁴

A weaker assumption than CONVEX that still has identifying power is simply that decision-makers' choices do not violate the weak axiom of revealed preference:

Assumption WARP (rationalizable choices). *Consider two budget functions B and B' and any unit i . If $d(i)$'s choice under B' is feasible under B , i.e. $z_{B'i} \geq B(\mathbf{x}_{B'i})$, then $(z_{Bi}, \mathbf{x}_{Bi}) = (z_{B'i}, \mathbf{x}_{B'i})$.*

I make the stronger assumption CONVEX for most of the identification results, but Assumption WARP still allows a version of many of them in which equalities become weak inequalities, indicating a degree of robustness with respect to departures from convexity (see Propositions 1 and 2 below). Note that the monotonicity assumption in CONVEX implies that choices will always satisfy $z = B(\mathbf{x})$, i.e. agents' choices will lay on their cost functions (despite Eq. 1 being an inequality, indicating “free-disposal”).

A.3 Examples from the general choice model in the overtime setting

To demonstrate the flexibility of the general choice model CONVEX, I below present some examples for the overtime setting. These examples are meant only to be illustrative, and each could apply to a different subset of units in the population. In these examples we continue to take the decision-maker for a given unit to be the firm employing that worker.⁵

Example 1: Substitution from bonus pay

⁴I note that SUTVA issues like those addressed in Appendix E could also occur in canonical bunching designs: for example if spouses choose their labor supply jointly, the introduction of a tax kink may cause one spouse to increase labor supply while the other decreases theirs.

⁵Appendix D discusses a further example in which the firm and worker bargain over this week's hours. This model can attenuate the wage elasticity of chosen hours since overtime pay gives the parties opposing incentives.

Let the firm's choice vector be $\mathbf{x} = (h, b)'$, where $b \geq 0$ indicates a bonus (or other fringe benefit) paid to the worker. Firms may find it optimal to offer bonuses to improve worker satisfaction and reduce turnover. Suppose firm preferences are: $\pi(z, h, b) = f(h) + g(z + b - \nu(h)) - z - b$, where z continues to denote wage compensation this week, $z + b - \nu(h)$ is the worker's utility with $\nu(h)$ a convex disutility from labor h , and $g(\cdot)$ increasing and concave. In this model firms will choose the surplus maximizing choice of hours $h_m := \operatorname{argmax}_h f(h) - \nu(h)$, provided that the corresponding optimal bonus is non-negative. Bonuses fully adjust to counteract overtime costs, and $h_0 = h_1 = h_m$.

Example 2: Off-the-clock hours and paid breaks

Suppose firms choose a pair $\mathbf{x} = (h, o)'$ with h hours worked and o hours worked "off-the-clock", such that $y(\mathbf{x}) = h - o$ are the hours for which the worker is ultimately paid. Evasion is harder the larger o is, which could be represented by firms facing a convex evasion cost $\phi(o)$, so that firm utility is $\pi(z, h, o) = f(h) - \phi(o) - z$.⁶ This model can also include some firms voluntarily offering paid breaks by allowing o to be negative.

Example 3: Complementaries between workers or weeks

Suppose the firm simultaneously chooses the hours $\mathbf{x} = (h, g)$ of two workers according to production that is isoelastic in a CES aggregate (g could also denote planned hours next week): $\pi(z, h, g) = a \cdot ((\gamma h^\rho + g^\rho)^{1/\rho})^{1+\frac{1}{\epsilon}} - z$ with γ a relative productivity shock. Let g^* denote the firm's optimal choice of hours for the second worker. Optimal h then maximizes $\pi(z, h, g^*)$ subject to $z = B_k(h)$, as if the firm faced a single-worker production function of $f(h) = a \cdot ((\gamma h^\rho + g^{*\rho})^{1/\rho})^{1+\frac{1}{\epsilon}}$. This function is more elastic than $a \cdot h^{1+\frac{1}{\epsilon}}$ provided that $\rho < 1 + 1/\epsilon$, attenuating the response to an increase in w implied by a given ϵ .⁷ Section 4.4 discusses how complementaries affect the final evaluation of the FLSA.

A.4 Observables in the kink bunching design

Lemma 1 outlines the core consequence of Assumption CONVEX for the relationship between observed Y_i and the potential outcomes introduced in the last section:

Lemma 1 (realized choices as truncated potential outcomes). *Under Assumptions CONVEX and CHOICE, the outcome observed given the constraint $z \geq \max\{B_{0i}(\mathbf{x}), B_{1i}(\mathbf{x})\}$*

⁶Note that the data observed in our sample are of hours of work $y(\mathbf{x})$ for which the worker is paid, when this differs from h . Appendix A describes how Equation 2 still holds, but for counterfactual values of hours *paid* $y = h - o$ rather than hours worked h . The bunching design lets us investigate treatment effects on paid hours, without observing off-the-clock hours or break time o .

⁷This expression overstates the degree of attenuation somewhat, since h_1 and h_0 maximize $f(h)$ above for different values g^* , which leads to a larger gap between h_0 and h_1 compared with a fixed g^* by the Le Chatelier principle (Milgrom and Roberts, 1996). However h_1/h_0 still increases on net given $\rho < 1 + 1/\epsilon$.

is:

$$Y_i = \begin{cases} Y_{0i} & \text{if } Y_{0i} < k \\ k & \text{if } Y_{1i} \leq k \leq Y_{0i} \\ Y_{1i} & \text{if } Y_{1i} > k \end{cases}$$

Proof. See Appendix B. \square

Lemma 1 says that the pair of counterfactual outcomes (Y_{0i}, Y_{1i}) is sufficient to pin down actual choice Y_i , which can be seen as an observation of one or the other potential outcome, or k , depending on how the potential outcomes relate to the kink point k .

Note that the “straddling” event $Y_{0i} \leq k \leq Y_{1i}$ from Lemma 1 can be written as $Y_{0i} \in [k, k + \Delta_i]$, where $\Delta_i := Y_{0i} - Y_{1i}$. Similarly, we can also write $Y_{1i} \leq k \leq Y_{0i}$ as $Y_i \in [k - \Delta_i, k]$. This forms the basic link between bunching and *treatment effects*.

Let $\mathcal{B} := P(Y_i = k)$ be the observable probability that the decision-maker chooses to locate exactly at $Y = k$. Proposition 1 gives the relationship between this bunching probability and treatment effects, which holds in a weakened form when CONVEX is replaced by WARP:

Proposition 1 (relation between bunching and Δ_i). *a) Under CONVEX and CHOICE: $\mathcal{B} = P(Y_{0i} \in [k, k + \Delta_i])$; b) under WARP and CHOICE: $\mathcal{B} \leq P(Y_{0i} \in [k, k + \Delta_i])$.*

Proof. See Appendix H. \square

Consider a random sample of observations of Y_i . Under i.i.d. sampling of Y_i , the distribution $F(y)$ of Y_i is identified.⁸ Let $F_1(y) = P(Y_{0i} \leq y)$ be the distribution function of the random variable Y_0 , and $F_1(y)$ the distribution function of Y_1 . From Lemma 1 it follows immediately that $F_0(y) = F(y)$ for all $y < k$, and $F_1(y) = F(y)$ for $y > k$. Thus observations of Y_i are also informative about the marginal distributions of Y_{0i} and Y_{1i} . Again, a weaker version of this also holds under WARP rather than CONVEX:

Proposition 2 (identification of truncated densities). *Suppose that F_0 and F_1 are continuously differentiable with derivatives f_0 and f_1 , and that F admits a derivative function $f(y)$ for $y \neq k$. Under WARP and CHOICE: $f_0(y) \leq f(y)$ for $y < k$ and $f_0(k) \leq \lim_{y \uparrow k} f(y)$, while $f_1(y) \leq f(y)$ for $y > k$ and $f_1(k) \leq \lim_{y \downarrow k} f(y)$, with equalities under CONVEX.*

Proof. See Appendix H. \square

As an example of how WARP alone (without CONVEX) can still be useful for identification, suppose that $\Delta_i = \Delta$ were known to be homogenous across units,⁹ and $f_0(y)$ were constant

⁸Note that in the overtime application sampling is actually at the firm level, which coincides with the level of decision-making units $d(i)$.

⁹One way to get homogenous treatment effects in levels in the overtime setting is to assume exponential production: $f(h) = \gamma(1 - e^{-h/\gamma})$ where $\gamma > 0$ and $h_{0it} - h_{1it} = \gamma \ln(1.5)$ for all units. The iso-elastic model instead gives homogeneous treatment effects for $\log(h)$.

across the interval $[k, k + \Delta]$, then by Propositions 1 and 2 we have that $\Delta \geq \mathcal{B}/f_0(k)$ under WARP and CHOICE.

A.5 Treatment effects in the bunching design

Proposition 1 establishes that bunching can be informative about features of the distribution of treatment effects Δ_i . This section discusses the interpretation of these treatment effects as well as some additional identification results omitted in the main text.

Unit i 's treatment effect $\Delta_i := Y_{0i} - Y_{1i}$ can be thought of as the causal effect of a counterfactual change from the choice set under B_1 to the choice set under B_0 . These treatment effects are “reduced form” in the sense that when the decision-maker has multiple margins of response \mathbf{x} to the incentives introduced by the kink, these may be bundled together in the treatment effect Δ_i (Appendix G.5 discusses this in the setting of Best et al. 2015). This clarifies a limitation sometimes levied against the bunching design, while also revealing a perhaps under-appreciated strength. On the one hand, it is not always clear “which elasticity” is revealed by bunching at a kink, complicating efforts to identify a elasticity parameter having a firm structural interpretation (Einav et al., 2017).

On the other hand, the bunching design can be useful for ex-post policy evaluation and even forecasting effects of small policy changes (as described in Section 4.4), without committing to a tightly parameterized underlying model of choice. This provides a response to the note of caution by Einav et al. (2017), which points out that alternative structural models calibrated from the bunching-design can yield very different predictions about counterfactuals. By focusing on the counterfactuals Y_{0i} and Y_{1i} , we can specify a *particular* type of counterfactual question that can be answered robustly across a broad class of models.

The “trick” of Lemma 1 is to express the observable data in terms of counterfactual choices, rather than of primitives of the utility function. The underlying utility function $u_i(z, \mathbf{x})$ is used only as an intermediate step in the logic, which only requires the nonparametric restrictions of convexity and monotonicity rather than knowing its functional form (or even what vector of choice variables \mathbf{x} underly a given agent's observed value of y). This greatly increases the robustness of the method to potential misspecification of the underlying choice model.

Additional identification results for the bunching design:

While Theorem 1 of Section 4 develops the treatment effect identification result used to evaluate the FLSA, Appendix G presents some further identification results for the bunching design that are not used in this paper, which can be considered alternatives to Theorem 1. This includes re-expressing canonical results from the literature in the general framework of this section, including the linear interpolation approach of Saez (2010), the polynomial ap-

proach of Chetty et al. (2011) and a “small-kink” approximation appearing in Saez (2010) and Kleven (2016). Appendix G also discusses alternative shape constraints to bi-log-concavity, including monotonicity of densities. I also give there a result in which a lower bound to a certain local average treatment effect is identified under WARP, without requiring convexity of preferences.

The buncher ATE when Assumption RANK fails:

This section picks up from the discussion in Section 4.3, but continues with the notation of this Appendix. When RANK fails (and $p = 0$ for simplicity), the bounds from Theorem 1 are still valid under BLC of Y_0 and Y_1 for the following averaged quantile treatment effect:

$$\frac{1}{\mathcal{B}} \int_{F_0(k)}^{F_1(k)} \{Q_0(u) - Q_1(u)\} du = \mathbb{E}[Y_{0i}|Y_{0i} \in [k, k + \Delta_0^*]] - \mathbb{E}[Y_{1i}|Y_{1i} \in [k - \Delta_1^*, k]], \quad (3)$$

where $\Delta_0^* := Q_0(F_1(k)) - Q_1(F_1(k)) = Q_0(F_1(k)) - k$ and $\Delta_1^* := Q_0(F_0(k)) - Q_1(F_0(k)) = k - Q_1(F_0(k))$. Thus, Δ_0^* is the value such that $F_0(k + \Delta_0^*) = F_0(k) + \mathcal{B}$, and Δ_1^* is the value such that $F_1(k - \Delta_1^*) = F_1(k) - \mathcal{B}$. The averaged quantile treatment effect of Eq. (3) yields a lower bound on the buncher ATE, as described in Fig. 3.

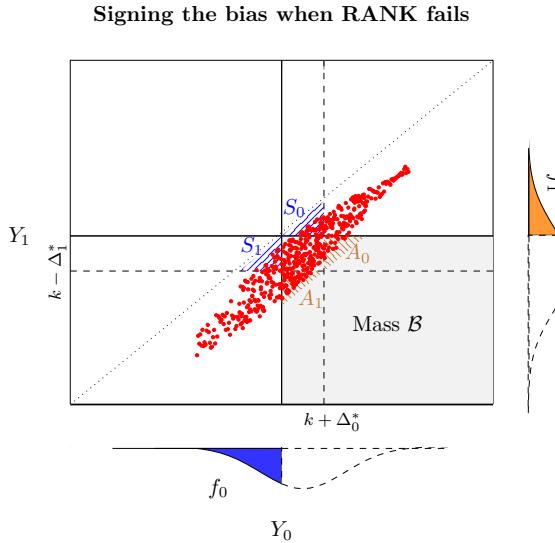


Figure 3: When Assumption RANK fails, the average $\mathbb{E}[Y_{0i}|Y_{0i} \in [k, k + \Delta_0^*]]$ will include the mass in the region S_0 , who are not bunchers (NE lines) but will be missing the mass in the region A_0 (NW lines) who are. This causes an under-estimate of the desired quantity $\mathbb{E}[Y_{0i}|Y_{1i} \leq k \leq Y_{0i}]$. Similarly, $\mathbb{E}[Y_{1i}|Y_{1i} \in [k - \Delta_1^*, k]]$ will include the mass in the region S_1 , who are not bunchers but will be missing the mass in A_1 , who are. This causes an over-estimate of the desired quantity $\mathbb{E}[Y_{1i}|Y_{1i} \leq k \leq Y_{0i}]$.

A.6 Policy changes in the bunching-design

This section presents the logic establishing Theorem 2 in the main text regarding the effects of changes to the policy generating a kink. Consider a bunching design setting in which the cost functions B_0 and B_1 can be viewed as members of family $B_i(\mathbf{x}; \rho, k)$ parameterized by a continuum of scalars ρ and k , where $B_{0i}(\mathbf{x}) = B_i(\mathbf{x}; \rho_0, k^*)$ and $B_{1i}(\mathbf{x}) = B_i(\mathbf{x}; \rho_1, k^*)$ for some $\rho_1 > \rho_0$ and value k^* of k . In the overtime setting ρ represents a wage-scaling factor, with $\rho = 1$ for straight-time and $\rho = 1.5$ for overtime:

$$B_i(y; \rho, k) = \rho w_i y - k w_i (\rho - 1) \quad (4)$$

where work hours y may continue to be a function $y(\mathbf{x})$ of a vector of choice variables to the firm. In this example, k controls the size of the lump-sum subsidy $k w_i (\rho - 1)$ that keeps $B_i(k; \rho, k)$ invariant as ρ is changed.

In the general setting, assume that ρ takes values in a convex subset of \mathbb{R} containing ρ_0 and ρ_1 , and that for any k and $\rho' > \rho$ the cost functions $B_i(\mathbf{x}; \rho, k)$ and $B_i(\mathbf{x}; \rho', k)$ satisfy the conditions of the bunching design framework from Section 4 (with the function $y_i(\mathbf{x})$ fixed across all ρ and k). That is, $B_i(\mathbf{x}; \rho', k) > B_i(\mathbf{x}; \rho, k)$ iff $y_i(\mathbf{x}) > k$ with equality when $y_i(\mathbf{x}) = k$, the functions $B_i(\cdot; \rho, k)$ are weakly convex and continuous, and $y_i(\cdot)$ is continuous. It is readily verified that Equation (4) satisfies these requirements with $y_i(h) = h$.¹⁰

For any value of ρ , let $Y_i(\rho, k)$ be agent i 's realized value of $y_i(\mathbf{x})$ when a choice of (z, \mathbf{x}) is made under the constraint $z \geq B_i(\mathbf{x}; \rho, k)$. A natural restriction in the overtime setting that is that the function $Y_i(\rho, k)$ does not depend on k , and some of the results below will require this. A sufficient condition for $Y_i(\rho, k) = Y_i(\rho)$ is a family of cost functions that are linearly separable in k , as we have in the overtime setting with Equation (4), along with quasi-linearity of preferences. Quasilinearity of preferences is a property of profit-maximizing firms when z represents a cost, and is thus a natural assumption in the overtime setting.

Assumption SEPARABLE (invariance of potential outcomes with respect to k). *For all i, ρ and k , $B_i(\mathbf{x}; \rho, k)$ is additively separable between k and \mathbf{x} (e.g. $b_i(\mathbf{x}, \rho) + \phi_i(\rho, k)$ for some functions b_i and ϕ_i), and for all i $u_i(z, \mathbf{x})$ can be chosen to be additively separable and linear in z .*

Additive separability of $B_i(\mathbf{x}; \rho, k)$ in k may be context specific: in the example from Best et al. (2015) described in Appendix G.5, quasi-linearity of preferences is not sufficient since the cost functions are not additively separable in k . To maintain clarity of exposition, I will keep k implicit in $Y_i(\rho)$ throughout the foregoing discussion, but the proofs make it clear when SEPARABLE is being used.

¹⁰As an alternative example, I construct in Appendix G.5 functions $B_i(\mathbf{x}; \rho, k)$ for the bunching design setting from Best et al. (2015). In that case, ρ parameterizes a smooth transition between an output and a profit tax, where k enters into the rate applied to the tax base for that value of ρ .

Below I state two intermediate results that allow us to derive expressions for the effects of marginal changes to ρ_1 or k on hours. Lemma 2 generalizes an existing result from Blomquist et al. (2015), and makes use of a regularity condition I introduce in the proof as Assumption SMOOTH.¹¹ Counterfactual bunchers $K_i^* = 1$ are assumed to stay at some fixed value k^* (40 in the overtime setting), regardless of ρ and k . Let $p(k) = p \cdot \mathbb{1}(k = k^*)$ denote the possible counterfactual mass at the kink as a function of k . Let $f_\rho(y)$ be the density of $Y_i(\rho)$, which exists by SMOOTH and is defined for $y = k^*$ as a limit (see proof).

Lemma 2 (bunching expressed in terms of marginal responsiveness). *Assume CHOICE, SMOOTH and WARP. Then:*

$$\mathcal{B} - p(k) \leq \int_{\rho_0}^{\rho_1} f_\rho(k) \mathbb{E} \left[-\frac{dY_i(\rho)}{d\rho} \middle| Y_i(\rho) = k \right] d\rho$$

with equality under CONVEX.

Proof. See Appendix H. □

The main tool in establishing Lemma 2 is to relate the integrand in the above to the rate at which kink-induced bunching goes away as the “size” of the kink goes to zero.

Lemma SMALL (small kink limit). *Assume CHOICE*, WARP, and SMOOTH. Then:*

$$\lim_{\rho' \downarrow \rho} \frac{P(Y_i(\rho') \leq k \leq Y_i(\rho)) - p(k)}{\rho' - \rho} = -f_\rho(k) \mathbb{E} \left[\frac{dY_i(\rho)}{d\rho} \middle| Y_i(\rho) = k \right]$$

Proof. See Appendix H. □

Note that the quantity $P(Y_i(\rho') \leq k \leq Y_i(\rho)) - p(k)$ is an upper bound on the bunching that would occur due to a kink between budget functions $B_i(\mathbf{x}; \rho, k)$ and $B_i(\mathbf{x}; \rho', k)$ (under WARP, with equality under CONVEX). As a result, Lemma SMALL shows that the uniform density approximation that has appeared in Saez (2010) and Kleven (2016) (stated in Appendix Theorem 8) for “small” kinks becomes exact in the limit that the two cost functions approach one another. The small kink approximation says that $\mathcal{B} \approx f_\rho(k) \cdot \mathbb{E}[Y_i(\rho) - Y_i(\rho')]$, where note that treatment effects can be written:

$$Y_i(\rho) - Y_i(\rho') = \frac{dY_i(\rho)}{d\rho} (\rho' - \rho) + O((\rho' - \rho)^2)$$

By Lemma 2, we can also see that the RHS in Lemma SMALL evaluated at $\rho = \rho_1$ is equal to the derivative of bunching as ρ_1 is increased, under CONVEX.

Lemma 2 is useful for identification results regarding changes to k when it is combined with a result from Kasy (2022), which considers how the distribution of a generic outcome

¹¹Blomquist et al. (2021) derive the special case of Lemma 2 with convex preferences over a scalar choice variable and $p = 0$, in the context of labor supply under piecewise linear taxation. I establish it here for the general bunching design model where in particular, the $Y_i(\rho)$ may depend on an underlying vector \mathbf{x} which are not observed by the econometrician. I also use different regularity conditions.

variable changes as heterogeneous units flow to different values of that variable in response to marginal policy changes.

Lemma 3 (continuous flows under a small change to ρ). *Under SMOOTH:*

$$\partial_\rho f_\rho(y) = \partial_y \left\{ f_\rho(y) \mathbb{E} \left[-\frac{dY_i(\rho)}{d\rho} \middle| Y_i(\rho) = y, K_i^* = 0 \right] \right\}$$

Proof. See Kasy (2022). □

The intuition behind Lemma 3 comes from the physical dynamics of fluids. When ρ changes, a mass of units will “flow” out of a small neighborhood around any y , and this mass is proportional to the density at y and to the average rate at which units move in response to the change. When the magnitude of this net flow varies with y , the change to ρ will lead to a change in the density there.

With ρ_0 fixed at some value, let us index observed Y_i and bunching \mathcal{B} with the superscript $[k, \rho_1]$ when they occur in a kinked policy environment with cost functions $B_i(\cdot; \rho_0, k)$ and $B_i(\cdot; \rho_1, k)$. Lemmas 2 and 3 together imply Theorem 2 (see Appendix B for proof), which in the notation of this section reads as:

1. $\partial_k \{\mathcal{B}^{[k, \rho_1]} - p(k)\} = f_1(k) - f_0(k)$
2. $\partial_k \mathbb{E}[Y_i^{[k, \rho_1]}] = \mathcal{B}^{[k, \rho_1]} - p(k)$
3. $\partial_{\rho_1} \mathcal{B}^{[k, \rho_1]} = -k f_{\rho_1}(k) \mathbb{E} \left[\frac{dY_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = k \right]$
4. $\partial_{\rho_1} \mathbb{E}[Y_i^{[k, \rho_1]}] = - \int_k^\infty f_{\rho_1}(y) \mathbb{E} \left[\frac{dY_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = y \right] dy$

Note: Assumption SEPARABLE is only necessary for Items 1-2 in Theorem 2, Item 3 holds without it and with $\frac{\partial Y_i(\rho, k)}{\partial \rho}$ replacing $\frac{dY_i(\rho)}{d\rho}$.

B Main proofs

B.1 Proof of Lemma 1

The proof proceeds in the following two steps:

- i) First, I show that $Y_{0i} \leq k$ implies that $Y_i = Y_{0i}$, and similarly $Y_{1i} \geq k$ implies that $Y_i = Y_{1i}$. This holds under CONVEX but also under the weaker assumption of WARP.
- ii) Second, I show that under CONVEX $Y_i < k \implies Y_i = Y_{0i}$ and $Y_i > k \implies Y_i = Y_{1i}$.

Item i) above establishes the first and third cases of Lemma 1. The only remaining possible case is that $Y_{1i} \leq k \leq Y_{0i}$. However, to finish establishing Lemma 1, we also need the reverse implication: that $Y_{1i} \leq k \leq Y_{0i}$ implies $Y_i = k$. This comes from taking the contrapositive of

each of the two claims in item ii).

Proof of i): Let $\mathcal{X}_{0i} = \{\mathbf{x} : y_i(\mathbf{x}) \leq k\}$ and $\mathcal{X}_{1i} = \{\mathbf{x} : y_i(\mathbf{x}) \geq k\}$. If $Y_{0i} \leq k$, then by CHOICE \mathbf{x}_{B_0} is in \mathcal{X}_0 . Since $B_k(\mathbf{x}) = B_0(\mathbf{x})$ for all $\mathbf{x} \in \mathcal{X}_0$, it follows that $z_{B_0i} \geq B_k(\mathbf{x}_{B_0i})$, i.e. Y_{0i} is feasible under B_k . Note that $B_{ki}(\mathbf{x}) \geq B_{0i}(\mathbf{x})$ for all \mathbf{x} . By WARP then $(z_{B_0i}, \mathbf{x}_{B_0i}) = (z_{B_ki}, \mathbf{x}_{B_ki})$. Thus $Y_i = y_i(\mathbf{x}_{B_k}) = y_i(\mathbf{x}_{B_0}) = Y_{0i}$. So $Y_{0i} \leq k \implies Y_i = Y_{0i}$. By the same logic we can show that $Y_{1i} \geq k \implies Y_i = Y_{1i}$.

Proof of ii): For any convex budget function $B(\mathbf{x})$, $(z_{Bi}, \mathbf{x}_{Bi}) = \operatorname{argmax}_{z, \mathbf{x}} \{u_i(z, \mathbf{x}) \text{ s.t. } z \geq B(\mathbf{x})\}$. If $u_i(z, \mathbf{x})$ is strictly quasi-concave, then the RHS exists and is unique since it maximizes u_i over the convex domain $\{(z, \mathbf{x}) : z \geq B(\mathbf{x})\}$. Furthermore, by monotonicity of $u_i(z, \mathbf{x})$ in z we may substitute in the constraint $z = B(\mathbf{x})$ and write

$$\mathbf{x}_{Bi} = \operatorname{argmax}_{\mathbf{x}} u_i(B(\mathbf{x}), \mathbf{x})$$

Suppose that $y_i(\mathbf{x}_{Bi}) \neq k$, and consider any $\mathbf{x} \neq \mathbf{x}_{Bi}$ such that $y_i(\mathbf{x}) \neq k$. Let $\tilde{\mathbf{x}} = \theta\mathbf{x} + (1 - \theta)\mathbf{x}^*$ where $\mathbf{x}^* = \mathbf{x}_{Bi}$ and $\theta \in (0, 1)$. Since $B(\mathbf{x})$ is convex in \mathbf{x} and $u_i(z, \mathbf{x})$ is weakly decreasing in z :

$$u_i(B(\tilde{\mathbf{x}}), \tilde{\mathbf{x}}) \geq u_i(\theta B(\mathbf{x}) + (1 - \theta)B(\mathbf{x}^*), \tilde{\mathbf{x}}) > \min\{u_i(B(\mathbf{x}), \mathbf{x}), u_i(B(\mathbf{x}^*), \mathbf{x}^*)\} = u_i(B(\mathbf{x}), \mathbf{x}) \quad (5)$$

where I have used CONVEX in the second step, and that \mathbf{x}^* is a maximizer in the third. This result implies that for any such $\mathbf{x} \neq \mathbf{x}^*$, if one draws a line between \mathbf{x} and \mathbf{x}^* , the function $u_i(B(\mathbf{x}), \mathbf{x})$ is strictly increasing as one moves towards \mathbf{x}^* . When \mathbf{x} is a scalar, this argument is used by Blomquist et al. (2015) (see Lemma A1 therein) to show that $u_i(B(\mathbf{x}), \mathbf{x})$ is strictly increasing to the left of \mathbf{x}^* , and strictly decreasing to the right of \mathbf{x}^* . Note that for any (binding) linear budget constraint $B(\mathbf{x})$, the result holds without monotonicity of $u_i(z, \mathbf{x})$ in z . This is useful for Theorem 1* in which some workers choose their hours.

For any function B , let $u_{Bi}(\mathbf{x}) = u_i(B(\mathbf{x}), \mathbf{x})$, and note that

$$u_{B_{ki}}(\mathbf{x}) = \begin{cases} u_{B_0i}(\mathbf{x}) & \text{if } \mathbf{x} \in \mathcal{X}_{0i} \\ u_{B_1i}(\mathbf{x}) & \text{if } \mathbf{x} \in \mathcal{X}_{1i} \end{cases}$$

Let \mathbf{x}_{ki} be the unique maximizer of $u_{B_{ki}}(\mathbf{x})$, where $Y_i = y_i(\mathbf{x}_{ki})$. Suppose that $Y_i < k$. Suppose furthermore that $Y_{0i} \neq Y_i$, with $Y_{0i} = y_i(\mathbf{x}_{0i})$ and \mathbf{x}_{0i} the maximizer of $u_{B_0i}(\mathbf{x})$. Note that we must have that $\mathbf{x}_{0i} \notin \mathcal{X}_{0i}$, because $B_0 = B_k$ in \mathcal{X}_{0i} so we can't have $u_{B_0i}(\mathbf{x}_{0i}) > u_{B_0i}(\mathbf{x}_{ki})$ (since \mathbf{x}_{ki} maximizes $u_{B_{ki}}(\mathbf{x})$). Thus $Y_{0i} > k$.

By continuity of $y_i(\mathbf{x})$, \mathcal{X}_{0i} is a closed set and \mathbf{x}_{ki} belongs to the interior of \mathcal{X}_{0i} . Thus, while \mathbf{x}_{0i} is not in \mathcal{X}_{0i} , there exists a point $\tilde{\mathbf{x}} \in \mathcal{X}_{0i}$ along the line between \mathbf{x}_{0i} to \mathbf{x}_{ki} .

Since $Y_i \neq k$ and $Y_{0i} \neq k$, Eq. (5) with $B = B_k$ then implies that $u_{B_k i}(\tilde{\mathbf{x}}) > u_{B_k i}(\mathbf{x}_{0i})$. Since $u_{B_0 i}(\mathbf{x}) = u_{B_k i}(\mathbf{x})$ for all \mathbf{x} in \mathcal{X}_{0i} , it follows that $u_{B_0 i}(\tilde{\mathbf{x}}) > u_{B_0 i}(\mathbf{x}_{0i})$. However, this contradicts the premise that \mathbf{x}_{0i} maximizes $u_{B_0 i}(\mathbf{x})$. Thus, $Y_i < k$ implies $Y_i = Y_{0i}$. Figure 4 depicts the logic visually. The proof that $Y_i > k$ implies $Y_i = Y_{1i}$ is analogous.

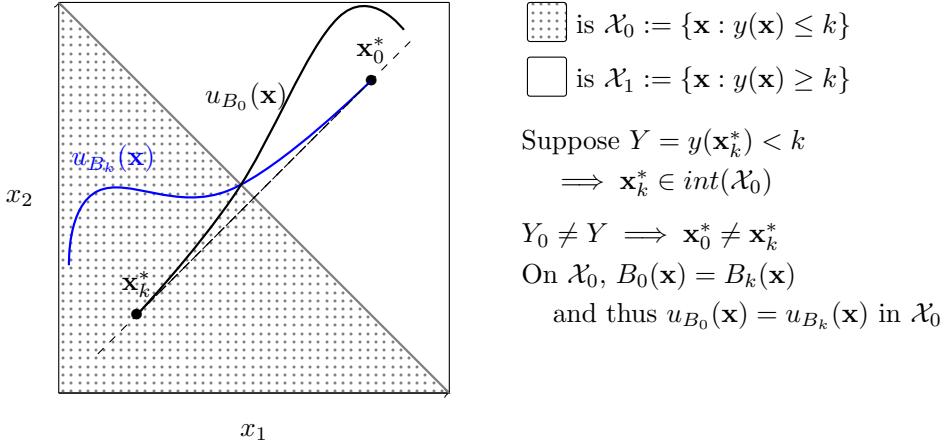


Figure 4: Depiction of the step establishing $(Y < k) \implies (Y = Y_0)$ in the proof of Lemma 1. In this example $z = (x_1, x_2)$ and $y(\mathbf{x}) = x_1 + x_2$. We suppress indices i for clarity. Proof is by contradiction. If $Y_0 \neq Y$, then $\mathbf{x}_k^* \neq \mathbf{x}_0^*$, where \mathbf{x}_k^* and \mathbf{x}_0^* are the unique maximizers of $u_{B_k}(\mathbf{x})$ and $u_{B_0}(\mathbf{x})$, respectively. By Equation 5, we have that the function $u_{B_0}(\mathbf{x})$, depicted heuristically as a solid black curve, is strictly increasing as one moves along the dotted line from \mathbf{x}_k^* towards \mathbf{x}_0^* . Similarly, the function $u_{B_0}(\mathbf{x})$, depicted as a solid blue curve, is strictly increasing as one moves in the opposite direction along the same line, from \mathbf{x}_0^* towards \mathbf{x}_k^* . By the assumption that $Y < k$, then using continuity of $y(\mathbf{x})$ it must be the case that \mathbf{x}_k^* lies in the interior of \mathcal{X}_0 , the set of \mathbf{x} 's that make $y(\mathbf{x}) \leq k$. This means that there is some interval of the dotted line that is within \mathcal{X}_0 . On this interval, the functions B_0 and B_k are equal, and thus so must be the functions u_{B_k} and u_{B_0} . Since the same function cannot be both strictly increasing and strictly decreasing, we have obtained a contradiction.

B.2 Proof of Theorem 1

Theorem 1 of Dümbgen et al. (2017) gives a characterization of bi-log concavity in terms of a random variable's CDF *and* its density. In our case this reads as follows: for $d \in \{0, 1\}$ and any t ,

$$1 - (1 - F_{d|K^*=0}(k))e^{-\frac{f_{d|K^*=0}(k)}{1-F_{d|K^*=0}(k)}t} \leq F_{d|K^*=0}(k+t) \leq F_{d|K^*=0}(k)e^{\frac{f_{d|K^*=0}(k)}{F_{d|K^*=0}(k)}t}$$

Defining $u = F_{0|K^*=0}(k+t)$, we can use the substitution $t = Q_{0|K^*=0}(u) - k$ to translate the above into bounds on the conditional quantile function of Y_{0i} , evaluated at u :

$$\frac{F_{0|K^*=0}(k)}{f_{0|K^*=0}(k)} \cdot \ln\left(\frac{u}{F_{0|K^*=0}(k)}\right) \leq Q_{0|K^*=0}(u) - k \leq -\frac{1 - F_{0|K^*=0}(k)}{f_{0|K^*=0}(k)} \cdot \ln\left(\frac{1-u}{1-F_{0|K^*=0}(k)}\right)$$

And similarly for Y_1 , letting $v = F_{1|K^*=0}(k - t)$:

$$\frac{1 - F_{1|K^*=0}(k)}{f_{1|K^*=0}(k)} \cdot \ln \left(\frac{1 - v}{1 - F_{1|K^*=0}(k)} \right) \leq k - Q_{1|K^*=0}(v) \leq -\frac{F_{1|K^*=0}(k)}{f_{1|K^*=0}(k)} \cdot \ln \left(\frac{v}{F_{1|K^*=0}(k)} \right)$$

Note that, given Assumption RANK and Lemma 1:

$$\begin{aligned} E[Y_{0i} - Y_{1i}|Y_i = k, K_i^* = 0] &= \frac{1}{\mathcal{B}^*} \int_{F_{0|K^*=0}(k)}^{F_{0|K^*=0}(k)+\mathcal{B}^*} \{Q_{0|K^*=0}(u) - Q_{0|K^*=0}(u)\} du \\ &= \frac{1}{\mathcal{B}^*} \int_{F_{0|K^*=0}(k)}^{F_{0|K^*=0}(k)+\mathcal{B}^*} \{Q_{0|K^*=0}(u) - k\} du + \frac{1}{\mathcal{B}^*} \int_{F_{1|K^*=0}(k)-\mathcal{B}^*}^{F_{1|K^*=0}(k)} \{k - Q_{1|K^*=0}(v)\} dv \end{aligned}$$

where $\mathcal{B}^* := P(h_{it} = k|K^* = 0)$. A lower bound for $E[Y_{0i} - Y_{1i}|Y_i = k, K_i^* = 0]$ is thus:

$$\begin{aligned} &\frac{F_{0|K^*=0}(k)}{f_{0|K^*=0}(k) \cdot \mathcal{B}^*} \int_{F_{0|K^*=0}(k)}^{F_{0|K^*=0}(k)+\mathcal{B}^*} \ln \left(\frac{u}{F_{0|K^*=0}(k)} \right) du + \frac{1 - F_{1|K^*=0}(k)}{f_{1|K^*=0}(k) \cdot \mathcal{B}^*} \int_{F_{1|K^*=0}(k)-\mathcal{B}^*}^{F_{1|K^*=0}(k)} \ln \left(\frac{1 - v}{1 - F_{1|K^*=0}(k)} \right) dv \\ &= g(F_{0|K^*=0}(k), f_{0|K^*=0}(k), \mathcal{B}^*) + h(F_{1|K^*=0}(k), f_{1|K^*=0}(k), \mathcal{B}^*) \end{aligned}$$

where

$$\begin{aligned} g(a, b, x) &:= \frac{a}{bx} \int_a^{a+x} \ln \left(\frac{u}{a} \right) du = \frac{a^2}{bx} \int_1^{1+\frac{x}{a}} \ln(u) du \\ &= \frac{a^2}{bx} \{u \ln(u) - u\} \Big|_1^{1+\frac{x}{a}} = \frac{a^2}{bx} \left\{ \left(1 + \frac{x}{a}\right) \ln \left(1 + \frac{x}{a}\right) - \frac{x}{a} \right\} \\ &= \frac{a}{bx} (a + x) \ln \left(1 + \frac{x}{a}\right) - \frac{a}{b} \end{aligned}$$

and

$$h(a, b, x) := \frac{1 - a}{bx} \int_{a-x}^a \ln \left(\frac{1 - v}{1 - a} \right) dv = \frac{(1 - a)^2}{bx} \int_1^{1+\frac{x}{1-a}} \ln(u) du = g(1 - a, b, x)$$

Similarly, an upper bound is:

$$\begin{aligned} &-\frac{1 - F_{0|K^*=0}(k)}{f_{0|K^*=0}(k)(\mathcal{B}^*)} \int_{F_{0|K^*=0}(k)}^{F_{0|K^*=0}(k)+\mathcal{B}^*} \ln \left(\frac{1 - u}{1 - F_{0|K^*=0}(k)} \right) du \\ &\quad - \frac{F_{1|K^*=0}(k)}{f_{1|K^*=0}(k)(\mathcal{B}^*)} \int_{F_{1|K^*=0}(k)-\mathcal{B}^*}^{F_{1|K^*=0}(k)} \ln \left(\frac{v}{F_{1|K^*=0}(k)} \right) dv \\ &= \tilde{g}(F_{0|K^*=0}(k), f_{0|K^*=0}(k), \mathcal{B}^*) + \tilde{h}(F_{1|K^*=0}(k), f_{1|K^*=0}(k), \mathcal{B}^*) \end{aligned}$$

where

$$\begin{aligned} \tilde{g}(a, b, x) &:= -\frac{1 - a}{bx} \int_a^{a+x} \ln \left(\frac{1 - u}{1 - a} \right) du = -\frac{(1 - a)^2}{bx} \int_{1-\frac{x}{1-a}}^1 \ln(u) du \\ &= \frac{(1 - a)^2}{bx} \{u - u \ln(u)\} \Big|_{1-\frac{x}{1-a}}^1 = \frac{1 - a}{b} + \frac{1 - a}{bx} (1 - a - x) \ln \left(1 - \frac{x}{1 - a}\right) \\ &= -g(1 - a, b, -x) \end{aligned}$$

and

$$\tilde{h}(a, b, x) := -\frac{a}{bx} \int_{a-x}^a \ln\left(\frac{v}{a}\right) dv = -\frac{a^2}{bx} \int_{1-\frac{x}{a}}^1 \ln(u) du = \tilde{g}(1-a, b, x) = -g(a, b, -x)$$

Given p , we relate the $K^* = 0$ conditional quantities to their unconditional analogues:

$$F_{0|K^*=0}(k) = \frac{F_0(k) - p}{1-p} \quad \text{and} \quad F_{1|K^*=0}(k) = \frac{F_1(k) - p}{1-p} \quad \text{and} \quad \mathcal{B}^* = \frac{\mathcal{B} - p}{1-p}$$

$$f_{0|K^*=0}(k) = \frac{f_0(k)}{1-p} \quad \text{and} \quad f_{1|K^*=0}(k) = \frac{f_1(k)}{1-p}$$

Let $F(h) = P(h_{it} \leq h)$ be the CDF of the data, and define $f(h) = \frac{d}{dh} P(h_{it} \leq h)$ for $h \neq k$. By Proposition 2 and the BLC assumption, the above quantities are related to observables as:

$$F_0(k) = \lim_{h \uparrow k} F(h) + p, \quad F_1(k) = F(k), \quad f_0(k) = \lim_{h \uparrow k} f(h), \quad \text{and} \quad f_1(k) = \lim_{h \downarrow k} f(h)$$

As shown by Dümbgen et al. (2017), BLC implies the existence of a continuous density function, which assures that the required density limits exist, and delivers Item 1. of the theorem.

To obtain the final result, note that the function $g(a, b, x)$ is homogeneous of degree zero. Thus $\Delta_k^* \in [\Delta_k^L, \Delta_k^U]$, with

$$\Delta_k^L := g(F_-(k), f_-(k), \mathcal{B} - p) + g(1 - F(k), f_+(k), \mathcal{B} - p)$$

$$\Delta_k^U := -g(1 - p - F_-(k), f_-(k), p - \mathcal{B}) - g(F(k) - p, f_+(k), p - \mathcal{B})$$

where $-$ and $+$ subscripts denote left and right limits. The bounds are sharp as CHOICE, CONVEX and RANK imply no further restrictions on the potential outcome distributions.

B.3 Proof of Theorem 2

This proof follows the notation of Appendix A. Throughout this proof we let $Y_i(\rho, k) = Y_i(\rho)$, given Assumption SEPARABLE. By Appendix A Lemmas 2 and 3 the effect of changing k on bunching is:

$$\begin{aligned} \partial_k \{\mathcal{B} - p(k)\} &= -\frac{\partial}{\partial k} \int_{\rho_0}^{\rho_1} f_\rho(k) \mathbb{E} \left[\frac{Y_i(\rho)}{d\rho} \middle| Y_i(\rho) = k \right] d\rho \\ &= -\int_{\rho_0}^{\rho_1} \frac{\partial}{\partial k} \left\{ f_\rho(k) \mathbb{E} \left[\frac{Y_i(\rho)}{d\rho} \middle| Y_i(\rho) = k \right] \right\} d\rho = \int_{\rho_0}^{\rho_1} \partial_\rho f_\rho(k) d\rho = f_1(k) - f_0(k) \end{aligned}$$

Turning now to the total effect on average hours.

$$\begin{aligned} \partial_k E[Y_i^{[k, \rho_1]}] &= \partial_k \{P(Y_i(\rho_0) < k) \mathbb{E}[Y_i(\rho_0) | Y_i(\rho_0) < k]\} + k \partial_k (\mathcal{B}^{[k, \rho_1]} - p(k)) + \mathcal{B}^{[k, \rho_1]} - p(k) \\ &\quad + \partial_k \{P(Y_i(\rho_1) > k) \mathbb{E}[Y_i(\rho_1) | Y_i(\rho_1) > k]\} \\ &= \partial_k \int_{-\infty}^k y \cdot f_{\rho_0}(y) dy + k(f_0(k) - f_1(k)) + \mathcal{B}^{[k, \rho_1]} - p(k) + \partial_k \int_k^\infty y \cdot f_{\rho_1}(y) dy \\ &= k f_0(k) + k \underline{(f_1(k) - f_0(k))} + \mathcal{B}^{[k, \rho_1]} - p(k) - k f_1(k) \end{aligned}$$

Meanwhile: $\partial_{\rho_1} \mathbb{E}[Y_i^{[k, \rho_1]}] = - \int_k^\infty f_{\rho_1}(y) \mathbb{E} \left[\frac{dY_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = y \right] dy$ follows directly from Lemma 2 and differentiating both sides with respect to ρ_1 , and thus

$$\begin{aligned} \partial_{\rho_1} E[Y_i^{[k, \rho_1]}] &= k \partial_{\rho_1} \mathcal{B}^{[k, \rho_1]} + \partial_{\rho_1} \{P(Y_i(\rho_1) > k) \mathbb{E}[Y_i(\rho_1)|Y_i(\rho_1) > k]\} = k \partial_{\rho_1} \mathcal{B}^{[k, \rho_1]} + \int_k^\infty y \cdot \partial_{\rho_1} f_{\rho_1}(y) \cdot dy \\ &= -k f_{\rho_1}(k) \mathbb{E} \left[\frac{Y_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = k \right] - \int_k^\infty y \cdot \partial_y \left\{ f_{\rho_1}(y) \mathbb{E} \left[\frac{dY_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = y \right] \right\} dy \\ &= -k f_{\rho_1}(k) \mathbb{E} \left[\frac{Y_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = k \right] + \underbrace{y f_{\rho_1}(y) \mathbb{E} \left[\frac{dY_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = y \right] \Big|_{y=k}}_{-\int_k^\infty f_{\rho_1}(y) \mathbb{E} \left[\frac{dY_i(\rho_1)}{d\rho} \middle| Y_i(\rho_1) = y \right] dy} \end{aligned}$$

where I have used Lemma 2 with the Leibniz rule (establishing Item 3 in Theorem 2) as well as Lemma 3 in the third step, and then integration by parts along with the boundary condition that $\lim_{y \rightarrow \infty} y \cdot f_{\rho_1}(y) = 0$, implied by Assumption SMOOTH.

C Additional empirical information and results

C.1 Sample restrictions

Beginning with the initial sample described in Column (2) of Table 1, I keep paychecks from workers who are paid on a weekly basis, and condition on paychecks that contain a record of positive hours for work, vacation, holidays, or sick leave, totaling fewer than 80 hours in a week.¹² I also drop observations from California, which has a daily overtime rule that is binding for a significant number of workers, and could confound the effects of the weekly FLSA rule.

Further, I focus on hourly workers. While the data include a field for the employer to input a salary, there is no guarantee that employers actually use this feature in the payroll software. Therefore, I use a combination of sampling restrictions to ensure I remove all non-hourly workers from the sample. First, I drop workers that ever have a salary on file with the payroll system. Second, I only keep workers at firms for whom *some* workers have a salary on file, the assumption being that employers either don't use the feature at all or use it for all of their salaried employees. I also drop paychecks from workers for whom hours are recorded as 40 in every week that they appear in the data,¹³ as it is possible that these workers are simply coded as working 40 hours despite being paid on a salary basis. I also drop workers who never receive overtime pay.

¹²This restriction removes about 2% of the sample after the other restrictions. While a genuine 80 hour workweek is possible, I consider these observations to likely correspond to two weeks of work despite the worker's pay frequency being coded as weekly.

¹³For the purposes of this restriction, I count the "40 hours" event as occurring when either hours worked or hours paid is equal to 40.

C.2 A test of the Trejo (1991) model of straight-time wage adjustment

One way to assess the role of the wage rigidity reported in Table 2 is to test directly whether straight-time wages and hours are plausibly related *at the weekly level* according to Equation (1). Given the kink in Eq. (1), we can perform such a test using the wage and hours reported on each paycheck, while making only weak differentiability assumptions on unobservables for identification.

Suppose that for some subset of units it , wages are actively adjusted to the hours they work according to Equation (1), in order to target some total earnings z_{it} . Denote the corresponding units by a latent variable $A_{it} = 1$. These units may come from workers with limited variation in their schedules in those weeks in which $h_{it} = h_i^*$ for some typical hours h_i^* according to which their wages were set by Eq. (1) at hiring. $A_{it} = 1$ units might instead have dynamic wages that adjust to week-by-week variation in their hours h_{it} . Let $A_{it} = 0$ denote units for whom the worker's wage is determined in some other way. Let $q(h) = P(A_{it} = 1|h_{it} = h)$ denote the proportion of these two groups at various points in the hours distribution. An extreme version of the fixed-job model of Trejo (1991) for example, would have $q(h) = 1$ for all h .

By the law of iterated expectations and some algebra we have that:

$$\begin{aligned}\mathbb{E} [\ln w_{it}|h_{it} = h] &= q(h) \{\mathbb{E} [\ln z_{it}|h_{it} = h, A_{it} = 1] - \ln (h + 0.5(h - 40)\mathbb{1}(h \geq 40))\} \\ &\quad - (1 - q(h))\mathbb{E} [\ln w_{it}|h_{it} = h, A_{it} = 0]\end{aligned}$$

The middle term above introduces a kink in the conditional expectation of the log of straight-time wages with respect to hours. If we assume that $\mathbb{E} [\ln z_{it}|h_{it} = h, A_{it} = 1]$, $\mathbb{E} [\ln w_{it}|h_{it} = h, A_{it} = 0]$ and $q(h)$ are all continuously differentiable in h , then the magnitude of this kink identifies $q(40)$, the proportion of active wage responders local to $h = 40$:

$$\lim_{h \downarrow 40} \frac{d}{dh} \mathbb{E} [\ln w_{it}|h_{it} = h] - \lim_{h \uparrow 40} \frac{d}{dh} \mathbb{E} [\ln w_{it}|h_{it} = h] = -\frac{1}{2} \cdot \frac{q(40)}{40}$$

These continuous differentiability assumptions are reasonable, if wage setting according to Equation (1) is the only force introducing non-smoothness in the relationship between wages and hours at 40. For instance, we assume that production technologies do not have any special features at 40 hours that would cause the distribution of target earnings levels z_{it} among the $A_{it} = 1$ units to itself have a kink around $h_{it} = 40$.

Figure 5 reports the results of fitting separate local linear functions to the CEF of log wages on either side of $h = 40$. We can reject the hypothesis that the fixed-job model applies to all employees at all times, near 40. However, the data appear to be consistent with a proportion $q(40) \approx 0.25$ of all paychecks close to 40 hours reflecting an hours/wage relationship governed by Equation (1). This can be rationalized by straight-wages being updated intermittently to reflect expected or anticipated hours, which vary in practice quite a bit between pay periods.

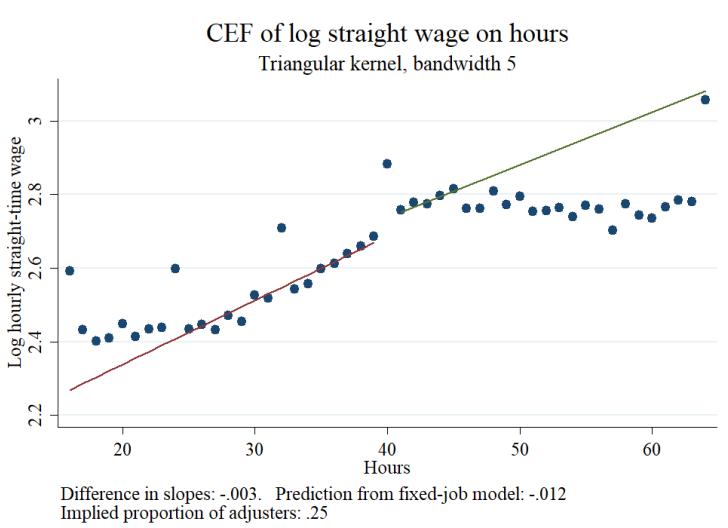


Figure 5: A test of the fixed-jobs model presented in Trejo (1991), based on the magnitude of the kink in the conditional expectation of log wages with respect to hours at 40 (see above). Regression lines fit on each side with a uniform kernel within 25 hours of the 40. This figure closely resembles Figure 5 of Bick et al. (2022) which uses CPS data for hourly workers.

C.3 Further characteristics of the sample

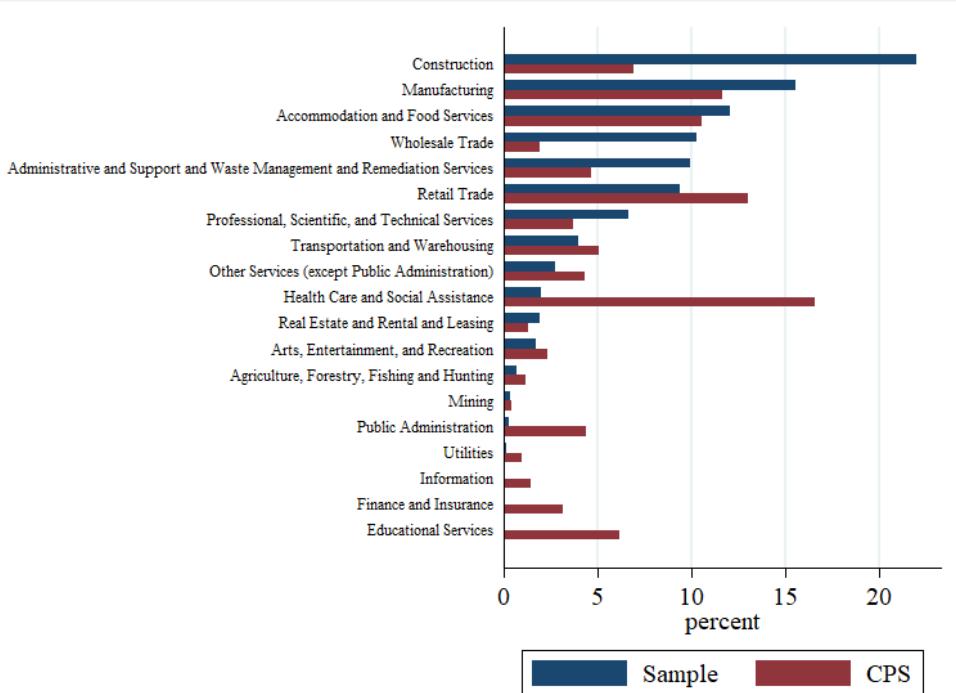


Figure 6: Industry distribution of estimation sample versus the Current Population Survey sample described in Section 3.

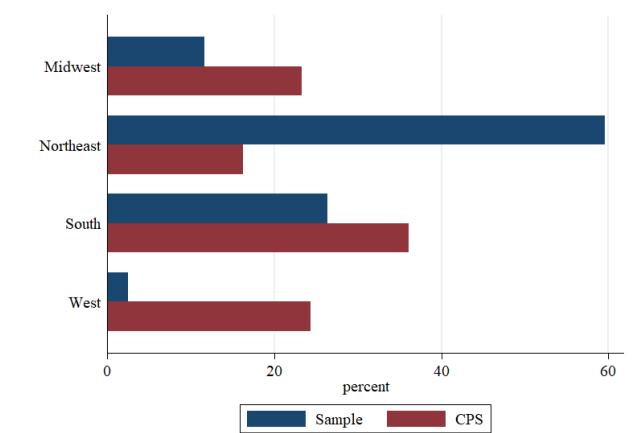


Figure 7: Geographical distribution of estimation sample versus the Current Population Survey sample described in Section 3.

Industry	Avg. OT hours	OT % hours	OT % pay	Industry share
Accommodation and Food Services	2.37	0.06	0.11	0.08
Administrative and Support	5.69	0.13	0.18	0.08
Agriculture, Forestry, Fishing and Hunting	3.76	0.11	0.15	0.00
Arts, Entertainment, and Recreation	3.87	0.10	0.13	0.00
Construction	3.09	0.07	0.10	0.20
Educational Services	1.83	0.05	0.07	0.00
Finance and Insurance	0.31	0.00	0.01	0.00
Health Care and Social Assistance	4.59	0.12	0.12	0.02
Information	1.67	0.04	0.06	0.00
Manufacturing	3.37	0.08	0.11	0.18
Mining	2.26	0.07	0.12	0.00
Other Services	2.61	0.06	0.09	0.02
Professional, Scientific, and Technical Services	2.91	0.07	0.10	0.06
Public Administration	2.36	0.05	0.08	0.00
Real Estate and Rental and Leasing	2.85	0.07	0.09	0.02
Retail Trade	2.83	0.07	0.10	0.08
Transportation and Warehousing	5.24	0.12	0.17	0.04
Utilities	3.80	0.08	0.11	0.00
Wholesale Trade	5.15	0.11	0.14	0.10
Total Sample	3.55	0.08	0.12	0.98

Table 1: Overtime prevalence by industry in the sample, including average number of OT hours per weekly paycheck, % OT hours among hours worked, % pay for hours work going to OT, and industry share of total hours in sample.

	(1) Work hours=40	(2) OT hours	(3) Total work hours	(4) Work hours=40	(5) OT hours
Tenure	0.000400 (0.95)	0.0515 (3.95)	0.0796 (3.31)		
Age	0.000690 (3.82)	0.00266 (0.74)	0.0250 (3.25)		
Female	0.0140 (2.08)	-1.322 (-9.07)	-1.943 (-6.08)		
Minimum wage worker	0.00121 (0.29)	-1.687 (-2.39)	-5.352 (-4.08)		
Firm just hired				-0.00572 (-2.95)	0.553 (5.78)
Date FE	Yes	Yes	Yes	Yes	Yes
Employer FE	Yes	Yes	Yes		
Worker FE				Yes	Yes
Observations	499619	499619	499619	628449	628449
R squared	0.229	0.264	0.260	0.387	0.515

t statistics in parentheses

Table 2: Columns (1)-(3) regress hours-related outcome variables on worker characteristics, with fixed effects for the date and employer. Standard errors clustered by firm. Columns (4)-(5) show that bunching and overtime hours among incumbent workers are both responsive to new workers being hired within a firm, even controlling for worker and day fixed effects. “Firm just hired” indicates that at least one new worker appears in payroll at the firm this week, and the new workers are dropped from the regression. “Minimum wage worker” indicates that the worker’s straight-time wage is at or below the maximum minimum wage in their state of residence for the quarter. Tenure and age are measured in years, and age is missing for some workers.

	(1) Total work hours	(2) Total work hours	(3) Total work hours
R squared	0.366	0.499	0.626
Date FE		Yes	
Worker FE		Yes	Yes
Employer x date FE	Yes		Yes
Observations	621011	628449	620854

t statistics in parentheses

Table 3: Decomposing variation in total hours. Worker fixed effects and employer by day fixed effects explain about 63% of the variation in total hours.

C.4 Additional treatment effect estimates and figures

	$p=0$		p from PTO	
	Bunching	Buncher ATE	Net Bunching	Buncher ATE
Accommodation and Food Services (N=69427)	0.036 [0.029, 0.044]	[0.937, 0.988] [0.734, 1.212]	0.036 [0.029, 0.044]	[0.937, 0.988] [0.734, 1.212]
Administrative and Support (N=49829)	0.062 [0.051, 0.074]	[1.625, 1.771] [1.313, 2.136]	0.009 [0.005, 0.013]	[0.251, 0.255] [0.143, 0.365]
Construction (N=136815)	0.139 [0.128, 0.149]	[2.759, 3.326] [2.341, 3.854]	0.029 [0.022, 0.035]	[0.612, 0.638] [0.442, 0.821]
Health Care and Social Assistance (N=13951)	0.051 [0.034, 0.069]	[1.412, 1.522] [0.570, 2.450]	0.005 [0.000, 0.010]	[0.146, 0.147] [-0.052, 0.348]
Manufacturing (N=112555)	0.137 [0.126, 0.148]	[2.098, 2.521] [1.894, 2.785]	0.018 [0.016, 0.021]	[0.307, 0.316] [0.255, 0.370]
Other Services (N=19263)	0.160 [0.132, 0.188]	[1.804, 2.240] [1.243, 2.996]	0.037 [0.024, 0.049]	[0.452, 0.478] [0.256, 0.693]
Professional, Scientific, Technical (N=47705)	0.136 [0.117, 0.155]	[2.281, 2.737] [1.862, 3.297]	0.010 [0.003, 0.016]	[0.178, 0.180] [0.060, 0.302]
Real Estate and Rental and Leasing (N=13498)	0.187 [0.141, 0.234]	[3.477, 4.478] [2.432, 6.053]	0.097 [0.060, 0.135]	[1.920, 2.215] [1.065, 3.316]
Retail Trade (N=56403)	0.129 [0.112, 0.146]	[3.694, 4.399] [2.447, 5.935]	0.032 [0.024, 0.040]	[0.969, 1.016] [0.550, 1.463]
Transportation and Warehousing (N=25926)	0.091 [0.070, 0.111]	[2.230, 2.530] [1.754, 3.127]	0.015 [0.009, 0.022]	[0.400, 0.409] [0.216, 0.602]
Wholesale Trade (N=66678)	0.126 [0.110, 0.141]	[2.751, 3.299] [2.321, 3.848]	0.046 [0.037, 0.055]	[1.068, 1.149] [0.765, 1.490]
All Industries (N=630217)	0.116 [0.112, 0.121]	[2.614, 3.054] [2.483, 3.217]	0.027 [0.024, 0.029]	[0.640, 0.666] [0.571, 0.740]

Table 4: Estimates of the buncher ATE by industry, based on $p = 0$ (left) or p estimated from paid time off (right). Estimates are reported only for industries having at least 10,000 observations. 95% bootstrap confidence intervals in gray, clustered by firm.

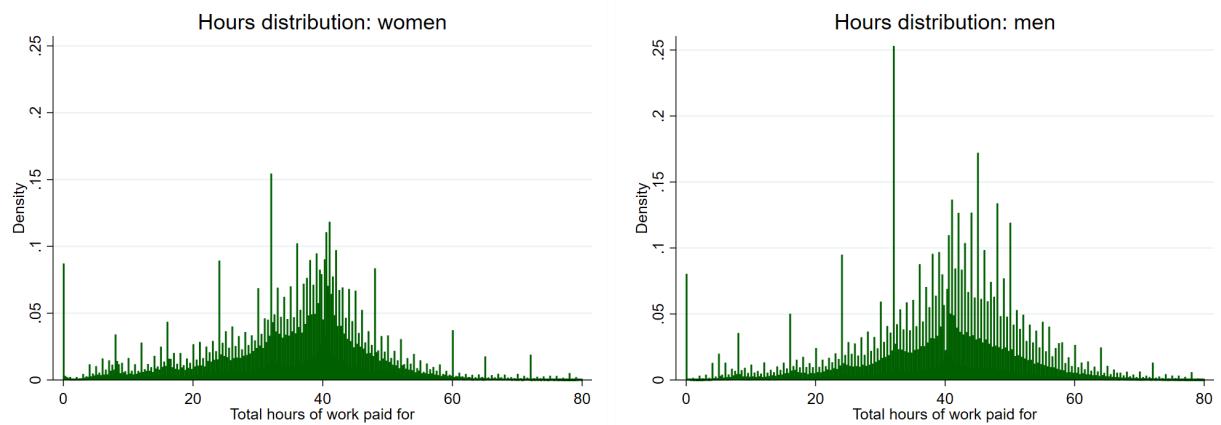


Table 6: Hours distribution by gender, conditional on different than 40 for visibility (size of point mass at 40 can be read from Figures 7 and 8).

	$p=0$		p from PTO	
	Bunching	Effect of the kink	Net Bunching	Effect of the kink
Accommodation and Food Services (N=69427)	0.036 [0.029, 0.044]	[-0.368, -0.248] [-0.450, -0.192]	0.036 [0.029, 0.044]	[-0.368, -0.248] [-0.450, -0.192]
Administrative and Support (N=49829)	0.062 [0.051, 0.074]	[-1.190, -0.681] [-1.424, -0.548]	0.009 [0.005, 0.013]	[-0.178, -0.101] [-0.256, -0.057]
Construction (N=136815)	0.139 [0.128, 0.149]	[-1.550, -1.121] [-1.771, -0.944]	0.029 [0.022, 0.035]	[-0.330, -0.219] [-0.422, -0.157]
Health Care and Social Assistance (N=13951)	0.051 [0.034, 0.069]	[-0.633, -0.320] [-1.020, -0.129]	0.005 [0.000, 0.010]	[-0.065, -0.030] [-0.155, 0.012]
Manufacturing (N=112555)	0.137 [0.126, 0.148]	[-1.167, -0.850] [-1.282, -0.766]	0.018 [0.016, 0.021]	[-0.162, -0.110] [-0.192, -0.090]
Other Services (N=19263)	0.160 [0.132, 0.188]	[-0.977, -0.811] [-1.300, -0.538]	0.037 [0.024, 0.049]	[-0.235, -0.176] [-0.345, -0.095]
Professional, Scientific, Technical (N=47705)	0.136 [0.117, 0.155]	[-1.192, -0.959] [-1.411, -0.767]	0.010 [0.003, 0.016]	[-0.090, -0.063] [-0.150, -0.021]
Real Estate and Rental and Leasing (N=13498)	0.187 [0.141, 0.234]	[-1.766, -1.466] [-2.303, -1.002]	0.097 [0.060, 0.135]	[-0.954, -0.725] [-1.378, -0.392]
Retail Trade (N=56403)	0.129 [0.112, 0.146]	[-1.685, -1.342] [-2.274, -0.908]	0.032 [0.024, 0.040]	[-0.434, -0.308] [-0.626, -0.175]
Transportation and Warehousing (N=25926)	0.091 [0.070, 0.111]	[-1.590, -0.998] [-1.935, -0.783]	0.015 [0.009, 0.022]	[-0.274, -0.166] [-0.406, -0.086]
Wholesale Trade (N=66678)	0.126 [0.110, 0.141]	[-2.122, -1.297] [-2.474, -1.088]	0.046 [0.037, 0.055]	[-0.776, -0.476] [-1.016, -0.333]
All Industries (N=630217)	0.116 [0.112, 0.121]	[-1.466, -1.026] [-1.542, -0.972]	0.027 [0.024, 0.029]	[-0.347, -0.227] [-0.386, -0.202]

Table 5: Estimates of the hours effect of the FLSA by industry, based on $p = 0$ (left) or p estimated from paid time off (right). Estimates are reported only for industries having at least 10,000 observations. 95% bootstrap confidence intervals in gray, clustered by firm. In the case of Accommodation and Food Services, $P(h_{it} = 40 | \eta_{it} > 0) > \mathcal{B}$, so I take the PTO-based estimate to be $p = 0$.

	$p=0$	p from non-changers	p from PTO
Net bunching:	0.090 [0.083, 0.098]	0.044 [0.041, 0.048]	0.011 [0.009, 0.012]
Buncher ATE	[1.507, 1.709] [1.387, 1.855]	[0.763, 0.814] [0.706, 0.877]	[0.187, 0.190] [0.150, 0.227]
Buncher ATE as elasticity	[0.093, 0.105] [0.086, 0.114]	[0.047, 0.050] [0.044, 0.054]	[0.012, 0.012] [0.009, 0.014]
Average effect of kink on hours	[-0.633, -0.489] [-0.688, -0.446]	[-0.319, -0.231] [-0.343, -0.213]	[-0.078, -0.054] [-0.094, -0.043]
Num observations	147953	147953	147953
Num clusters	352	352	352

Table 7: Results of the bunching estimator among women.

	$p=0$	p from non-changers	p from PTO
Net bunching:	0.124 [0.119, 0.129]	0.060 [0.058, 0.063]	0.031 [0.028, 0.034]
Buncher ATE	[3.074, 3.635] [2.777, 3.991]	[1.560, 1.701] [1.407, 1.869]	[0.828, 0.868] [0.717, 0.986]
Buncher ATE as elasticity	[0.190, 0.224] [0.171, 0.246]	[0.096, 0.105] [0.087, 0.115]	[0.051, 0.053] [0.044, 0.061]
Average effect of kink on hours	[-1.867, -1.271] [-2.060, -1.149]	[-0.921, -0.604] [-1.015, -0.545]	[-0.482, -0.311] [-0.549, -0.269]
Num observations	482264	482264	482264
Num clusters	524	524	524

Table 8: Results of the bunching estimator among men.

	$p=0$	p from non-changers	p from PTO
Net bunching:	0.116 [0.112, 0.120]	0.057 [0.055, 0.058]	0.027 [0.024, 0.030]
Treatment effect			
Linear density	2.794 [2.636, 2.952]	1.360 [1.287, 1.432]	0.644 [0.568, 0.719]
Monotonic density	[2.497, 3.171] [2.356, 3.353]	[1.215, 1.544] [1.153, 1.629]	[0.575, 0.731] [0.516, 0.805]
BLC buncher ATE	[2.614, 3.054] [2.493, 3.205]	[1.324, 1.435] [1.264, 1.501]	[0.640, 0.666] [0.574, 0.736]
Num observations	630217	630217	630217
Num clusters	566	566	566

Table 9: Treatment effects in levels with comparison to alternative shape constraints. Rows “Linear density” and “Monotonic density” assume homogenous treatment effects (see Figure 10 for analogous estimates that assume the isoelastic model).

	$p=0$	p from non-changers	p from PTO
Net bunching:	0.116 [0.112, 0.120]	0.057 [0.055, 0.058]	0.027 [0.024, 0.030]
Treatment effect			
Linear density	0.173 [0.163, 0.183]	0.084 [0.079, 0.088]	0.040 [0.035, 0.044]
Monotonic density	[0.154, 0.196] [0.145, 0.207]	[0.075, 0.095] [0.071, 0.100]	[0.035, 0.045] [0.032, 0.050]
BLC buncher ATE	[0.161, 0.188] [0.154, 0.198]	[0.082, 0.088] [0.078, 0.093]	[0.039, 0.041] [0.035, 0.045]
Num observations	630217	630217	630217
Num clusters	566	566	566

Table 10: Treatment effects expressed as elasticities, after applying each shape constraint to the distribution of log hours rather than the distribution of hours. Rows “Linear density” and “Monotonic density” thus assume constant treatment effects in logs, as in the isoelastic model.

	$p=0$	p from non-changers	p from PTO
Buncher ATE as elasticity	[0.161, 0.188] [0.153, 0.198]	[0.082, 0.088] [0.077, 0.093]	[0.039, 0.041] [0.035, 0.046]
Average effect of FLSA on hours	[-1.466, -1.329] [-1.541, -1.260]	[-0.727, -0.629] [-0.769, -0.593]	[-0.347, -0.294] [-0.385, -0.262]
Avg. effect among directly affected	[-2.620, -2.375] [-2.743, -2.259]	[-1.453, -1.258] [-1.532, -1.189]	[-0.738, -0.624] [-0.814, -0.560]
Double-time, average effect on hours	[-2.604, -0.950] [-2.716, -0.904]	[-1.239, -0.492] [-1.293, -0.464]	[-0.580, -0.241] [-0.639, -0.215]

Table 11: Estimates of policy effects (replicating Table 4) ignoring the potential effects of changes to straight-time wages.

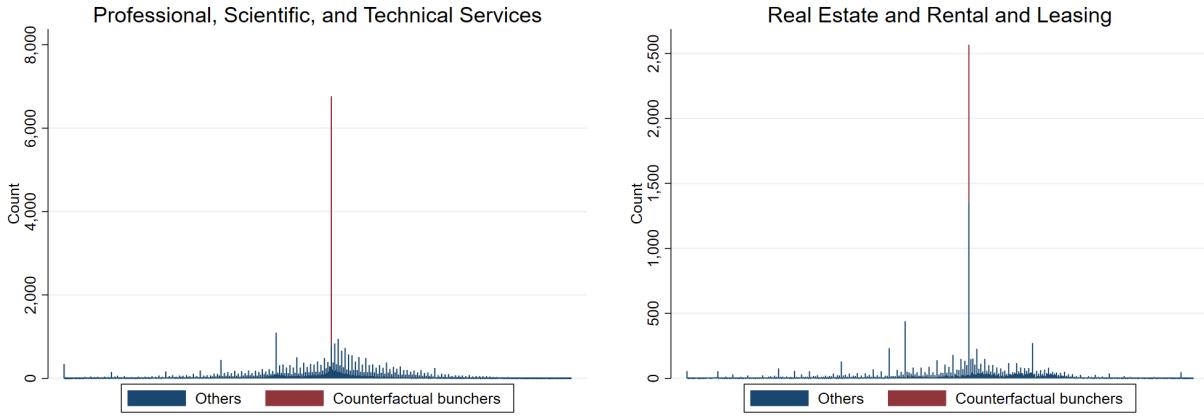


Figure 8: Hours distribution for an industry with a low treatment effect (left), and a high one (right). Both industries exhibit a comparable amount of raw bunching (14% and 19% respectively, see Table 5). In Professional, Scientific, and Technical Services, much more of the observable bunching is estimated to be counterfactual bunching, using the PTO-based method. Furthermore, the density of hours is higher just to the right of 40, meaning that the remaining bunching can be explained by a very small responsiveness of hours to the FLSA.

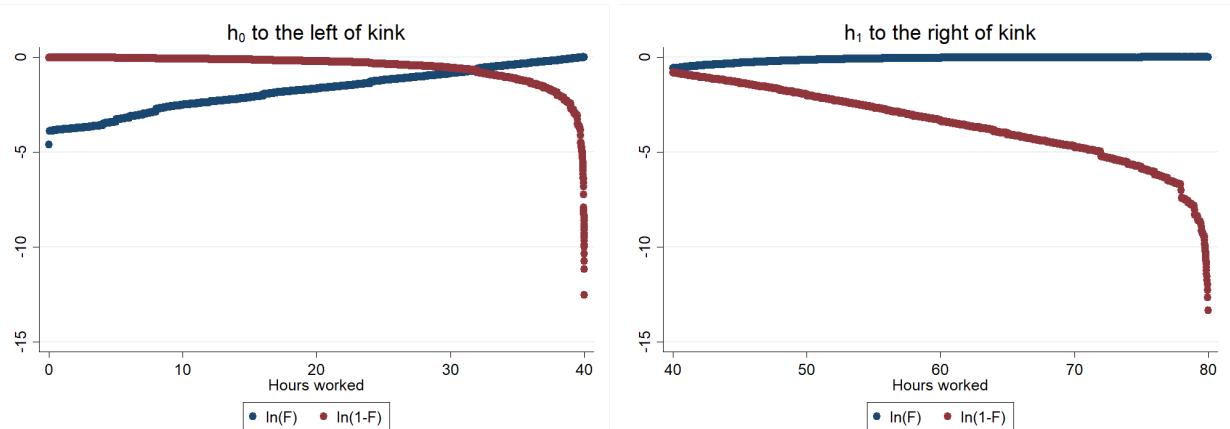


Figure 9: Validating the assumption of bi-log-concavity away from the kink. The left panel plots estimates of $\ln F_0(h)$ and $\ln(1 - F_0(h))$ for $h < 40$, based on the empirical CDF of observed hours worked. Similarly the right panel plots estimates of $\ln F_1(h)$ and $\ln(1 - F_1(h))$ for $h > k$, where I've conditioned the sample on $Y_i < 80$. Bi-log-concavity requires that the four functions plotted be concave globally.

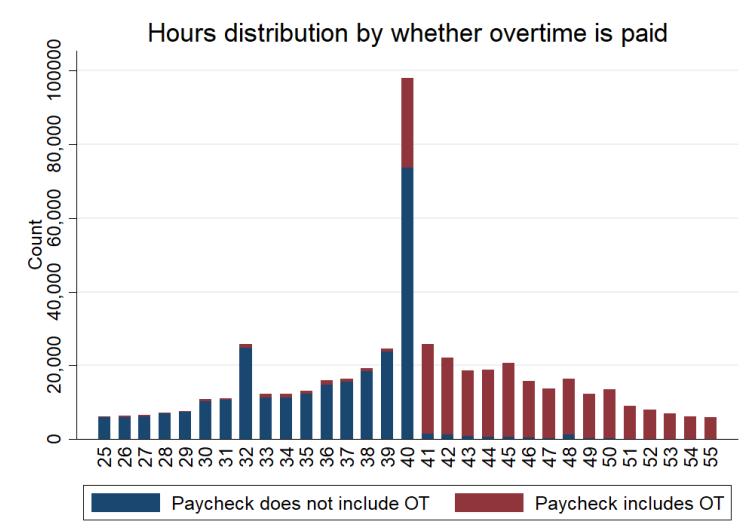


Figure 10: Histogram of hours worked pooling all paychecks in sample, with one hour bins. Blue mass in the stacks indicate that the paycheck included no overtime pay, while red indicates that the paycheck does include overtime pay.

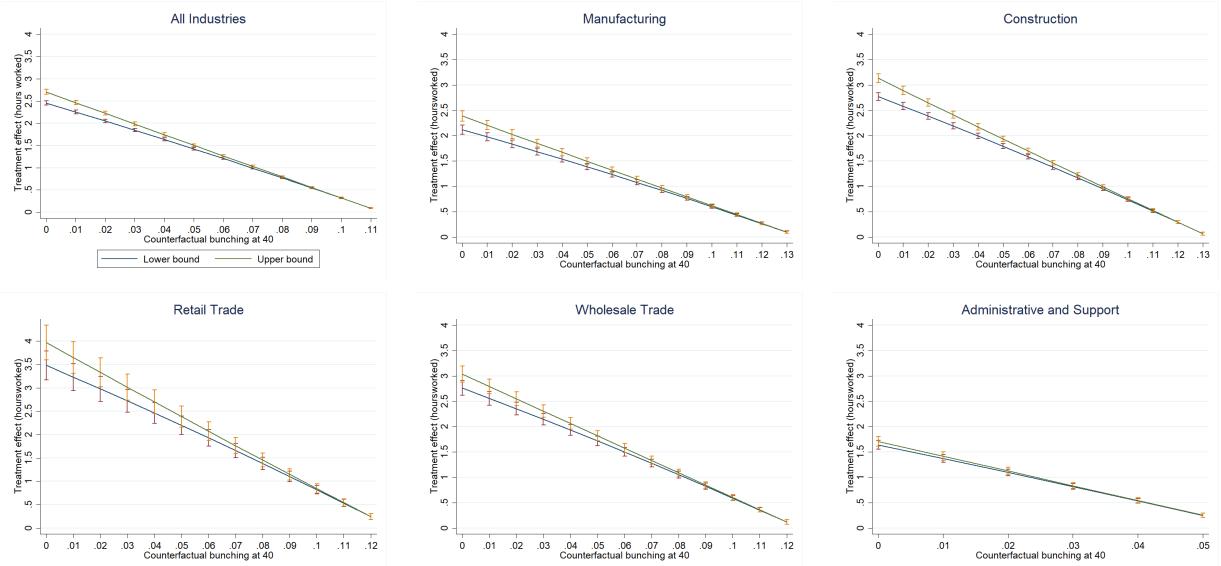


Figure 11: Estimates of the buncher ATE Δ_k^* as a function of p , pooled across industries and by each of the largest major industries.

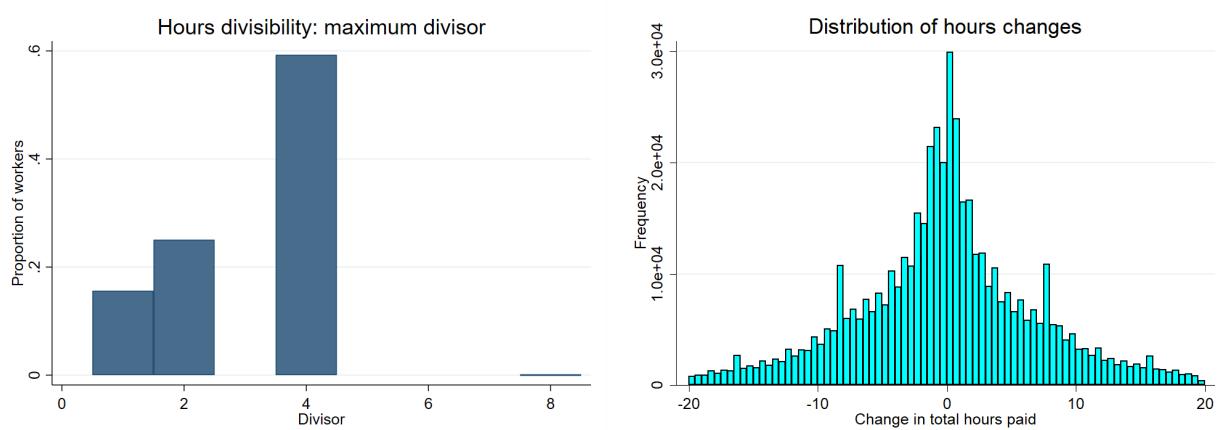


Figure 12: *Left:* distribution of the largest integer $m = 1 \dots 10$ that maximizes the proportion of worker i 's paychecks for which hours are divisible by m . This can be thought of as the granularity of hours reporting for worker i . *Right:* distribution of changes in total hours between subsequent pay periods (truncated at -20 and 20)

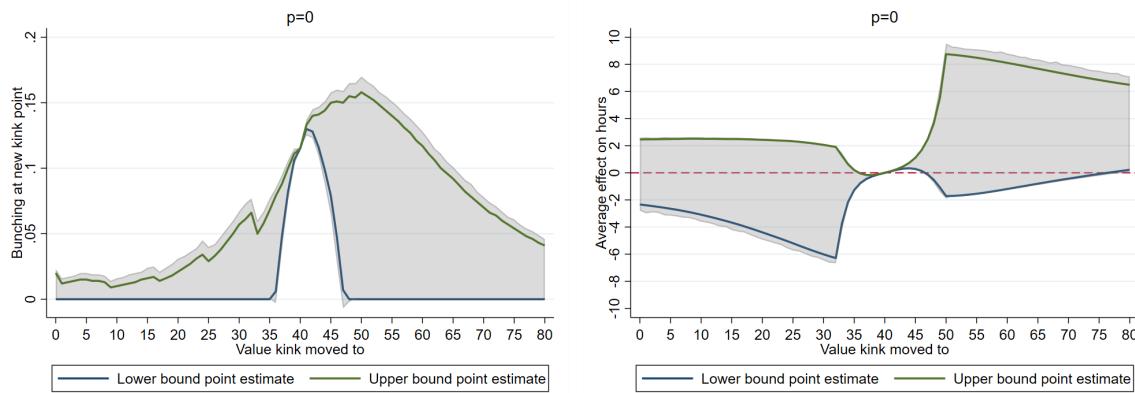
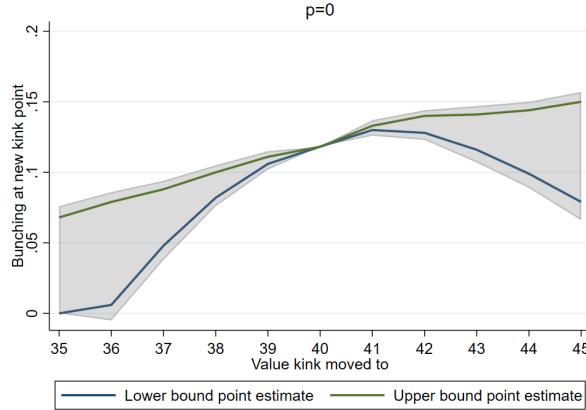
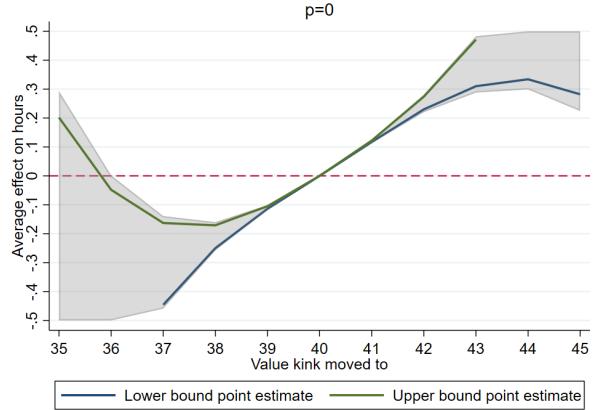


Figure 13: Estimates of the bunching (left panel) and average effect on hours (right panel) were k changed to any value from 0 to 80, assuming $p = 0$. Pointwise bootstrapped 95% confidence intervals, cluster bootstrapped by firm, are shaded gray. Bounds are not informative far from 40.

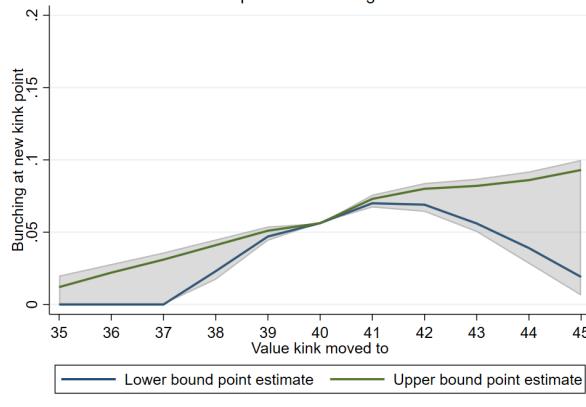
Bunching at new kink



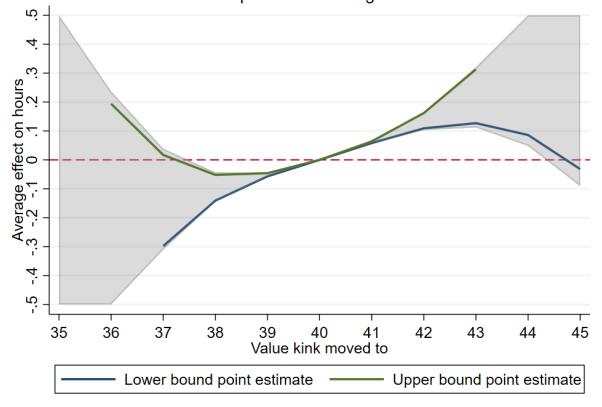
Average effect on hours



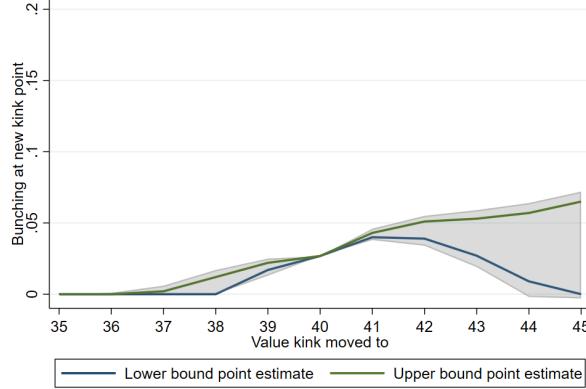
p from non-changers



p from non-changers



p from PTO



p from PTO

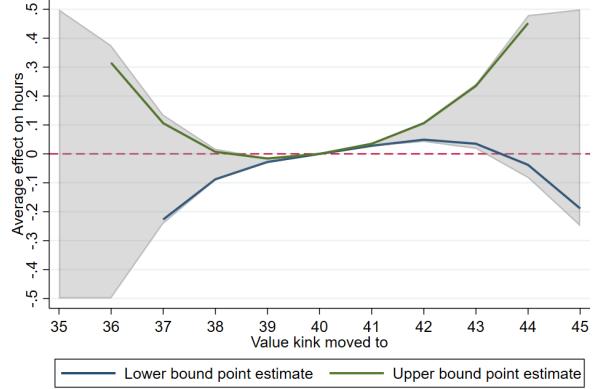


Figure 14: Bounds for the bunching that would exist at standard hours k if it were changed from 40 (left panel), as well as for the impact on average hours (right panel). Bounds of the effect on hours are clipped to the interval $[-0.5, 0.5]$ for visibility. Pointwise bootstrapped 95% confidence intervals, cluster bootstrapped by firm, are shaded gray.

C.5 Estimates from the iso-elastic model

This section estimates bounds on ϵ from the iso-elastic model described in Section 4.2, under the assumption that the distribution of $h_{0it} = \eta_{it}^{-\epsilon}$ is bi-log-concave (and linear as in Saez, 2010 as a special case). If h_{0it} is BLC, bounds on ϵ can be deduced from the fact that

$$F_0(40 \cdot 1.5^{-\epsilon}) = F_0(40) + \mathcal{B} = P(h_{it} \leq 40)$$

where $F_0(h) := P(h_{0it} \leq h)$ and the RHS of the above is observable in the data. $40 \cdot 1.5^{-\epsilon}$ is the location of this “marginal buncher” in the h_0 distribution. In particular,

$$\epsilon = -\ln(Q_0(F_0(40) + \mathcal{B})/40)/(\ln(1.5))$$

where $Q_0 := F_0^{-1}$ is guaranteed to exist by BLC (Dümbgen et al., 2017). In particular:

$$\epsilon \in \left[\frac{\ln \left(1 - \frac{1-F_0(40)}{40f(40)} \ln \left(1 - \frac{\mathcal{B}}{1-F_0(40)} \right) \right)}{-\ln(1.5)}, \frac{\ln \left(1 + \frac{F_0(40)}{40f(40)} \ln \left(1 + \frac{\mathcal{B}}{F_0(40)} \right) \right)}{-\ln(1.5)} \right]$$

where $F_0(k) = \lim_{h \uparrow 40} F(h)$ and $f_0(k) = \lim_{h \uparrow 40} f(h)$ are identified from the data. The bounds on ϵ estimated in this way are $\epsilon \in [-.210, -.167]$ in the full sample, with all bunching \mathcal{B} attributed to the kink ($p = 0$).

Since BLC is preserved when the random variable is multiplied by a scalar, BLC of h_{0it} implies BLC of $h_{1it} := \eta_{it}^{-\epsilon} \cdot 1.5^\epsilon$ as well. This implication can be checked in the data to the right of 40, since $\eta_{it}^{-\epsilon} \cdot 1.5^\epsilon$ is observed there. BLC of h_{1it} implies a second set of bounds on ϵ , because:

$$F_1(40 \cdot 1.5^\epsilon) = F_1(40) - \mathcal{B} = P(h_{it} < 40)$$

and the RHS is again observable in the data, where $F_1(h) := P(h_{1it} \leq h)$. Here $40 \cdot 1.5^\epsilon$ is the location of a second “marginal buncher” – for which $h_0 = 40$ – in the h_1 distribution. Now we have:

$$\epsilon \in \left[\frac{\ln \left(1 + \frac{F_1(40)}{40f_1(40)} \ln \left(1 - \frac{\mathcal{B}}{F_1(40)} \right) \right)}{\ln(1.5)}, \frac{\ln \left(1 - \frac{1-F_1(40)}{40f_1(40)} \ln \left(1 + \frac{\mathcal{B}}{1-F_1(40)} \right) \right)}{\ln(1.5)} \right]$$

where $F_1(k) = F(k)$ and $f_1(k) := \lim_{h \downarrow 40} f(h)$ are identified from the data. Empirically, these bounds are estimated as $\epsilon \in [-.179, -.141]$. Taking the intersection of these bounds with the range $\epsilon \in [-.210, -.168]$ estimated previously, we have that $\epsilon \in [-.179, -.168]$.¹⁴ The identified set is reduced from a length of .043 to .012, a factor of nearly 4. This underscores the importance of using the data from *both* sides of the kink for identification. Since a linear density satisfies BLC, the identification assumption of Saez, 2010, that the density of h_0 is

¹⁴Note that this interval differs slightly from the identified set of the buncher ATE as elasticity for $p = 0$ in Table 4. The latter quantity averages the effect in levels over bunchers and rescales: $\frac{1}{40 \ln(1.5)} \mathbb{E}[h_{0it}(1 - 1.5^\epsilon) | h_{it} = 40]$, but the two are approximately equal under $1.5^\epsilon \approx 1 + .5\epsilon$ and $\ln(1.5) \approx .5$.

linear, picks a point within the identified set under BLC. Table 9 verifies that this is born out in estimation (with results are expressed there as level effects rather than an elasticity).

As discussed in Section 4.2, a value of $\epsilon \approx -.175$ is difficult to reconcile with a realistic view of revenue production with respect to hours. Note that if instead of the isoelastic model, production were instead described by a more general separable and homogeneous production function like

$$\pi_{it}(z, h) = a_{it} \cdot f(h) - z$$

then treatment effects are $\Delta_{it} = g(1/\eta_{it}) - g(1.5/\eta_{it})$, where $g(m) := (f')^{-1}(m)$ yields the hours h at which $f'(h) = m$. We can then use the fundamental theorem of calculus to express this as $(h_{1it} - h_{0it})/h_{0it} = 1.5\bar{\epsilon}_{it} - 1$ where $\bar{\epsilon}_{it}$ is a unit-specific weighted average of the inverse elasticity of production between $1.5\eta_{it}$ and η_{it} : $\bar{\epsilon}_{it} := \int_{\eta_{it}^{-1}}^{1.5\eta_{it}^{-1}} \lambda(m) \cdot \epsilon(g(m)) \cdot dm$, and $\lambda(m) = \frac{1/m}{\ln 1.5}$ is a positive function integrating to one. Here $\bar{\epsilon}_{it}$ plays the role of an “effective” elasticity parameter that determines the size of treatment effects when the production function is $f(h)$. This suggests that simply generalizing the functional form $f(h)$ is not sufficient to reconcile a realistic picture of production with the data, since the observed bunching still maps to a local average elasticity of $f(h)$. However, the general choice model that allows multiple margins of choice \mathbf{x} can.

C.6 Results of the employment effect calculation

Taking my preferred estimate that hourly workers work approximately 1/3 of an hour less per week on average because of the rule, hours per worker are reduced by just under 1%. If we assume the same sized effect occurs for covered salaried workers, and ignore scale effects of the overtime rule on the total number of labor hours in FLSA-eligible jobs, this suggests employment among such jobs is 1% higher than it would be without the overtime premium. This serves as an upper bound, since overall total hours worked may decrease due to overtime regulation.

Following Hamermesh (1993), assume that the percentage change in employment decomposes as:

$$\Delta \ln E|_{EH} - \eta \cdot \Delta \ln LC \cdot \frac{\eta}{\alpha - \eta} \quad (6)$$

where η is constant-output demand elasticity for labor, α is a labor supply elasticity. Following Hamermesh (1993) I use $\Delta \ln LC = 0.7\%$ based on Ehrenberg and Schumann (1982), calibrated assuming that 80% of labor costs come from wages with overtime representing 2% of total hours. $\Delta \ln E|_{EH}$ is the quantity implied by my estimates: the percentage change in employment that would occur were the total number of worker-hours EH unchanged. Taking a preferred estimate of the average effect of the FLSA as reported in Table 4 to be about 1/3 of an hour, I use a value of $\Delta \ln E|_{EH} = \frac{1/3}{40} \approx 0.9\%$.

		η			
		-0.15	-0.3	-0.5	
		0	0.76	0.64	0.50
α		0.1	0.80	0.70	0.56
		0.5	0.85	0.79	0.68

Table 12: Back-of-the-envelope employment effects based on the average reduction in hours estimated via the bunching design and Equation (6), as a function of the demand elasticity for labor (rather than capital) η , and labor supply elasticity α . The bold entry reflects “best-guess” values of η and α .

“Best-guess” values for the other parameters used by Hamermesh, 1993 are $\eta = -0.3$ and $\alpha = 0.1$, based on a review of empirical estimates. This yields 0.17 percentage points for the substitution term $\eta \cdot \Delta \ln LC \cdot \frac{\eta}{\alpha - \eta}$, suggesting that the effect of the FLSA is attenuated from roughly 0.87 percentage points to about a 0.70 percentage point net increase in employment. I assume that the FLSA overtime rule applies to a total of 100 million workers, based on 80 million hourly workers combined with an estimated 20 million covered salary workers Kimball and Mishel (2015). Assuming the same percentage increase in employment applies to hourly workers and covered salary workers, the above estimate corresponds to 700,000 jobs created. Generating a negative overall employment response by assuming higher substitution to capital requires $\eta = -1.25$, well outside of empirical estimates.

D Incorporating workers that set their own hours

This section considers the robustness of the empirical strategy from Section 4 to a case where some workers are able to choose their own hours. In this case, a simple extension of the model leads to the bounds on the buncher ATE remaining valid, but it is only directly informative about the effects of the FLSA among workers who have their hours chosen by the firm. In this section I follow the notation from the main text where h_{it} indicate the hours of worker i in week t .

Suppose that some workers are able to choose their hours each week without restriction (“worker-choosers”), and that for the remaining workers (“firm-choosers”) their employers set their hours. In general we can allow who chooses hours for a given worker to depend on the period, so let $W_{it} = 1$ indicate that i is a worker-chooser in period t . Additionally, we continue to allow counterfactual bunchers for whom counterfactual hours satisfy $h_{0it} = h_{1it} = 40$, regardless of who chooses them. I replace Assumption CONVEX from Section 4 to allow agents to *either* dislike pay (firm-choosers), or like pay (worker-choosers):

Assumption CONVEX* (convex preferences, monotonic in either direction). *For each i, t and function $B(\mathbf{x})$, choice is $(z_{Bi}, \mathbf{x}_{Bi}) = \text{argmax}_{z, \mathbf{x}} \{u_i(z, \mathbf{x}) : z \geq B(\mathbf{x})\}$ where*

$u_i(z, \mathbf{x})$ is:

- strictly increasing in z , if $W_{it} = 1$
- strictly decreasing in z , if $W_{it} = 0$

and satisfies $u_i(\theta z + (1 - \theta)z^*, \theta \mathbf{x} + (1 - \theta)\mathbf{x}^*) > \min\{u_i(z, \mathbf{x}), u_i(z^*, \mathbf{x}^*)\}$ for any $\theta \in (0, 1)$ and points $(z, \mathbf{x}), (z^*, \mathbf{x}^*)$ such that $y_i(\mathbf{x}) \neq k$ and $y_i(\mathbf{x}^*) \neq k$.

For generality, I here use weaker notion of convexity of preferences from Assumption CONVEX in Appendix A. It is implied by strict quasi-concavity of $u_i(z, \mathbf{x})$.

Note: This setup is general enough to also allow a stylized bargaining-inspired model in which choices maximize a weighted sum of quasilinear worker and firm utilities. For example, suppose that for any pay schedule $B(h)$:

$$h = \underset{h}{\operatorname{argmax}} \beta(f(h) - z) + (1 - \beta)(z - \nu(h)) \quad \text{with } z = B(h) \quad (7)$$

where $f(h) - z$ is firm profits with concave production f , $z - \nu(h)$ is worker utility with a convex disutility of labor $\nu(h)$, and $\beta \in [0, 1]$ governs the weight of each party in the negotiation (this corresponds to Nash bargaining in which outside options are strictly inferior to all h for both parties, and utility is log-linear in z). Rearranging the maximand of Equation (7) as $(1 - 2\beta)z + \{\beta f(h) - (1 - \beta)\nu(h)\}$, we can observe that this setting delivers outcomes as-if chosen by a single agent with quasi-concave preferences, as $\beta f(h) - (1 - \beta)\nu(h)$ is concave. For Assumption CONVEX from Section 4 to hold with the assumed direction of monotonicity in pay z , we would require that $\beta > 1/2$ for all worker-firm pairs: informally, that firms have more say than workers do in determining hours. However the more general CONVEX* holds regardless of the distribution of β over worker-firm pairs. If $\beta_{it} < 1/2$, paycheck it will look like a worker-chooser, and if $\beta_{it} > 1/2$ paycheck it will look like a firm-chooser.

In the generalized model of CONVEX*, bunching is prima-facie evidence that firm-choosers exist, because there is no prediction of bunching among worker-choosers provided that potential outcomes are continuously distributed. By contrast, k is a “hole” in the worker-chooser hours distribution: intuitively, if a worker is willing to work 40 hours then they will also find it worthwhile to work more, given the sudden increase in their wage. Indeed under regularity conditions all of the data local to 40 are from firm-choosers (or counterfactual bunchers). To make this claim precise, assume that for worker-choosers, hours are the only margin of response (i.e. their utility depends on \mathbf{x} only thought $y(\mathbf{x})$), and let $IC_{0it}(y)$ and $IC_{1it}(y)$ be the worker’s indifference curves passing through h_{0it} and h_{1it} , respectively. I assume these indifference curves are twice Lipschitz differentiable, and let $M_{it} := \sup_y \max\{|IC''_{0it}(y)|, |IC''_{1it}(y)|\}$, where IC'' indicates second derivatives.

Proposition 3. Suppose that the joint distribution of h_{0it} and h_{1it} admits a continuous density conditional on $K_{it}^* = 0$, and that for any worker-chooser IC_{0it} and IC_{1it} are differentiable with M_{it}/w_{it} having bounded support. Then, under CHOICE and CONVEX*:

- $P(h_{it} = k \text{ and } K_{it}^* = 0) = P(h_{1it} \leq k \leq h_{0it} \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 0)$
- $\lim_{h \uparrow k} f(h) = P(W_{it} = 0) \lim_{h \uparrow k} f_{0|W=0}(h)$
- $\lim_{h \downarrow k} f(h) = P(W_{it} = 0) \lim_{h \downarrow k} f_{1|W=0}(h)$

Proof. See Appendix H. □

The first bullet of Proposition 3 says that all active bunchers are also firm-choosers, and have potential outcomes that straddle the kink. The second and third bullets state that the density of the data as hours approach 40 from either direction is composed only of worker-choosers. This result on density limits requires the stated regularity condition on M_{it}/w_{it} , which prevents worker indifference curves from becoming too close to themselves featuring a kink (plus a requirement that straight-time wages w_{it} be bounded away from zero).

Given the first item in Proposition 3, the buncher ATE introduced in Section 4 only includes firm-choosers:

$$\mathbb{E}[h_{0it} - h_{1it}|h_{it} = 40, K_{it}^* = 0] = \mathbb{E}[h_{0it} - h_{1it}|h_{it} = 40, K_{it}^* = 0, W_{it} = 0]$$

Accordingly, I assume rank invariance among the firm-chooser population only:

Assumption RANK* (near rank invariance and counterfactual bunchers). *The following are true:*

1. $P(h_{0it} = k) = P(h_{1it} = k) = p$
2. $Y_{it} = k \text{ iff } (h_{0it} \in [k, k + \Delta_0^*] \text{ and } W_{it} = 0) \text{ iff } (h_{1it} \in [k - \Delta_1^*, k] \text{ and } W_{it} = 0), \text{ for some } \Delta_0^*, \Delta_1^*$

where p continues to denote $P(K_{it}^* = 1)$.

We may now state a version of Theorem 2 that conditions all quantities on $W = 0$, provided that we assume bi-log concavity of h_0 and h_1 conditional on $W = 0$ and $K = 0$.

Theorem 1* (bi-log-concavity bounds on the buncher ATE, with worker-choosers). *Assume CHOICE, CONVEX* and RANK* hold. If both h_{0it} and h_{1it} are bi-log concave conditional on the event ($W_{it} = 0$ and $K_{it}^* = 0$), then:*

$$\mathbb{E}[h_{0it} - h_{1it}|h_{it} = k, K_{it}^* = 0] \in [\Delta_k^L, \Delta_k^U]$$

where

$$\Delta_k^L = g(F_{0|W=0, K^*=0}(k), f_{0|W=0, K^*=0}(k), \mathcal{B}^*) + g(1 - F_{1|W=0, K^*=0}(k), f_{1|W=0, K^*=0}(k), \mathcal{B}^*)$$

and

$$\Delta_k^U = -g(1 - F_{0|W=0, K^*=0}(k), f_{0|W=0, K^*=0}(k), -\mathcal{B}^*) - g(F_{1|W=0, K^*=0}(k), f_{1|W=0, K^*=0}(k), -\mathcal{B}^*)$$

where $\mathcal{B}^* = P(h_{it} = k | W_{it} = 0, K_{it}^* = 0)$ and

$$g(a, b, x) = \frac{a}{bx} (a + x) \ln \left(1 + \frac{x}{a} \right) - \frac{a}{b}$$

The bounds are sharp.

Proof. See Appendix H. \square

Identification with worker-choosers

Theorem 1* does not immediately yield identification of the buncher-ATE bounds Δ_k^L and Δ_k^U , as we need to estimate each of the arguments to the function g . As shown in the proof of Theorem 1*, the bounds can be rewritten in terms of p , the identified quantities \mathcal{B} , $P(W_{it} = 0) \lim_{y \uparrow k} f_{0|W=0}(y)$ and $P(W_{it} = 0) \lim_{y \uparrow k} f_{1|W=0}(y)$, and two unidentified probabilities: $P(h_{0it} < k \text{ and } h_{it} = h_{0it} \text{ and } W_{it} = 1)$ and $P(Y_{1it} > k \text{ and } h_{it} = h_{1it} \text{ and } W_{it} = 1)$.

To illustrate the unidentified quantities, Figure 15 depicts an example of a joint distribution of (h_0, h_1) that includes worker-choosers and satisfies Assumption RANK*. The x-axis is h_0 , and the y-axis is h_1 , with the solid lines indicating 40 hours and the dotted diagonal line depicting $h_1 = h_0$. The dots show a hypothetical joint-distribution of the potential outcomes, with the (red) dots south of the 45-degree line representing firm-choosers, and the (blue and orange) points above representing worker-choosers. Blue x's indicate worker-choosers who choose their value of h_0 , while orange circles indicate worker-choosers who choose their value of h_1 . The red dot at $(40, 40)$ represents a mass of counterfactual bunchers.

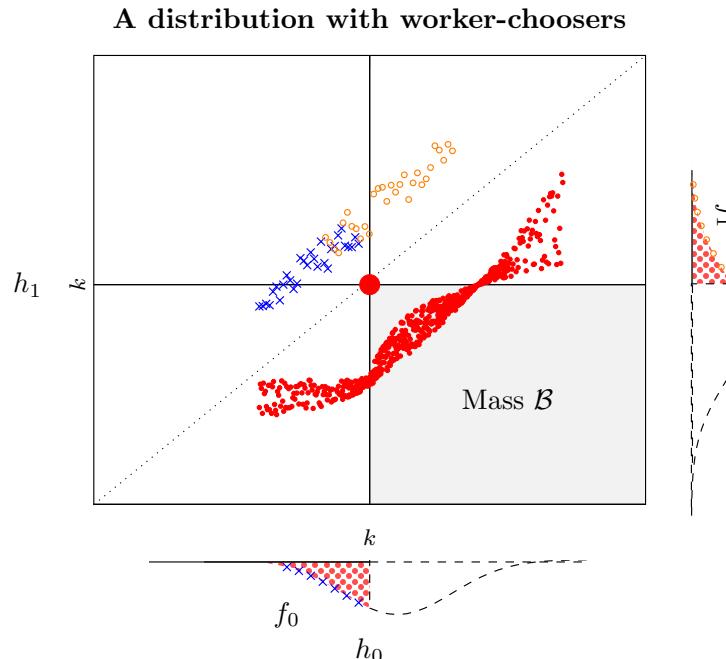


Figure 15: The joint distribution of (h_{0it}, h_{1it}) , for a distribution including worker-choosers and satisfying assumption RANK*, cf. Figure 5. See text for description.

Observed to the the econometrician is the point mass at 40 as well as the truncated marginal distributions depicted at the bottom and the right of the figure, respectively. The observable $P(h_{it} \leq h)$ for $h < 40$ doesn't exactly identify $P(h_{0it} \leq h)$ because some blue x's are missing: these are worker-choosers for whom $h_1 > 40 > h_0$ and choose to work overtime at their h_1 value. Thus they show up in the data at $h > 40$ even though they have $h_0 < 40$. Similarly, some orange circles do not appear in the observations above 40: these are worker-choosers for whom $h_1 > 40 > h_0$ and choose to work their h_0 value, not working overtime. The probabilities $P(h_{it} < 40 \text{ and } W_{it} = 0)$ and $P(h_{it} > 40 \text{ and } W_{it} = 0)$ can thus only be estimated with some error, with the size of the error depending on the mass of worker-choosers in the northwest quadrant of Figure 15. However, in practice this has little impact on the results, as the bounds Δ_k^L and Δ_k^U are not very sensitive to the values of the CDF inputs $F_{0|W=0,K^*=0}(k)$ and $F_{1|W=0,K^*=0}(k)$. The bounds mostly depend on the density estimates and the size of the bunching mass, given their empirical values. Thus Theorem 1* still partially identifies the buncher LATE among firm-choosers, to a good approximation.

However, a further caveat of Theorem 1* is worth mentioning. An evaluation of the FLSA would ideally account for worker-choosers (who are working longer hours as a result of the policy) when averaging treatment effects. However, the proportion of worker-choosers and the size of their hours increases are not identified. Using the buncher ATE to estimate the overall ex-post effect of the FLSA—as described in Section 4.4—may overstate its overall average net hours reduction. However, the survey evidence mentioned in Section 2 suggests that the set of worker-choosers is relatively small, mitigating this concern.

E Interdependencies among hours within the firm

In this section I consider the impact that interdependencies between the hours of different units may have on the estimates, reflected in the third term of Equation (8) from Section 4.4. First, I develop some structure to guide our intuition for this term, and then present some empirical evidence that it is likely to be small (recall that it is taken to be zero in the final results assessing the FLSA).

The basic issue is as follows: when a single firm chooses hours jointly among multiple units—either across different workers or across multiple weeks, or both—this term may be nonzero and contribute to the overall effect of the FLSA. In my potential outcomes donation, this represents a violation of the non-interference condition of the stable unit treatment value assumption (SUTVA), when using the treatment effects Δ_{it} to assess the average impact of the FLSA on hours. If firms maximize profits given a production function that is not linearly separable across workers or across weeks, such violations may occur.

To simplify the notation, suppose that SUTVA violations may occur across workers within

a firm in a single week, suppressing the time index t and focusing on a single firm. As in Section 4.4 let \mathbf{h}_{-i} denote the vector of actual (observed) hours for all workers aside from i within i 's firm. These hours are chosen according to the kinked cost schedule introduced by the FLSA. Let $\mathbf{h}_{0i}(\cdot)$ denote the hours that the firm would choose for worker i if they had to pay i' straight-wage w_i for all of i 's hours, as a function of the hours profile of the other workers in the firm (suppressing dependence on straight-wages in this section). Define $\mathbf{h}_{1i}(\cdot)$ analogously with $1.5w_i$. In this notation, the potential outcomes defined in Section 4 are $h_{0i} = \mathbf{h}_{0i}(\mathbf{h}_{-i})$ and $h_{1i} = \mathbf{h}_{1i}(\mathbf{h}_{-i})$. As in Section 4.4 let $(h_i^*, \mathbf{h}_{-i}^*)$ denote the hours profile that would occur absent the FLSA, so that the average ex-post effect of the FLSA is $\mathbb{E}[h_i - h_i^*]$.

Even if there are SUTVA violations, treatment effects $\Delta_i = \mathbf{h}_{0i}(\mathbf{h}_{-i}) - \mathbf{h}_{1i}(\mathbf{h}_{-i})$ remain meaningful as a reduced-form average labor demand elasticity, in which the wage of worker i is changed but with \mathbf{h}_{-i} held fixed. Furthermore, bunching is still informative about identify the buncher ATE: bunching will not occur unless $\Delta_i > 0$ from some units near the kink such that $h_{0i} \in [k, k + \Delta_i]$. The question is whether the treatment effects Δ remain a good guide to the overall effect of the FLSA, given that it may also change \mathbf{h}_{-i} for a given worker i .

For concreteness, let us now suppose that hours are chosen to maximize profits with a joint-production function $F(\mathbf{h})$, where \mathbf{h} is a vector of the hours this week across all workers in the firm. We then have that $(h_i, \mathbf{h}_{-i}) = \text{argmax} \left\{ F(\mathbf{h}) - \sum_j B_{kj}(h_j) \right\}$, where the sum is across workers j and $B_{kj}(h) := w_j h + .5w_j \mathbb{1}(h > 40)(h - 40)$. Similarly $(h_i^*, \mathbf{h}_{-i}^*) = \text{argmax} \left\{ F(\mathbf{h}) - \sum_j w_j h_j \right\}$. Whether $\mathbf{h}_{0i}(\mathbf{h}_{-i})$ is smaller or larger than h_i^* (with a fixed set of employees) will depend upon whether i 's hours are complements or substitutes in production with those of each of their colleagues, and with what strength. It is natural to expect that either case might occur. Consider for example a production function in which workers are divided into groups $\theta_1 \dots \theta_M$ corresponding to different occupations, and:

$$F(\mathbf{h}) = \prod_{m=1}^M \left(\left(\sum_{i \in \theta_m} a_i \cdot h_i^{\rho_m} \right)^{1/\rho_m} \right)^{\alpha_m} \quad (8)$$

where a_i is an individual productivity parameter for worker i . The hours of workers within an occupation enter as a CES aggregate with substitution parameter ρ_m , which then combine in a Cobb-Douglas form across occupations with exponents α_m . For this production function, the hours of two workers i and j belonging to different occupations are always complements in production: i.e. $\partial_{h_i} F(\mathbf{h})$ is increasing in h_j . When i and j belong to the same occupation θ_m , it can be shown that worker i and j 's hours are substitutes—i.e. $\partial_{h_i} F(\mathbf{h})$ is *decreasing* in h_j —when $\alpha_m \leq \rho_m$.

Thus both substitution and complementarity in hours can plausibly coexist within a firm, and it is difficult to sign theoretically the overall contribution of interdependencies on our parameter of interest θ (c.f. Eq. (8)). Given that neither occupations nor tasks are observed

in the data, it is also difficult to obtain direct evidence even with the aid of functional-form assumptions like Eq. (8). I therefore turn to an indirect empirical test of whether these effects are likely to play a significant role in θ .

An ideal test of interdependencies between hours within a firm would leverage random individual-level shocks to a worker’s hours, and look for a response in the hours of their colleagues. A worker taking sick-pay—thus reducing their hours of work—represents a compelling candidate as its timing may be uncorrelated with that of firm-level shocks (after controlling for seasonality). Figure 16 uses an event study design to show that in weeks when a worker receives a positive number of sick-pay hours, their individual hours worked for that week decline by about 8 hours on average. Yet I fail to find evidence of a corresponding change in the hours of others in the same firm. This suggests that short term variation in the hours of a worker’s colleagues does not tend to translate into contemporaneous changes in their own (for example, if the firm were dividing a fixed number of hours across workers). Figure 17 produces similar results when replacing the two-wage-fixed specification of Figure 16 with an “imputation”-based approach similar to Borusyak et al. (2021) and Gardner (2021).

Table 13 shows another piece of evidence: that my overall effect estimates are similar between small, medium, and large firms. If firms were to compensate for overtime hours reductions by “giving” some hours to similar workers who would otherwise be working less than 40, for instance, then we would expect this to play a larger role in firms where there are a large number of substitutable workers—causing a bias that increases with firm size. However, in Table 13 below, the confidence intervals for all three firm size categories overlap, in my preferred specification of estimating p using variation in PTO.

	$p=0$		p from PTO	
	Bunching	Effect of the kink	Net Bunching	Effect of the kink
Small firms	0.198 [0.189, 0.208]	[-1.525, -1.455] [-1.676, -1.299]	0.027 [0.023, 0.031]	[-0.231, -0.171] [-0.274, -0.139]
	0.103 [0.095, 0.110]	[-1.123, -0.786] [-1.237, -0.710]	0.030 [0.025, 0.035]	[-0.337, -0.224] [-0.407, -0.178]
Medium firms	0.050 [0.047, 0.054]	[-0.768, -0.468] [-0.861, -0.414]	0.024 [0.021, 0.028]	[-0.371, -0.224] [-0.444, -0.180]

Table 13: Estimates of the ex-post effect of the kink by firm size. “Small” firms have between 1 and 25 workers in my estimation sample, “Medium” have 26 to 50, and “Large” have more than 50. Note that the estimated net bunching caused by the FLSA is similar across firm sizes (right), despite the raw bunching observed in the data differing by firm size category.

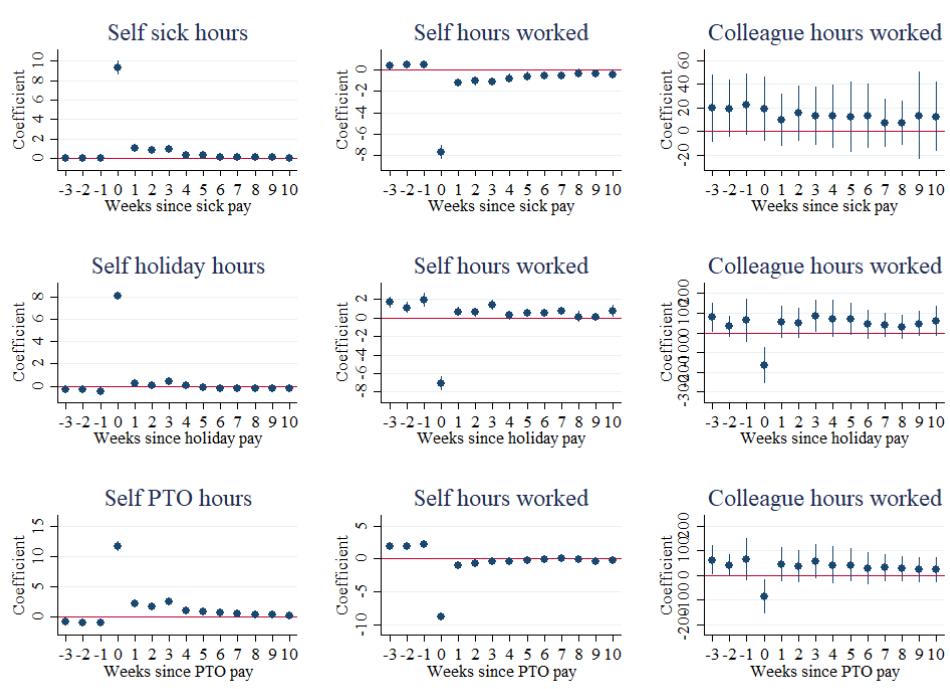


Figure 16: Event study coefficients β_j and 95% confidence intervals across an instance of a worker receiving pay for non-work hours (either sick pay, holiday pay, or paid time off ‘PTO’). Confidence intervals are constructed by non-parametric bootstrap clustered by firm. Estimating equation is $y_{it} = \mu_t + \lambda_i + \sum_{j=-3}^{10} \beta_j D_{it,j} + u_{it}$, where $D_{it,j} = 1$ if worker i in week t has a positive number of a given type of non-work hours j weeks ago (after a period of at least three weeks in which they did not), λ_i are worker fixed effects, and μ_t are calendar week effects. Rows correspond to choices of the non-work pay type: either sick, holiday, PTO. Columns indicate choices of the outcome y_{it} . “Colleague hours worked” sums the hours of work in t across all workers other than i in i ’s firm. The timing of both holiday and PTO hours appears to be correlated across workers, leading to a decrease in the working hours of i ’s colleagues in weeks in which i takes either holiday or PTO pay (center-right and bottom-right graphs). However I cannot reject that colleague work hours are unrelated to an instance of sick pay: before, during and after it occurs (top-right). Meanwhile i ’s hours of work reduce by about 8 hours on average during an instance of sick pay (top-center). This suggests that there is no contemporaneous reallocation of i ’s forgone work hours to their colleagues.

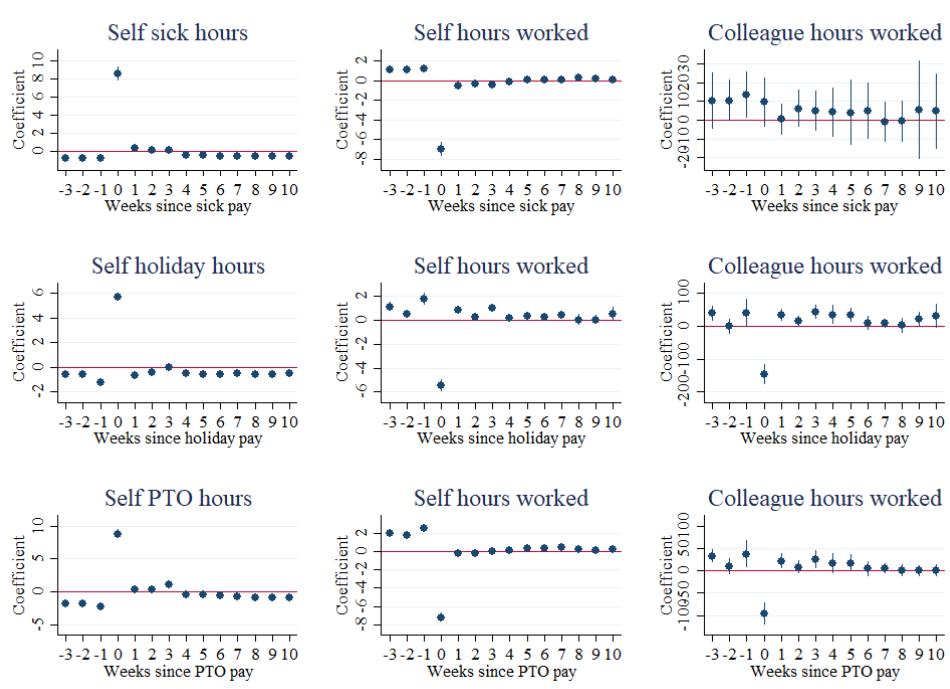


Figure 17: This figure replaces the two-way-fixed-effects estimator used in Figure 16 with an “imputation” approach similar to Borusyak et al. (2021) and Gardner (2021). Results are very similar to those in Figure 16. Specifically, I call all observations that are not between 3 weeks before and 10 weeks after a spell of non-work hours “clean controls”, and estimate a first regression $y_{it} = \mu_t + \lambda_i + \epsilon_{it}$ using these observations only. This regression includes all paychecks for workers that never have the corresponding type of non-work hour (sick pay, holiday pay, or PTO), but also a subset of paychecks for nearly all workers who do have a spell of non-work hours at some point (allowing me to estimate their fixed effect λ_i). Given the $\hat{\mu}_t$ and $\hat{\lambda}_i$, I compute $\tilde{y}_{it} = y_{it} - \hat{\mu}_t - \hat{\lambda}_i$ among units that are not clean controls (i.e. those between -3 and 10 weeks after the start of a spell), and estimate a second regression $\tilde{y}_{it} = \sum_{j=-3}^{10} \beta_j D_{it,j} + e_{it}$ on these units only (dropping a small number of workers i for whom there were no clean-control observations). 95% confidence intervals are constructed by non-parametric bootstrap clustered by firm.

F Modeling the determination of wages and “typical” hours

F.1 A simple model with exogenous labor supply

Each firm faces a labor supply curve $N(z, h)$, indicating the labor force N it can maintain if it offers total compensation z to each of its workers, when they are each expected to work h hours per week. The firm chooses a pair (z^*, h^*) based on the cost-minimization problem:

$$\min_{z, h, K, N} N \cdot (z + \psi) + rK \text{ s.t. } F(Ne(h), K) \geq Q \text{ and } N \leq N(z, h) \quad (9)$$

where the labor supply function is increasing in z while decreasing in h , $e(h)$ represents the “effective labor” from a single worker working h hours, and ψ represents non-wage costs per worker. The quantity ψ can include for example recruitment effort and training costs, administrative overhead and benefits that do not depend on h . Concavity of $e(h)$ captures

declining productivity at longer hours, for example from fatigue or morale effects. The function F maps total effective labor $Ne(h)$ and capital into level of output or revenue that is required to meet a target Q , and r is the cost of capital K . For simplicity, workers within a firm are here identical and all covered by the FLSA.

To understand the properties of the solution to Equation (9), let us examine two illustrative special cases.

Special case 1: an exogenous competitive straight-time wage (the “fixed-wage model”)

Much of the literature on hours determination has taken the hourly wage as a fixed input to the choice of hours, and assumed that at that wage the firm can hire any number of workers, regardless of hours. This can be motivated as a special case of Equation (9) in which there is perfect competition on the straight-time wage, i.e. $N(z, h) = \bar{N}\mathbb{1}(w_s(z, h) \geq w)$ for some large number \bar{N} and wage w exogenous to the firm, where the function $w_s(\cdot)$ is defined in Equation (1). Then Equation (9) reduces to:

$$\min_{N, h, K} N \cdot (hw + \mathbb{1}(h > 40)(w/2)(h - 40) + \psi) + rK \text{ s.t. } F(Ne(h), K) \geq Q \quad (10)$$

By limiting the scope of labor supply effects in the firm’s decision, Equation (10) is well-suited to illustrating the competing forces that shape hours choice on the production side: namely the fixed costs ψ on the one hand and the concavity of $e(h)$ on the other. Were ψ equal to zero with $e(h)$ strictly concave globally, a firm solving Equation (10) would always find it cheaper to produce a given level of output with more workers working less hours each. On the other hand, were ψ positive and e weakly convex, it would always be cheapest to hire a single worker to work all of the firm’s hours. In general, fixed costs and declining hours productivity introduce a tradeoff that leads to an interior solution for hours.¹⁵

Equation (10) introduces a kink into the firm’s costs as a function of hours, much as short-run wage rigidity does in my dynamic analysis. However, the assumption that the firm can demand any number of hours at a set straight-time wage rate is harder to defend when thinking about firms long-run expectations, a point emphasized by Lewis (1969). Equilibrium considerations will also tend to run against the independence of hourly wages and hours - a mechanism explored in Appendix F.2.

Special case 2: iso-elastic functional forms (the “fixed-job model”)

By placing some functional form restrictions on Equation (9), we can obtain a closed-form expression for (z^*, h^*) . In particular, when both labor supply and $e(h)$ are iso-elastic, production is separable between capital and labor and linear in the latter, and firms set the

¹⁵In the fixed-wage special case, these two forces along with the wage are in fact sufficient to pin down hours, which do not depend on the production function F or the chosen output level Q . See e.g. Cahuc and Zylberberg (2014) for the case in which $e(h)$ is iso-elastic.

output target Q to maximize profits, Proposition 4 characterizes the firm's choice of earnings and hours:

Proposition 4. *When i) $e(h) = e_0 h^\eta$ and $N(z, h) = N_0 z^{\beta_z} h^{\beta_h}$; ii) $F(L, K) = L + \phi(K)$ for some function ϕ ; and iii) Q is chosen to maximize profits, the (z^*, h^*) that solve Equation (9) are:*

$$h^* = \left[\frac{\psi}{e_0} \cdot \frac{\beta}{\beta - \eta} \right]^{1/\eta} \quad \text{and} \quad z^* = \psi \cdot \frac{\beta_z}{\beta_z + 1} \frac{\eta}{\beta - \eta}$$

where $\beta := \frac{|\beta_h|}{\beta_z + 1}$, provided that $\psi > 0$, $\eta \in (0, \beta)$, $\beta_h < 0$ and $\beta_z > 0$. Hours and compensation are both decreasing in $|\beta_h|$ and increasing in β_z .

Proof. See Appendix H. □

The proposition shows that the hours chosen depend on labor supply via $\beta = \frac{|\beta_h|}{1+\beta_z}$, which gauges how elastic labor supply is with respect to hours relative to earnings. The more sensitive labor supply is to a marginal increase in hours as compared with compensation, the higher β will be and lower the optimal number of hours. The proof of Proposition 4 also shows that the general model with $N(z, h)$ differentiable (unlike in Special Case 1) can support an interior solution for hours even without fixed costs $\psi = 0$. Proposition 4 provides an example of the *fixed-job* model: in the absence of perfect competition on the straight-wage, anticipated hours h^* , total pay z^* , and employment $N^* := N_0 \cdot (z^*)^{\beta_z} (h^*)^{\beta_h}$ are unaffected by the FLSA overtime rule, in this simple static model.

F.2 Endogenizing labor supply in an equilibrium search model

The last section treated the labor supply function $N(z, h)$ as exogenous, but in general it might be viewed as an equilibrium object that reflects both worker preferences over income/leisure and the competitive environment for labor. It is conceivable that equilibrium forces would lead to a labor supply function like that of the fixed-wage model, in which the FLSA has an effect on the hours set at hiring.

In this section, I show that the prediction of the fixed-job model that the FLSA has little to no effect on h^* or z^* is robust to embedding Equation (9) into an extension of the Burdett and Mortensen (1998) model of equilibrium with on-the-job search.¹⁶ In the context of the search model, the only effect of the overtime rule on the distribution of h^* is mediated through the minimum wage, which rules out some of the (z^*, h^*) pairs that would occur in the unregulated equilibrium. In a numerical calibration, this effect is quite small, suggesting that equilibrium effects play only a minor role in how the FLSA overtime rule impacts anticipated

¹⁶This remains true even in the perfectly competitive limit of the model, the basic reason being that workers choose to accept jobs on the basis of their known total earnings z^* , rather than the straight-time wage.

hours or straight-time wages. This motivates the strategy in Section 4.4, in which z^* and h^* are treated as fixed when considering the impact of the FLSA on straight-wages.

F.2.1 The model

I focus on a minimal extension of Burdett and Mortensen (1998) that takes firms to be homogeneous in their technology and workers to be homogeneous in their tastes over the tradeoff between income and working hours.¹⁷ Let there be a large number N_w of workers and large number N_f of firms, and define $m = N_w/N_f$.¹⁸ Formally, we model this as a continuum of workers with mass m , and continuum of firms with unit mass. Firms choose a value of pay z and hours h to apply to all of their workers. Each period, there is an exogenous probability λ that any given worker receives a job offer, drawn uniformly from the set of all firms. Employed workers accept a job offer when they receive an earnings-hours package that they prefer to the one they currently hold, where preferences are captured by a utility function $u(z, h)$ taken to be homogeneous across workers and strictly quasiconcave, where $u_z > 0$ and $u_h < 0$. If a worker is not currently employed, they leave unemployment for a job offer if $u(z, h) \geq u(b, 0)$, where b represents a reservation earnings level required to incent a worker to enter employment. Workers leave the labor market with probability δ each period, and an equal number enters the non-employed labor force.

Before we turn to earnings-hours posting decision of firms, we can already derive several relationships that must hold for the earnings-hours distribution in a steady state equilibrium. First note that the share unemployed v of the workforce must be $v = \frac{\delta}{\delta+\lambda}$, since mass $m(1-v)\delta$ enters unemployment each period, and $m\lambda v$ leaves (taking for granted here that firms only post job offers that are preferred to unemployment, which is indeed a feature of the actual equilibrium). Let's say that job (z, h) is "inferior" to (z', h') when $u(z, h) \leq u(z', h')$. Let P_{ZH} be the firm-level distribution over offers (Z_j, H_j) , and define

$$F(z, h) := P_{ZH}(u(Z_j, H_j) \leq u(z, h)) \quad (11)$$

to be the fraction of firms offering inferior job packages to (z, h) . The separation rate of workers at a firm choosing (z, h) is thus: $s(z, h) = \delta + \lambda(1 - F(z, h))$. To derive the recruitment of new workers to a given firm each period, we define the related quantity $G(z, h)$ – the fraction of employed workers that are at inferior firms to (z, h) . In a steady state, note that $G(z, h)$

¹⁷The model presented here bears similarity to that of Hwang et al. (1998), which also considers search equilibrium with non-wage amenities such as hours. My model generalizes the preferences of workers to be possibly non-quasilinear, which allows my model to support hours dispersion in equilibrium, even with identical firms. In their model, by contrast, firms are allowed to be heterogeneous but all firms with the same production technology would offer the same quantity of hours.

¹⁸Here we largely follow the notation of the presentation of the Burdett & Mortensen model by Manning (2003).

must satisfy

$$\underbrace{m(1-v) \cdot G(z, h)(\delta + \lambda(1 - F(z, h)))}_{\text{mass of workers leaving set of inferior firms}} = \underbrace{mv\lambda F(z, h)}_{\text{mass of workers entering set of inferior firms}}$$

since the number of workers at firms inferior to (z, h) is assumed to stay constant. To get the RHS of the above, note that workers only enter the set of firms inferior to (z, h) from unemployment, and not from firms that they prefer. This expression allows us to obtain the recruitment function $R(z, h)$ to a firm offering (z, h) . Recruits will come from inferior firms and from unemployment, so that

$$\begin{aligned} R(z, h) &= \lambda m ((1-v)G(z, h) + v) \\ &= \lambda mv \left(\frac{\lambda F(z, h)}{\delta + \lambda(1 - F(z, h))} + 1 \right) \\ &= m \left(\frac{\delta \lambda}{\delta + \lambda(1 - F(z, h))} \right) \end{aligned}$$

Combining with the separation rate, we obtain the steady-state labor supply function facing each firm:

$$N(z, h) = R(z, h)/s(z, h) = \frac{m\delta\lambda}{(\delta + \lambda(1 - F(z, h))^2} \quad (12)$$

Eq. (12) is analogous to the baseline Burdett and Mortensen model without hours, with the quantity $F(z, h)$ playing the role of the firm-level CDF of wages from the baseline model.

Now we turn to how the form of $F(z, h)$ in general equilibrium. We take the profits of firms to be

$$\pi(z, h) = N(z, h)(p(h) - z) = m\delta\lambda \cdot \frac{p(h) - z}{(\delta + \lambda(1 - F(z, h))^2} \quad (13)$$

where the function $p(h)$ corresponds to net revenue per worker $e(h) - \psi$, with $e(h)$ being a weakly concave and increasing “effective labor” function with $e(0) = 0$, and z recurring non-wage costs per worker. To simplify some of the exposition, we will emphasize the simplest case of $p(h) = p \cdot h$, such that worker hours are perfectly substitutable across workers.

In equilibrium, the identical firms each playing a best response to $F(z, h)$, and thus all choices of (z, h) in the support of P_{ZH} must yield the same level of profits π^* . This gives an expression for $F(z, h)$ over all (z, h) in the support of P_{ZH} , in terms of π^* :

$$F(z, h) = 1 + \frac{\delta}{\lambda} - \sqrt{\frac{m\delta}{\lambda} \cdot \frac{p(h) - z}{\pi^*}} \quad (14)$$

where we subtract the positive square root since the negative square root cannot deliver a real number less than or equal to unity for $F(z, h)$. Note that Eq. (14) only needs to hold at (z, h) that are actually chosen by firms in equilibrium

It follows from Eqs. (14) and (12) that we can rank firms in equilibrium by $F(z, h)$ and therefore by size according to the quantity $z - p(h)$. Note that since Eq. (12) is continuously differentiable in (z, h) , we can rule out mass points in P_{ZH} by an argument paralleling that in Burdett and Mortensen (1998). Suppose $P_{ZH}(z, h) = \delta > 0$ for some (z, h) . Then any firm located at (z, h) and earning positive profits could increase their profits further by offering a sufficiently small increase in compensation (or reduction in hours, or a combination of both). Since $F(z + \delta_z, h) = F(z, h) + \delta$ to first order, there exists a small enough δ_z such that $\pi(z + \delta_z, h) > \pi(z, h)$ by Eq. (13).

To fully characterize the equilibrium P_{ZH} , I first argue that for a strictly quasiconcave utility function u , workers cannot be indifferent between more than two points (z, h) that share a value of $z - p(h)$ (see Figure 18 below). This implies that offers in the support of P_{ZH} lie along a one dimensional path through \mathbb{R}^2 . Consider for example the case of perfect hours substitutability: $p(h) = ph$, and imagine moving from a given point (z, h) in the support of P_{ZH} continuously along a line that keeps $z - ph$ and hence $F(z, h)$ constant. Since $F(z, h)$ is constant along this line, we must have that either worker utility is constant or that P_{ZH} has no additional mass along the line. However, we cannot be moving along an indifference curve of $u(z, h)$, as strict convexity of preferences implies that the marginal rate of substitution between compensation and hours can equal p (or more generally $p'(h)$, which is non-increasing) at no more than a single point for a single level of utility. Thus, P_{ZH} puts a positive density on at most one point along each isoquant of $z - p(h)$, and must have positive density on each isoquant within some connected interval. Given this, we can parametrize the points in support of P_{ZH} by a single scalar $t \in [0, 1]$, such that $\text{supp}(P_{ZH}) = \{(z(t), h(t))\}_{t \in [0, 1]}$ and $t = F(z(t), h(t))$.

Now observe that each $(z(t), h(t))$ must pick out the point along its respective isoquant of $z - p(h)$ which delivers the highest utility to workers, i.e.:

$$(z(t), h(t)) = \underset{z, h}{\text{argmax}} u(z, h) \text{ s.t. } z - p(h) = \eta(t) \quad (15)$$

where $\eta(t) = \frac{\pi^* \lambda}{m \delta} (1 - \frac{t}{1 + \delta/\lambda})^2$ is the value of $p(h(t)) - z(t)$ such that $F(z(t), h(t)) = t$ according to Eq.(14), viewed as a function of t . If instead we had $u(z(t), h(t)) < \max_{(z, h): z - p(h) = F^{-1}(t)} u(z, h)$, then any firm located at $(z(t), h(t))$ could profitably deviate to the argmax while keeping profits per worker constant but increasing their labor supply by attracting workers from $(z(t), h(t))$. The first order condition for this problem implies that $(z(t), h(t))$ maintains a marginal rate of substitution of $p'(h(t))$ (p in the baseline case) between compensation and hours at all t , while the slope of the path $(z(t), h(t))$ can be derived from the implicit function theorem:

$$\frac{z'(t)}{h'(t)} = - \left. \frac{u_{hh}(z, h) + p''(h)u_z(z, h) + p'(h)u_{zh}(z, h)}{p'(h)u_{zz}(z, h) + u_{zh}(z, h)} \right|_{(z, h) = (z(t), h(t))}$$

The curve AB shown in Figure 18 depicts the path $\{(z(t), h(t))\}_{t \in [0, 1]}$ for a case in which

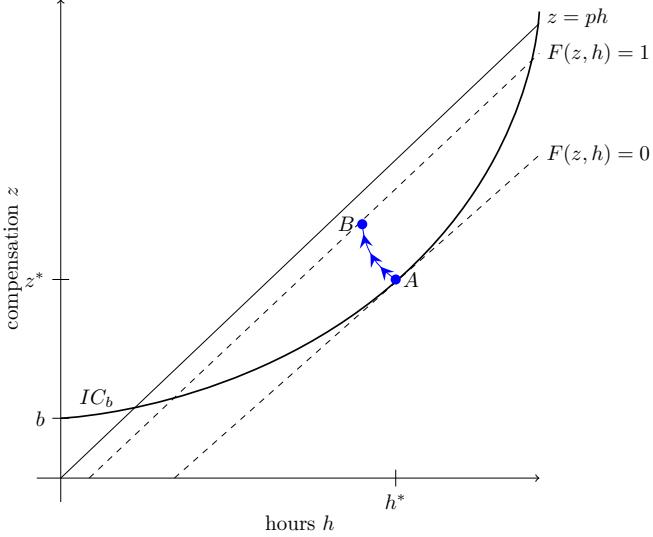


Figure 18: The support of the equilibrium distribution of compensation-hours offers (z, h) lies along the arrowed (blue) curve AB . Figure shows the case of perfect hours substitutability $p(h) = ph$. Plain curve IC_b is the indifference curve passing through the unemployment point $(b, 0)$. The least desirable firm in the economy lies at the pair (z^*, h^*) , labeled by A , where IC_b has a slope of p . The other points chosen by firms are found by starting at point A and moving in the direction of higher utility, while maintaining a marginal rate of substitution of p between hours and earnings. This path intersects the line of solutions to $F(z, h) = 1$ given Eq. (14) at point B . Note that this line still lies below the zero profit line $z = ph$, as firms make positive profit. Curve AB is shown for a general non-quasilinear, non-homothetic utility function (see text for details).

preferences are neither homothetic nor quasilinear, for example: $u(z, h) = \frac{z^{1-\gamma}}{1-\gamma} - \beta \frac{h^{1+1/\epsilon}}{1+1/\epsilon}$. If preferences were instead homothetic then AB would be a straight line pointing to the northwest from A . In the numerical calibration, I take preferences to follow the non-quasilinear Stone-Geary functional form.¹⁹ If preferences were quasilinear in income (for example the above with $\gamma = 0$), then AB would be a vertical line rising from point A and there would be no hours dispersion in equilibrium (as in Hwang et al., 1998).

To pin down the initial point A , we note that it must lie on the indifference curve passing through the unemployment point $(b, 0)$, labeled as IC_b in Figure 18. If it were to the northwest of the IC_b curve, a firm located there could increase profits by offering a lower value of $z - p(h)$, since given that $F(z(0), h(0)) = 0$ their steady state labor supply already only recruits from unemployment. However, they cannot offer a pair that lies to the southeast of IC_b , since they could never attract workers from unemployment. I assume that the marginal rate of substitution between compensation and hours is less than $p'(0)$ at $(z, h) = (b, 0)$ (such that there are gains from trade) and increases continuously with h , eventually passing $p'(h)$ at

¹⁹A CES generalization of Stone-Geary preferences also results in a straight line AB : $u(z, h) = [\theta(z - \gamma_z)^\lambda + (1 - \theta)(\gamma_h - h)^\lambda]^{1/\lambda}$. It is also possible to obtain a non-linear path AB while maintaining constant elasticity of substitution between earnings and leisure. The work of Sato (1975) on production functions suggests utility functions satisfying $\frac{u_z(z, h)}{u_h(z, h)} = \left(\frac{z - \gamma_z}{h - \gamma_h}\right)^{\frac{1}{1-\lambda}} \phi(u(c, h))$ where ϕ is any positive function.

some point h^* . This point is unique given strict quasiconcavity of $u(\cdot)$. Then, let z^* be the earnings value such that workers are indifferent between (z^*, h^*) and unemployment $(b, 0)$, which represents a reservation level of utility required to enter employment.

Finally, we can also express $F(z, h)$ as a function of $(z^*, h^*) = (z(0), h(0))$. Using that $F(z^*, h^*) = 0$ and $\pi^* = \pi(z^*, h^*)$, we can rewrite Equation (14) as:

$$F(z, h) = \left(1 + \frac{\delta}{\lambda}\right) \left[1 - \sqrt{\frac{p(h) - z}{p(h^*) - z^*}}\right] \quad (16)$$

The firms at point B in Figure 18 thus solve $z - p(h) = \left(\frac{\delta}{\delta + \lambda}\right)^2 (z^* - p(h^*))$. Equilibrium profits are $\pi^* = m(p(h^*) - z^*) \cdot \frac{\lambda/\delta}{(1+\lambda/\delta)^2}$. By Eq. (16) we can also work out that $\eta(t) = \left(1 - \frac{t}{1+\delta/\lambda}\right)^2 (ph^* - z^* - \psi)$. Note that in equilibrium, there exists dispersion not only in both earnings and in hours (provided preferences are not quasi-linear), but also in effective hourly wages, as the ratio $z(t)/h(t)$ is also strictly increasing with t . Note that π^* goes to zero in the limit that $\lambda/\delta \rightarrow \infty$. In this limit dispersion over hours, earnings, and hourly earnings all disappear as the line AB collapses to a single point on the zero profit line $z = p(h)$.²⁰

F.2.2 Effects of FLSA policies

Now consider the introduction of a minimum wage, which introduces a floor on the hourly wage $w := z/h$. I assume that the point (z^*, h^*) does not satisfy the minimum wage, so that the minimum wage binds and rules out part of the unregulated support of P_{ZH} . The left panel of Figure 19 depicts the resulting equilibrium, in which the initial point $(z(0), h(0))$ moves to the point marked A' , at which the marginal rate of substitution between compensation and hours is $p'(h)$, but the compensation-hours pair just meets the minimum wage. This compresses the distribution P_{ZH} compared with the unregulated equilibrium from Figure 18, which now follows a subset of the original path AB . In a stochastic dominance sense, all jobs see a reduction in hours and an increase in total compensation (and hence a compounded effect on hourly wages) when a minimum wage is introduced or increased.

The right panel of Figure 19 shows how equilibrium is further affected if in addition to a binding minimum wage, premium pay is required at a higher minimum wage $1.5\bar{w}$ for hours in excess of 40, provided that the point A' lies at an hours value that is greater than 40. In this case, $(z(0), h(0))$ will lie at point A'' , at which the marginal rate of substitution between compensation and hours is equal to h' , and compensation is equal to the minimum-compensation function under both the minimum wage and overtime policies: $\underline{w}(h) := \underline{w}h + 1(h > 40)(h - 40)\underline{w}/2$.

²⁰Note that there is no contradiction here as the argument that point A must be on IC_b relies on $F(z(0), h(0)) = 0$, which is implied by no mass points in P_{ZH} , in turn implied by firms making positive profit.

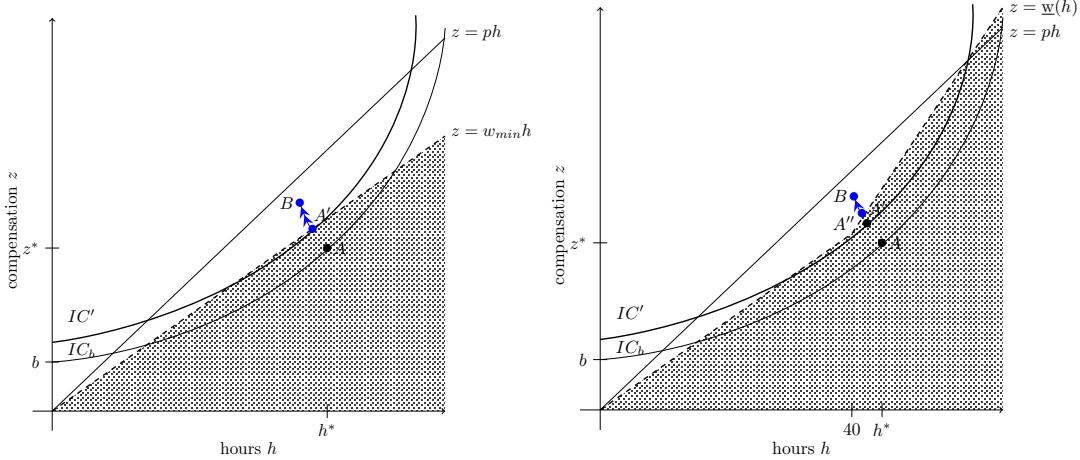


Figure 19: Left panel shows the support of the equilibrium distribution of compensation-hours offers (z, h) under a binding minimum wage. The compensation hours pairs that do not meet $\underline{w}h$ are indicated by the shaded region. The lowest-wage offer in the economy moves from point A in the unregulated equilibrium to the point A' on the minimum wage line $z = \underline{w}h$ at which the marginal rate of substitution between compensation and hours equals p . This is equal to the point at which curve AB from Figure 18 crosses the minimum wage line. Curve $A'B$ traces the remainder of curve AB . The compensation-hours offers are thus more compressed and the new distribution of earnings stochastically dominates the distribution from the unregulated equilibrium, while the opposite is true of hours. Right panel shows how this effect is augmented when overtime premium pay for hours in excess of 40 is required, and A' lies at greater than 40 hours. In this case the support of P_{ZH} begins at point A'' , which lies on the kinked minimum wage function $\underline{w}(h)$.

F.2.3 Calibration

To allow wealth effects in worker utility while facilitating calibration based on existing empirical studies, we assume worker utility is Stone-Geary:

$$u(z, h) = \beta \log(z - \gamma_z) + (1 - \beta) \log(\gamma_h - h)$$

This simple specification allows a closed form solution to the path $(z(t), h(t))$, given by the following Proposition, which follows directly from the optimization problem (15), while also working out the initial point $(z(0), h(0))$ in each policy regime. Using this, I can then calibrate the model to consider the effects of FLSA policies on earnings and hours.

Proposition. *Under Stone-Geary preferences and linear production $p(h) = ph - \psi$, the equilibrium offer distribution is a uniform distribution over $\{(z(t), h(t))\}_{t \in [0,1]}$, where:*

$$\begin{pmatrix} z(t) \\ h(t) \end{pmatrix} = \begin{pmatrix} p\beta\gamma_h + (1 - \beta)\gamma_z - \beta\psi - \beta\eta(t) \\ \beta\gamma_h + \frac{1-\beta}{p}(\gamma_z + \psi) + \frac{(1-\beta)}{p}\eta(t) \end{pmatrix}$$

where $\eta(t) = \left(1 - \frac{t}{1+\delta/\lambda}\right)^2 \cdot (ph(0) - z(0) - \psi)$. The initial point $(z(0), h(0))$ is

1. $h(0) = \gamma_h - \left(\frac{(b-\gamma_c)(1-\beta)}{p\beta}\right)^\beta \gamma_h^{1-\beta}$ and $z(0) = z^* = \gamma_z + \left(\frac{p\beta\gamma_h}{1-\beta}\right)^{1-\beta} ((b - \gamma_c)(1 - \beta))^\beta$ in the unregulated equilibrium

2. $h(0) = (\frac{p\beta}{1-\beta}\gamma_h + \gamma_z)(\underline{w} - \frac{p\beta}{1-\beta})^{-1}$ and $z(0) = \underline{w}h(0)$ with a binding minimum wage of \underline{w} (binding in the sense that $z^* < \underline{w}h^*$)
3. $h(0) = (\frac{p\beta}{1-\beta}\gamma_h + \gamma_z + 20\underline{w})(1.5\underline{w} - \frac{p\beta}{1-\beta})^{-1}$ and $z(0) = 1.5\underline{w}h(0) - 20\underline{w}$ with a minimum wage of \underline{w} and time-and-a-half overtime pay after 40 hours, if the expression for $h(0)$ in item 2. is greater than 40

Moments with respect to the worker distribution can be evaluated for any measurable function $\phi(z, h)$ as:

$$E_{workers}[\phi(Z_i, H_i)] = \left(1 + \frac{\lambda}{\delta}\right) \int_0^1 \phi(z(t), h(t)) \cdot \left(1 + \frac{\lambda}{\delta}(1-t)\right)^{-2} dt$$

I calibrate the model focusing on a lower-wage labor market where productivity is a constant $p = \$15$. I allow non-wage costs of $\psi = \$100$ a week, with the value based on estimates of benefit costs in the low-wage labor market.²¹ I take $b = \$250$ corresponding to unemployment benefits that can be accrued at zero weekly hours of work.²² We calibrate the factor λ/δ using estimates from Manning (2003) using the proportion of recruits from unemployment. Using Manning's estimates from the US in 1990 of about 55% of recruits coming from unemployment, this implies a value of $\lambda/\delta \approx 3$ in the baseline Burdett and Mortensen, 1998 model.

To calibrate the preference parameters, I first pin down β from estimates of the marginal propensity to reduce earnings after random lottery wins (Imbens et al. 2001; Cesarini et al. 2017). Both of these studies report a value in the neighborhood of $\beta = 0.85$. I take a value of $\gamma_z = \$200$ as the level of consumption at which the marginal willingness to work is infinite, and take $\gamma_h = 50$ hours of work per week. I choose this value according to a rule-of-thumb as the average hours among full-time workers in the US (42.5), divided by β .²³

Given these values, we can compute moments of functions of the joint distribution of compensation and hours using the Proposition and numerical evaluation of the integrals. Table 14 reports worker-level means of hours, weekly compensation, and the hourly wage z/h , as well as employment and profits per worker averaged across the firm distribution. In the unregulated equilibrium, the lowest-compensated workers work about 49 hours a week earning about \$300, while the highest-compensated workers work about 46 hours and earn more than \$550. This equates to a more than doubling of the hourly wage, which is about \$6 for the $t = 0$ workers and over \$12 for the $t = 1$ workers. For each of the first three variables, the mean is much closer to the $t = 1$ value than the $t = 0$ value, which follows

²¹Specifically, I take a benefit cost of \$2.43 per hour worked for the 10th percentile of wages in 2019: BLS ECEC, multiplied by the average weekly hours worked of 42.5 from the 2018 CPS summary, which results in $102.425 \approx 100$.

²²We use the UI replacement rate for single adults 2 months after unemployment from the OECD. Taking this for individuals at 2/3 of average income (the lowest available in this table), and then use a BLS figure for average income at the 10% percentile of 22,880 , we have $b \approx \$22,880 \cdot 0.6/52.25 = \263

²³ Cesarini et al. (2017) point out that when γ_c and no-earned income, optimal hours choice is $\beta\gamma_h$. By comparison, these authors calibrate γ_h to be about 35 hours in the Swedish labor market.

from the higher- t firms having more employees. The convexity of the labor supply function across values of t is apparent from the firm size row: the largest firm is about 16 times as large as the smallest, while the average firm size is four times larger than the $t = 0$ firms. The final row reports weekly profits per worker: the average worker captures more than half of the employer surplus for the $t = 0$ worker, whose weekly compensation is comparable to the employer's profit for that worker.

	Unregulated equilibrium			$\underline{w} = 7.25$	$\underline{w} = 7.25$ & OT	$\underline{w} = 12$ & OT
	t=0	t=1	mean	mean	mean	mean
weekly hours	48.85	45.71	46.34	46.18	46.11	45.51
weekly earnings	297.36	564.68	511.22	524.31	530.93	581.78
hourly wage	6.09	12.35	11.06	11.37	11.53	12.78
firm size / smallest	1.00	16.00	4.00	4.00	4.00	4.00
weekly profit per worker	335.46	20.97	146.76	119.81	106.18	1.49

Table 14: Results from the calibration. The parameter $t \in [0, 1]$ indicates firm rank in desirability from the perspective of workers. Means for weekly hours, weekly earnings, and hourly wages are computed with respect to the worker distribution, while firm size and profits per worker is averaged with respect to the firm distribution.

The third column of Table 14 adds a minimum wage set at the current federal rate of \$7.25. This provides a small increase of about 30 cents to the average hourly wage, which now begins at \$7.25 for $t = 0$ rather than \$6.06. Note that the minimum wage provides spillovers by reallocating firm mass up the entire wage ladder, beyond the mechanical effect of increasing the wages of those previously below 7.25. Average hours worked are decreased slightly due to the minimum wage, by about ten minutes per week. As this effect is mediated by a wealth effect in labor supply, we can expect it to be small unless worker preferences deviate significantly from quasi-linearity with respect to income. With $\beta = .85$, this effect is reasonably modest but non-negligible. In the fourth column, we see that the combination of the minimum wage and overtime premium has little effect beyond the direct effect of the minimum wage: hourly earnings increase another 15 cents and hours worked go down by another 0.07. Finally, we see that increasing the minimum wage to \$12 has much larger effects: adding another dollar to average wages and reducing working time by a bit more than half an hour per week. Given the fixed costs assumed in this calibration, a \$12 minimum wage causes employers to run on extremely thin margins for these workers (who have \$15 an hour productivity). However, note that in this model a minimum wage causes neither an increase nor decrease in aggregate non-employment, as this is governed in the steady state only by λ/δ . Thus, the average absolute firm size is unchanged across the policy environments.

G Additional identification results for the bunching design

This section presents several additional sufficient conditions for point or partial identification in the bunching design, beyond Theorem 1 from the main text. In this section, I continue with the notation Y_i rather than h_{it} as in Appendix A. For simplicity, I in this section assume that Y_0 and Y_1 admit a density everywhere so there is no counterfactual bunching at the kink. However, the results here can be applied given a known $p = P(Y_{0i} = Y_{1i} = k)$, as in Section 4.3, by trimming p from the observed bunching and re-normalizing the distribution $F(y)$.

I first consider parametric assumptions when treatment effects are assumed homogeneous, recasting some existing results from the literature into my generalized framework. Then I turn to nonparametric restrictions that also assume homogeneous treatment effects, before stating some results with heterogeneous treatments.

G.1 A generalized notion of homogeneous treatment effects

Recall that in the isoelastic model, treatment effects are homogeneous across units after a log transformation of the choice variable y . In order to formalize and generalize results from the literature that have focused on the isoelastic model, let begin with a generalized notion of homogenous treatment effects. For any strictly increasing and differentiable transformation $G(\cdot)$, let us define for each unit i :

$$\delta_i^G := G(Y_{0i}) - G(Y_{1i})$$

The iso-elastic model common in the bunching-design literature predicts that while Δ_i is heterogeneous across i , δ_i^G is homogeneous when G is taken to be the natural logarithm function. In this case Δ_i^G is proportional to a reduced form elasticity measuring the percentage change in $y_i(\mathbf{x})$ when moving from constraint B_{1i} to B_{0i} . In particular, in the simplest case of a bunching design in which B_0 and B_1 are linear functions of y with slopes ρ_0 and ρ_1 respectively, and utility follows the iso-elastic quasi-linear form of Equation (4), we have:

$$\delta_i^G = \delta := |\epsilon| \cdot \ln(\rho_1/\rho_0)$$

for all units i , when G is taken to be the natural logarithm.

Note that under CHOICE and CONVEX the result of Lemma 1 holds with $G(\cdot)$ applied to each of Y_i , Y_{0i} , and Y_{1i} , since G is strictly increasing. When δ_i^G is homogeneous for some G with common value δ , we thus have that $\mathcal{B} = P(G(Y_{0i}) \in [G(k), G(k) + \delta])$ by Proposition 1. Since $G(\cdot)$ is strictly increasing, we can still write the bunching condition in terms of counterfactual “levels” Y_{0i} as

$$\mathcal{B} = P(Y_{0i} \in [k, k + \Delta]) \text{ where } \Delta = G^{-1}(G(k) + \delta) - k \quad (17)$$

For example, $\Delta = k(e^\delta - 1)$ in the iso-elastic model. The parameter Δ is equal to the parameter Δ_0^* introduced in Section 4.3, since $\delta_i^G = \delta$ implies rank invariance between Y_{0i} and Y_{1i} . Δ can be seen as a pseudo-parameter plays the same role as Δ would in a setup in which we assumed a constant treatment effects in levels $\Delta_i = \Delta$. If it can be pinned down, it will also be possible to identify δ . Nevertheless, it will be important to keep track of the function G when δ_i^G is assumed homogeneous. For instance, homogeneous $\delta_i^G = \delta$ implies that $f_0^G(G(k) + \delta) = f_1^G(G(k))$ but not that $f_0(k + \Delta) = f_1(k)$, where f_d^G is the density of $G(d_i)$ for each $d \in \{0, 1\}$.

G.2 Parametric approaches with homogeneous treatment effects

The approach introduced by Saez 2010 assumes that the density $f_0(y)$ is linear on the bunching interval $[k, k + \Delta]$. This affords point-identification of ϵ in an iso-elastic utility model. We can use the notation above to provide the following generalization of this result:

Proposition 5 (identification by linear interpolation, à la Saez 2010). *If $\delta_i^G = \delta$ for some G , $F_1(y)$ and $F_0(y)$ are continuously differentiable, and $f_0(y)$ is linear on the interval $[k, k + \Delta]$, then with CONVEX, CHOICE:*

$$\mathcal{B} = \frac{1}{2} (G^{-1}(G(k) + \delta) - k) \left\{ \lim_{y \uparrow k} f(y) + \frac{G'(G^{-1}(G(k) + \delta))}{G'(k)} \lim_{y \downarrow k} f(y) \right\}$$

Proof. See Section H. □

In particular, given the iso-elastic model with budget slopes ρ_0 and ρ_1 :

$$\mathcal{B} = \frac{\Delta}{2} \left\{ \lim_{y \uparrow k} f(y) + \frac{k}{k + \Delta} \lim_{y \downarrow k} f(y) \right\} = \frac{k}{2} \left(\left(\frac{\rho_0}{\rho_1} \right)^\epsilon - 1 \right) \left(\lim_{y \uparrow k} f(y) + \left(\frac{\rho_0}{\rho_1} \right)^{-\epsilon} \lim_{y \downarrow k} f(y) \right) \quad (18)$$

which serves as the main estimating equation from Saez (2010) (and can be solved for ϵ by the quadratic formula). The empirical approach of Saez (2010) can thus be seen as applying a result justified in a much more general model than the iso-elastic utility function assumed therein, provided that the researcher is willing to assume homogeneous treatment effects (possibly after some known transformation G , and/or conditional on observables).²⁴ Note that the linearity assumption of Proposition 5 could be falsified by visual inspection: it implies that right and left limits of the derivative of the density of Y_i at the kink are equal.

A more popular approach, following Chetty et al. (2011), is to use a global polynomial approximation to $f_0(y)$, which interpolates $f_0(y)$ inwards from both directions across the

²⁴Note that if we had instead assumed that $f_0^G(y)$ is linear (on the interval $[G(k), G(k) + \delta^G]$), then we simply replace $f(y)$ by $f^G(y)$ in the above and let G be the identity function, which can be readily solved for δ^G with the simpler expression $\delta^G = \mathcal{B}/\frac{1}{2} \{ \lim_{y \uparrow k} f^G(y) + \lim_{y \downarrow k} f^G(y) \}$.

missing region of unknown width Δ . This technique has the added advantage of accommodating diffuse bunching, for which the relevant \mathcal{B} is the total “excess-mass” around k rather than a perfect point mass at k . I focus here on the simplest case in which bunching is exact, as in the overtime setting. The polynomial approach can be seen as a special case of the following result:

Proposition 6 (identification from global parametric fit, à la Chetty et al. 2011).

Suppose $f_0(y)$ exists and belongs to a parametric family $g(y; \theta)$, where $f_0(y) = g(y; \theta_0)$ for some $\theta_0 \in \Theta$, and that $\delta_i^G = \delta$ for some G and CONVEX and CHOICE hold. Then, if:

1. $g(y; \theta)$ is an analytic function of y on the interval $[k, k + \Delta]$ for all $\theta \in \Theta$, and
2. $g(y; \theta_0) > 0$ for all $y \in [k, k + \Delta]$,

it follows that Δ (and hence δ) is identified as $\Delta(\theta_0)$, where for any θ , $\Delta(\theta)$ is the unique Δ such that $\mathcal{B} = \int_k^{k+\Delta} g(y; \theta) dy$, and θ_0 satisfies

$$f(y) = \begin{cases} g(y; \theta_0) & y < k \\ g(y + \Delta(\theta_0); \theta_0) & y > k \end{cases} \quad (19)$$

Proof. See Section H. □

The standard approach of fitting a high-order polynomial to $f_0(y)$ can satisfy the assumptions of Proposition 6, since polynomial functions are analytic everywhere. Proposition 6 yields an identification result that can justify an estimation approach similar to one often made in the literature, based on Chetty et al. (2011).²⁵ However, it requires taking seriously the idea that $f_0(y) = g(y; \theta_0)$, treating the approach as parametric rather than as a series approximation to a nonparametric density $f_0(y)$. This assumption is very strong. Indeed, assuming that $g(y; \theta_0)$ follows a polynomial exactly has even more identifying power than is exploited by Proposition 6. In particular, if we also have that $f_1(y) = g(y; \theta_1)$ then we could use data on either side of the kink to identify by θ_0 and θ_1 , which would allow identification of the average treatment effect with complete treatment effect heterogeneity.

G.3 Nonparametric approaches with homogeneous treatment effects

The additional assumptions from the preceding section have allowed for point-identification of causal effects under an assumption of homogenous treatment effects. These assumptions have taken the form of parametric restrictions on the density of counterfactual choices Y_{0i} in the missing region $[k, k + \Delta]$: that this density is linear or more generally fits a parametric family

²⁵The estimation technique proposed by Chetty et al. (2011) ignores the shift term $\Delta(\theta)$ in Equation (19), a limitation discussed by Kleven (2016). This is perhaps less problematic in typical settings where the bunching is somewhat diffuse around the kink, in contrast to the overtime setting in which bunching is exact, and the slope of the density is far from zero near 40. A more robust estimation procedure for parametric bunching designs could be based on iterating on Equation (19) after updating $\Delta(\theta)$, until convergence. This presents an interesting topic for future research.

of analytic functions. As has been argued in Blomquist and Newey (2017), these parametric assumptions drive all of the identification, an undesirable feature from the standpoint of robustness to departures from them. I now explore some non-parametric assumptions about $f_0(y)$ that yield bounds on Δ in a model with homogeneous treatment effects.

For example, monotonicity of $f_0(y)$ has been suggested by Blomquist and Newey (2017) as an alternative assumption in the context of the iso-elastic model. A result based on monotonicity that allows heterogeneous treatment effects is presented in Section G.4. However, monotonicity may be restrictive if the density of Y_0 has a mode near the kink point. In this case, local log-concavity of $f_0(y)$ may be a more attractive assumption (concavity or convexity would be another). ²⁶ Note that log-concavity is a stronger version of the bi-log-concavity assumption used in the main text, but still nests many common parametric distributions such as the uniform, normal, exponential extreme value and logistic. For simplicity, this result assumes homogeneous treatment effects in levels (rather than after applying a function G).

Proposition 7 (bounds from log-concavity). *Suppose that $\Delta_i = \Delta$ and that $f_0(y)$ is log-concave in the interval $y \in [k, k + \Delta]$ and continuously differentiable at k and $k + \Delta$. Then, under CONVEX and CHOICE:*

$$\Delta \in [\Delta^L, \Delta^U]$$

where

$$\Delta^U = \mathcal{B} \cdot \frac{\ln(f_+) - \ln(f_-)}{f_+ - f_-} \quad \text{and} \quad \Delta^L = \left(\frac{f_-}{f'_-} - \frac{f_+}{f'_+} \right) \ln \left(\frac{\mathcal{B} + \frac{f_-^2}{f'_-} - \frac{f_+^2}{f'_+}}{\frac{f_-}{f'_-} - \frac{f_+}{f'_+}} \right) + \frac{f_+}{f'_+} \ln f_+ - \frac{f_-}{f'_-} \ln f_-$$

where $f'_- := \lim_{y \uparrow k} f'(y)$ and $f'_+ := \lim_{y \downarrow k} f'(y)$

Proof. See Figure 20. Derivation of expressions available by request. \square

Intuition for Proposition 7 is provided in Figure 20. If $f_0(y)$ is log convex rather than log-concave in the missing region, then the bounds Δ^L and Δ^U can simply be swapped. Or, if we suppose that f_0 is either log-concave or log-convex locally: $\Delta \in [\min\{\Delta^U, \Delta^L\}, \max\{\Delta^U, \Delta^L\}]$.

G.4 Alternative identification strategies with heterogeneous treatment effects

An argument made in Saez 2010 and Kleven and Waseem (2013) uses a uniform density assumption to allow heterogeneous treatment in the bunching-design. If a kink is very small, then this might be justified as an approximation given smoothness of $f(\Delta, y)$, since Δ_i will be “small” for all i . Below I state an analog of this result in the generalized bunching design

²⁶Log concavity has previously been assumed as a shape restriction in the context of bunching by Diamond and Persson (2016), though to study the effects of manipulation on other variables, rather than for the effect of incentives on the variable being manipulated.

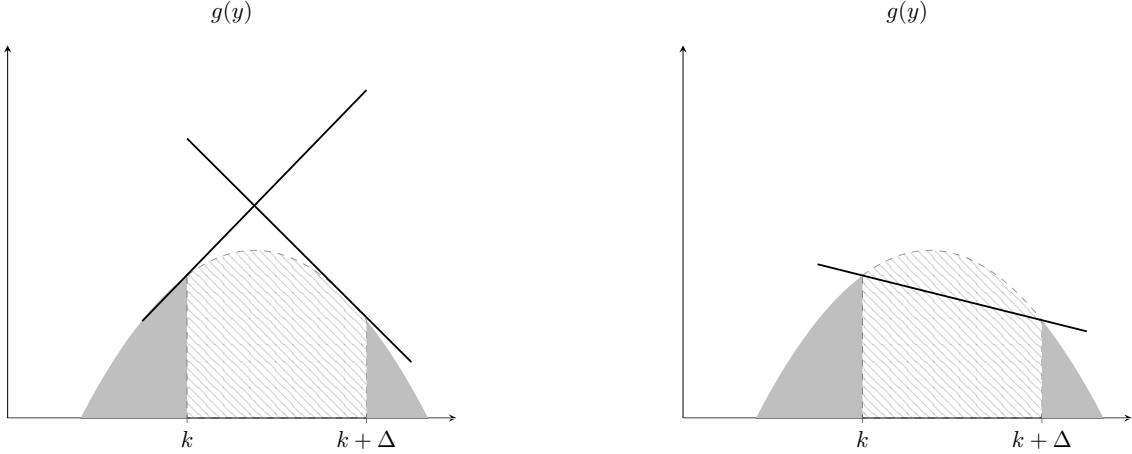


Figure 20: The left and right panels of this figure depict intuition for the lower and upper bounds on Δ in Proposition 7. In both panels, the hatched region is the missing region $[k, k + \Delta]$ which contains known mass B . The function plotted is $g(y)$, the logarithm of $f_0(y)$. Outside of the missing region, this function is known. Concavity of $g(y)$ provides both upper and lower bounds for the values of $g(y)$ inside the missing region, yielding the analytic bounds in Proposition 7.

framework of this paper. The result will make use of the following Lemma, which states that treatment effects must be positive at the kink:

Lemma POS (positive treatment effect at the kink). *Under WARP and CHOICE, $P(\Delta_i \geq 0 | Y_{0i} = k) = P(\Delta_i \geq 0 | Y_{1i} = k) = 1$.*

Proof. See proof of Lemma 1, which rules out the events $Y_{0i} \leq k < Y_{1i}$ and $Y_{0i} < k \leq Y_{1i}$. \square

Proposition 8 (identification of an ATE under uniform density approximation). *Let Δ_i and Y_{0i} admit a joint density $f(\Delta, y)$ that is continuous in y at $y = k$. For each Δ assume that $f(\Delta, Y_0) = f(\Delta, k)$ for all Y_0 in the region $[k, k + \Delta]$. Under Assumptions WARP and CHOICE*

$$\mathbb{E} [\Delta_i | Y_{0i} = k] \geq \frac{B}{\lim_{y \uparrow k} f(y)},$$

with equality under CONVEX.

Proof. Note that

$$\begin{aligned} B \leq P(Y_{0i} \in [k, k + \Delta_i]) &= \int_0^\infty d\Delta \int_k^{k+\Delta} dy \cdot f(\Delta, y) = \int_0^\infty f(\Delta, k) \Delta d\Delta \\ &= f_0(k) P(\Delta_i \geq 0 | Y_{0i} = k) \mathbb{E} [\Delta_i | Y_{0i} = k, \Delta \geq 0] \\ &\leq \lim_{y \uparrow k} f(y) \cdot \mathbb{E} [\Delta_i | Y_{0i} = k] \end{aligned}$$

using Lemma POS in the last step. The inequalities are equalities under CONVEX. \square

Lemma SMALL in Appendix B formalizes the idea that the uniform density approximation from Proposition 8 becomes exact in the limit of a “small” kink.

We can also produce a result based on monotonicity, allowing heterogeneous treatment effects. Let $\tau_0 := E[\Delta_i | Y_{0i} = k]$ and $\tau_1 := E[\Delta_i | Y_{1i} = k]$.

Proposition 9 (monotonicity with heterogeneous treatment effects). *Assume CONVEX and CHOICE, and suppose the joint density $f_0(\Delta, y)$ of Δ_i and Y_{0i} and the joint density $f_1(\Delta, y)$ of Δ_i both exist. Suppose first that $f_0(\Delta, y)$ is weakly increasing on the interval $y \in [k, k + \Delta]$ for all Δ in the support of Δ_i . Then*

$$\tau_1 \geq \frac{\mathcal{B}}{\lim_{y \downarrow k} f(y)} \quad \text{and} \quad \tau_0 \leq \frac{\mathcal{B}}{\lim_{y \uparrow k} f(y)}$$

Alternatively, if $f_1(\Delta, y)$ is weakly decreasing on the interval $y \in [k - \Delta, k]$ for each Δ , then

$$\tau_0 \geq \frac{\mathcal{B}}{\lim_{y \uparrow k} f(y)} \quad \text{and} \quad \tau_1 \leq \frac{\mathcal{B}}{\lim_{y \downarrow k} f(y)}$$

Proof. Note that $f_1(\Delta, y) = f_0(\Delta, y + \Delta)$ for any y, Δ , and hence $f_0(y, \Delta)$ is increasing (decreasing) on $[k, k + \Delta]$ whenever $f_1(y, \Delta)$ is increasing (decreasing) on $[k - \Delta, k]$. Then:

$$\begin{aligned} \mathcal{B} &= \int_0^\infty d\Delta \int_k^{k+\Delta} dy \cdot f_0(\Delta, y) \leq \int_0^\infty \Delta f_0(\Delta, k) d\Delta = f_0(k)\tau_0 \\ \mathcal{B} &= \int_0^\infty d\Delta \int_{k-\Delta}^k dy \cdot f_1(\Delta, y) \geq \int_0^\infty \Delta f_1(\Delta, k) d\Delta = f_1(k)\tau_0 \end{aligned}$$

for example in the first case, where we have used Lemma POS. The reverse case is analogous \square

This result implies that when treatment effects are statistically independent of Y_0 (for example when they are homogenous): $\Delta_i \perp Y_{0i}$, we have that $\mathbb{E}[\Delta_i] = \tau_0 = \tau_1 \in \left[\frac{\mathcal{B}}{\max\{f_-, f_+\}}, \frac{\mathcal{B}}{\min\{f_-, f_+\}} \right]$.

Other approaches to identification with heterogeneous treatment effects are possible when the researcher observes covariates X_i that are unaffected by a counterfactual shift between B_1 and B_0 . For example, assuming that $E[X_i | Y_{0i} = y]$ or $E[X_i | Y_{1i} = y]$ are Lipschitz continuous with a known constant leads to a lower bound on maximum of τ_0 and τ_1 from an observed discontinuity of $E[X_i | Y_i = y]$ at $y = k$. Another strategy for using covariates would be to model the potential outcomes Y_{0i} and Y_{1i} as functions of them. If we are willing to suppose that

$$Y_{0i} = g_0(X_i) + U_{0i} \quad \text{and} \quad Y_{1i} = g_1(X_i) + U_{1i}$$

with U_{1i} and U_{0i} each statistically independent of X_i , then the censoring of the distributions of Y_{0i} and Y_{1i} in Lemma 1 can be “undone”, following the results of Lewbel and Linton (2002).²⁷ This would allow estimation of the unconditional average treatment effect as well as quantile treatment effects at all levels. However, the assumption that U_0 and U_1 are independent of X is quite strong.

²⁷Lewbel and Linton (2002) establish identification of $g(x)$ and $F_U(u)$ in a model where the econometrician observes censored observations of $Y = g(X) + U$. Given knowledge of the distribution of X , the estimated marginal distributions of U_1 and U_2 , and the estimated function $g(x)$ the researcher could estimate the distributions $F_1(y) = P(Y_{1i} \leq y)$ and $F_0(y) = P(Y_{0i} \leq y)$ by deconvolution, and then estimate causal effects.

G.5 Two bunching design settings from the literature

Below I discuss two examples from the literature that illustrate the general kink bunching design framework described in Section 4. The first is the classic labor supply example, where convexity of preferences arises from increasing opportunity costs of time allocated to labor. In the second example, firms are again the decision makers but now the “running variable” y is a function of two underlying choice variables \mathbf{x} .

Example 1: Labor supply with taxation

Individuals have preferences $\tilde{u}_i(c, y)$ over consumption c , and labor earnings y , where ϵ_i is a vector of parameters capturing heterogeneity over the disutility of labor, labor productivity, etc. The agent’s budget constraint is $c \leq y - B(y)$ where $B(y)$ is income tax as a function of pre-tax earnings y . $\tilde{u}_i(c, y)$ is taken to be strictly quasi-concave in (c, y) for each i as the opportunity cost of leisure rises with greater earnings, and monotonically increasing in consumption. Define $z = y - c$ to be tax liability, and let $u_i(z, y) = \tilde{u}_i(y - z, y)$ which is monotonically decreasing in tax. Individuals now choose a value of y (e.g. by adjusting hours of work, number of jobs, or misreporting) given a progressive tax schedule $B_k(y) = \tau_0 y + 1(y \geq k)(\tau_1 - \tau_0)(y - k)$, with the kink arising from an increase in marginal tax rates from τ_0 to $\tau_1 > \tau_0$ at $y = k$. The budget functions are $B_0(y) = \tau_0 y$, $B_1(y) = \tau_1 y - (\tau_1 - \tau_0)k$, and the kinked budget constraint can be written $z \geq B_k(y) = \max\{B_0(y), B_1(y)\}$.

Example 2: Minimum tax schemes

Best et al. (2015) study a feature of corporate taxation in Pakistan in which firms pay the maximum of a tax on output and a tax on reported profits:

$$B(r, \hat{w}) = \max\{\tau_\pi(r - \hat{w}), \tau_r r\}$$

where r is firm revenue, \hat{w} is reported costs, and $\tau_r < \tau_\pi$. Under the profit tax, firms have incentive to reduce their tax liability by inflating the value \hat{w} above their true costs of production $w_i(r)$. One can write tax liability as a piecewise function in which the tax regime depends on reported profits as a fraction of output: $y = \frac{r - \hat{w}}{r} = 1 - \frac{\hat{w}}{r}$:

$$B(r, \hat{w}) = \begin{cases} \tau_r r & \text{if } y \leq \tau_r / \tau_\pi \\ \tau_\pi(r - \hat{w}) & \text{if } y > \tau_r / \tau_\pi \end{cases}$$

This function has a “kink” in both r and \hat{w} when $y(r, \hat{w}) = k = \tau_r / \tau_\pi$. In this setting, $B_0(r, \hat{w}) = \tau_r r$, corresponding to a tax on output while $B_1(r, \hat{w}) = \tau_\pi(r - \hat{w})$ describes a tax on (reported) profits. Both functions are linear, and hence weakly convex, in the vector (r, \hat{w}) . The functions B_{0i} , B_{1i} and y_i are all common across firms.

Assume that firm i chooses the pair $\mathbf{x} = (r, \hat{w})$ according to preferences $u_i(z, \mathbf{x})$, which are strictly decreasing in tax liability z and strictly quasiconcave in (z, r, \hat{w}) . In Best et al. (2015), preferences are for example taken to be in a baseline model:

$$u_i(z, r, \hat{w}) = r - w_i(r) - g_i(\hat{w} - w_i(r)) - z \quad (20)$$

where $g_i(\cdot)$ represents costs of tax evasion by misreporting costs. This specification of $u_i(z, r, \hat{w})$ is strictly quasi-concave provided that the production and evasion cost functions $w_i(\cdot)$ and $g_i(\cdot)$ are strictly convex.

With such preferences, the presence of the minimum tax kink can be expected to lead to a firm response among both margins of \mathbf{x} : r and \hat{w} . In particular, consider a linear approximation to $\Delta_i = Y_i(0) - Y_i(1)$ for a buncher with $Y_{0i} \approx k$, keeping the i implicit:

$$\begin{aligned} \Delta &\approx \frac{dy(r, \hat{w})}{\hat{w}} \Big|_{(r_0, \hat{w}_0)} \Delta_{\hat{w}} + \frac{dy(r, \hat{w})}{r} \Big|_{(r_0, \hat{w}_0)} \Delta_r \\ &= \frac{\hat{w}_0}{r_0^2} \Delta_r - \frac{1}{r_0} (\Delta_{w(r)} + \Delta_{(\hat{w}-w(r))}) \\ &\approx \frac{\hat{w}_0}{r_0^2} \Delta_r - \frac{1}{r_0} (w'(r_0) \Delta_{ri} + \Delta_{(\hat{w}-w(r))}) \\ &= \frac{1}{r_0} \{(1 - Y_0 - w'(r_0)) \Delta_r \Delta_{\hat{w}}\} \approx \frac{1}{r} \{-k \Delta_r - \Delta_{(\hat{w}-w)}\} \\ &\approx \frac{1}{r_0} \left\{ -\frac{\tau_r}{\tau_\pi} \cdot r \epsilon^r \frac{d(1 - \tau_E)}{\tau_E} - \Delta_{\hat{w}i} \right\} = \frac{\tau_r^2}{\tau_\pi} \epsilon^r - \frac{\Delta_{(\hat{w}-w)}}{r_0} \end{aligned} \quad (21)$$

where ϵ^r is the elasticity of firm revenue with respect to the net of effective tax rate $1 - \tau_E$. In this case, when crossing from the output to reported profits regime $\frac{d(1 - \tau_E)}{\tau_E} = -\tau_r$, implying the final expression (see Best et al. 2015 for definition of τ_E). We have also used the optimality condition that $w'(r_0) = 1$. Expression (21) shows that the response to the minimum tax kink is almost entirely driven by a response on the difference between reported and actual costs: $\hat{w}_i - w_i(r)$. This is because τ_r is less than 1%, so the first term ends up not contributing meaningfully in practice (it scales as the square of τ_r). In this empirical setting, it is thus possible to interpret the bunching response as a response to one of the components of \mathbf{x} , despite \mathbf{x} being a vector.

We can also situate the setting of Best et al. (2015) in terms of a continuum of cost functions, as in Section A.6. In particular, let $\rho \in [0, 1]$ and define

$$B(r, \hat{w}; \rho, k) = \frac{\tau_r}{1 - \rho(1 - k)} (y - \rho c)$$

Then $B_0(r, \hat{w}) = B(r, \hat{w}; 0)$ and $B_1(r, \hat{w}; \tau_r/\tau_\pi) = B(r, \hat{w}; 1, \tau_r/\tau_\pi)$. It can be verified that for any $\rho' > \rho$ and k , $B(r, \hat{w}; \rho', k) > B(r, \hat{w}; \rho, k)$ iff $y_i(r, \hat{w}) > k$, with equality when $y_i(r, \hat{w}) = k$. The path from $\rho_0 = 0$ to $\rho_1 = 1$ passes through a continuum of tax policies in which the tax base gradually incorporates reported costs, while the tax rate on that tax base also increases continuously with ρ .

H Additional proofs

H.1 Proof of Propositions 1 and 2

Consider Proposition 1. Item i) in the proof of Lemma 1 establishes that under CHOICE and WARP $Y_i = k$ implies $Y_{1i} \leq k \leq Y_{0i}$, since taking contrapositives we have that $(Y_i \geq k \text{ and } Y_i \leq k)$ implies $Y_{1i} \leq k \leq Y_{0i}$. We have also seen from item ii) that under CHOICE and CONVEX $Y_{1i} \leq k \leq Y_{0i}$ also implies $Y_i = k$, thus $Y_{1i} \leq k \leq Y_{0i}$ and $Y_i = k$ are equivalent. Note that by adding $\Delta_i = Y_{0i} - Y_{1i}$ to both sides of the inequality $Y_{1i} \leq k$, we have that $Y_{0i} \leq k + \Delta_i$. Combining with the other inequality that $Y_{0i} \geq k$, we can thus rewrite the event $Y_{1i} \leq k \leq Y_{0i}$ as $Y_{0i} \in [k, k + \Delta_i]$ (or equivalently to $Y_{1i} \in [k - \Delta_i, k]$). We thus have that $\mathcal{B} \leq P(Y_{0i} \in [k, k + \Delta])$ under CHOICE and WARP, and that $\mathcal{B} = P(Y_i = k) = P(Y_{1i} \leq k \leq Y_{0i})$ under CHOICE and CONVEX.

Now consider Proposition 2. By item i) in the proof of Proposition 1, we have that under WARP and CHOICE $Y_{0i} \leq k \implies Y_i = Y_{0i}$. Thus, for any $\delta > 0$ and $y < k$: $Y_{0i} \in [y - \delta, y]$ implies that $Y_i \in [y - \delta, y]$ and hence $P(Y_{0i} \in [y - \delta, y]) \leq P(Y_i \in [y - \delta, y])$. This implies that $f_0(y) - f(y) = \lim_{\delta \downarrow 0} \frac{P(Y_{0i} \in [y - \delta, y]) - P(Y_i \in [y - \delta, y])}{\delta} \leq 0$, i.e. that $f(y) \geq f_0(y)$. An analogous argument holds for Y_1 , where we consider the event $Y_{1i} \in [y, y + \delta]$ any $y > k$. Under CONVEX instead of WARP, the inequalities are all equalities, by Lemma 1.

H.2 Proof of Lemma 2

Let $\Delta_i^k(\rho, \rho') := Y_i(\rho, k) - Y_i(\rho', k)$ for any $\rho, \rho' \in [\rho_0, \rho_1]$ and value of k .

Assumption SMOOTH (regularity conditions). *The following hold:*

1. $P(\Delta_i^k(\rho, \rho') \leq \Delta, Y_i(\rho, k) \leq y)$ is twice continuously differentiable at all $(\Delta, y) \neq (0, k^*)$, for any $\rho, \rho' \in [\rho_0, \rho_1]$ and k .
2. $Y_i(\rho, k) = Y(\rho, k, \epsilon_i)$, where ϵ_i has compact support $E \subset \mathbb{R}^m$ for some m . $Y(\cdot, k, \cdot)$ is continuously differentiable on all of $[\rho_0, \rho_1] \times E$, for every k .
3. there possibly exists a set $\mathcal{K}^* \subset E$ such that $Y(\rho, k, \epsilon) = k^*$ for all $\rho \in [\rho_0, \rho_1]$ and $\epsilon \in \mathcal{K}^*$. The quantity $\mathbb{E} \left[\frac{\partial Y_i(\rho, k)}{\partial \rho} \middle| Y_i(\rho, k) = y, \epsilon_i \notin \mathcal{K}^* \right]$ is continuously differentiable in y for all y including k^* .

In the remainder of this proof I keep k be implicit in the functions $Y_i(\rho, k)$ and $\Delta_i^k(\rho, \rho')$, as it will remained fixed. Item 1 of SMOOTH excludes the point $(0, k^*)$ on the basis that we may expect point masses at $Y_i(\rho) = k^*$, as in the overtime setting. Following Section 4, item 3 imposes that all such ‘‘counterfactual bunchers’’ have zero treatment effects, while also introducing a further condition that will be used later in Lemma 3. Let K_i^* be an indicator for $\epsilon_i \in \mathcal{K}^*$ and denote $p = P(K_i^* = 1)$. Item 1 implies that the density $f_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y)$ is continuous in y whenever $y \neq k^*$ or $\Delta \neq 0$, so I define $f_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, k^*) = \lim_{y \rightarrow k^*} f_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y)$

for any ρ, ρ' and Δ . Similarly, we can define the marginal density $f_\rho(y)$ of $Y_i(\rho)$ at k^* to be $\lim_{y \rightarrow k^*} f_\rho(y)$ for any ρ .

By item 1 of Assumption SMOOTH, the marginal $F_\rho(y) := P(Y_i(\rho) \leq y)$ is differentiable away from $y = k$ with derivative $f_\rho(y)$. From the proof of Theorem 1 it follows that $\mathcal{B} \leq F_{\rho_1}(k) - F_{\rho_0}(k) + p(k)$ with equality under CONVEX, and thus:

$$\begin{aligned} \mathcal{B} - p(k) &\leq F_{\rho_1}(k) - F_{\rho_0}(k) \\ &= \int_{\rho_0}^{\rho_1} \frac{d}{d\rho} F_\rho(k) d\rho \\ &= \int_{\rho_0}^{\rho_1} \lim_{\delta \downarrow 0} \frac{F_{\rho+\delta}(k) - F_\rho(k)}{\delta} d\rho \\ &= \int_{\rho_1}^{\rho_0} \lim_{\delta \downarrow 0} \frac{P(Y_i(\rho+\delta) \leq k \leq Y_i(\rho)) - p(k)}{\delta} d\rho \\ &= \int_{\rho_1}^{\rho_0} f_\rho(k) \mathbb{E} \left[-\frac{Y_i(\rho)}{d\rho} \middle| Y_i(\rho) = k \right] d\rho \end{aligned}$$

where the third equality applies the identity $1 = P(Y_{0i} \leq k) + P(Y_i(\rho) \leq k \leq Y_i(\rho+\delta)) + P(Y_{1i} > k)$ under CHOICE and WARP (this follows from item i) of the proof of Lemma 1) to the pair of choice constraints $B(\rho)$ and $B(\rho+\delta)$, noting that $P(Y_i(\rho) < k) = F_\rho(k) - p(k)$. The final equality uses Lemma SMALL.

H.3 Proof of Lemma SMALL

Throughout this proof we let f_W denote the density of a generic random variable or random vector W_i , if it exists. Write $\Delta_i(\rho, \rho') = \Delta_i(\rho, \rho', \epsilon_i)$ where $\Delta_i(\rho, \rho', \epsilon) := Y(\rho, \epsilon) - Y(\rho', \epsilon)$.

$$\begin{aligned} \lim_{\rho' \downarrow \rho} \frac{P(Y_i(\rho) \leq k \leq Y_i(\rho')) - p(k)}{\rho' - \rho} &= \lim_{\rho' \downarrow \rho} \frac{P(Y_i(\rho) \in [k, k + \Delta(\rho, \rho')_i]) - p(k)}{\rho' - \rho} \\ &= \lim_{\rho' \downarrow \rho} \frac{P(Y_i(\rho) \in (k, k + \Delta(\rho, \rho')_i])}{\rho' - \rho} \\ &= \lim_{\rho' \downarrow \rho} \frac{1}{\rho' - \rho} \int_0^\infty d\Delta \int_k^{k+\Delta} dy \cdot f_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y) \\ &= \lim_{\rho' \downarrow \rho} \int_0^\infty d\Delta \int_k^{k+\Delta} dy \cdot \frac{f_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, k) + (y - k)r_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, k, y)}{\rho' - \rho} \end{aligned} \tag{22}$$

where we have used that by item 1 the joint density of $\Delta_i(\rho, \rho')$ and $Y_i(\rho)$ exists for any ρ, ρ' and is differentiable and $r_{\Delta(\rho, \rho'), Y(\rho)}$ is a first-order Taylor remainder term satisfying

$$\lim_{y \downarrow k} |r_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y)| = |r_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, k)| = 0$$

for any Δ .

I now show that the whole term corresponding to this remainder is zero. First, note that:

$$\begin{aligned} \left| \lim_{\rho' \downarrow \rho} \int_0^\infty d\Delta \int_k^{k+\Delta} dy \cdot \frac{(y - k)r_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y)}{\rho' - \rho} \right| &= \lim_{\rho' \downarrow \rho} \left| \int_0^\infty d\Delta \int_k^{k+\Delta} dy \cdot \frac{(y - k)r_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y)}{\rho' - \rho} \right| \\ &\leq \lim_{\rho' \downarrow \rho} \int_0^\infty d\Delta \int_k^{k+\Delta} dy \cdot \left| \frac{(y - k)r_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y)}{\rho' - \rho} \right| \\ &\leq \lim_{\rho' \downarrow \rho} \int_0^\infty d\Delta \frac{\Delta}{\rho' - \rho} \int_k^{k+\Delta} dy \cdot |r_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, y)| \end{aligned}$$

where I've used continuity of the absolute value function and the Minkowski inequality. Define $\xi(\rho, \rho') = \sup_{\epsilon \in E} \Delta(\rho, \rho', \epsilon)$. The strategy will be show that $\lim_{\rho' \downarrow \rho} \xi(\rho, \rho') = 0$, and then since $r_{\Delta_i(\rho, \rho'), Y_i(\rho)}(\Delta, y) = 0$ for any $\Delta > \xi(\rho, \rho')$ and all y (since the marginal density $f_{\Delta(\rho, \rho')}(\Delta)$ would be zero for such Δ). With $\xi(\rho, \rho')$ so-defined:

$$\begin{aligned} \text{RHS of above} &\leq \lim_{\rho' \downarrow \rho} \int_0^{\xi(\rho, \rho')} d\Delta \frac{\xi(\rho, \rho')}{\rho' - \rho} \int_k^{k+\xi(\rho, \rho')} dy \cdot |r_{\Delta_i(\rho, \rho'), Y_i(\rho)}(\Delta, y)| \\ &= \lim_{\rho' \downarrow \rho} \frac{\xi(\rho, \rho')}{\rho' - \rho} \cdot \lim_{\rho' \downarrow \rho} \int_0^{\xi(\rho, \rho')} d\Delta \int_0^{\xi(\rho, \rho')} dy \cdot |r_{\Delta_i(\rho, \rho'), Y_i(\rho)}(\Delta, k + y)| \end{aligned} \quad (23)$$

where in the second step I have assumed that each limit exists (this will be demonstrated below). Let us first consider the inner integral of the above: $\int_k^{k+\xi(\rho, \rho')} dy \cdot |r_{\Delta_i(\rho, \rho'), Y_i(\rho)}(\Delta, y)|$, for any Δ . Supposing that $\lim_{\rho' \downarrow \rho} \xi(\rho, \rho') = 0$, it follows that this inner integral evaluates to zero, by the Leibniz rule and using that $r_{\Delta_i(\rho, \rho'), Y_i(\rho)}(\Delta, k) = 0$. Thus the entire second limit is equal to zero.

Now I prove that $\lim_{\rho' \downarrow \rho} \xi(\rho, \rho') = 0$ and that $\lim_{\rho' \downarrow \rho} \frac{\xi(\rho, \rho')}{\rho' - \rho}$ exists. First, note that continuous differentiability of $Y(\rho, \epsilon_i)$ implies $Y_i(\rho)$ is continuous for each i so $\lim_{\rho' \downarrow \rho} \Delta_i(\rho, \rho') = 0$ point-wise in ϵ . We seek to turn this point-wise convergence into uniform convergence over ϵ , i.e. that $\lim_{\rho' \downarrow \rho} \sup_{\epsilon \in E} \Delta(\rho, \rho', \epsilon) = \sup_{\epsilon \in E} \lim_{\rho' \downarrow \rho} \Delta(\rho, \rho', \epsilon) = \sup_{\epsilon \in E} 0 = 0$. The strategy will be to use equicontinuity of the sequence and compactness of E . Consider any such sequence $\rho_n \xrightarrow{n} \rho$ from above, and let $f_n(\epsilon) := Y(\rho, \epsilon) - Y(\rho_n, \epsilon)$ and $f(\epsilon) = \lim_{n \rightarrow \infty} f_n(\epsilon) = 0$. Equicontinuity of the sequence $f_n(\epsilon)$ says that for any $\epsilon, \epsilon' \in E$ and $e > 0$, there exists a $\delta > 0$ such that $\|\epsilon - \epsilon'\| < \delta \implies |f_n(\epsilon) - f_n(\epsilon')| < e$.

This follows from continuous differentiability of $Y(\rho, \epsilon)$. Let $M = \sup_{\rho \in [\rho_0, \rho_1], \epsilon \in E} |\nabla_{\rho, \epsilon} Y(\rho, \epsilon)|$. M exists and is finite given continuity of the gradient and compactness of $[\rho_0, \rho_1] \times E$. Then, for any two points $\epsilon, \epsilon' \in E$ and any $\rho \in [\rho_0, \rho_1]$:

$$|Y(\rho, \epsilon) - Y(\rho, \epsilon')| = \left| \int_{\epsilon'}^{\epsilon} \nabla_{\epsilon} Y(\rho, \epsilon) \cdot d\epsilon \right| \leq \int_{\epsilon'}^{\epsilon} |\nabla_{\epsilon} Y(\rho, \epsilon) \cdot d\epsilon| \leq M \int_{\epsilon'}^{\epsilon} \|d\epsilon\| \leq M \|\epsilon - \epsilon'\|$$

where $d\epsilon$ is any path from ϵ to ϵ' and I have used the definition of M and Cauchy-Schwarz in the second inequality. The existence of a uniform Lipschitz constant M for $Y(\rho, \epsilon)$ implies a uniform equicontinuity of $Y(\rho, \epsilon)$ of the form that for any $e > 0$ and $\epsilon, \epsilon' \in E$, there exists a $\delta > 0$ such that $\|\epsilon - \epsilon'\| < \delta \implies \sup_{\rho \in [\rho_0, \rho_1]} |Y(\rho, \epsilon) - Y(\rho, \epsilon')| < e/2$, since we can simply take $\delta = e/(2M)$. This in turn implies that whenever $\|\epsilon - \epsilon'\| < \delta$:

$$\begin{aligned} |Y(\rho, \epsilon) - Y(\rho_n, \epsilon) - \{Y(\rho, \epsilon') - Y(\rho_n, \epsilon')\}| &= |Y(\rho, \epsilon) - Y(\rho, \epsilon') - \{Y(\rho_n, \epsilon) - Y(\rho_n, \epsilon')\}| \\ &\leq |Y(\rho, \epsilon) - Y(\rho, \epsilon')| + |Y(\rho_n, \epsilon) - Y(\rho_n, \epsilon')| \leq e, \end{aligned}$$

our desired result. Together with compactness of E , equicontinuity implies that $\lim_{n \rightarrow \infty} \sup_{\epsilon \in E} f_n(\epsilon) = \sup_{\epsilon \in E} \lim_{n \rightarrow \infty} f_n(\epsilon) = 0$.

We apply an analogous argument for $\lim_{\rho' \downarrow \rho} \frac{\xi(\rho, \rho')}{\rho' - \rho}$, where now $f_n(\epsilon) = \frac{Y(\rho, \epsilon) - Y(\rho_n, \epsilon)}{\rho_n - \rho}$. For this case it's easier to work directly with the function $\frac{Y(\rho, \epsilon) - Y(\rho_n, \epsilon)}{\rho_n - \rho}$, showing that it is Lipschitz

in deviations of ϵ uniformly over $n \in \mathbb{N}, \epsilon \in E$.

$$\begin{aligned} \left| \frac{Y(\rho, \epsilon) - Y(\rho_n, \epsilon)}{\rho_n - \rho} - \frac{Y(\rho, \epsilon') - Y(\rho_n, \epsilon')}{\rho_n - \rho} \right| &= \frac{1}{\rho_n - \rho} \left| \int_{\epsilon'}^{\epsilon} \nabla_{\epsilon} Y(\rho, \epsilon) \cdot \mathbf{d}\epsilon - \int_{\epsilon'}^{\epsilon} \nabla_{\epsilon} Y(\rho_n, \epsilon) \cdot \mathbf{d}\epsilon \right| \\ &\leq \frac{1}{\rho_n - \rho} \left(\left| \int_{\epsilon'}^{\epsilon} \nabla_{\epsilon} Y(\rho, \epsilon) \cdot \mathbf{d}\epsilon \right| + \left| \int_{\epsilon'}^{\epsilon} \nabla_{\epsilon} Y(\rho_n, \epsilon) \cdot \mathbf{d}\epsilon \right| \right) \\ &\leq \frac{2M}{\rho_n - \rho} \int_{\epsilon'}^{\epsilon} \|\mathbf{d}\epsilon\| \leq \frac{2M}{\rho_n - \rho} \|\epsilon - \epsilon'\| \end{aligned}$$

This implies equicontinuity of $\frac{Y(\rho, \epsilon) - Y(\rho_n, \epsilon)}{\rho_n - \rho}$ with the choice $\delta = e(\rho_n - \rho)/(2M)$. As before, equicontinuity and compactness of E allow us to interchange the limit and the supremum, and thus:

$$\begin{aligned} \lim_{n \rightarrow \infty} \frac{\xi(\rho, \rho_n)}{\rho_n - \rho} &= \lim_{n \rightarrow \infty} \frac{\sup_{\epsilon \in E} \{Y(\rho, \epsilon) - Y(\rho_n, \epsilon)\}}{\rho_n - \rho} = \lim_{n \rightarrow \infty} \sup_{\epsilon \in E} \frac{Y(\rho, \epsilon) - Y(\rho_n, \epsilon)}{\rho_n - \rho} \\ &= \sup_{\epsilon \in E} \lim_{n \rightarrow \infty} \frac{Y(\rho, \epsilon) - Y(\rho_n, \epsilon)}{\rho_n - \rho} = \sup_{\epsilon \in E} \frac{\partial Y(\rho, \epsilon)}{\partial \rho} := M' < \infty \end{aligned}$$

where finiteness of M' follows from it being defined as the supremum of a continuous function over a compact set. This establishes that the first limit in Eq. (23) exists and is finite, completing the proof that it evaluates to zero.

Given that the second term in Eq. (22) is zero, we can simplify the remaining term as:

$$\begin{aligned} \lim_{\rho' \downarrow \rho} \frac{P(Y_i(\rho) \leq k \leq Y_i(\rho')) - p(k)}{\rho' - \rho} &= \lim_{\rho' \downarrow \rho} \frac{1}{\rho' - \rho} \int_0^\infty f_{\Delta(\rho, \rho'), Y(\rho)}(\Delta, k) \Delta d\Delta \\ &= f_\rho(k) \lim_{\rho' \downarrow \rho} \frac{1}{\rho' - \rho} P(\Delta_i(\rho, \rho') \geq 0 | Y_i(\rho) = k) \\ &\quad \cdot \mathbb{E} [\Delta_i(\rho, \rho') | Y_i(\rho) = k, \Delta_i(\rho, \rho') \geq 0] \\ &= f_\rho(k) \lim_{\rho' \downarrow \rho} \frac{1}{\rho' - \rho} \mathbb{E} [\Delta_i(\rho, \rho') | Y_i(\rho) = k, \Delta_i(\rho, \rho')] \\ &= f_\rho(k) \mathbb{E} \left[\lim_{\rho' \downarrow \rho} \frac{\Delta_i(\rho, \rho')}{\rho' - \rho} \middle| Y_i(\rho) = k \right] \\ &= f_\rho(k) \mathbb{E} \left[-\frac{Y_i(\rho)}{d\rho} \middle| Y_i(\rho) = k \right] \end{aligned}$$

where I have used Lemma POS and then finally the dominated convergence theorem. To see that we may use the latter, note that $\frac{dY_i(\rho)}{d\rho} = \frac{\partial Y(\rho, \epsilon_i)}{\partial \rho} < M$ uniformly over all $\epsilon_i \in E$, and $\mathbb{E} [M | Y_i(\rho) = k] = M < \infty$.

H.4 Proof of Appendix D Proposition 3

Note: this proof follows the notation of Y_i from Appendix A, rather than h_{1it} from Appendix D and the main text. Begin with the following observations:

- $(Y < k) \implies (Y_0 = Y)$ and $(Y > k) \implies (Y_1 = Y)$ both follow from convexity of preferences, and linearity of the cost functions B_1 and B_0 . From these two it also follows that $(Y_1 \leq k \leq Y_0) \implies (Y = k)$. See proof of Theorem 1, which treats this case.
- For firm-choosers: $(Y_0 < k) \implies (Y = Y_0)$, since the cost function B_0 coincides with B_k for $y \leq k$, and is higher otherwise. Similarly $(Y_1 > k) \implies (Y = Y_1)$. Together these also imply that $(Y = k) \implies (Y_1 \leq k \leq Y_0)$.
- By analogous logic, for worker-choosers: $(Y_0 \geq k) \implies (Y = Y_1)$, and $(Y_1 \leq k) \implies (Y = Y_0)$ using that their utility functions are strictly increasing in c . Together these also imply that $Y_1 \leq k \leq Y_0$ can only occur if $Y_0 = Y_1 = k$.

Now consider the claims of the Proposition:

- $P(Y_{it} = k \text{ and } K_{it}^* = 0) = P(Y_{1it} \leq 40 \leq Y_{0it} \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 0)$
- $\lim_{y \uparrow 40} f(y) = P(W_{it} = 0) \lim_{y \uparrow 40} f_{0|W=0}(y)$
- $\lim_{y \downarrow 40} f(y) = P(W_{it} = 0) \lim_{y \downarrow 40} f_{1|W=0}(y)$

First claim:

$$\begin{aligned} P(Y_{it} = k \text{ and } K_{it}^* = 0) &= P(Y_{it} = k \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 0) + P(Y_{it} = k \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 1) \\ &= P(Y_{1it} \leq 40 \leq Y_{0it} \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 0) + 0 \end{aligned}$$

where for the first term I've used that when $W_{it} = 0$, $(Y_{it} = k) \iff (Y_{1it} \leq 40 \leq Y_{0it})$ following Theorem 1. For the second, I've used that by the absolute continuity assumption: $P(Y_{0it} = k \text{ or } Y_{1it} = k | K_{it}^* = 0) = 0$, so:

$$\begin{aligned} P(Y_{it} = k \text{ and } K_{it}^* = 0) &= P(Y_{it} = k \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 1 \text{ and } Y_{0it} < k) \\ &\quad + P(Y_{it} = k \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 1 \text{ and } Y_{0it} > k) \\ &= P(Y_{it} = k \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 1 \text{ and } Y_{0it} < k \text{ and } Y_{1it} = k) \\ &\quad + P(Y_{it} = k \text{ and } K_{it}^* = 0 \text{ and } W_{it} = 1 \text{ and } Y_{0it} > k \text{ and } Y_{1it} = k) \\ &= 0 + 0 = 0 \end{aligned}$$

where I've used that $W_{it} = 1$ and $Y_{0it} < k$ implies that $Y_{it} = Y_{0it}$ if $Y_{1it} < k$, and $Y_{it} \in \{Y_{0it}, Y_{1it}\}$ if $Y_{1it} > k$ to eliminate the first term. The second term uses that $Y_1 \leq k \leq Y_0$ can only occur when $Y_0 = Y_1 = k$.

Second claim:

$$\begin{aligned} \lim_{y \uparrow k} f(y) &= \lim_{y \uparrow k} \frac{d}{dy} P(Y_{it} \leq y) \\ &= \lim_{y \uparrow k} \frac{d}{dy} P(Y_{it} \leq y \text{ and } W_{it} = 0) + \lim_{y \uparrow k} \frac{d}{dy} P(Y_{it} \leq y \text{ and } W_{it} = 1) \end{aligned}$$

The first term is equal to $P(W_{it} = 0) \lim_{y \uparrow k} f_{0|W=0}(y)$, and I now show that the second is equal to zero:

$$\begin{aligned} & \lim_{y \uparrow k} \frac{d}{dy} P(Y_{it} \leq y \text{ and } W_{it} = 1) \\ &= \lim_{y \uparrow k} \frac{d}{dy} P(Y_{0it} \leq y \text{ and } Y_{it} = Y_{0it} \text{ and } W_{it} = 1) \\ &= \lim_{y \uparrow k} \frac{d}{dy} P(Y_{0it} \leq y \text{ and } \{u(B_0(Y_{0it}), Y_{0it}) \geq u_{it}(B_1(y), y) \text{ for all } y > k\} \text{ and } W_{it} = 1) \end{aligned}$$

For it 's utility under B_k at Y_{0it} to be greater than that attainable at any $y > k$, the indifference curve IC_{0it} passing through Y_{0it} must lie above $B_{1it}(y) = w_{it}y + \frac{w_{it}}{2}(y - k)$ for all $y > k$. Using that IC_{0it} passes through the point $(w_{it}Y_{0it}, Y_{0it})$ with derivative w_{it} there (by the first-order condition for an optimum), we may write it as

$$\begin{aligned} IC_{0it}(y) &= w_{it}Y_{0it} + \int_{Y_{0it}}^y IC'_{0it}(y') dy' = w_{it}Y_{0it} + \int_{Y_{0it}}^y \left\{ w_{it} + \int_{Y_{0it}}^{y'} IC''_{0it}(y'') dy'' \right\} dy' \\ &\leq w_{it}y + \int_{Y_{0it}}^y M(y' - Y_{0it}) dy = w_{it}y + \frac{1}{2}(y - Y_{0it})^2 M_{it} \end{aligned}$$

using that IC_{0it} is twice differentiable. Now $IC_{0it}(y) \geq B_{1it}(y)$ for $y > k$ implies that

$$\frac{w_{it}}{M_{it}}(y - k) \leq (y - Y_{0it})^2$$

Taking for example $y = 80 - Y_{0it}$, such that $y - k = y - Y_{0it}$, we have that $Y_{0it} \leq k - \frac{w_{it}}{M_{it}}$. Thus:

$$\begin{aligned} & \lim_{y \uparrow k} \frac{d}{dy} P(Y_{it} \leq y \text{ and } Y_{it} > Y_{0it} \text{ and } W_{it} = 1) \\ &\leq \lim_{y \uparrow k} \lim_{\delta \downarrow 0} \frac{1}{\delta} P(Y_{0it} \in (y - \delta, y] \text{ and } Y_{0it} \leq k - \frac{w_{it}}{M_{it}} \text{ and } W_{it} = 1) \\ &\leq \lim_{y \uparrow k} \lim_{\delta \downarrow 0} \frac{1}{\delta} P(Y_{0it} \in (y - \delta, y] \text{ and } \frac{w_{it}}{M_{it}} \leq k - y + \delta \text{ and } W_{it} = 1) \\ &\leq \lim_{y \uparrow k} \lim_{\delta \downarrow 0} \frac{1}{\delta} P(\frac{w_{it}}{M_{it}} \leq k - y + \delta \text{ and } W_{it} = 1) \\ &\leq \lim_{\delta \downarrow 0} \frac{1}{\delta} P\left(\frac{w_{it}}{M_{it}} \leq \delta \text{ and } W_{it} = 1\right) \\ &= f_{w/m|W=1}(0) = 0 \end{aligned}$$

where we may interchange the limits given that $\frac{w_{it}}{M_{it}}$ conditional on $W_{it} = 1$ admits a density $f_{w/m|W=1}$ that is bounded in a neighborhood around 0. This, and that $f_{w/m|W=1}(0) = 0$ follows from the assumption that the distribution of M_{it}/w_{it} is bounded.

We have now proved the second claim, that $\lim_{y \uparrow k} f(y) = P(W_{it} = 0) \lim_{y \uparrow k} f_{0|W=0}(y)$.

Third claim: Analogous logic to the second claim, using the bounded 2nd derivative of IC_{1it} .

H.5 Proof of Appendix D Theorem 1*

Note: this proof follows the notation of Y_i from Appendix A, rather than h_{1it} from Appendix D and the main text. Let $T_i = 1$ be a shorthand for firm-choosers who are not counterfactual bunchers, i.e. the event $K_{it}^* = 0$ and $W_{it} = 0$.

By Theorem 1 of Dümbgen et al., 2017: for $d \in \{0, 1\}$ and any t , bi-log concavity implies that:

$$1 - (1 - F_{d|T=1}(k))e^{-\frac{f_{d|T=1}(k)}{1-F_{d|T=1}(k)}t} \leq F_{d|T=1}(k+t) \leq F_{d|T=1}(k)e^{\frac{f_{d|T=1}(k)}{F_{d|T=1}(k)}t}$$

Defining $u = F_{0|T=1}(k+t)$, we can use the substitution $t = Q_{0|T=1}(u) - k$ to translate the above into bounds on the conditional quantile function of Y_{0i} , evaluated at u :

$$\frac{F_{0|T=1}(k)}{f_{0|T=1}(k)} \cdot \ln\left(\frac{u}{F_{0|T=1}(k)}\right) \leq Q_{0|T=1}(u) - k \leq -\frac{1 - F_{0|T=1}(k)}{f_{0|T=1}(k)} \cdot \ln\left(\frac{1-u}{1-F_{0|T=1}(k)}\right)$$

And similarly for Y_1 , letting $v = F_{1|T=1}(k-t)$:

$$\frac{1 - F_{1|T=1}(k)}{f_{1|T=1}(k)} \cdot \ln\left(\frac{1-v}{1-F_{1|T=1}(k)}\right) \leq k - Q_{1|T=1}(v) \leq -\frac{F_{1|T=1}(k)}{f_{1|T=1}(k)} \cdot \ln\left(\frac{v}{F_{1|T=1}(k)}\right)$$

By RANK, we have that $Y_i = k \iff F_{0|T=1}(Y_{0i}) \in [F_{0|T=1}(k), F_{0|T=1}(k) + \mathcal{B}^*] \iff F_{1|T=1}(Y_{1i}) \in [F_{1|T=1}(k) - \mathcal{B}^*, F_{1|T=1}(k)]$ where $\mathcal{B}^* := P(Y_i = k | T = 1)$, and thus:

$$E[Y_{0i} - Y_{1i} | Y_i = k, T_i = 0] = \frac{1}{\mathcal{B}^*} \int_{F_{0|T=1}(k)}^{F_{0|T=1}(k) + \mathcal{B}^*} \{Q_{0|T=1}(u) - k\} du + \frac{1}{\mathcal{B}^*} \int_{F_{1|T=1}(k) - \mathcal{B}^*}^{F_{1|T=1}(k)} \{k - Q_{1|T=1}(v)\} dv$$

A lower bound for $E[Y_{0i} - Y_{1i} | Y_i = k, T_i = 0]$ is thus:

$$\begin{aligned} & \frac{F_{0|T=1}(k)}{f_{0|T=1}(k)(\mathcal{B}^*)} \int_{F_{0|T=1}(k)}^{F_{0|T=1}(k) + \mathcal{B}^*} \ln\left(\frac{u}{F_{0|T=1}(k)}\right) du + \frac{1 - F_{1|T=1}(k)}{f_{1|T=1}(k)(\mathcal{B}^*)} \int_{F_{1|T=1}(k) - \mathcal{B}^*}^{F_{1|T=1}(k)} \ln\left(\frac{1-v}{1-F_{1|T=1}(k)}\right) dv \\ &= g(F_{0|T=1}(k), f_{0|T=1}(k), \mathcal{B}^*) + h(F_{1|T=1}(k), f_{1|T=1}(k), \mathcal{B}^*) \end{aligned}$$

where as in Theorem 1: $g(a, b, x) = \frac{a}{bx} (a+x) \ln(1+\frac{x}{a}) - \frac{a}{b}$ and $h(a, b, x) = g(1-a, b, x)$. Similarly, an upper bound is:

$$\begin{aligned} & -\frac{1 - F_{0|T=1}(k)}{f_{0|T=1}(k)(\mathcal{B}^*)} \int_{F_{0|T=1}(k)}^{F_{0|T=1}(k) + \mathcal{B}^*} \ln\left(\frac{1-u}{1-F_{0|T=1}(k)}\right) du \\ & \quad - \frac{F_{1|T=1}(k)}{f_{1|T=1}(k)(\mathcal{B}^*)} \int_{F_{1|T=1}(k) - \mathcal{B}^*}^{F_{1|T=1}(k)} \ln\left(\frac{v}{F_{1|T=1}(k)}\right) dv \\ &= \tilde{g}(F_{0|T=1}(k), f_{0|T=1}(k), \mathcal{B}^*) + \tilde{h}(F_{1|T=1}(k), f_{1|T=1}(k), \mathcal{B}^*) \end{aligned}$$

where again $\tilde{g}(a, b, x) = -g(1-a, b, -x)$ and $\tilde{h}(a, b, x) = -g(a, b, -x)$. We have then that $E[Y_{0i} - Y_{1i} | Y_i = k, T_i = 0] \in [\Delta_k^L, \Delta_k^U]$, where:

$$\begin{aligned} \Delta_k^L &= g(F_{0|T=1}(k), f_{0|T=1}(k), \mathcal{B}^*) + g(1 - F_{1|T=1}(k), f_{1|T=1}(k), \mathcal{B}^*) \\ &= g(P(Y_{0i} \leq k \text{ and } T_i = 1), P(T_i = 1)f_{0|T=1}(k), P(Y_i = k \text{ and } T_i = 1)) \\ &\quad + g(P(Y_{1i} > k \text{ and } T_i = 1), P(T_i = 1)f_{1|T=1}(k), P(Y_i = k \text{ and } T_i = 1)) \end{aligned}$$

$$\begin{aligned}
\Delta_k^U &= -g(1 - F_{0|T=1}(k), f_{0|T=1}(k), -\mathcal{B}^*) - g(F_{1|T=1}(k), f_{1|T=1}(k), -\mathcal{B}^*) \\
&= -g(P(Y_{0i} > k \text{ and } T_i = 1), P(T_i = 1)f_{0|T=1}(k), -P(Y_i = k \text{ and } T_i = 1)) \\
&\quad - g(P(Y_{1i} \leq k \text{ and } T_i = 1), P(T_i = 1)f_{1|T=1}(k), -P(Y_i = k \text{ and } T_i = 1))
\end{aligned}$$

where I've used that the function $g(a, b, x)$ is homogeneous of degree zero and multiplied each argument by $P(T_i = 1)$. The bounds are sharp as CHOICE, CONVEX and RANK imply no further restrictions on the marginal potential outcome distributions.

Next, note that:

$$\begin{aligned}
\lim_{y \uparrow k} f(y) &= \lim_{y \uparrow k} \frac{d}{dy} P(Y_{0i} \leq y \text{ and } W_i = 0) = \lim_{y \uparrow k} \frac{d}{dy} P(Y_{0i} \leq y \text{ and } W_i = 0 \text{ and } K_i^* = 0) \\
&= P(T_i = 1) \cdot \lim_{y \uparrow k} \frac{d}{dy} P(Y_{0i} \leq y | T_i = 1) = P(T_i = 1) \cdot f_{0|T=1}(k)
\end{aligned}$$

$$\begin{aligned}
\lim_{y \downarrow k} f(y) &= -\lim_{y \downarrow k} \frac{d}{dy} P(Y_{1i} \geq y \text{ and } W_i = 0) = -\lim_{y \downarrow k} \frac{d}{dy} P(Y_{1i} \geq y \text{ and } W_i = 0 \text{ and } K_i^* = 0) \\
&= P(T_i = 1) \cdot -\lim_{y \downarrow k} \frac{d}{dy} P(Y_{1i} \geq y | T_i = 1) = P(T_i = 1) \cdot f_{1|T=1}(k)
\end{aligned}$$

$$\mathcal{B}-p = P(Y_i = k \text{ and } K_i^* = 0) = P(Y_i = k \text{ and } K_i^* = 0 \text{ and } W_i = 0) = P(Y_i = k \text{ and } T_i = 1)$$

As shown by Dümbgen et al., 2017, BLC implies the existence of a continuous density function, which assures that these density limits exist and are equal to the corresponding potential outcome densities above. Thus, the quantities $P(Y_i = k \text{ and } T_i = 1)$, $P(T_i = 1) \cdot f_{0|T=1}(k)$ and $P(T_i = 1) \cdot f_{1|T=1}(k)$ are all point-identified from the data.

Now we turn to the CDF arguments of Δ_k^L and Δ_k^U . Note that the desired quantities can be written

- $P(Y_{0i} \leq k \text{ and } T_i = 1) = P(Y_{0i} < k \text{ and } T_i = 1) = P(Y_{0i} < k \text{ and } W_i = 0)$
- $P(Y_{1i} > k \text{ and } T_i = 1) = P(Y_{1i} > k \text{ and } W_i = 0)$
- $P(Y_{0i} > k \text{ and } T_i = 1) = P(Y_{0i} > k \text{ and } W_i = 0)$
- $P(Y_{1i} \leq k \text{ and } T_i = 1) = P(Y_{1i} < k \text{ and } T_i = 1) = P(Y_{1i} < k \text{ and } W_i = 0)$

Let

$$A := P(Y_{0i} < k \text{ and } Y_i = Y_{0i} \text{ and } W_i = 1) \quad \text{and} \quad B := P(Y_{1i} > k \text{ and } Y_i = Y_{1i} \text{ and } W_i = 1)$$

The desired quantities are related to observables via A and B :

- $P(Y_i < k) = P(Y_{0i} < k \text{ and } W_i = 0) + A$
- $P(Y_i > k) = P(Y_{1i} > k \text{ and } W_i = 0) + B$
- $P(Y_i \leq k) - p = P(Y_i \leq k \text{ and } K_i^* = 0) = P(Y_i \leq k \text{ and } T_i = 1) + A = P(Y_{1i} \leq k \text{ and } W_i = 0) + A$

- $P(Y_i \geq k) - p = P(Y_i \geq k \text{ and } K_i^* = 0) = P(Y_i \geq k \text{ and } T_i = 1) + B = P(Y_{0i} > k \text{ and } W_i = 0) + B$

The four CDF arguments appearing in Δ_k^L and Δ^U are thus identified up to the correction terms A and B . A simple sufficient condition for $A = B = 0$ is that there are no worker-choosers.

H.6 Proof of Appendix F.1 Proposition 4

The first order conditions with respect to z and h are:

$$\lambda F_L(L, K)e(h) = \phi + \frac{\beta_Y(z, h) + 1}{\beta_Y(z, h)}z$$

and

$$\lambda F_L(L, K)e(h)(\eta(h)/\beta_h(z, h) + 1) = z + \phi$$

where $L = N(z, h)e(h)$, $\eta(h) := e'(h)h/e(h)$, $\beta_h(z, h) := N_h(z, h)h/N(z, h)$ and $\beta_z(z, h) := N_z(z, h)Y/N(z, h)$ are elasticity functions and λ is a Lagrange multiplier. I have assumed that the functions $|\beta_h|$, β_h , and η are strictly positive and finite globally. Combining the two equations, we have that an interior solution must satisfy either: $z = \frac{\phi \frac{\eta}{\beta_h}}{1 - \frac{\beta_z+1}{\beta_z} \frac{\beta_h+\eta}{\beta_h}}$ (Case 1), or that the denominator of the above is zero: $\frac{\beta_h}{\beta_h+\eta} = \frac{\beta_z+1}{\beta_z}$ (Case 2), where the dependence of β_z and β_h has been left implicit. Defining $\beta(z, h) = |\beta_h(z, h)|/(\beta_z(z, h) + 1)$, we can rewrite the condition for Case 2 as $\beta(z, h) = \eta(h)$.

With $\phi = 0$, we must be in Case 2 for any $z > 0$ to have positive profits, and not that positivity of z requires $\beta < \eta$ in case one. On the other hand if $\phi > 0$ we cannot have Case 1 provided that $\eta/\beta_h > 0$. Now specialize to the conditions set out in the Proposition: that $F_L = 1$, $\lambda = 1$ (profit maximization), and β_h , β_z and η are all constants. Then $z = \frac{\phi \frac{\eta}{\beta_h}}{1 - \frac{\beta_z+1}{\beta_z} \frac{\beta_h+\eta}{\beta_h}} = \phi \cdot \frac{\beta_z}{\beta_z+1}$ and the first order condition for hours becomes

$$e(h) = \phi + \phi \frac{\eta}{\beta - \eta}$$

which simplifies to $h = \left[\frac{\phi}{e_0} \cdot \frac{\beta}{\beta - \eta} \right]^{1/\eta}$.

H.7 Proof of Appendix Proposition 5

By constant treatment effects, $f_1^G(y) = f_0^G(y + \delta)$ and note that both $f_0^G(k)$ and $f_1^G(k)$ are identified from the data. These can be transformed into densities for Y_{0i} and Y_{1i} via $f_d(y) = G'(y)f_d^G(G(y))$ for $d \in \{0, 1\}$. With $f_0(y)$ linear on the interval $[k, k + \Delta]$, the integral $\int_k^{k+\Delta} f_0(y)dy$ evaluates to $\mathcal{B} = \frac{\Delta}{2}(f_0(k) + f_0(k + \Delta))$. Although $f_0(k) = \lim_{y \uparrow k} f(y)$ by CONT, $f_0(k + \Delta)$ is not immediately observable. However:

$$f_0(k + \Delta) = f_0(G^{-1}(G(k) + \delta)) = G'(k + \Delta)f_0^G(G(k) + \delta)$$

and furthermore by constant treatment effects:

$$f_0^G(G(k) + \delta) = f_1^G(G(k)) = (G'(k))^{-1} f_1(k) = (G'(k))^{-1} \lim_{y \downarrow k} f(y)$$

Combining these equations, we have the result.

H.8 Proof of Appendix Proposition 6

We seek a Δ such that for some θ_0 :

$$\mathcal{B} = \int_{\tilde{k}}^{k+\Delta} g(y; \theta_0) dy \quad (24)$$

and

$$f(y) = \begin{cases} g(y; \theta_0) & y < k \\ g(y + \Delta; \theta_0) & y > k \end{cases} \quad (25)$$

and

$$g(y; \theta_0) > 0 \text{ for all } y \in [k, k + \Delta] \quad (26)$$

Recall from Equation (17) that $\Delta = G^{-1}(G(k) + \delta) - k$ and hence $\delta = G(k + \Delta) - G(k)$. Thus if we find a unique Δ satisfying the two equations, we have found a unique value of δ : the true value of the homogenous effect δ^G .

Suppose we have two candidate values $\Delta' > \Delta$. For them to both satisfy (24), we would need $\Delta' = \Delta(\theta')$ and $\Delta = \Delta(\theta)$, where $\theta, \theta' \in \Theta$ and $\Delta(\theta_0)$ is the unique Δ satisfying Eq. (24) for a given θ_0 , which is unique for each permissible value θ_0 by the positivity condition (26). To satisfy (25), we would also need

$$g(y; \theta) = \begin{cases} f(y) & y < k \\ f(y - \Delta(\theta)) & y > k + \Delta(\theta) \end{cases} \quad g(y; \theta') = \begin{cases} f(y) & y < k \\ f(y - \Delta(\theta')) & y > k + \Delta(\theta') \end{cases} \quad (27)$$

Since $g(y; \theta)$ is a real analytic function for any $\theta \in \Theta$, the function $h_{\theta\theta'}(y) := g(y; \theta) - g(y; \theta')$ is real analytic. An implication of this is that if $h_{\theta\theta'}(y)$ vanishes on the interval $[0, \tilde{k}]$, as it must by Equation (27), it must vanish everywhere on \mathbb{R} . Thus for any $y > k + \Delta(\theta)$:

$$g(y + \Delta(\theta') - \Delta(\theta); \theta) = g(y + \Delta(\theta') - \Delta(\theta); \theta') = g(y; \theta)$$

So $g(y; \theta)$ is periodic with period $\Delta(\theta') - \Delta(\theta)$. Since g is non-negative, it cannot integrate to unity globally, and thus cannot be the same function as $f_0(y)$.

H.9 Details of calculations for policy estimates

H.9.1 Ex-post evaluation of time-and-a-half after 40

$$\mathbb{E}[Y_{0i} - Y_i] = (\mathcal{B} - p)\mathbb{E}[Y_{0i} - k | Y_i = k, K_i^* = 0] + p \cdot 0 + P(Y_{1i} > k)\mathbb{E}[Y_{0i} - Y_{1i} | Y_i > k]$$

Consider the first term

$$(\mathcal{B} - p)E[Y_{0i} - k|Y_i = k, K_i^* = 0] = (1 - p)\mathcal{B}^* \cdot \frac{1}{\mathcal{B}^*} \int_{F_{0|K^*=0}(k)}^{F_{0|K^*=0}(k)+\mathcal{B}^*} \{Q_{0|K^*=0}(u) - k\} du$$

where $\mathcal{B}^* := P(Y_i = k|K^* = 0) = \frac{\mathcal{B}-p}{1-p}$. Bounds for the rightmost quantity are given by bi-log-concavity of Y_{0i} , just as in Theorem 1. In particular:

$$\begin{aligned} (\mathcal{B} - p)E[Y_{0i} - k|Y_i = k, K_i^* = 0] &\geq (1 - p)\mathcal{B}^* \cdot \frac{F_{0|K^*=0}(k)}{f_{0|K^*=0}(k)(\mathcal{B}^*)} \int_{F_{0|K^*=0}(k)}^{F_{0|K^*=0}(k)+\mathcal{B}^*} \ln\left(\frac{u}{F_{0|K^*=0}(k)}\right) du \\ &= (1 - p)\mathcal{B}^* \cdot g(F_{0|K^*=0}(k), f_{0|K^*=0}(k), \mathcal{B}^*) \\ &= (\mathcal{B} - p) \cdot g(F_-, f_-, \mathcal{B} - p) \end{aligned}$$

and

$$\begin{aligned} (\mathcal{B} - p)E[Y_{0i} - k|Y_i = k, K_i^* = 0] &\leq -(1 - p)\mathcal{B}^* \cdot \frac{1 - F_{0|K^*=0}(k)}{f_{0|K^*=0}(k)(\mathcal{B}^*)} \int_{F_{0|K^*=0}(k)}^{F_{0|K^*=0}(k)+\mathcal{B}^*} \ln\left(\frac{1-u}{1-F_{0|K^*=0}(k)}\right) du \\ &= (1 - p)\mathcal{B}^* \cdot g'(F_{0|K^*=0}(k), f_{0|K^*=0}(k), \mathcal{B}^*) \\ &= -(\mathcal{B} - p) \cdot g(1 - p - F_-, f_+, p - \mathcal{B}) \end{aligned}$$

where as before $g(a, b, x) = \frac{a}{bx} (a + x) \ln(1 + \frac{x}{a}) - \frac{a}{b}$ and $g'(a, b, x) = -g(1 - a, b, -x)$.

Now consider the second term of $\mathbb{E}[Y_{0i} - Y_{1i}]$: $P(Y_{1i} > k)\mathbb{E}[Y_{0i} - Y_{1i}|Y_i > k]$. Taking as a lower bound an assumption of constant treatment effects in levels: $P(Y_{1i} > k)\mathbb{E}[Y_{0i} - Y_{1i}|Y_i > k] \geq P(Y_{1i} > k)\Delta_k^L$.

For an upper bound, we assume that $\mathbb{E}\left[\frac{dY_i(\rho)}{d\rho} \frac{\rho}{Y_i(\rho)} \middle| Y_i(\rho') = y, K_i^* = 0\right] = \mathcal{E}$ for all ρ, ρ' and y . Consider then the buncher ATE in logs:

$$\begin{aligned} \mathbb{E}[\ln Y_{0i} - \ln Y_{1i}|Y_i = k, K_i^* = 0] &= \mathbb{E}[\ln Y_{0i} - \ln Y_{1i}|Y_{0i} \in [k, Q_{0|K^*=0}(F_{1|K^*=0})], K_i^* = 0] \\ &= \int_{\rho_0}^{\rho_1} d\rho \cdot \mathbb{E}\left[\frac{dY_i(\rho)}{d\rho} \frac{1}{Y_i(\rho)} \middle| Y_{0i} \in [k, k + \Delta_0^*], K_i^* = 0\right] \\ &= \int_{\rho_0}^{\rho_1} d\ln\rho \cdot \frac{1}{\mathcal{B}^*} \int_k^{k+\Delta_0^*} dy \cdot f_0(y) \cdot \mathbb{E}\left[\frac{dY_i(\rho)}{d\rho} \frac{\rho}{Y_i(\rho)} \middle| Y_{0i} = y, K_i^* = 0\right] \\ &= \mathcal{E} \int_{\rho_0}^{\rho_1} d\ln\rho = \mathcal{E} \ln(\rho_1/\rho_0) \end{aligned}$$

with the notation that $\Delta_0^* := Q_{0|K^*=0}(F_{1|K^*=0}) - k$. Moreover:

$$\begin{aligned} \mathbb{E}[Y_{0i} - Y_{1i}|Y_i > k] &= \int_{\rho_0}^{\rho_1} d\rho \cdot \mathbb{E}\left[\frac{dY_i(\rho)}{d\rho} \middle| Y_{1i} > k, K_i^* = 0\right] \\ &= P(Y_{1i} > k)^{-1} \int_{\rho_0}^{\rho_1} d\ln\rho \cdot \int_k^\infty y \cdot f_1(y) \cdot \mathbb{E}\left[\frac{dY_i(\rho)}{d\rho} \frac{\rho}{Y_i(\rho)} \middle| Y_{1i} = y, K_i^* = 0\right] dy \\ &= \mathcal{E} \cdot \mathbb{E}[Y_{1i}|Y_{1i} > k] \int_{\rho_0}^{\rho_1} d\ln\rho = \mathcal{E} \ln(\rho_1/\rho_0) \cdot \mathbb{E}[Y_{1i}|Y_{1i} > k] \end{aligned}$$

Thus in the isoelastic model

$$E[Y_{0i} - Y_i] = (\mathcal{B} - p) E[Y_{0i} - k | Y_i = k, K_i^* = 0] + \mathbb{E}[Y_{1i} | Y_{1i} > k] \cdot P(Y_{1i} > k) \mathbb{E}[\ln Y_{0i} - \ln Y_{1i} | Y_i = k, K_i^* = 0]$$

and an upper bound is

$$\delta_k^U \cdot E[Y_i | Y_i > k] - (\mathcal{B} - p) \cdot g(1 - p - F_-, f_+, p - \mathcal{B})$$

where δ_k^U is an upper bound to the buncher ATE in logs $\mathbb{E}[\ln Y_{0i} - \ln Y_{1i} | Y_i = k, K_i^* = 0]$.

H.9.2 Moving to double time

I make use of the first step deriving the expression for $\partial_{\rho_1} E[Y_i^{[k, \rho_1]}]$ in Theorem 2, namely that:

$$\partial_{\rho_1} E[Y_i^{[k, \rho_1]}] = k \partial_{\rho_1} \mathcal{B}^{[k, \rho_1]} + \partial_{\rho_1} \{P(Y_i(\rho_1) > k) \mathbb{E}[Y_i(\rho_1) | Y_i(\rho_1) > k]\}$$

Thus:

$$\begin{aligned} E[Y_i^{[k, \rho_1]}] - E[Y_i^{[k, \bar{\rho}_1]}] &= - \int_{\rho_1}^{\bar{\rho}_1} \partial_\rho E[Y_i^{[k, \rho]}] d\rho = - \int_{\rho_1}^{\bar{\rho}_1} \left\{ k \partial_\rho \mathcal{B}^{[k, \rho]} + \partial_\rho \{P(Y_i(\rho) > k) \mathbb{E}[Y_i(\rho) | Y_i(\rho) > k]\} \right\} d\rho \\ &= -k(\mathcal{B}^{[k, \bar{\rho}_1]} - \mathcal{B}^{[k, \rho_1]}) + P(Y_i(\rho_1) > k) \mathbb{E}[Y_i(\rho_1) | Y_i(\rho_1) > k] - P(Y_i(\bar{\rho}_1) > k) \mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k] \\ &= -k(\mathcal{B}^{[k, \bar{\rho}_1]} - \mathcal{B}^{[k, \rho_1]}) + \{P(Y_i(\rho_1) > k) - P(Y_i(\bar{\rho}_1) > k)\} \cdot \mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k] \\ &\quad + P(Y_i(\rho_1) > k) (\mathbb{E}[Y_i(\rho_1) | Y_i(\rho_1) > k] - \mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k]) \\ &= (\mathbb{E}[Y_{1i} | Y_{1i} > k] - k)(\mathcal{B}^{[k, \bar{\rho}_1]} - \mathcal{B}^{[k, \rho_1]}) + P(Y_{1i} > k) (\mathbb{E}[Y_{1i} | Y_{1i} > k] - \mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k]) \\ &\leq (\mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k] - k)(\mathcal{B}^{[k, \bar{\rho}_1]} - \mathcal{B}^{[k, \rho_1]}) + P(Y_{1i} > k) \mathbb{E}[Y_i(\rho_1) - Y_i(\bar{\rho}_1) | Y_{1i} > k] \\ &\leq (\mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k] - k)(\mathcal{B}^{[k, \rho_1]} - p) + P(Y_{1i} > k) \mathbb{E}[Y_i(\rho_1) - Y_i(\bar{\rho}_1) | Y_{1i} > k] \\ &\leq (\mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k] - k)(\mathcal{B}^{[k, \rho_1]} - p) + P(Y_{1i} > k) \mathbb{E}[Y_{0i} - Y_{1i} | Y_{1i} > k] \\ &\approx (\mathbb{E}[Y_{1i} | Y_{1i} > k] - k)(\mathcal{B}^{[k, \rho_1]} - p) + P(Y_{1i} > k) \mathbb{E}[Y_{0i} - Y_{1i} | Y_{1i} > k] \\ &\leq (\mathbb{E}[Y_{1i} | Y_{1i} > k] - k)(\mathcal{B}^{[k, \rho_1]} - p) + P(Y_{1i} > k) E[Y_i | Y_i > k] \cdot \delta_k^U \end{aligned}$$

In the iso-elastic model, making use instead of the final expression for $\partial_{\rho_1} E[Y_i^{[k, \rho_1]}]$ in Thm. 2:

$$\begin{aligned} E[Y_i^{[k, \rho_1]}] - E[Y_i^{[k, \bar{\rho}_1]}] &= - \int_{\rho_1}^{\bar{\rho}_1} \partial_\rho E[Y_i^{[k, \rho_1]}] d\rho = \int_{\rho_1}^{\bar{\rho}_1} d\rho \int_k^\infty f_\rho(y) \mathbb{E}\left[\frac{dY_i(\rho)}{d\rho} \middle| Y_i(\rho) = y\right] dy \\ &= \int_{\rho_1}^{\bar{\rho}_1} d\ln \rho \int_k^\infty f_\rho(y) y \cdot \mathbb{E}\left[\frac{dY_i(\rho)}{d\rho} \frac{\rho}{Y_i(\rho)} \middle| Y_i(\rho) = y\right] dy \\ &\geq \mathcal{E} \int_{\rho_1}^{\bar{\rho}_1} d\ln \rho \int_k^\infty f_\rho(y) y \cdot dy = \mathcal{E} \int_{\rho_1}^{\bar{\rho}_1} d\ln \rho \cdot P(Y_i(\rho) > k) \mathbb{E}[Y_i(\rho) | Y_i(\rho) > k] \\ &\geq \mathcal{E} \ln(\bar{\rho}_1 / \rho_1) \cdot P(Y_i(\bar{\rho}_1) > k) \mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k] \\ &= \mathcal{E} \ln(\bar{\rho}_1 / \rho_1) \cdot \{P(Y_{1i} > k) \mathbb{E}[Y_{1i} | Y_{1i} > k] + (P(Y_i(\bar{\rho}_1) > k) \mathbb{E}[Y_i(\bar{\rho}_1) | Y_i(\bar{\rho}_1) > k] - P(Y_{1i} > k) \mathbb{E}[Y_{1i} | Y_{1i} > k])\} \\ &= \mathcal{E} \ln(\bar{\rho}_1 / \rho_1) \cdot \left\{ P(Y_{1i} > k) \mathbb{E}[Y_{1i} | Y_{1i} > k] - (E[Y_i^{[k, \rho_1]}] - E[Y_i^{[k, \bar{\rho}_1]}]) + k(\mathcal{B}^{[k, \bar{\rho}_1]} - \mathcal{B}^{[k, \rho_1]}) \right\} \end{aligned}$$

where in the fourth step I've used that $Y_i(\rho)$ is decreasing in ρ with probability one, which follows from SEPARABLE and CONVEX. So:

$$\begin{aligned} E[Y_i^{[k, \rho_1]}] - E[Y_i^{[k, \bar{\rho}_1]}] &\geq \frac{\mathcal{E} \ln(\bar{\rho}_1/\rho_1)}{1 + \mathcal{E} \ln(\bar{\rho}_1/\rho_1)} \cdot \{P(Y_{1i} > k) \mathbb{E}[Y_{1i}|Y_{1i} > k] + k(\mathcal{B}^{[k, \bar{\rho}_1]} - \mathcal{B}^{[k, \rho_1]})\} \\ &\geq \frac{\mathcal{E} \ln(\bar{\rho}_1/\rho_1)}{1 + \mathcal{E} \ln(\bar{\rho}_1/\rho_1)} \cdot P(Y_{1i} > k) \mathbb{E}[Y_{1i}|Y_{1i} > k] \end{aligned}$$

H.9.3 Effect of a change to the kink point on bunching

Using that $p(k^*) = p$ and $p(k') = 0$:

$$\begin{aligned} \mathcal{B}^{[k', \rho_1]} - \mathcal{B}^{[k^*, \rho_1]} &= (\mathcal{B}^{[k', \rho_1]} - p(k')) - (\mathcal{B}^{[k^*, \rho_1]} - p(k^*)) - p = -p + \int_{k^*}^{k'} dk \cdot \partial_k (\mathcal{B}^{[k', \rho_1]} - p(k)) \\ &= -p + \int_{k^*}^{k'} dk \cdot (f_1(k) - f_0(k)) = -p + F_1(k') - F_1(k^*) - F_0(k') + F_0(k^*) \\ &= P(k^* < Y_{1i} \leq k') - P(k^* < Y_{0i} \leq k') - p \\ &= P(k^* < Y_i \leq k') - P(k^* < Y_{0i} \leq k') - p \end{aligned}$$

if $k' > k^*$. Similarly, if $k' < k^*$:

$$\mathcal{B}^{[k', \rho_1]} - \mathcal{B}^{[k^*, \rho_1]} = P(k' \leq Y_{0i} < k^*) - P(k' \leq Y_{1i} < k^*) - p = P(k' \leq Y_i < k^*) - P(k' \leq Y_{1i} < k^*) - p$$

The lemma in the next section gives identified bounds on the counterfactual quantity that appears in the expression in each case.

H.9.4 Average effect of a change to the kink point on hours

$$\begin{aligned} E[Y_i^{[k', \rho_1]}] - E[Y_i^{[k^*, \rho_1]}] &= \int_{k^*}^{k'} \partial_k E[Y_i^{[k, \rho_1]}] dk = \int_{k^*}^{k'} \{\mathcal{B}^{[k, \rho_1]} - p(k)\} dk \\ &= k (\mathcal{B}^{[k, \rho_1]} - p(k)) \Big|_{k^*}^{k'} - \int_{k^*}^{k'} k \cdot \partial_k \{\mathcal{B}^{[k, \rho_1]} - p(k)\} dk \\ &= k' \mathcal{B}^{[k', \rho_1]} - k^* (\mathcal{B} - p) - \int_{k^*}^{k'} y (f_1(y) - f_0(y)) dy \\ &= (k' - k^*) \mathcal{B}^{[k', \rho_1]} + k^* (\mathcal{B}^{[k', \rho_1]} - \mathcal{B}) + p k^* - \int_{k^*}^{k'} y (f_1(y) - f_0(y)) dy \end{aligned}$$

For $k' > k^*$, this is equal to

$$\begin{aligned} (k' - k^*) \mathcal{B}^{[k', \rho_1]} + k^* (\mathcal{B}^{[k', \rho_1]} - (\mathcal{B} - k)) + P(k^* < Y_{0i} \leq k') (\mathbb{E}[Y_{0i}|k^* < Y_{0i} \leq k'] - P(k^* < Y_{1i} \leq k') (\mathbb{E}[Y_{1i}|k^* < Y_{1i} \leq k']) \\ = (k' - k^*) \mathcal{B}^{[k', \rho_1]} + P(k^* < Y_{0i} \leq k') (\mathbb{E}[Y_{0i}|k^* < Y_{0i} \leq k'] - k^*) - P(k^* < Y_{1i} \leq k') (\mathbb{E}[Y_{1i}|k^* < Y_{1i} \leq k'] - k^*) \\ = (k' - k^*) \mathcal{B}^{[k', \rho_1]} + P(k^* < Y_{0i} \leq k') (\mathbb{E}[Y_{0i}|k^* < Y_{0i} \leq k'] - k^*) - P(k^* < Y_i \leq k') (\mathbb{E}[Y_i|k^* < Y_i \leq k'] - k^*) \end{aligned}$$

The first term represents the mechanical effect from the bunching mass under k' being transported from k^* to k' , and can be bounded given the bounds for $\mathcal{B}^{[k', \rho_1]} - \mathcal{B}^{[k^*, \rho_1]}$ in the last

section. The last term is point identified from the data, while the middle term can be bounded using bi-log concavity of Y_{0i} conditional on $K^* = 0$. Similarly, when $k' < k^*$, the effect on hours is:

$$(k' - k^*)\mathcal{B}^{[k', \rho_1]} + P(k' \leq Y_{0i} < k^*)(k^* - \mathbb{E}[Y_{0i}|k' \leq Y_{0i} < k^*]) - P(k' \leq Y_{1i} < k^*)(k^* - \mathbb{E}[Y_{1i}|k' \leq Y_{1i} < k^*]) \\ = (k' - k^*)\mathcal{B}^{[k', \rho_1]} + P(k' \leq Y_i < k^*)(k^* - \mathbb{E}[Y_i|k' \leq Y_i < k^*]) - P(k' \leq Y_{1i} < k^*)(k^* - \mathbb{E}[Y_{1i}|k' \leq Y_{1i} < k^*])$$

with the middle term point identified from the data and last term bounded by bi-log concavity of Y_{1i} conditional on $K^* = 0$. The analytic bounds implied by BLC in each case are given by the Lemma below.

Lemma. Suppose Y_i is a bi-log concave random variable with CDF $F(y)$. Let $F_0 := F(y_0)$ and $f_0 = f(y_0)$ be the CDF and density, respectively, evaluated at a fixed y_0 . Then, for any $y' > y_0$:

$$A \leq P(y_0 \leq Y_i \leq y') (\mathbb{E}[Y_i|y_0 \leq Y_i \leq y'] - y_0) \leq B$$

and for any $y' < y_0$:

$$B \leq P(y' \leq Y_i \leq y_0) (y_0 - \mathbb{E}[Y_i|y' \leq Y_i \leq y_0]) \leq A$$

where $A = g(F_0, f_0, F_L(y'))$ and $B = g(1 - F_0, f_0, 1 - F_U(y'))$, with

$$F_L(y') = 1 - (1 - F_0)e^{-\frac{f_0}{1-F_0}(y-y_0)}, \quad F_U(y') = F_0e^{\frac{f_0}{F_0}(y'-y_0)}$$

and

$$g(a, b, c) = \begin{cases} \frac{ac}{b} \left(\ln \left(\frac{c}{a} \right) - 1 \right) + \frac{a^2}{b} & \text{if } c > 0 \\ \frac{a^2}{b} & \text{if } c \leq 0 \end{cases}$$

In either of the two cases $\max\{0, F_L(y')\} \leq F(y') \leq \min\{1, F_U(y')\}$.

Proof. As shown by Dümbgen et al., 2017, bi-log concavity of Y_i implies not only that $f(y)$ exists, but that it is strictly positive, and we may then define a quantile function $Q = F^{-1}$ such that $Q(F(y)) = y$ and $y = Q(F(y))$. Theorem 1 of Dümbgen et al., 2017 also shows that for any y' :

$$\underbrace{1 - (1 - F_0)e^{-\frac{f_0}{1-F_0}(y-y_0)}}_{:=F_L(y')} \leq F(y') \leq \underbrace{F_0e^{\frac{f_0}{F_0}(y'-y_0)}}_{:=F_U(y')}$$

We can re-express this as bounds on the quantile function evaluated at any $u' \in [0, 1]$:

$$\underbrace{y_0 + \frac{F_0}{f_0} \ln \left(\frac{u}{F_0} \right)}_{Q_L(u')} \leq Q(u') \leq \underbrace{y_0 - \frac{1 - F_0}{f_0} \ln \left(\frac{1 - u}{1 - F_0} \right)}_{Q_U(u')}$$

Write the quantity of interest as:

$$P(y_0 \leq Y_i \leq y') (\mathbb{E}[Y_i|y_0 \leq Y_i \leq y'] - y_0) = \int_{y_0}^{y'} (y - y_0) f(y) dy = \int_{F_0}^{F(y')} (Q(u) - y_0) du$$

Given that $Q(u) \geq y_0$, the integral is increasing in $F(y')$. Thus an upper bound is:

$$\begin{aligned}
P(y_0 \leq Y_i \leq y') (\mathbb{E}[Y_i|y_0 \leq Y_i \leq y'] - y_0) &\leq \int_{F_0}^{F_U(y')} (Q_U(u) - y_0) du \\
&= -\frac{1-F_0}{f_0} \int_{F_0}^{F_U(y')} \ln\left(\frac{1-u}{1-F_0}\right) du \\
&= \frac{(1-F_0)^2}{f_0} \int_1^{\frac{1-F_U(y')}{1-F_0}} \ln(v) dv \\
&= \frac{(1-F_0)(1-F_U(y'))}{f_0} \left(\ln\left(\frac{1-F_U(y')}{1-F_0}\right) - 1 \right) + \frac{(1-F_0)^2}{f_0}
\end{aligned}$$

where we've made the substitution $v = \frac{1-u}{1-F_0}$ and used that $\int \ln(v)dv = v(\ln(v) - 1)$. Inspection of the formulas for F_U and F_L reveal that $F_U \in (0, \infty)$ and $F_L \in (-\infty, 1)$. In the event that $F_U(y') \geq 1$, the above expression is undefined but we can replace $F_U(y')$ with one and still obtain valid bounds:

$$P(y_0 \leq Y_i \leq y') (\mathbb{E}[Y_i|y_0 \leq Y_i \leq y'] - y_0) \leq -\frac{(1-F_0)^2}{f_0} \int_0^1 \ln(v) dv = \frac{(1-F_0)^2}{f_0}$$

where we've used that $\int_0^1 \ln(v)dv = -1$.

Similarly, a lower bound is:

$$\begin{aligned}
P(y_0 \leq Y_i \leq y') (\mathbb{E}[Y_i|y_0 \leq Y_i \leq y'] - y_0) &\geq \int_{F_0}^{F_L(y')} (Q_L(u) - y_0) du = \frac{F_0}{f_0} \int_{F_0}^{F_L(y')} \ln\left(\frac{u}{F_0}\right) du \\
&= \frac{F_0^2}{f_0} \int_1^{F_L(y')/F_0} \ln(v) du \\
&= \frac{F_0 F_L(y')}{f_0} \left(\ln\left(\frac{F_L(y')}{F_0}\right) - 1 \right) + \frac{F_0^2}{f_0}
\end{aligned}$$

where we've made the substitution $v = \frac{u}{F_0}$. If $F_L(y') \leq 0$, then we replace with zero to obtain

$$P(y_0 \leq Y_i \leq y') (\mathbb{E}[Y_i|y_0 \leq Y_i \leq y'] - y_0) \geq -\frac{F_0^2}{f_0} \int_0^1 \ln(v) du = \frac{F_0^2}{f_0}$$

When $y' < y$, write the quantity of interest as:

$$P(y' \leq Y_i \leq y_0) (y_0 - \mathbb{E}[Y_i|y' \leq Y_i \leq y_0]) = \int_{y'}^{y_0} (y_0 - y) f(y) dy = \int_{F(y')}^{F_0} (y_0 - Q(u)) du$$

This integral is decreasing in $F(y')$, so an upper bound is:

$$\begin{aligned}
P(y' \leq Y_i \leq y_0) (y_0 - \mathbb{E}[Y_i|y' \leq Y_i \leq y_0]) &\leq \int_{F_L(y')}^{F_0} (y_0 - Q_L(u)) du = -\frac{F_0}{f_0} \int_{F_L(y')}^{F_0} \ln\left(\frac{u}{F_0}\right) du \\
&= -\frac{F_0^2}{f_0} \int_{F_L(y')/F_0}^1 \ln(v) du \\
&= \frac{F_0 F_L(y')}{f_0} \left(\ln\left(\frac{F_L(y')}{F_0}\right) - 1 \right) + \frac{F_0^2}{f_0}
\end{aligned}$$

or simply F_0^2/f_0 when $F_L(y') \leq 0$, and a lower bound is:

$$\begin{aligned}
P(y' \leq Y_i \leq y_0) (y_0 - \mathbb{E}[Y_i | y' \leq Y_i \leq y_0]) &\geq \int_{F_U(y')}^{F_0} (y_0 - Q_U(u)) du \\
&= \frac{1 - F_0}{f_0} \int_{F_U(y')}^{F_0} \ln\left(\frac{1 - u}{1 - F_0}\right) du \\
&= -\frac{(1 - F_0)^2}{f_0} \int_{\frac{1 - F_U(y')}{1 - F_0}}^1 \ln(v) dv \\
&= \frac{(1 - F_0)(1 - F_U(y'))}{f_0} \left(\ln\left(\frac{1 - F_U(y')}{1 - F_0}\right) - 1 \right) + \frac{(1 - F_0)^2}{f_0}
\end{aligned}$$

or simply $(1 - F_0)^2/f_0$ when $F_U(y') \geq 1$. \square

In estimation, I censor intermediate CDF bound estimates based on the above lemma at zero and one. These constraints are not typically binding so I ignore the effect of this on asymptotic normality of the final estimators, when constructing confidence intervals.

H.10 Details of calculating wage correction terms

For the ex-post effect of the kink

Suppose that straight-time wages w^* are set according to Equation (1) for all workers, where h^* are their anticipated hours. The straight-wages that would exist absent the FLSA w_0^* , yield the same total earnings z^* , so:

$$w_0^* h^* = w^*(h^* + (\rho_1 - 1)(h^* - k)\mathbb{1}(h^* > k))$$

where $k = 40$ and $\rho_1 = 1.5$. The percentage change is thus

$$(w_0^* - w^*)/w^* = \frac{(\rho_1 - 1)(h^* - k)\mathbb{1}(h^* > k)}{h^* + (\rho_1 - 1)(h^* - k)\mathbb{1}(h^* > k)}$$

If h_{0it} is constant elasticity in the wage with elasticity \mathcal{E} , then we would expect

$$\frac{h_{0it} - h_{0it}^*}{h_{0it}} = 1 - \left(1 + \frac{(\rho_1 - 1)(h^* - k)\mathbb{1}(h^* > k)}{h^* + (\rho_1 - 1)(h^* - k)\mathbb{1}(h^* > k)}\right)^\mathcal{E}$$

Taking $h_{0it} \approx h_{1it} \approx h^*$ and integrating along the distribution of h_{1it} , we have:

$$\mathbb{E}[h_{0it} - h_{0it}^*] \approx \mathbb{E} \left[\mathbb{1}(h_{it} > k) h_{it} \left(1 - \left(1 + \frac{(\rho_1 - 1)(h_{it} - k)}{h_{it} + (\rho_1 - 1)(h_{it} - k)}\right)^\mathcal{E} \right) \right]$$

which will be negative provided that $\mathcal{E} < 0$. The total ex-post effect of the kink is:

$$\mathbb{E}[h_{it} - h_{0it}^*] = \mathbb{E}[h_{it} - h_{0it}] + \mathbb{E}[h_{0it} - h_{0it}^*]$$

For a move to double-time

The straight-wages w_2^* that would exist with double time, for workers with $h^* > k$, that yield the same total earnings z^* as the actual straight wages w^* satisfy:

$$w_2^*(k + (\bar{\rho}_1 - 1)(h^* - k)) = w^*(k + (\rho_1 - 1)(h^* - k))$$

where $\bar{\rho}_1 = 2$. The percentage change is thus

$$(w_2^* - w^*)/w^* = \frac{k + (\rho_1 - 1)(h^* - k)}{k + (\bar{\rho}_1 - 1)(h^* - k)} - 1$$

Let \bar{h}_{0i} be hours under a straight-time wage of w_2^* . By a similar calculation thus:

$$\mathbb{E}[\bar{h}_i^{[\bar{\rho}_1, k]} - h_{it}^{[\bar{\rho}_1, k]}] \approx \mathbb{E} \left[\mathbb{1}(h_{it} > k) h_{it} \left(\left(\frac{k + (\rho_1 - 1)(h^* - k)}{k + (\bar{\rho}_1 - 1)(h^* - k)} \right)^\varepsilon - 1 \right) \right]$$

The total effect of a move to double-time is:

$$\mathbb{E}[\bar{h}_{it}^{[\bar{\rho}_1, k]} - h_{it}] = \mathbb{E}[\bar{h}_{it}^{[\bar{\rho}_1, k]} - h_{it}^{[\bar{\rho}_1, k]}] + \mathbb{E}[h_{it}^{[\bar{\rho}_1, k]} - h_{it}]$$

The above definitions are depicted visually in Figure 21 below.

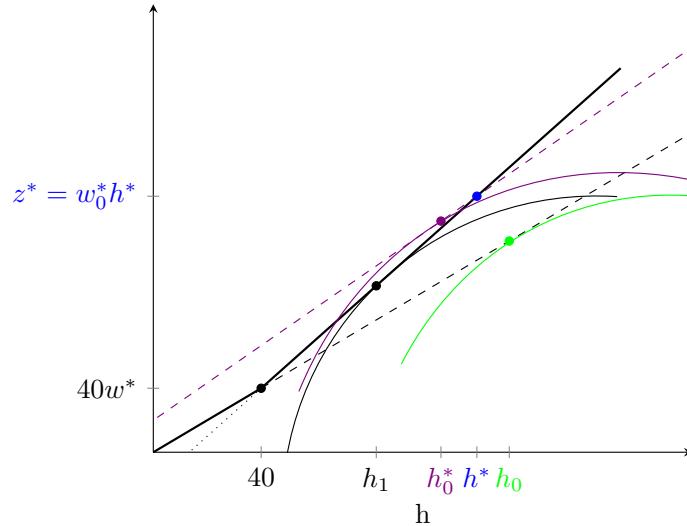


Figure 21: Depiction of h^* , h_0 , h_0^* and h_1 for a single fixed worker that works overtime at h_1 hours this week. Their realized wage w^* has been set to yield earnings z^* based on anticipated hours h^* given the FLSA kink. In a world without the FLSA, the worker's wage would instead be $w_0^* = z^*/h^*$, and this week the firm would have chosen h_0^* hours, where the worker's marginal productivity this week is w_0^* (in the benchmark model). Note: while (z^*, h^*) is chosen jointly with employment and on the basis of anticipated productivity, choice of h_0^* is instead constrained by the contracted purple pay schedule (with the worker already hired) and on the basis of updated productivity. h_1 may differ from h^* for this same reason. In the numerical calculation h^* is approximated by h_1 – which corresponds to productivity variation being small and h^* being a credible choice given the FLSA. If credibility (the firm not wanting to renege too far on hours after hiring) were a constraint on the choice of (z^*, h^*) in the no-FLSA counterfactual, then h^* would be smaller without the FLSA, but I consider this “second-order” and do not attempt a correction here.

Changing the location of the kink

Let $\mathcal{B}_w^{[k]}$ denote bunching with the kink at location k and (a distribution of) wages denoted by w . Then the effect of moving k on bunching is

$$\mathcal{B}_{w'}^{[k']} - \mathcal{B}_w^{[k^*]} = (\mathcal{B}_w^{[k']} - \mathcal{B}_w^{[k^*]}) + (\mathcal{B}_{w'}^{[k']} - \mathcal{B}_w^{[k']})$$

where w' are the wages that would occur with bunching at the new kink point k' . The first term has been estimated by the methods described above, with the second term representing a correction due to wage adjustment. Taking $Y_{0i} \approx Y_{1i} \approx h^*$, the straight-time wages w^* set according to Equation (1) that would change are those between k' and k^* . Consider the case $k' < k^*$. We expect wages to fall, as the overtime policy becomes more stringent, and $(\mathcal{B}_{w'}^{[k']} - \mathcal{B}_w^{[k']})$ is only nonzero to the extent that the increase in Y_0 and Y_1 changes the mass of each in the range $[k', k^*]$. With the range $[k', k^*]$ to the left of the mode of Y_{0i} , it is most plausible that this mass will decrease. Similarly, for Y_{1i} , it is most likely that this mass will decrease, making the overall sign of $(\mathcal{B}_{w'}^{[k']} - \mathcal{B}_w^{[k']})$ ambiguous. However, since most of the adjustment should occur for workers who are typically found between k and k' , we would not expect either term to be very different from zero.

Now consider the effect of average hours:

$$\mathbb{E}[Y_{w'}^{[k']} - Y_w^{[k^*]}] = \mathbb{E}[Y_w^{[k']} - Y_w^{[k^*]}] + \mathbb{E}[Y_{w'}^{[k']} - Y_w^{[k']}]$$

For a reduction in k , we would expect wages w' to be lower with $k = k'$ and hence the second term positive. This will attenuate the effects that are bounded by the methods above, holding the wages fixed at their realized levels.

Consider first the case of $k' < k^*$. Let w' be wages under the new kink point k' , and assuming they adjust to keep total earnings z^* constant, wages w' will change if h^* is between k and k' as: $w'(k' + 0.5(h^* - k')) = w^*h^*$, and the percentage change for these workers is thus

$$(w' - w^*)/w^* = \frac{h^*}{k' + 0.5(h^* - k')} - 1$$

$$\mathbb{E}[Y_{w'}^{[k']} - Y_w^{[k']}] \approx \mathbb{E} \left[\mathbb{1}(k' < Y_i < k^*) Y_i \left(\left(\frac{Y_i}{k' + 0.5(Y_i - k')} \right)^\varepsilon - 1 \right) \right]$$

In the case of $k' > k^*$, we will have wages change as: $w'h^* = w^*(k^* + 0.5(h^* - k^*))$ if h^* is between k and k' . The percentage change for these workers is thus

$$(w' - w^*)/w^* = \frac{k^* + 0.5(h^* - k^*)}{h^*} - 1$$

$$\mathbb{E}[Y_{w'}^{[k']} - Y_w^{[k']}] \approx \mathbb{E} \left[\mathbb{1}(k^* < Y_i < k') Y_i \left(\left(\frac{k^* + 0.5(Y_i - k^*)}{Y_i} \right)^\varepsilon - 1 \right) \right]$$

References

- BEST, M. C., BROCKMEYER, A., KLEVEN, H. J., SPINNEWIJN, J. and WASEEM, M. (2015). "Production vs Revenue Efficiency With Limited Tax Capacity: Theory and Evidence From Pakistan". *Journal of Political Economy* 123 (6), p. 48.
- BICK, A., BLANDIN, A. and ROGERSON, R. (2022). "Hours and Wages". *The Quarterly Journal of Economics* (forthcoming).
- BLOMQUIST, S., KUMAR, A., LIANG, C.-Y. and NEWHEY, W. (2021). "On Bunching and Identification of the Taxable Income Elasticity". *Journal of Political Economy* 129 (8).
- BLOMQUIST, S., KUMAR, A., LIANG, C.-Y. and NEWHEY, W. K. (2015). "Individual heterogeneity, non-linear budget sets and taxable income". *The Institute for Fiscal Studies Working Paper* CWP21/15.
- BLOMQUIST, S. and NEWHEY, W. (2017). "The Bunching Estimator Cannot Identify the Taxable Income Elasticity". *The Institute for Fiscal Studies Working Paper* CWP40/17.
- BORUSYAK, K., JARAVEL, X. and SPIESS, J. (2021). "Revisiting Event Study Designs: Robust and Efficient Estimation". *arXiv:2108.12419 [econ]*.
- BURDETT, K. and MORTENSEN, D. T. (1998). "Wage Differentials, Employer Size, and Unemployment". *International Economic Review* 39 (2), p. 257.
- CAHUC, P. and ZYLBERBERG, A. (2014). *Labor economics*. 2nd. Cambridge, Mass.: MIT Press.
- CESARINI, D., LINDQVIST, E., NOTOWIDIGDO, M. J. and ÖSTLING, R. (2017). "The Effect of Wealth on Individual and Household Labor Supply: Evidence from Swedish Lotteries". *American Economic Review* 107 (12), pp. 3917–46.
- CHETTY, R., FRIEDMAN, J. N., OLSEN, T. and PISTAFERRI, L. (2011). "Adjustment Costs, Firm Responses, and Micro vs. Macro Labor Supply Elasticities: Evidence from Danish Tax Records." *Quarterly Journal of Economics* 126 (2), pp. 749–804.
- DIAMOND, R. and PERSSON, P. (2016). *The Long-term Consequences of Teacher Discretion in Grading of High-stakes Tests*. w22207. Cambridge, MA: National Bureau of Economic Research, w22207.
- DÜMBGEN, L., KOLESNYK, P. and WILKE, R. A. (2017). "Bi-log-concave distribution functions". *Journal of Statistical Planning and Inference* 184, pp. 1–17.
- EHRENBERG, R. and SCHUMANN, P. (1982). *Longer hours or more jobs? : an investigation of amending hours legislation to create employment*. New York State School of Industrial and Labor Relations, Cornell University.
- EINAV, L., FINKELSTEIN, A. and SCHRIMPFF, P. (2017). "Bunching at the kink: Implications for spending responses to health insurance contracts". *Journal of Public Economics* 146, pp. 27–40.

- GARDNER, J. (2021). "Two-stage differences in differences". *Working Paper*, p. 34.
- HAMERMESH, D. S. (1993). *Labor demand*. Princeton, NJ: Princeton Univ. Press.
- HWANG, H., MORTENSEN, D. T. and REED, W. R. (1998). "Hedonic Wages and Labor Market Search". *Journal of Labor Economics* 16 (4), pp. 815–847.
- IMBENS, G. W., RUBIN, D. B. and SACERDOTE, B. I. (2001). "Estimating the Effect of Unearned Income on Labor Earnings, Savings, and Consumption: Evidence from a Survey of Lottery Players". *The American Economic Review* 91 (4), p. 25.
- KASY, M. (2022). "Who wins, who loses? Identification of the welfare impact of changing wages". *Journal of Econometrics* 226 (1), pp. 1–26.
- KIMBALL, W. and MISHEL, L. (2015). "Estimating the Number of Workers Directly Benefiting from the Proposed Increase in the Overtime Salary Threshold". *Economic Policy Institute Technical Paper*, p. 12.
- KLEVEN, H. J. and WASEEM, M. (2013). "Using Notches to Uncover Optimization Frictions and Structural Elasticities: Theory and Evidence from Pakistan*". *The Quarterly Journal of Economics* 128 (2), pp. 669–723.
- KLEVEN, H. J. (2016). "Bunching". *Annual Review of Economics* 8, pp. 435–464.
- LEWBEL, A. and LINTON, O. (2002). "Nonparametric Censored and Truncated Regression". *Econometrica* 70 (2), pp. 765–779.
- LEWIS, H. G. (1969). "Employer Interest in Employee Hours of Work". *Unpublished paper*.
- MANNING, A. (2003). *Monopsony in Motion: Imperfect Competition in Labor Markets*. Princeton: Princeton University Press.
- MILGROM, P. and ROBERTS, J. (1996). "The LeChatelier Principle". *American Economic Review* 1 (86), pp. 173–179.
- SAEZ, E. (2010). "Do Taxpayers Bunch at Kink Points?" *American Economic Journal: Economic Policy* 2 (3), pp. 180–212.
- SATO, R. (1975). "The Most General Class of CES Functions". *Econometrica* 43 (5), p. 999.
- TREJO, B. S. J. (1991). "The Effects of Overtime Pay Regulation on Worker Compensation". *American Economic Review* 81 (4), pp. 719–740.

Chapter 2: A Vector Monotonicity Assumption for Multiple Instruments

A Vector Monotonicity Assumption for Multiple Instruments

Leonard Goff^{a,b*}

^aDepartment of Economics, University of Calgary
2500 University Dr. N.W. Calgary, AB T2N1N4, Canada

^bDepartment of Economics, University of Georgia
620 S Lumpkin St, Athens, GA 30602, USA

This version: August 31, 2022

Abstract

When a researcher combines multiple instrumental variables for a single binary treatment, the monotonicity assumption of the local average treatment effects (LATE) framework can become restrictive: it requires that all units share a common direction of response even when separate instruments are shifted in opposing directions. What I call *vector monotonicity*, by contrast, simply assumes treatment uptake to be monotonic in all instruments, capturing the intuitive notion of “no defiers” for each. I characterize the class of causal parameters that are point identified under vector monotonicity, when the instruments are binary. This class includes, for example, the average treatment effect among units that are in any way responsive to the collection of instruments, or those that are responsive to any given subset of them. I propose a simple estimator for the identified treatment effect parameters that differs from the popular two-stage least squares estimator. Empirical applications revisit the labor market returns to college, and the labor supply effects of family size.

*Email: leonard.goff@ucalgary.ca. I am grateful to Josh Angrist, Simon Lee, Suresh Naidu and Bernard Salanié for insightful feedback throughout this project, which was a chapter of my PhD dissertation. I also thank Isaiah Andrews, Jushan Bai, Junlong Feng, Peter Hull, Jack Light, José Luis Montiel Olea, Serena Ng, Vitor Possebom and Alex Torgovitsky for helpful comments and discussion. I thank attendees of the Columbia econometrics colloquium, the 2019 Young Economists Symposium, the 2019 Empirics and Methods in Economics Conference, the 2022 Georgia Econometrics Workshop, and the 2022 North American Meeting of the Econometric Society for their feedback. I thank the Bureau of Labor Statistics for data access. Any errors are my own.

1 Introduction

The local average treatment effects (LATE) framework of Imbens and Angrist (1994) allows for causal inference with arbitrary heterogeneity in treatment effects, but in doing so makes an important assumption about homogeneity in selection behavior. This takes the form of the LATE *monotonicity* assumption, which is often quite natural to make when the researcher has a single instrumental variable at their disposal. When multiple instruments are combined however, this monotonicity assumption can become hard to justify—a point that has recently been emphasized by Mogstad et al. (2021).

This paper considers a natural alternative assumption, which is that monotonicity holds on an instrument-by-instrument basis: what I call *vector monotonicity*. Vector monotonicity (VM) captures the notion that each instrument has an impact on treatment uptake in a direction that is common across units. This direction need not be known ex-ante by the researcher, but is often implied by economic theory. For example, two instruments for college enrollment might be: i) proximity to a college; and ii) affordability of nearby colleges. Vector monotonicity assumes that increasing either instrument induces some individuals towards going to college, while discouraging none. This contrasts with traditional LATE monotonicity, which requires either proximity or affordability to dominate in the selection behavior of all units.

To appreciate this, let us code each of the above instruments for college as a binary variable (“far” / “close” and “cheap” / “expensive”). The LATE monotonicity assumption says that a counterfactual change to the proximity and/or tuition instruments can either move some students into college attendance, or some students out, but not both. In particular, this requires that all units who would go to college when it is far but cheap would also go to college if it was close and expensive, or that the reverse is true. This implication will generally fail to hold if individuals differ in how much each of the instruments “matters” to them: for example, if some students are primarily sensitive to distance while others are primarily sensitive to tuition. Vector monotonicity instead says something quite natural in this context: proximity to a college weakly encourages college attendance, regardless of price, and lower tuition weakly encourages college attendance, regardless of distance.

I provide a simple approach to estimating causal effects under vector monotonicity. I show that in a setting with a binary treatment and any number of binary instruments satisfying VM, average treatment effects can be identified for subgroups of the population if and only if

that subgroup satisfies a certain condition.¹ The condition is met by, for example, the group of all units that move into treatment when any fixed subset of the instruments are switched “on”. As special cases, this includes e.g. those units that respond to a movement of a single particular instrument, or those units that have any variation whatsoever in counterfactual treatment status given the available instruments. I then propose a two-step estimator for the identified causal parameters. The estimator is equally scalable in terms of its computational burden as is the popular two-stage least squares estimator (as the number of instruments increases), despite the rapid proliferation of potential selection patterns compatible with VM. The estimator is implemented in a companion Stata package **ivcombine**.²

In a closely related paper, Mogstad, Torgovitsky and Walters (2021) (henceforth MTW) discuss the above difficulties for LATE monotonicity with multiple instruments, and introduce an alternative assumption they call *partial monotonicity* (PM). PM is similar to VM but is weaker: PM allows the direction of “compliance” for each instrument to depend on the values of the other instruments. In a second paper, Mogstad, Torgovitsky and Walters (2022) extend the marginal treatment effects framework (see e.g. Heckman and Vytlacil 2005) for use under partial monotonicity. This orients attention to a broad class of causal parameters that are often only partially identified by IV methods, typically requiring continuous instruments and/or parametric assumptions for point identification. By contrast, I maintain the stronger monotonicity assumption VM and characterize the class of causal parameters that are then point identified even with discrete instruments and without any auxiliary assumptions. I show that VM differs from PM only by the addition of a testable condition, that this restriction is quite naturally satisfied, and that VM adds useful identifying power beyond that of PM alone.

The estimator I propose in this paper can be seen as an alternative to two-stage least squares (2SLS), which has been the most popular method to make use of multiple instruments in applied work. 2SLS is known to identify a convex combination of local average treatment effects under the standard LATE assumptions (if the first stage is correctly specified to recover the true propensity score function), but this implication does not hold generally under VM or PM: the weights on various complier groups can be negative. MTW derive additional testable conditions which are sufficient for the 2SLS estimand to deliver positive weights under PM, but

¹I show how general discrete instruments can be accommodated by re-expressing them as a larger number of binary instruments, while preserving vector monotonicity. Section 3 also gives some point identification results under VM that make no assumptions about the support of the instruments.

²The package is available from <https://github.com/leonardgoff/ivcombine>.

the number of conditions to be verified grows combinatorially with the number of instruments. Targeting a particular treatment effect parameter that is identified under VM avoids the need for such tests altogether.

This paper thus contributes by providing a straightforward means to combine multiple instruments to estimate useful treatment effect parameters, requiring only vector monotonicity and the standard assumptions of instrument validity. While a growing body of work has considered alternatives to the basic LATE model of Imbens and Angrist (1994),³ this literature has not emphasized the distinction between separate instruments, when more than one is available. Using multiple instruments together has the benefit of enlarging the population of “compliers” for which treatment effects can be estimated, improving external validity and potentially aiding with issues of weak identification.

In Section 2 I discuss the basic instrumental variables (IV) setup and definitions, comparing vector monotonicity to the traditional monotonicity assumption and MTW’s proposal of partial monotonicity. Section 3 provides an introduction to identification under VM, giving leading examples of treatment effect parameters that are point identified. In Section 4, I show that VM partitions the population into well-defined “response groups” in a setting with any number of binary instruments (nesting results from MTW for the two-instrument case). I then use this taxonomy of response groups in Section 5 to fully characterize the family of identified parameters, which leads to the estimator proposed in Section 6. Section 7 reports results from two empirical applications.

2 Setup

To begin, let us consider a basic setup in which a researcher has multiple instrumental variables for a single binary treatment. Within this framework, I contrast vector monotonicity with traditional LATE monotonicity and the partial monotonicity assumption of Mogstad, Torgovitsky and Walters (2021).

Consider a setting with a binary treatment variable D , scalar outcome variable Y , and vector

³In the case of a binary treatment, Gautier and Hoderlein (2011), Lewbel and Yang (2016) and Gautier (2020) consider explicit selection models, while Chaisemartin (2017) shows that a weaker notion than monotonicity can be sufficient to give a causal interpretation to LATE estimands. LATE monotonicity is also generally not assumed by nonseparable triangular models with endogeneity (e.g. Imbens and Newey 2009, Torgovitsky 2015, D’Haultfœuille and Février 2015, Gunsilius 2020, Feng 2020), which typically impose some version of monotonicity in unobserved heterogeneity instead of with respect to instrument values. Lee and Salanié (2018) relax monotonicity in a setting with multivalued treatment and continuous instruments, generalizing results from the local instrumental variables approach (see e.g. Heckman and Vytlacil 2005).

$Z = (Z_1, Z_2, \dots, Z_J)'$ of J instrumental variables that can take values in set $\mathcal{Z} \subseteq (\mathcal{Z}_1 \times \mathcal{Z}_2 \times \dots \times \mathcal{Z}_J)$, where we let \mathcal{Z}_j denote the set of values that instrument Z_j can take.⁴ A typical value in \mathcal{Z} will be denoted with the boldface notation \mathbf{z} , with z_j denoting the component corresponding to the j^{th} instrument.

Definition 1 (potential outcomes and treatments). Let $D_i(\mathbf{z})$ denote the counterfactual treatment status of unit i when the vector of instruments takes value \mathbf{z} , and $Y_i(d, \mathbf{z})$ the outcome that would occur with treatment status $d \in \{0, 1\}$ and value \mathbf{z} for the instruments.

The following assumption states that the available instrumental variables are *valid*:

Assumption 1 (exclusion and independence). i) $Y_i(d, \mathbf{z}) = Y_i(d)$ for all $\mathbf{z} \in \mathcal{Z}, d \in \{0, 1\}$; and ii)

$$(Y_i(1), Y_i(0), \{D_i(\mathbf{z})\}_{\mathbf{z} \in \mathcal{Z}}) \perp (Z_{1i}, \dots, Z_{Ji})$$

The first part of Assumption 1 states that the instruments satisfy the exclusion restriction that potential outcomes do not depend on them once treatment status is fixed. The second part of Assumption 1 states that the instruments are statistically independent of potential outcomes and potential treatments.⁵ In practice, it is common to maintain a version of this independence assumption that holds only conditional on a set of observed covariates. For ease of exposition, I implicitly condition on any such covariates and then discuss incorporating them into estimation in Section 6.

2.1 Notions of monotonicity

It is well-known that when treatment effects are heterogeneous, Assumption 1 alone is not sufficient for instrument variation to identify treatment effects. The seminal LATE model of Imbens and Angrist (1994) introduces the additional assumption of monotonicity:

Assumption IAM (traditional LATE monotonicity). For all $\mathbf{z}, \mathbf{z}' \in \mathcal{Z}$: $D_i(\mathbf{z}) \geq D_i(\mathbf{z}')$ for all i or $D_i(\mathbf{z}) \leq D_i(\mathbf{z}')$ for all i .

⁴ \mathcal{Z} may be a strict subset of $(\mathcal{Z}_1 \times \mathcal{Z}_2 \times \dots \times \mathcal{Z}_J)$ when certain combinations of instrument values are ruled out on conceptual grounds, e.g. Z_1 indicates a mothers' first two births being girls and Z_2 indicates them both being boys.

⁵It's worth noting that whether or not to use multiple instruments may not be "optional", in the sense that if a collection of instruments are valid, this does not imply that a subset of the instruments is also valid. Consider two instruments that both affect treatment uptake $D_i(z_1, z_2)$. If instrument 1 were used alone, the counterfactual selection function would become $\tilde{D}_i(z_1) := D_i(z_1, Z_{2i})$, which may be correlated with Z_{1i} if the instruments are not independent of one another, even if $D_i(z_1, z_2) \perp (Z_{1i}, Z_{2i})$.

I follow the terminology of MTW and refer to this as Assumption IAM (for “Imbens and Angrist monotonicity”) going forward, rather than “LATE monotonicity” as above.

My proposed assumption of *vector monotonicity* captures monotonicity as the notion that “increasing” the value of any one instrument weakly encourages (or discourages) all units to take treatment, regardless of the values of the other instruments:⁶

Assumption 2 (vector monotonicity). *There exists an ordering \geq_j on \mathcal{Z}_j for each $j \in \{1 \dots J\}$ such that for all $\mathbf{z}, \mathbf{z}' \in \mathcal{Z}$, if $\mathbf{z} \geq \mathbf{z}'$ component-wise according to the orderings $\{\geq_j\}_{j=1}^J$, then $D_i(\mathbf{z}) \geq D_i(\mathbf{z}')$ for all i .*

The assumption of partial monotonicity introduced by MTW is weaker than both IAM and VM. Let (z_j, \mathbf{z}_{-j}) denote a given value in \mathcal{Z} as the combination of a value $z_j \in \mathcal{Z}_j$ for the j^{th} instrument and value $\mathbf{z}_{-j} \in \mathcal{Z}_{-j}$ for the other instruments, where \mathcal{Z}_{-j} denotes the set of possible values for all instruments aside from Z_j .

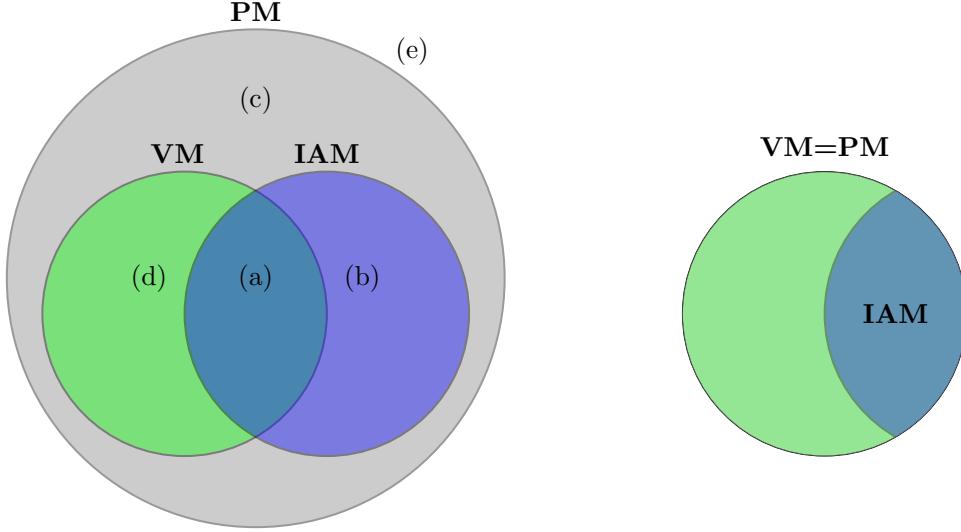
Assumption PM (partial monotonicity). *Consider any $j \in \{1 \dots J\}$, $z_j, z'_j \in \mathcal{Z}_j$, and $\mathbf{z}_{-j} \in \mathcal{Z}_{-j}$ such that $(z_j, \mathbf{z}_{-j}) \in \mathcal{Z}$ and $(z'_j, \mathbf{z}_{-j}) \in \mathcal{Z}$. Then either $D_i(z_j, \mathbf{z}_{-j}) \geq D_i(z'_j, \mathbf{z}_{-j})$ for all i or $D_i(z_j, \mathbf{z}_{-j}) \leq D_i(z'_j, \mathbf{z}_{-j})$ for all i .*

Under PM, there exists an ordering on the points $z \in \mathcal{Z}_j$ such that $D_i(z, \mathbf{z}_{-j})$ is weakly increasing along that order, for any fixed choice of j and \mathbf{z}_{-j} . The crucial restriction made by VM beyond PM is that under VM, this ordering must be *the same* across all values of $\mathbf{z}_{-j} \in \mathcal{Z}_{-j}$ for a given j . For example, close proximity to a college encourages going to college, whether or not nearby colleges are cheap. By contrast, PM could capture a situation in which college proximity encourages attendance when nearby colleges are cheap, but discourages attendance when they are expensive.

While VM is technically stronger than PM, the additional restriction made by VM over PM is empirically testable, by inspecting the propensity score function. Call \mathcal{Z} *connected* when for any two $\mathbf{z}, \mathbf{z}' \in \mathcal{Z}$ there exists a sequence of vectors $\mathbf{z}_1, \dots, \mathbf{z}_M \in \mathcal{Z}$ with $\mathbf{z}_1 = \mathbf{z}$, $\mathbf{z}_M = \mathbf{z}'$ and each \mathbf{z}_m and \mathbf{z}_{m-1} differing on only one component.⁷

⁶VM is referred to as “actual monotonicity” by Mogstad et al. (2021), when each \geq_j is the standard ordering on real numbers. Mountjoy (2022) imposes a version of VM in a setting with a multivalued treatment and continuous instruments. An alternative but equivalent characterization of VM is given in Appendix E, which is useful in some proofs.

⁷This rules out cases where \mathcal{Z} is disjoint with respect to such chains of single-instrument switches, for example in a case of two binary instruments if \mathcal{Z} consists only of the points $(0, 0)$ and $(1, 1)$.



Without restriction on the propensity score When propensity score is monotonic

Figure 1: Left panel shows an ex-ante comparison of Imbens & Angrist monotonicity (IAM), vector monotonicity (VM), and partial monotonicity (PM), before the propensity score function is known. Right panel depicts the relationship when the propensity score is component-wise monotonic: PM and VM become identical, with IAM a special case. Examples from a linear latent-index model of cases (a)-(e) are discussed in Appendix D.

Proposition 1. *Suppose PM and Assumption 1 hold, and \mathcal{Z} is connected. Then VM holds if and only if $\mathbb{E}[D_i|Z_i = \mathbf{z}]$ is component-wise monotonic in \mathbf{z} , for some fixed ordering \succeq_j on each \mathcal{Z}_j .*

Proof. See Appendix E. □

By contrast, PM (like IAM) is compatible with any propensity score function. Note that if Assumption 1 holds conditional on covariates X_i , Proposition 1 also need only hold with respect to the *conditional* propensity score $\mathbb{E}[D_i|Z_i = \mathbf{z}, X_i = x]$ (see Section 7).

Since IAM implies PM, it also follows from Proposition 1 that if IAM and Assumption 1 hold and $\mathbb{E}[D_i|Z_i = \mathbf{z}]$ is component-wise monotonic in \mathbf{z} , then VM holds. Thus if a researcher has verified that the propensity score function is monotonic, VM becomes a strictly weaker assumption than IAM. The overall relationship between Assumptions IAM, VM and PM is depicted graphically in Figure 1.

Finally, I note that a sufficient condition for the restriction from PM to VM is the existence of groups that are responsive to each instrument alone. For example, suppose Alice only cares about proximity, and Bob only cares about tuition, with selection functions:

$$D_{alice}(z_1, z_2) = \mathbb{1}(z_2 = close) \quad \text{and} \quad D_{bob}(z_1, z_2) = \mathbb{1}(z_1 = cheap)$$

If Alice and Bob are both present in the population, PM then requires that the directions of response that Alice and Bob exhibit for the two instruments hold (weakly) for all other units in the population, implying VM.⁸ The existence of both Alice and Bob also violates IAM, since the points (expensive, close) and (cheap, far) cannot be ordered for all i .

3 Identification under vector monotonicity: motivation

This section begins our analysis of identification under VM by providing some simple but general point identification results for treatment effects, which do not require any assumptions about the support of the instruments. Sections 4 and 5 then refine and extend these results in the case of binary instruments, by characterizing explicitly the response groups that can exist in the population under VM, and fully delimiting which kinds of conditional treatment effects (and potential outcome means) can be identified.

Let us say that $\mathbf{z}' \geq_D \mathbf{z}$ if $D_i(\mathbf{z}') \geq D_i(\mathbf{z})$ for all i . Under IAM, the relation \geq_D constitutes a total order on \mathcal{Z} , because for any $\mathbf{z}, \mathbf{z}' \in \mathcal{Z}$ either $\mathbf{z}' \geq_D \mathbf{z}$ or $\mathbf{z} \geq_D \mathbf{z}'$. In the returns-to-schooling example, a possible such order might be the following, where we depict that $\mathbf{z}' \geq_D \mathbf{z}$ with an arrow from \mathbf{z} to \mathbf{z}' :

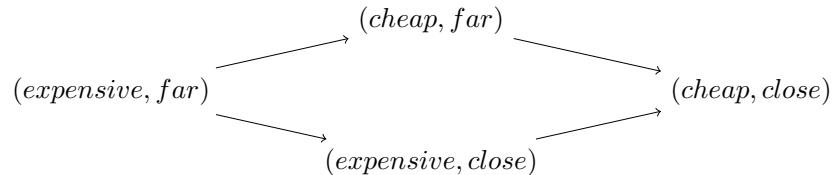
$$\begin{array}{ccccccc} \bullet & \xrightarrow{\hspace{2cm}} & \bullet & \xrightarrow{\hspace{2cm}} & \bullet & \xrightarrow{\hspace{2cm}} & \bullet \\ (\text{expensive, far}) & & (\text{cheap, far}) & & (\text{expensive, close}) & & (\text{cheap, close}) \end{array}$$

However, another such ordering (which could be equally plausible ex ante) is:

$$\begin{array}{ccccccc} \bullet & \xrightarrow{\hspace{2cm}} & \bullet & \xrightarrow{\hspace{2cm}} & \bullet & \xrightarrow{\hspace{2cm}} & \bullet \\ (\text{expensive, far}) & & (\text{expensive, close}) & & (\text{cheap, far}) & & (\text{cheap, close}) \end{array}$$

In either case, the two points (expensive, close) and (cheap, far) must be ranked by \geq_D .

By contrast, PM and VM both ascribe *partial* orders on \mathcal{Z} in which only some pairs z, z' are ranked. In particular, $\mathbf{z}' \geq_D \mathbf{z}$ if $z'_j \geq_j z_j$ for all $j = 1 \dots J$. In the returns to schooling example, the natural such partial order can be represented as:



⁸In particular, $D_{alice}(1, z_2) > D_{alice}(0, z_2)$ for all $z_2 \in \mathcal{Z}_2$ implies through PM that $D_i(1, z_2) \geq D_i(0, z_2)$ for all i and $z_2 \in \mathcal{Z}_2$, and similarly Bob implies that $D_i(z_1, 1) \geq D_i(z_1, 0)$ for all i and $z_1 \in \mathcal{Z}_1$.

Let \mathbf{l} and \mathbf{u} be any two points in \mathcal{Z} such that $\mathbf{u} \geq_D \mathbf{l}$. Without loss of generality, take each instrument Z_j to be numeric with values increasing along \geq_j , so that we may simply refer to \mathbf{u} and \mathbf{l} as instrument values such that $\mathbf{u} \geq \mathbf{l}$ component-wise. It is straightforward to show that the standard identification argument of Imbens and Angrist (1994) (made there in the case of IAM) applies here to the pair \mathbf{l}, \mathbf{u} , i.e.

$$\mathbb{E}[Y_i(1) - Y_i(0)|D_i(\mathbf{u}) > D_i(\mathbf{l})] = \frac{\mathbb{E}[Y_i|Z_i = \mathbf{u}] - \mathbb{E}[Y_i|Z_i = \mathbf{l}]}{\mathbb{E}[D_i|Z_i = \mathbf{u}] - \mathbb{E}[D_i|Z_i = \mathbf{l}]} \quad (1)$$

By Equation (1), the local average treatment effect among the group of individuals who take treatment at \mathbf{u} but not at \mathbf{l} is identified from the observables (Y_i, D_i, Z_i) .⁹

When there are multiple instruments, a natural generalization of the above result is to allow \mathbf{u} and \mathbf{l} to only change a subset of the instruments, leaving other instruments at their realized values. To this end, let \mathbf{u} allow $\mathbf{u}(\mathbf{z})$ and $\mathbf{l}(\mathbf{z})$ to be known *functions* of Z_i , with the property that $\mathbf{u}(\mathbf{z}) \geq \mathbf{l}(\mathbf{z})$ component-wise for all values of \mathbf{z} . For example, let $\mathbf{u}(Z_i)$ be the function that replaces the first component of Z_i with a 1, and let $\mathbf{l}(Z_i)$ be the function that replaces the first component of Z_i with a 0, so that:

$$\mathbf{u}(Z_i) = (1, Z_{2i}, \dots, Z_{ji})' \quad \mathbf{l}(Z_i) = (0, Z_{2i}, \dots, Z_{ji})'$$

An extension of result (1) establishes that we can identify the average treatment effect among units who shift into treatment when the first instrument is changed from 0 to 1. For general functions $\mathbf{u}(\cdot)$ and $\mathbf{l}(\cdot)$, define $\mathcal{C}^{u,l} := \{i : D_i(\mathbf{u}(Z_i)) > D_i(\mathbf{l}(Z_i))\}$. Then:

$$\mathbb{E}[Y_i(1) - Y_i(0)|i \in \mathcal{C}^{u,l}] = \frac{\int dF(\mathbf{z}) \cdot \{\mathbb{E}[Y_i|Z_i = \mathbf{u}(\mathbf{z})] - \mathbb{E}[Y_i|Z_i = \mathbf{l}(\mathbf{z})]\}}{\int dF(\mathbf{z}) \cdot \{\mathbb{E}[D_i|Z_i = \mathbf{u}(\mathbf{z})] - \mathbb{E}[D_i|Z_i = \mathbf{l}(\mathbf{z})]\}}, \quad (2a)$$

where the denominator yields the size of the group $\mathcal{C}^{u,l}$:

$$P(i \in \mathcal{C}^{u,l}) = \int dF(\mathbf{z}) \cdot \{\mathbb{E}[D_i|Z_i = \mathbf{u}(\mathbf{z})] - \mathbb{E}[D_i|Z_i = \mathbf{l}(\mathbf{z})]\}, \quad (2b)$$

⁹Equation (1) can be seen by noting that since for all i , $D_i(\mathbf{u}) - D_i(\mathbf{l})$ is either equal to 0 or 1, we have that

$$\mathbb{E}[\{Y_i(1) - Y_i(0)\} \{D_i(\mathbf{u}) - D_i(\mathbf{l})\}] = P(D_i(\mathbf{u}) > D_i(\mathbf{l})) \cdot \mathbb{E}[Y_i(1) - Y_i(0)|D_i(\mathbf{u}) > D_i(\mathbf{l})]$$

while by independence condition of Assumption 1 $P(D_i(\mathbf{u}) > D_i(\mathbf{l})) = \mathbb{E}[D_i(\mathbf{u})] - \mathbb{E}[D_i(\mathbf{l})] = \mathbb{E}[D_i|Z_i = \mathbf{u}] - \mathbb{E}[D_i|Z_i = \mathbf{l}]$. Finally, independence also shows that $\mathbb{E}[Y_i|Z_i = \mathbf{u}] - \mathbb{E}[Y_i|Z_i = \mathbf{l}] = \mathbb{E}[\{Y_i(1) - Y_i(0)\} \{D_i(\mathbf{u}) - D_i(\mathbf{l})\}]$. See also Sun and Wüthrich (2022), who introduce a general notion of “pairwise validity” for particular pairs of two instrument values like \mathbf{l}, \mathbf{u} .

and $F(\mathbf{z})$ is the CDF of the instruments.¹⁰ If \mathcal{Z} is discrete, then each $\int dF(\mathbf{z})$ above becomes a sum over $\sum_{\mathbf{z} \in \mathcal{Z}} P(Z_i = \mathbf{z})$.

Equations (2a) and (2b) are quite general, but they do not provide much intuition about the groups of individuals i in $\mathcal{C}^{u,l}$ that “respond” to a shift from \mathbf{l} to \mathbf{u} . Nor do they tell us whether there are additional treatment effect parameters that might be identified under VM. The remainder of this paper specializes to the case in which the instruments are binary, and in that setting gives an explicit characterization of the set $\mathcal{C}^{u,l}$ in terms of latent groups in the population. This leads to a general identification result (Theorem 1), which also provides a more natural guide to estimation than Eq. (2). Theorem 2 then establishes that the class of conditional treatment effect parameters covered by this result are the *only* ones that can be point identified absent further assumptions.

3.1 Some examples from the family of identified parameters

This section highlights some of the identified parameters under VM (which can be seen via Eq. (2)), and discusses their interpretation in the returns to schooling setting mentioned in the introduction. For simplicity, suppose that all instruments are binary, and consider the following leading examples of conditional average treatment effects of the form $\mathbb{E}[Y_i(1) - Y_i(0)|i \in \mathcal{C}^{u,l}]$ in Table 1:

Parameter name	$\mathcal{C}^{u,l}$
<i>ACL</i>	$D_i(1, 1 \dots 1) > D_i(0, 0 \dots 0)$
<i>SLATE_J</i>	$D_i((1 \dots 1), \mathbf{z}_{-\mathcal{J}}) > D_i((0 \dots 0), \mathbf{z}_{-\mathcal{J}}))$
<i>SLATT_J</i>	$D_i(\mathbf{z}) > D_i((0 \dots 0), \mathbf{z}_{-\mathcal{J}})$
<i>SLATU_J</i>	$D_i((1 \dots 1), \mathbf{z}_{-\mathcal{J}}) > D_i(\mathbf{z})$
<i>PTE_j(\mathbf{z}_{-j})</i>	$D_i(1, \mathbf{z}_{-j}) > D_i(0, \mathbf{z}_{-j})$

Table 1: Examples of identified treatment effect parameters under VM (see text for details).

I call the first item in Table 1 the “all-compliers LATE” (ACL). The ACL is the average treatment effect among all units who change their treatment uptake in any way in response to the instruments, and is the largest subgroup of the population for which treatment effects can be generally point identified from instrument variation. In the returns to schooling example, the ACL can equivalently be described as the average treatment effect among individuals who would go to college were it close and cheap, but would not were it far and expensive.

¹⁰To see this, note that $\mathbb{E}[Y_i(1) - Y_i(0)|i \in \mathcal{C}^{u,l}] = \int dF(\mathbf{z}|i \in \mathcal{C}^{u,l}) \cdot \mathbb{E}[Y_i(1) - Y_i(0)|D_i(\mathbf{u}(\mathbf{z})) > D_i(\mathbf{l}(\mathbf{z}))]$, while $dF(\mathbf{z}|i \in \mathcal{C}^{u,l}) = dF(\mathbf{z})/P(i \in \mathcal{C}^{u,l}) \cdot P(i \in \mathcal{C}^{u,l}|Z_i = \mathbf{z})$, and $P(i \in \mathcal{C}^{u,l}|Z_i = \mathbf{z}) = \mathbb{E}[D_i|Z_i = \mathbf{u}(\mathbf{z})] - \mathbb{E}[D_i|Z_i = \mathbf{l}(\mathbf{z})]$.

On the other end of the spectrum, the final row of Table 1 gives the most disaggregated type of identified parameter, what might be called a *partial treatment effect* $PTE_j(\mathbf{z}_{-j})$. This is the average treatment effect among individuals that move into treatment when a single instrument j is shifted from zero to one, while the other instrument values are held fixed at some explicit vector of values \mathbf{z}_{-j} . An example is the average treatment effect among individuals who go to college if it is close and cheap, but do not if it is far and cheap.

The remaining parameters in Table 1 constitute a middle ground between the granular PTE's and the very broad averaging of the ACL. For example, a *set local average treatment effect*, or $SLATE_{\mathcal{J}}$, captures the average treatment effect among units that move into treatment when all instruments in some fixed set \mathcal{J} are changed from 0 to 1, with the other instruments not in \mathcal{J} fixed at each unit's realized values. When \mathcal{J} contains just one instrument index, SLATE recovers treatment effects among those who would “comply” with variation in that single instrument. For example, $SLATE_{\{2\}}$ is the average treatment effect among individuals who don't go to college if it is far, but do if it is close given their realized value of the tuition instrument.¹¹ This parameter may for example be of interest to policymakers considering whether to expand a community college to a new campus.

The treatment effect parameters $SLATT_{\mathcal{J}}$ and $SLATU_{\mathcal{J}}$ in the final two rows of Table 1 are similar to $SLATE_{\mathcal{J}}$ but additionally condition on units' realized treatment status. For example $SLATT_{\{1,2\}}$ with our two instruments averages over individuals who do go to college, but wouldn't have gone were it far and expensive.¹² SLATT and SLATU can also be used to partially identify the average treatment effect among the treated or untreated, when potential outcomes are bounded, following logic detailed below for the ATE. Section 6 discusses estimation of the parameters listed in Table 1.

Identified sets for ATE, ATT, and ATU

Note that all of the parameters given in Table 1 omit “always-takers” and ‘never-takers’: any individuals for whom $\inf_{z \in \mathcal{Z}} D_i(z) = 1$ and $\sup_{z \in \mathcal{Z}} D_i(z) = 0$, respectively. Thus, these parame-

¹¹Note that a single-instrument SLATE like $SLATE_{\{2\}}$ does *not* generally correspond to using Z_2 alone as an instrument, e.g. $Cov(Y, Z_2)/Cov(D, Z_2)$, since this estimand does not control for variation in Z_1 that is correlated with Z_2 . For this reason, linear 2SLS does not generally estimate a convex combination of treatment effects under VM. However, this is also true under “fully-saturated” 2SLS, in which the instruments are interacted with one another (Mogstad et al., 2021).

¹²Note that with a single binary instrument, $SLATT_{\{1\}}$ coincides with $ACL = SLATE_{\{1\}}$, as $\mathbb{E}[Y_i(1) - Y_i(0)|D_i = 1, G_i = \text{complier}] = \mathbb{E}[Y_i(1) - Y_i(0)|Z_i = 1, \text{complier}] = \mathbb{E}[Y_i(1) - Y_i(0)|\text{complier}]$, using Assumption 1. However, when the group \mathcal{G}^c consists of more than one group, the “all-compliers” version of $SLATT$ generally differs from ACL .

ters do not speak directly to some policy-relevant quantities, such as the overall (unconditional) average treatment effect.

However, when Y_i has bounded support, we can use parameters such as those in Table 1 to generate sharp bounds in the spirit of Manski (1990). For example, consider the overall average treatment effect, which we can write as:

$$ATE := \mathbb{E}[Y_i(1) - Y_i(0)] = p_a\Delta_a + p_n\Delta_n + (1 - p_t - p_a)ACL \quad (3)$$

where Δ_a is the average treatment effect among always-takers, and p_a is the proportion of the population that are always-takers (and analogously for Δ_n and p_n).

With binary instruments having full rectangular support $\mathcal{Z} = \{0, 1\}^J$, the data at $\bar{Z} := (1, 1, \dots, 1)'$ and $\underline{Z} := (0, 0, \dots, 0)'$ are informative about always- and never-takers. For example, $p_a = \mathbb{E}[D_i | Z_i = \underline{Z}]$ and $p_n = \mathbb{E}[1 - D_i | Z_i = \bar{Z}]$. Now we suppose that $Y_i(0), Y_i(1) \in [\underline{Y}, \bar{Y}]$ for some \underline{Y} and \bar{Y} and all i . Then:

$$p_a\Delta_a \in [\mathbb{E}[Y_i D_i | Z_i = \underline{Z}] - p_a \cdot \bar{Y}, \quad \mathbb{E}[Y_i D_i | Z_i = \underline{Z}] - p_a \cdot \underline{Y}]$$

and

$$p_n\Delta_n \in [\underline{Y} \cdot p_n - \mathbb{E}[Y_i(1 - D_i) | Z_i = \bar{Z}], \quad \bar{Y} \cdot p_n - \mathbb{E}[Y_i(1 - D_i) | Z_i = \bar{Z}]]$$

Combining these two equations with (3) leads to bounds on the ATE, given the point-identified value of the ACL.

4 Response groups under vector monotonicity

In this section I show that when the instruments are binary the assumption of vector monotonicity partitions the population of interest into a set of groups that naturally generalize the familiar taxonomy of “always-takers”, “never-takers”, and “compliers” from the case of a single binary instrument. Understanding the structure of these groups paves the way for the identification results of Section 5.

Let us note first that specializing to binary instruments is with no loss of generality, if one begins with finite discrete instruments satisfying vector monotonicity. Such discrete instruments can be re-expressed as a larger number of binary instruments in a way that preserves VM:

Proposition 2. Let Z_1 be a discrete variable with M ordered points of support $z_1 < z_2 < \dots < z_M$, and $Z_2 \dots Z_J$ be other instrumental variables. Define $\tilde{Z}_{mi} := \mathbb{1}(Z_{1i} \geq z_m)$. If the vector $Z = (Z_1, \dots, Z_J)$ satisfies Assumption VM on a connected \mathcal{Z} then so does the vector $(\tilde{Z}_2, \dots, \tilde{Z}_M, Z_2, \dots, Z_J)$.

Proof. See Appendix E. □

Applying Proposition 2 iteratively allows one to begin with discrete instruments in a given empirical setting, and then replace them with a set of binary instruments that still satisfy VM. This is done by introducing one binary instrument per value for any instrument that was initially discrete (omitting the lowest value for each initial instrument).

To simplify notation, let G_i represent an individual's entire vector of counterfactual treatments $\{D_i(\mathbf{z})\}_{\mathbf{z} \in \mathcal{Z}}$. For example, with a single binary instrument $G_i = \text{always-taker}$ indicates that $D_i(0) = D_i(1) = 1$. I refer to G_i as unit i 's “response group”.¹³ Response groups partition individuals in the population based on upon their selection behavior over all counterfactual values of the instruments. We can think of VM as a restriction on the support \mathcal{G} of G_i , limiting the types of response groups that can coexist in the population.

Without loss of generality, let us continue to normalize the value labeled “1” for each binary instrument to be the direction in which potential treatments are increasing according to VM. These “up” values might be known ex-ante, on the basis of theoretical predictions about choice behavior, but they are also empirically identified from the propensity score function (see Proposition 1).

4.1 A single binary instrument

With one binary instrument, VM and IAM coincide. \mathcal{G} then contains three familiar groups (see e.g. Angrist and Pischke 2008): “compliers” (for whom $D_i(1) > D_i(0)$), “always-takers” (for whom $D_i(1) = D_i(0) = 1$) and “never-takers” (for whom $D_i(1) = D_i(0) = 0$).

4.2 Two binary instruments

In the case of two binary instruments, \mathcal{G} contains six distinct response groups under VM. These groups have been described by MTW, who give them the names reported in Table 2 below.

¹³This language follows that of Lee and Salanié (2020). Heckman and Pinto (2018) refer to such groups as *response-types* or *strata*.

A Z_1 complier, for example, would in the returns to college example go to college if and only if college is cheap, regardless of whether it is close. A Z_2 complier, by contrast, would go to college if and only if college is close, regardless of whether it is cheap. A reluctant complier is “reluctant” in the sense that they require college to be both cheap and close to attend, while an eager complier goes to college so long as it is either cheap or close. Never and always takers are defined in the same way as they are under IAM: $\max_{\mathbf{z} \in \mathcal{Z}} D_i(\mathbf{z}) = 0$ and $\min_{\mathbf{z} \in \mathcal{Z}} D_i(\mathbf{z}) = 1$, respectively.¹⁴

Name	$\mathbf{z} = (0, 0)'$	$\mathbf{z} = (0, 1)'$	$\mathbf{z} = (1, 0)'$	$\mathbf{z} = (1, 1)'$
never takers	N	N	N	N
always takers	T	T	T	T
Z_1 compliers	N	N	T	T
Z_2 compliers	N	T	N	T
eager compliers	N	T	T	T
reluctant compliers	N	N	N	T

Table 2: The six response groups under VM with two binary instruments. Here “T” denotes treatment ($D_i(\mathbf{z}) = 1$) and “N” no-treatment ($D_i(\mathbf{z}) = 0$).

A natural question is whether the sizes $p_g := P(G_i = g)$ of the six groups in Table 2 can be quantified empirically. In general, only two of them are point identified. Let $P(\mathbf{z}) := \mathbb{E}[D_i | Z_i = \mathbf{z}]$ be the propensity score function. From the definitions in Table 2, it is clear that $p_{n.t} = 1 - P(1, 1)$ and $p_{a.t.} = P(0, 0)$. We cannot point identify the remaining group occupancies, but as a consequence of Eq. (2b) we can identify certain linear combinations of them, e.g. $P(1, 0) - P(0, 0) = p_{Z_1} + p_{eager}$, $P(0, 1) - P(0, 0) = p_{Z_2} + p_{eager}$, and $P(1, 1) - P(0, 1) = p_{Z_1} + p_{reluctant}$. Positivity of each p_g implies bounds, however: for example, $\{P(1, 0) - P(0, 0)\} - \{P(1, 1) - P(0, 1)\} \leq p_{eager} \leq \min\{P(0, 1) - P(0, 0), P(1, 0) - P(0, 0)\}$.

4.3 Multiple binary instruments

In this section I show how the two-instrument case generalizes to a case where the researcher has any number J of binary instruments. This provides a foundation for our characterization of identified treatment effect parameters in Section 5.

¹⁴The examples from Table 1 can be understood in terms of these groups. For example, the group of individuals included in $SLATE_{\{2\}}$ are Z_2 compliers, eager compliers with high tuition rates ($Z_{1i} = 0$), and reluctant compliers with low tuition rates ($Z_{1i} = 1$). The ACL includes all individuals who are not always-takers or never-takers. For a discrete instrumental variable mapped to multiple binary instruments by Proposition 2, the LATE among units moved into treatment between any two of its values will also be an example of a SLATE. For example, if Z_1 has support $z_1 < z_2 < z_3 < z_4$, the average treatment effect among individuals for which $D_i(\mathbf{z}_4, Z_{-1,i}) > D_i(\mathbf{z}_2, Z_{-1,i})$ corresponds to $SLATE_{\mathcal{J}}$ with $\mathcal{J} = \{\tilde{Z}_3, \tilde{Z}_4\}$. SLATE thus allows the practitioner to flexibly condition upon response to individual or joint variation in the instruments.

I focus on the baseline case in which the space of conceivable instrument values is rectangular: $\mathcal{Z} = \{0, 1\}^J$ (see Section 7 for an example not of this form). We wish to characterize the subset of the 2^{2^J} distinct possible mappings between vectors of instrument values and treatment that satisfy VM. The number of such response groups G_i as a function of J is equal to the number of isotone boolean functions on J variables, which is known to follow the so-called Dedekind sequence (Kisielewicz, 1988):¹⁵

$$3, 6, 20, 168, 7581, 7828354\dots$$

For example, there are 20 response groups when $J = 3$. Let Ded_J denote the J^{th} number in the Dedekind sequence. We have seen above the response groups corresponding to $\text{Ded}_1 = 3$ and $\text{Ded}_2 = 6$.

Now let us consider an arbitrary finite J . One group that always satisfies VM is composed of those units for whom $D_i(\mathbf{z}) = 0$ for all values $\mathbf{z} \in \mathcal{Z}$: “never-takers”. Each of the remaining response groups can be associated with a collection of minimal combinations of instruments that are sufficient for that unit to take treatment: for example, in a setting with three instruments, one response group would be the units that take treatment if either $Z_1 = 1$, or if $Z_2 = Z_3 = 1$. Note that by vector monotonicity, any unit in this group must then also take treatment if $Z_1 = Z_2 = Z_3 = 1$. Another group of units might take treatment *only* if $Z_1 = Z_2 = Z_3 = 1$. This group is more “reluctant” than the former. The group of always-takers are the least “reluctant”: they require no instruments to equal one in order for them to take treatment.

Continuing this logic, we can associate each response group (aside from never-takers) with a family F of subsets $S \subseteq \{1 \dots J\}$ of the instrument labels, where a unit takes treatment when all instruments in some $S \in F$ are equal to one. However, we need only consider families for which no element S of the family is a subset of some other S' : so-called *Sperner families* (see e.g. Kleitman and Milner 1973). Families that are not Sperner would be redundant under VM. For example, if F consists of the set $\{2, 3\}$ and $\{1, 2, 3\}$, then given VM the set $\{1, 2, 3\}$ could be dropped from F without affecting the implied selection function $D_i(\mathbf{z})$.

Definition 2 (response group for a Sperner family). *For any Sperner family F , let $g(F)$*

¹⁵An analytical expression for the Dedekind numbers is given by Kisielewicz (1988), but only the first eight have been calculated numerically due to the computational burden of doing so. While the Dedekind numbers explode quite rapidly, they do so much more slowly than the total number 2^{2^J} of boolean functions of J variables. For example while $3/4 = 75\%$ of conceivable response groups for $J = 1$ satisfy VM, only $20/256 \approx 7.8\%$ do for $J = 3$, and just $7581/4294967296 \approx 1.7 * 10^{-4}$ do for $J = 5$. Thus the “bite” of VM is increasing with J , ruling out an increasing fraction of conceivable selection patterns.

denote the response group described by the following treatment rule: units take treatment if and only if $z_j = 1$ for all j in S , for at least one of the sets S in F .

All together, the response groups satisfying VM with J binary instruments are as follows: i) the never-takers group, and ii) $Ded_J - 1$ further groups $g(F)$ corresponding to each of the distinct Sperner families F of instrument labels (one such family is always the null-set, which corresponds to the always-takers).

In the simplest example of the above, when $J = 1$, vector monotonicity coincides with PM and IAM, and the Sperner families corresponding to this single instrument are simply the null set and the singleton $\{1\}$: corresponding to always-takers and compliers, respectively. Together with never-takers, we have the familiar three groups from LATE analysis with a single binary instrument.

For $J = 2$, the five groups (apart from never takers) described in Table 2 map to Sperner families as follows:

F	name of G_F
\emptyset	“always takers”
$\{1\}$	“ Z_1 compliers”
$\{2\}$	“ Z_2 compliers”
$\{1\}, \{2\}$	“eager compliers”
$\{1, 2\}$	“reluctant compliers”

With $J = 3$ there are 19 Sperner families, each indicated below within bold brackets:

$$\begin{aligned} & \{\emptyset\}, \{\{1\}\}, \{\{2\}\}, \{\{3\}\}, \{\{1, 2\}\}, \{\{1, 3\}\}, \{\{2, 3\}\}, \{\{1, 2, 3\}\}, \\ & \{\{1\}, \{2\}\}, \{\{2\}, \{3\}\}, \{\{1\}, \{3\}\}, \{\{1\}, \{2\}, \{3\}\}, \{\{1, 2\}, \{3\}\}, \{\{1, 3\}, \{2\}\}, \{\{2, 3\}, \{1\}\}, \\ & \{\{1, 2\}, \{1, 3\}\}, \{\{1, 2\}, \{2, 3\}\}, \{\{1, 3\}, \{2, 3\}\}, \{\{1, 2\}, \{1, 3\}, \{2, 3\}\} \end{aligned}$$

For instance, an individual with G_i corresponding to $\{\{1, 2\}, \{1, 3\}, \{2, 3\}\}$ takes treatment so long as any two of the instruments take a value of one.

A key difference between VM and IAM for identification is that under VM, the selection functions corresponding to the various response groups are not all linearly independent of one

another. Only 2^J such functions can be independent (though Ded_J is strictly larger than 2^J for $J > 1$). Let $\mathcal{G}^c := \mathcal{G}/\{\text{a.t., n.t.}\}$ denote the set of $\text{Ded}_J - 2$ response groups compatible with Assumption VM, aside from the never-takers and always takers. All of the groups in \mathcal{G}^c can be thought of as generalized “compliers” of some kind: units that vary treatment uptake in *some* way across possible instrument values.

Denote by $\mathcal{D}_g(\mathbf{z})$ the potential treatments function $D_i(\mathbf{z})$ that is common to all units sharing a value g of G_i . A natural basis for the set of selection functions $\{\mathcal{D}_g(\mathbf{z})\}_{g \in \mathcal{G}^c}$ can be formed by considering the treatment selection functions for response groups $g(S)$ corresponding to Sperner families that consist of a single set $S \subseteq \{1 \dots J\}$:

$$\mathcal{D}_{g(S)}(\mathbf{z}) = \prod_{j \in S} z_j = \mathbb{1}(z_j = 1 \text{ for all } j \text{ in } S)$$

I refer to such response groups $g(S)$ as *simple*.¹⁶ For $J = 2$, the selection functions for the simple response groups are:

$$\mathcal{D}_{Z_1}(\mathbf{z}) = z_1 \quad \mathcal{D}_{Z_2}(\mathbf{z}) = z_2 \quad \mathcal{D}_{reluctant}(\mathbf{z}) = z_1 z_2$$

The selection function for the remaining group, eager compliers, can then be obtained as:

$$\mathcal{D}_{eager}(\mathbf{z}) = z_1 + z_2 - z_1 z_2 = \mathcal{D}_{Z_1}(\mathbf{z}) + \mathcal{D}_{Z_2}(\mathbf{z}) - \mathcal{D}_{reluctant}(\mathbf{z})$$

We can express this linear dependency by the matrix M_J in the system:

$$\begin{pmatrix} \mathcal{D}_{Z_1}(\mathbf{z}) \\ \mathcal{D}_{Z_2}(\mathbf{z}) \\ \mathcal{D}_{reluctant}(\mathbf{z}) \\ \mathcal{D}_{eager}(\mathbf{z}) \end{pmatrix} = \underbrace{\begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & -1 \end{pmatrix}}_{:= M_2} \begin{pmatrix} \mathcal{D}_{Z_1}(\mathbf{z}) \\ \mathcal{D}_{Z_2}(\mathbf{z}) \\ \mathcal{D}_{reluctant}(\mathbf{z}) \end{pmatrix} \quad (4)$$

For general J , we define the matrix M_J from the analogous system of equations:

$$\{\mathcal{D}_{g(F)}(\mathbf{z})\}_{F: g(F) \in \mathcal{G}^c} = M_J \{\mathcal{D}_{g(S)}(\mathbf{z})\}_{S \subseteq \{1 \dots J\}, S \neq \emptyset},$$

¹⁶Note that a similar construction plays a central role in Lee and Salanié, 2018.

which holds for all $\mathbf{z} \in \mathcal{Z}$. The rows of matrix M_J are indexed by Sperner families (corresponding to the groups in \mathcal{G}^c), and the columns by the simple Sperner families (for non-null S). The entries of M_J are given explicitly by the following expression:

Proposition 3. $[M_J]_{F,S'} = \sum_{f \in s(F,S')} (-1)^{|f|+1}$ where $s(F,S') := \left\{ f \subseteq F : \left(\bigcup_{S \in f} S \right) = S' \right\}$.

Proof. See Appendix E. □

5 Identification under VM with binary instruments

In this section I define and characterize the class of causal parameters that are point identified under vector monotonicity, and a full-support condition on the instruments. This section maintains a setup of J binary instruments with $\mathcal{Z} = \{0,1\}^J$.

5.1 Parameters of interest and identification

Generalizing from Section 3, let us begin by considering conditional averages of either potential outcome, after possible transformation by a function f . For $d \in \{0,1\}$, let

$$\theta_c^{fd} := \mathbb{E}[f(Y_i(d)) | C_i = 1]$$

where $C_i = c(G_i, Z_i)$ is any function $c : \mathcal{G} \times \mathcal{Z} \rightarrow \{0,1\}$ of individual i 's response group and their realization of the instruments. Intuitively, the event $C_i = 1$ will indicate that unit i belongs to a particular subgroup of generalized ‘‘compliers’’. Most of the discussion will center on the class of average treatment effect parameters:

$$\Delta_c := \mathbb{E}[Y_i(1) - Y_i(0) | C_i = 1] = \theta_c^{y1} - \theta_c^{y0}$$

with $f(y) = y$ the identity function. Functions of the form $c(g, z)$ are most general type of conditioning event that depends on primitives of the basic IV model, without depending directly on potential outcomes.¹⁷ For example, the definition Δ_c nests treatment effect parameters familiar both from the LATE (Imbens and Angrist, 1994) and marginal treatment effects (Heckman and Vytlacil, 2005) literatures. For instance, with a single binary instrument the LATE sets $c(g, \mathbf{z}) = \mathbb{1}(g = \text{complier})$. The class of estimands considered in Section 3 that condition on

¹⁷However, no restrictions are imposed on the joint distribution of $(Y_i(1), Y_i(0), G_i)$, so this model is compatible with G_i being arbitrarily correlated with potential outcomes or with treatment effects, as in Roy-type models.

$i \in \mathcal{C}^{u,l}$ are also an example of this general form, since

$$D_i(u(Z_i)) > D_i(l(Z_i)) \iff c(G_i, Z_i) = 1 \text{ where } c(g, \mathbf{z}) = \mathcal{D}_g(u(\mathbf{z})) - \mathcal{D}_g(l(\mathbf{z}))$$

I now characterize the family of $c(g, \mathbf{z})$ under which identification of θ_c^{fd} is possible. In particular, a necessary and sufficient condition will be what I call “Property M”:

Definition 3 (Property M). Suppose there are J binary instruments. We say the function $c(g, \mathbf{z})$ satisfies Property M if for every $g \in \mathcal{G}^c$:

$$c(g, \mathbf{z}) = \sum_{S \subseteq \{1, \dots, J\}, S \neq \emptyset} [M_J]_{F(g), S} \cdot c(g(S), \mathbf{z}),$$

while $c(always-takers, \mathbf{z}) = c(never-takers, \mathbf{z}) = 0$, for all $\mathbf{z} \in \mathcal{Z}$.

where the matrix M_J is defined in Proposition 3. I’ll also say that a parameter θ_c^{fd} or Δ_c “satisfies Property M” if its underlying function $c(g, \mathbf{z})$ does.

While Property M is somewhat abstract, the following result connects it to a generalization of the groups $\mathcal{C}^{u,l}$ from Section 3:

Proposition 4. A function $c : \mathcal{G} \times \mathcal{Z} \rightarrow \{0, 1\}$ satisfies Property M if and only if

$$c(g, \mathbf{z}) = \sum_{k=1}^K \{\mathcal{D}_g(u_k(\mathbf{z})) - \mathcal{D}_g(l_k(\mathbf{z}))\}$$

for some $K \leq J/2$, where $u_k(\cdot)$ and $l_k(\cdot)$ are functions $\mathcal{Z} \rightarrow \mathcal{Z}$ such that $u_k(\mathbf{z}) \geq l_k(\mathbf{z})$ component-wise while $l_k(\mathbf{z}) \geq u_{k+1}(\mathbf{z})$ component-wise, for all k and $\mathbf{z} \in \mathcal{Z}$.

Proof. See Appendix E. □

This yields a natural interpretation of average treatment effects that satisfy Property M, which is that they can be written $\Delta_c = \mathbb{E} \left[Y_i(1) - Y_i(0) \mid \bigcup_{k=1}^K \{i : D_i(u_k(Z_i)) > D_i(l_k(Z_i))\} \right]$ for some sequence of functions u_k and l_k having the properties stated in Proposition 4. It also shows that each of the parameters introduced in Table 1, e.g. the ACL and the SLATE’s, satisfy Property M when $\mathcal{Z} = \{0, 1\}^J$.¹⁸

¹⁸Some further examples of identified parameters from those mentioned in Table 1 can be constructed using a closure property of the set of c satisfying Property M. Let \mathcal{C} denote the set of $c : \mathcal{G} \times \mathcal{Z} \rightarrow \{0, 1\}$ that satisfy Property M, and let $c_a(g, \mathbf{z})$ and $c_b(g, \mathbf{z})$ be two functions in \mathcal{C} . Then it is straightforward to show that $c_a(g, \mathbf{z}) - c_b(g, \mathbf{z}) \in \mathcal{C}$ iff $c_b(g, \mathbf{z}) \leq c_a(g, \mathbf{z})$ for all $\mathbf{z} \in \mathcal{Z}, g \in \mathcal{G}^c$. We can use this observation to generate parameters that condition on the “complement” of the complier group for Δ_{c_b} within the larger complier group

Causal parameters that satisfy Property M are identified under VM with binary instruments, provided the various instruments provide sufficient independent variation in treatment uptake. A simple sufficient condition for this is that the instruments have full (rectangular) support. This assumption is stronger than necessary, but simplifies presentation (Appendix A gives a generalization).

Assumption 3 (full support). $P(Z_i = \mathbf{z}) > 0$ for all $\mathbf{z} \in \{0, 1\}^J$

An alternative expression of Assumption 3 is useful for stating the constructive identification result below. For an arbitrary ordering of the $k := 2^J - 1$ non-empty subsets $S \subseteq \{1 \dots J\}$, define the random vector $\Gamma_i = (Z_{S_1 i} \dots Z_{S_k i})'$ from products of the Z_{ji} for j within each subset S , i.e. $Z_{Si} = \prod_{j \in S} Z_{ji}$. That is, each element of Γ_i indicates the treatment status of a particular simple response group $g(S)$, given Z_i . Let Σ be the covariance matrix of Γ_i .

Lemma 1. *Assumption 3 holds if and only if Σ has full rank.*

Proof. See Appendix E. □

Lemma 1 shows that full support of the instruments is equivalent to there being linearly independent variation in treatment takeup among all of the simple response groups.

Theorem 1 provides an explicit estimand for parameters θ_c^{fd} when the function c satisfies Property M:

Theorem 1. *Under Assumptions 1-3 (independence & exclusion, VM, and full support), for any c satisfying Property M and any measurable function $f(Y)$:*

$$\theta_c^{fd} = (-1)^{d+1} \frac{\mathbb{E}[f(Y_i)h(Z_i)\mathbb{1}(D_i = d)]}{\mathbb{E}[h(Z_i)D_i]},$$

and

$$P(C_i = 1) = E[h(Z_i)D_i]$$

provided that $P(C_i = 1) > 0$, where $h(Z_i) = \lambda' \Sigma^{-1}(\Gamma_i - \mathbb{E}[\Gamma_i])$ and

$$\lambda = (\mathbb{E}[c(g(S_1), Z_i)], \dots, \mathbb{E}[c(g(S_k), Z_i)])'$$

for Δ_{ca} . For example with $J = 2$, consider the average treatment effect among individuals who are counted in the ACL but not in $SLATE_{\{1\}}$: $\mathbb{E}[\Delta_i | G_i \in \mathcal{G}^c \text{ but } \{D_i(1, Z_{2i}) = D_i(0, Z_{2i})\}]$. This could represent the average treatment effect among individuals that would not respond to a counterfactual reduction in college tuition alone, but would respond if both tuition and proximity were shifted in concert.

Proof. See Appendix E. □

It follows immediately from Theorem 1 that conditional average treatment effects $\Delta_c = \theta_c^{y1} - \theta_c^{y0}$ satisfying Property M are identified, and the expression simplifies to:

$$\Delta_c = \mathbb{E}[h(Z_i)Y_i]/\mathbb{E}[h(Z_i)D_i]$$

Note that as the numerator of Δ_c depends on Z_i and Y_i only and the denominator depends on Z_i and D_i only, identification of Δ_c would hold in a “split-sample” setting where Y_i and D_i are not necessarily linked in the same dataset.

Now I show that Theorem 1 has a converse: *any* identified Δ_c must satisfy Property M. In this sense, Property M is both a necessary and sufficient condition for identification. To state this result in a simple way, let us consider so-called “IV-like estimands” introduced by Mogstad et al. (2018), which are any cross moment $\mathbb{E}[s(D_i, Z_i)Y_i]$ between Y_i and a function of treatment and instruments. Let \mathcal{P}_{DZ} denote the joint distribution of D and \mathbf{z} (which is identified from the data). Then:

Theorem 2. Suppose θ_c^{fd} is identified by a finite set of IV-like estimands and \mathcal{P}_{DZ} , provided that Assumptions 1-3 hold and $P(C_i = 1) > 0$. Then θ_c^{fd} satisfies Property M.

Proof. See Appendix E.¹⁹ □

An algebraic intuition for Theorem 1

The basic logic behind Theorem 1 can be appreciated by focusing on the average treatment effect parameters Δ_c , and observing that by Assumption 1 and the law of iterated expectations they can be written as a weighted average over response-group specific average treatment effects $\Delta_g := \mathbb{E}[Y_i(1) - Y_i(0)|G_i = g]$:

$$\Delta_c = \sum_{g \in \mathcal{G}} \left\{ \frac{P(G_i = g)\mathbb{E}[c(g, Z_i)]}{\sum_{g' \in \mathcal{G}} P(G_i = g')\mathbb{E}[c(g', Z_i)]} \right\} \cdot \Delta_g \quad (5)$$

where the weight on Δ_g is proportional to the quantity $\mathbb{E}[c(g, Z_i)]$ (as well as $P(G_i = g)$). Now consider a general type of estimand in which a single scalar “instrument” $h(Z_i)$ is constructed

¹⁹In saying that a parameter θ is *identified* by some set of empirical estimands, I mean that the set of values of θ that are compatible with the empirical estimands is a singleton, regardless of the distribution of the latent variables $(G_i, Y_i(1), Y_i(0))$ —for all \mathcal{P}_{DZ} satisfying maintained assumptions, namely that: i) the marginal distribution of Z_i satisfies Assumption 3; ii) $P(C_i > 0)$, and iii) VM.

from the vector of instruments Z_i according to some function h , and then used in a simple linear IV regression.²⁰ Some similar algebra shows that under Assumption 1:

$$\frac{Cov(Y_i, h(Z_i))}{Cov(D_i, h(Z_i))} = \sum_{g \in \mathcal{G}} \frac{P(G_i = g) \cdot Cov(\mathcal{D}_g(Z_i), h(Z_i))}{\sum_{g' \in \mathcal{G}} P(G_i = g') \cdot Cov(\mathcal{D}_{g'}(Z_i), h(Z_i))} \cdot \Delta_g \quad (6)$$

This reveals that such estimands also uncover a weighted average of the Δ_g , where the weight placed on each response group g is governed by the covariance between $\mathcal{D}_g(Z_i)$ and $h(Z_i)$. Comparing Equations (5) and 6, we see that a simple IV estimand using $h(Z_i)$ can identify Δ_c if the function h is chosen in such a way that $Cov(\mathcal{D}_g(Z_i), h(Z_i)) = \mathbb{E}[c(g, Z_i)]$ for all the response groups g . However, since the covariance operator is linear, the linear dependencies examined in Section 4.3 translate into a set of linear restrictions among these weights, captured by the matrix M_J . Property M guarantees that the vector of $\mathbb{E}[c(g(F), Z_i)]$ across Sperner families F belongs to the column-space of the matrix M_J , whatever the distribution of Z_i . What remains to secure identification is then simply to independently “tune” the covariances for each simple response group, which is possible given full rectangular support of the instruments and is achieved by the specific choice of $h(Z_i)$ in Theorem 1.

The role of Property M in Theorem 1 can thus be thought of as emerging from there being under VM more response groups in \mathcal{G}^c than there are independent pairs of points in the support of the instruments. This contrasts with IAM, under which with J binary instruments both are generally equal to $2^J - 1$. As a result, it is possible under IAM to identify for any single such response group g^* the average treatment effect $\Delta_{g^*} := \mathbb{E}[Y_i(1) - Y_i(0)|G_i = g^*]$ (and hence also obtain any desired convex combination of the Δ_{g^*}). However, under VM the corresponding choice $c(g, \mathbf{z}) = \mathbb{1}(g = g^*)$ fails to satisfy Property M, and we cannot identify the Δ_g individually in general.²¹

5.2 Extensions and further results on identification

This section outlines some further results related to identification under VM.

1) Identification is less powerful under PM alone. The assumption of vector monotonicity does

²⁰Special cases of this form include 2SLS: $h(\mathbf{z}) = \mathbb{E}[D_i|Z_i = \mathbf{z}]$, and Wald-type estimands: $h(\mathbf{z}) = \frac{\mathbb{1}(Z_i = \mathbf{z})}{P(Z_i = \mathbf{z})} - \frac{\mathbb{1}(Z_i = \mathbf{z}')}{P(Z_i = \mathbf{z}')}}$.

²¹We can see this in a simple example with $J = 2$ and $g = Z_1$ complier. In this case Property M would require that $c(\text{eager complier}, \mathbf{z}) = c(Z_1 \text{ complier}, \mathbf{z}) + c(Z_2 \text{ complier}, \mathbf{z}) - c(\text{reluctant complier}, \mathbf{z})$, i.e. that $0 = 1 + 0 - 1$, cf Eq. (4).

have additional identifying power beyond that of PM. Consider the *ACL*, which by Theorem 1 is identified under VM. In the $J = 2$ case, it is possible to see that the *ACL* is not identified under PM only, by explicit enumeration of the possible response groups.

Proposition 5. *When $J = 2$, if PM holds but neither VM nor IAM hold, the ACL is not point identified from knowledge of any set of IV-like estimands and \mathcal{P}_{DZ} .*

Proof. See Appendix E. □

2) *Linear dependency among the instruments:* Assumption 3 is stronger than is strictly necessary for identification, since linear dependencies between products of the instruments pose no problem if the corresponding “weights” in Δ_c do not need be tuned independently from one another. In Appendix A, I give a version of Assumption 3 and generalization of Theorem 1 that can accommodate instrument support restrictions and/or non-rectangular \mathcal{Z} (for instance after applying Proposition 2 to map discrete instruments to binary ones).

3) *Conditional distributions of the potential outcomes* By choosing $f(Y) = \mathbb{1}(Y \leq y)$ in Theorem 1 for some value y in the support of Y_i , we can identify the CDF of each potential outcome at y conditional on $C_i = 1$ as: $F_{Y(d)|C=1}(y) = (-1)^{d+1} \frac{\mathbb{E}[h(Z_i)\mathbb{1}(D_i=d)\mathbb{1}(Y_i \leq y)]}{\mathbb{E}[h(Z_i)D_i]}$ (note that unlike identification of Δ_c this requires observing (Y_i, Z_i, D_i) all in the same sample). This allows for the identification of $C_i = 1$ conditional quantile treatment effects, bounds on the distribution of treatment effects (Fan and Park, 2010), or distributional treatment effects: $F_{Y(1)|C=1}(y) - F_{Y(0)|C=1}(y)$ as $\frac{\mathbb{E}[h(Z_i)\mathbb{1}(Y_i \leq y)]}{\mathbb{E}[h(Z_i)D_i]}$.

6 Estimation

This section proposes a simple two-step estimator for the family of identified causal parameters introduced in Section 5, focusing on the conditional average treatment effects Δ_c . Appendix C discusses its limiting distribution, which is asymptotically normal and converges at the parametric rate.

Theorem 1 establishes that Δ_c satisfying Property M are equal to a ratio of two population expectations involving the function $h(Z_i)$. A natural plug-in estimator simply replaces these expectations with their sample counterparts, and $h(Z_i)$ by a first-step estimate. Following $h(Z_i) = (\Gamma_i - \mathbb{E}[\Gamma_i])'\mathbb{E}[(\Gamma_i - \mathbb{E}[\Gamma_i])(\Gamma_i - \mathbb{E}[\Gamma_i])']^{-1}\lambda$ from Theorem 1, define $\hat{H} = n\tilde{\Gamma}(\tilde{\Gamma}'\tilde{\Gamma})^{-1}\hat{\lambda}$,

where $\tilde{\Gamma}$ is a $n \times k$ design matrix with entries $\tilde{\Gamma}_{il} = Z_{S_l i} - \frac{1}{n} \sum_{j=1}^n Z_{S_l j}$, where S_l is the l^{th} subset according to some arbitrary ordering of the $k := 2^J - 1$ non-empty such subsets of $\{1 \dots J\}$. Note that the rows of $\tilde{\Gamma}$ correspond to observations of the vector Γ_i introduced in Section 5.1, residualized with respect to the sample mean. The vector $\hat{\lambda}$ is a sample estimator of $\lambda = (\mathbb{E}[c(g(S_1), Z_i)], \dots \mathbb{E}[c(g(S_k), Z_i)])'$, given in Table 3 for the leading examples from Table 1.

Given the vector \hat{H} as defined above, consider the estimator $\hat{\Delta}_c = (\hat{H}'D)^{-1}(\hat{H}'Y)$, where Y and D are $n \times 1$ vectors of observations of Y_i and D_i , respectively. An application of the Frisch-Waugh-Lovell theorem allows us to rewrite $\hat{\Delta}_c$ as $\hat{\rho}(\hat{\lambda})$,²² where

$$\hat{\rho}(\lambda) := ((0, \lambda')(\Gamma'\Gamma)^{-1}\Gamma'D)^{-1}(0, \lambda')(\Gamma'\Gamma)^{-1}\Gamma'Y \quad (7)$$

and we introduce the $n \times 2^J$ matrix Γ that adds a column of ones before $\tilde{\Gamma}$ and skips the demeaning of each $Z_{S_l i}$. Note that existence of the population analog of $(\Gamma'\Gamma)^{-1}$ is guaranteed by Assumption 3. When Assumption 3 does not hold but identification is still possible (see Appendix A), the matrices $\tilde{\Gamma}$ and Γ may be defined in the same way but using only sets S within a smaller collection $\mathcal{F} \subset \{S \subseteq \{1, 2, \dots, J\}, S \neq \emptyset\}$. For example, when using construction of Proposition 2 that maps discrete to binary instruments, \mathcal{F} can be taken to include all sets of the final binary instruments that do not contain distinct \tilde{Z} from the same original discrete instrument. In all cases, let \mathcal{F} index the elements of Γ_i .

Comparison with 2SLS: Note that the estimator $\hat{\Delta}_c$ has a similar form to a “fully-saturated” 2SLS estimator that includes an indicator for each value of Z_i in the first stage. Indeed, that estimator can be written in the form of $\hat{\rho}(\lambda)$ in which each component of λ is a sample covariance between D_i and the corresponding component of Γ_i .²³ The weights $\widehat{\text{Cov}}(D_i, \Gamma_i)$ that 2SLS uses to aggregate over linear projection coefficients $(\Gamma'\Gamma)^{-1}\Gamma'V$ for $V \in \{D, Y\}$ are determined by the joint distribution of D_i and Z_i , and in general may not be positive. By contrast, $\hat{\Delta}_c$ uses weights $\hat{\lambda}_S = \hat{\mathbb{E}}[c(g(S), Z_i)]$ chosen to match the desired parameter of interest, guaranteeing that the estimator recovers a well-defined causal parameter under VM.²⁴ It bears emphasizing

²²To see this, note that as a vector of linear projection coefficients, $(\tilde{\Gamma}'\tilde{\Gamma})^{-1}\tilde{\Gamma}'D$ is the same as the final k components of the vector $(\Gamma'\Gamma)^{-1}\Gamma'D$, and thus $\lambda'(\tilde{\Gamma}'\tilde{\Gamma})^{-1}\tilde{\Gamma}'D = (0, \lambda')(\Gamma'\Gamma)^{-1}\Gamma'D$ (and similarly for Y).

²³ $\hat{\rho}_{2sls} = ((0, \widehat{\text{Cov}}(D_i, \Gamma_i))'(\Gamma'\Gamma)^{-1}\Gamma'D)^{-1}(0, \widehat{\text{Cov}}(D_i, \Gamma_i))'(\Gamma'\Gamma)^{-1}\Gamma'Y \xrightarrow{P} \sum_{g \in \mathcal{G}^c} \frac{P(G_i=g) \cdot \text{Cov}(D_i, \mathcal{D}_g(Z_i))}{\sum_{g'} P(G_i=g') \cdot \text{Cov}(D_i, \mathcal{D}_{g'}(Z_i))} \cdot \Delta_g$.

²⁴Despite not being consistent for any causal parameter Δ_c in general, $\hat{\rho}_{2sls}$ may have lower variance than $\hat{\Delta}_c$

that with a large number of instruments, $\hat{\Delta}_c$ is no more “expensive” than 2SLS: both involve computing a pair of linear projections with the same number 2^J of terms (despite the fact that the richness of possible selection behavior is more complex under VM than under IAM, scaling as Ded_J rather than 2^J).

Estimating λ : Under regularity conditions (see Theorem 3 in Appendix C), we have that for any consistent estimator $\hat{\lambda}$ of a $\lambda \in \mathbb{R}^{|\mathcal{F}|}$:

$$\hat{\rho}(\hat{\lambda}) \xrightarrow{p} \sum_{g \in \mathcal{G}^c} \frac{P(G_i = g)[M_J \lambda]_g}{\sum_{g' \in \mathcal{G}^c} P(G_i = g')[M_J \lambda]_{g'}} \cdot \Delta_g$$

Matching the RHS of the above to particular estimands Δ_c that satisfy Property M is achieved by choosing $\hat{\lambda}$ appropriately for that Δ_c . Table 3 gives consistent estimators $\hat{\lambda}$ for ACL, SLATE, SLATT, SLATU and PTE. Note that in the case of the ACL and PTEs, $\hat{\lambda}$ does not depend on the data and thus no “first-step” is necessary in estimation. The estimator $\hat{\Delta}_c = \hat{\rho}(\hat{\lambda})$ is implemented in the Stata package **ivcombine**.

Parameter	Estimator $\hat{\lambda}$ of population λ
ACL	$(1, 1, \dots, 1)'$
SLATE $_{\mathcal{J}}$	$\hat{\lambda}_S = \mathbb{1}(\mathcal{J} \cap S \neq \emptyset) \hat{P}(Z_{S-\mathcal{J}, i} = 1)$
SLATT $_{\mathcal{J}}$	$\hat{\lambda}_S = \mathbb{1}(\mathcal{J} \cap S \neq \emptyset) \hat{P}(Z_{S, i} = 1)$
SLATU $_{\mathcal{J}}$	$\hat{\lambda}_S = \mathbb{1}(\mathcal{J} \cap S \neq \emptyset) \hat{P}(Z_{S-\mathcal{J}, i}(1 - Z_{\mathcal{J}, i}) = 1)$
PTE $_j(\mathbf{z}_{-j})$	$\hat{\lambda}_S = \mathbb{1}(S = \mathbf{z}_{-j, 1} \cup \{j\})$

Table 3: Estimators $\hat{\lambda}$ for leading treatment effect parameters. $S - \mathcal{J}$ denotes the set difference $\{j : j \in S, j \notin \mathcal{J}\}$, $\mathbf{z}_{-j, 1}$ denotes the set of instruments that are equal to one in \mathbf{z}_{-j} , and $\hat{P}(E_i) = n^{-1} \sum_{i=1}^n \mathbb{1}(E_i)$ for any event E_i .

Estimation of the ACL from a single Wald ratio: The population estimand corresponding to the all-compliers LATE takes on a particularly simple form. In particular, the ACL is equal to the following single “Wald ratio”:

$$\rho_{\bar{Z}, \underline{Z}} := \frac{\mathbb{E}[Y_i | Z_i = \bar{Z}] - \mathbb{E}[Y_i | Z_i = \underline{Z}]}{\mathbb{E}[D_i | Z_i = \bar{Z}] - \mathbb{E}[D_i | Z_i = \underline{Z}]} \quad (8)$$

where $\bar{Z} = (1, 1, \dots, 1)'$ and $\underline{Z} = (0, 0, \dots, 0)'$, provided that $P(Z_i = \bar{Z}) > 0$ and $P(Z_i = \underline{Z}) > 0$, and the denominator is non-zero. This can be shown via the following corollary to Theorem 1, which expresses Δ_c in terms of conditional expectation functions of each of Y_i and D_i on the

as an estimator of Δ_c . In Section C.1 I discuss regaining some of that lost efficiency through regularization, and give simulation evidence of its usefulness.

instruments:²⁵

Corollary to Theorem 1. *Under the Assumptions of Theorem 1:*

$$\Delta_c = \frac{\sum_{\mathbf{z} \in \mathcal{Z}} \left(\sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} \lambda_S A_{S,\mathbf{z}} \right) \mathbb{E}[Y_i | Z_i = \mathbf{z}]}{\sum_{\mathbf{z} \in \mathcal{Z}} \left(\sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} \lambda_S A_{S,\mathbf{z}} \right) \mathbb{E}[D_i | Z_i = \mathbf{z}]}$$

where λ_S is as defined in Theorem 1 and $A_{S,\mathbf{z}} = \sum_{\substack{f \subseteq \mathbf{z}_0 \\ (\mathbf{z}_1 \cup f) = S}} (-1)^{|f|}$, with $(\mathbf{z}_1, \mathbf{z}_0)$ a partition of the indices $j \in \{1 \dots J\}$ that take a value of zero or one in \mathbf{z} , respectively.

Proof. See Appendix E. □

By Eq. (8), a very simple consistent estimator of the ACL is thus:

$$\hat{\rho}_{\bar{Z}, \underline{Z}} := \frac{\hat{\mathbb{E}}[Y_i | Z_i = \bar{Z}] - \hat{\mathbb{E}}[Y_i | Z_i = \underline{Z}]}{\hat{\mathbb{E}}[D_i | Z_i = \bar{Z}] - \hat{\mathbb{E}}[D_i | Z_i = \underline{Z}]} \quad (9)$$

It turns out that $\hat{\rho}_{\bar{Z}, \underline{Z}}$ is in fact numerically equivalent in finite sample to $\hat{\rho}((1, 1, \dots, 1)')$ from Eq. (7).²⁶ In situations where there is non-zero but small support on the points \bar{Z} and \underline{Z} , we may thus expect that $\hat{\rho}((1, 1, \dots, 1)')$ may perform poorly as an estimator of ACL in small samples, since it effectively ignores all of the data for which $Z_i \notin \{\underline{Z}, \bar{Z}\}$. This issue is mentioned by Frölich (2007) in the context of IAM, in which case $\hat{\rho}_{\bar{Z}, \underline{Z}}$ is also consistent for the ACL with finite \mathcal{Z} .²⁷ Appendix C develops and investigates the performance of a data-driven regularization procedure to ameliorate this problem, while also showing asymptotic normality of the estimator with or without such regularization. Appendix C also reports a simulation study that shows the regularization procedure can indeed be helpful in practice, and demonstrates the potential pitfalls of estimating 2SLS.

Covariates: In Appendix B, I describe how covariates can be accommodated in estimation when instrument independence holds only after conditioning on observed variables X , i.e.:

$$\{(Y_i(1), Y_i(0), G_i) \perp Z_i\} | X_i \quad (10)$$

²⁵In particular the identity $\sum_{f \subseteq S} (-1)^{|f|} = 0$ for any $S \neq \emptyset$ annihilates all but two of the components of $(0, \lambda')A$.

²⁶To see this, note that the vector H of H_i solves the system of equations $\Gamma' H_i = (1 \dots 1)'$. Among vectors that are in the column space of Γ , H is the unique such solution, given that the design matrix Γ has full column rank. One can readily verify that $\Gamma' H = (1, 1, \dots, 1)$ with the choice $H_i = \frac{1(Z_i=(1\dots1))}{P(Z_i=(0\dots0))} - \frac{1(Z_i=(0\dots0))}{P(Z_i=(0\dots0))}$, and that this $H = \Gamma\eta$ with $\eta = (1/\hat{P}(Z_i = (1 \dots 1))), 0, \dots, 0, -1/\hat{P}(Z_i = (0 \dots 0))$.

²⁷An analogous result to Eq. (8) holds under IAM with finite instruments, where in that case we take any $\bar{Z} \in \text{argmax}_z \mathbb{E}[D_i | Z_i = z]$ and $\underline{Z} \in \text{argmin}_z \mathbb{E}[D_i | Z_i = z]$, and define $\mathcal{G}^c := \{g \in \mathcal{G} : \mathbb{E}[\mathcal{D}_g(Z_i)] \in (0, 1)\}$.

The main result there is that while conditional average treatment effects $\Delta_c(x) := \mathbb{E}[Y_i(1) - Y_i(0)|C_i = c, X_i = x]$ can be identified for each x in the support of X_i , the unconditional Δ_c can be easier to estimate. Estimation can proceed nonparametrically, but a particularly simple case occurs when the conditional expectation functions $\mathbb{E}[Y_i|Z_i = \mathbf{z}, X_i = x]$ and $\mathbb{E}[D_i|Z_i = \mathbf{z}, X_i = x]$ are each additively separable between \mathbf{z} and x , and linear in x . Then, a simple consistent estimator of Δ_c is:

$$\hat{\Delta}_c = \left((0, \hat{\lambda}') (\Gamma' \mathcal{M}_X \Gamma)^{-1} \Gamma' \mathcal{M}_X D \right)^{-1} (0, \hat{\lambda}') (\Gamma' \mathcal{M}_X \Gamma)^{-1} \Gamma' \mathcal{M}_X Y \quad (11)$$

where \mathcal{M}_X is a projection onto the null space of the design matrix of X_i . In this case, the only change required to the unconditional estimator $\hat{\Delta}_c$ is to augment the linear projections of Y_i and D_i onto the instruments with X_i as additional linear regressors. I implement this estimator in the first of the two empirical applications below.

7 Empirical applications

In this section I apply the results of this paper to two well-known settings in which multiple instruments have been used: i) the labor market returns to college; and ii) the effects of children on labor supply. The two empirical applications highlight different aspects of identification and estimation under VM. In the first setting, I use covariates to increase the plausibility of Assumption 1. In the second, the instruments do not have rectangular support so I use identification results that relax Assumption 3.

In both empirical settings, existing literature has based IV methods on the traditional IAM notion of monotonicity (or on an assumption of homogeneous treatment effects). I instead base estimates on identification results that hold under VM and find new evidence of heterogeneity in treatment effects across response groups. This complements existing results that assume IAM—which is a less plausible assumption in either context.

7.1 The returns to college

I use the dataset from Carneiro, Heckman and Vytlacil (2011) (henceforth CHV) constructed from the 1979 National Longitudinal Survey of Youth. The sample consists of 1,747 white males in the U.S., first interviewed in 1979 at ages that ranged from 14 to 22, and then again annually. The outcome of interest Y_i is the log of individual i 's wage in 1991, and treatment

$D_i = 1$ indicates i attended at least some college. As in CHV, treatment effects are expressed in roughly per-year equivalents by dividing by four.

CHV consider four separate instruments for schooling. In a baseline setup, I use the two binary instruments from our running example: tuition and proximity. A second setup then adds the remaining two instruments from CHV, which capture local labor market conditions when a student is in high school. The first two instruments are defined as follows: $Z_{2i} = 1$ indicates the presence of a public college in i 's county of residence at age 14, while $Z_{1i} = 1$ indicates that average tuition rates local to i 's residence around age 17 falls below the sample median, which corresponds to about \$2,170 in 1993 dollars. This represents one particular choice of how tuition can be discretized into a binary variable, but note that the methods in this paper could also be used with tuition recast as a discrete variable with a rich set of tuition levels.

While VM is a natural assumption for the tuition and proximity instruments, a conditional version of instrument validity is more plausible than Assumption 1. Conditioning on observables can help control for unobserved heterogeneity that may be correlated with location during teenage years. I follow CHV and as control variables a student's corrected Armed Forces Qualification Test score, mother's years of education, number of siblings, "permanent" local earnings in county of residence at 17, "permanent" unemployment in county of residence at 17, earnings in county of residence in 1991, and unemployment in state of residence in 1991, along with an indicator for urban residence at 17 and cohort dummies (see CHV for variable definitions and construction). Also following CHV, I include as components of X_i the squares of continuous control variables, relaxing the assumption of strict linearity in each.²⁸ I implement estimator (11), and compute standard errors by applying the delta method to the system of estimated regression equations (allowing for heteroscedasticity and cross-correlation between the equations).²⁹

²⁸The above variables represent the union of variables that CHV use in their first stage and outcome equation, with one exception: I drop years of experience in 1991 since it may itself be affected by schooling, as MTW do as well in their empirical application. In the two instrument setup, I also add to X_i the two "unused" instruments from CHV and their squares: long-run local earnings in county of residence at 17 and long run unemployment in state of residence at 17.

²⁹Note that while Appendix C Theorem 3 provides an expression for the asymptotic variance of $\hat{\rho}(\hat{\lambda})$ that could be estimated, this result does not cover the case with covariates. Also, as the distribution of Z_i is fairly well balanced across the four cells of \mathcal{Z} in this application, I do not implement the regularization procedure proposed in Appendix C.

7.1.1 Results from baseline setup with two instruments

The left panel of Table 4 reports a cross tabulation of the two instruments. As noted, the observations are relatively evenly distributed across the four cells. The instruments are positively correlated, with a Pearson correlation coefficient of about 0.13.

		Distribution of the instruments		Mean fitted propensity scores			
		$Z_2 = \text{"close"}$		Z_2			
		0	1	$Z_1 = \text{"cheap"}$	far	expensive	close
0		469	401		0.451	0.509	
1		361	516		0.487	0.530	

Table 4: Left: number of observations having each pair of values of the instruments, with total sample size $N = 1,747$. Right: fitted propensity scores estimated by OLS, evaluated at the sample mean of the X_i variables.

The right panel of Table 4 reports predictions from the estimated conditional propensity score function $\mathbb{E}[D_i|Z_i = \mathbf{z}, X_i = x]$ estimated via a linear regression of D_i on the instruments (and their interaction) as well as X_i , then evaluated at the mean of X_i . This allows us to take the (*expensive, far*) cell 45.1% as an estimate of the overall proportion of never-takers in the population, while the share of never-takers is estimated to be 47.0%. The remaining roughly 8% of the population are generalized “compliers” consisting of the tuition (Z_1), proximity (Z_2), eager and reluctant compliers. From the table we can also see that $P(D_i(\text{expensive, close}, x) > D_i(\text{expensive, far}, x)) \approx 5.7\%$, and $P(D_i(\text{cheap, far}, x) > D_i(\text{expensive, far}, x)) \approx 3.6\%$. Combining these figures and the response group definitions from Section 4, we see that between 1.5% and 3.6% of the population are eager compliers, while no more than 2.1% are reluctant compliers. Similarly, no more than 3.6% are tuition compliers, and between 2.1% and 5.7% are proximity compliers. Overall, the data are compatible with a roughly even split between the four groups, but it is also possible that proximity compliers account for more than half of all generalized compliers.

We now turn to treatment effects. Figure 2 reports estimates of several of the parameters introduced in Section 5, alongside fully-saturated 2SLS for comparison. Consider first the ACL: the point estimate of 0.14 indicates that having attended a year of college increases 1991 wages of all compliers by roughly 14% on average. This estimate is within the range of roughly -0.1 to 0.3 of the marginal treatment effect (MTE) function estimated by CHV under the assumption of IAM, and is similar to their point estimate of the average treatment on the treated under a

parametric normal selection model. The 2SLS estimate from Figure 2 yields a similar value at 0.12. Note that given the limited sample size none of the estimates are quite significant at even the 90% level. I focus discussion on the point estimates for the sake of illustration with this important caveat.

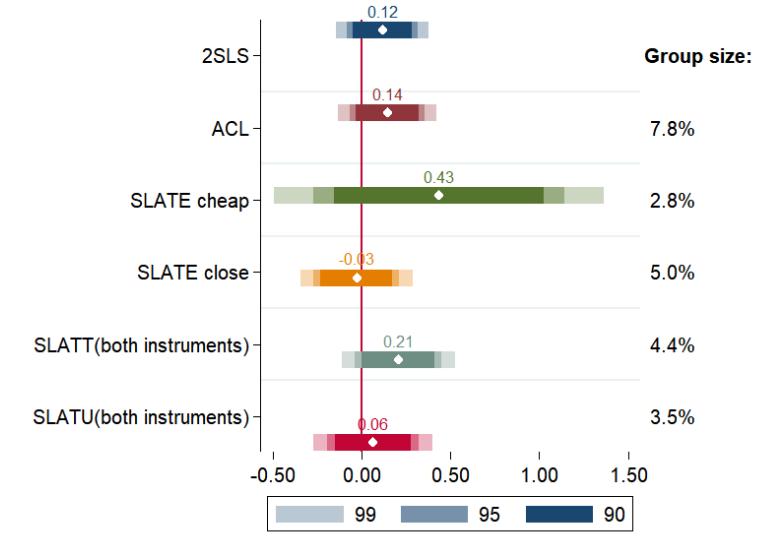


Figure 2: Estimates of various causal parameters identified under VM with two instruments, alongside fully-saturated 2SLS for comparison. Bars indicate 95% confidence intervals, and “Group Size” refers to the identified quantity $P(C_i = 1)$ for each parameter.

The point estimates from the remaining rows in Figure 2 suggest that the ACL aggregates over substantial heterogeneity in the population. For example, the tuition SLATE suggests that a year of college has no average effect on the wages of individuals who move into treatment if and only if a college is nearby, given local affordability. Recall that this group includes proximity compliers, eager compliers for whom college is expensive, and reluctant compliers for whom it is cheap. On the other hand, the SLATE for tuition is about three times as large as the ACL. These results are suggestive that the average treatment effect among tuition compliers is larger than it is among proximity compliers, however the sign of the difference is not identified.³⁰ Note finally that the point estimates for *SLATU* and *SLATT* suggest that among the compliers averaged over by the ACL, those who in fact go to college have greater treatment effects on average than those who do not, which is consistent with some students selecting on the basis of their heterogeneous future gains (as in a Roy-type model).

³⁰I note that in the $J = 2$ case, if Δ_g and corresponding group size p_g is known for one group $g \in \mathcal{G}^c$ ex- ante, then the remaining three group specific treatment effects and group sizes can be point identified.

7.1.2 Results with all four instruments

I now add the additional two instruments from CHV, to increase comparability and emphasize the scalability of my proposed methods to multiple instruments.

Accordingly, let Z_{3i} indicate that local earnings in i 's county of residence at 17 is below the sample median, and Z_{4i} indicates that unemployment in i 's state of residence at 17 is above the sample 25% percentile (this threshold is chosen as it yields a stronger predictor of college attendance as compared with using the median). The two local labor market variables and their squares are removed from the vector of controls X_i .

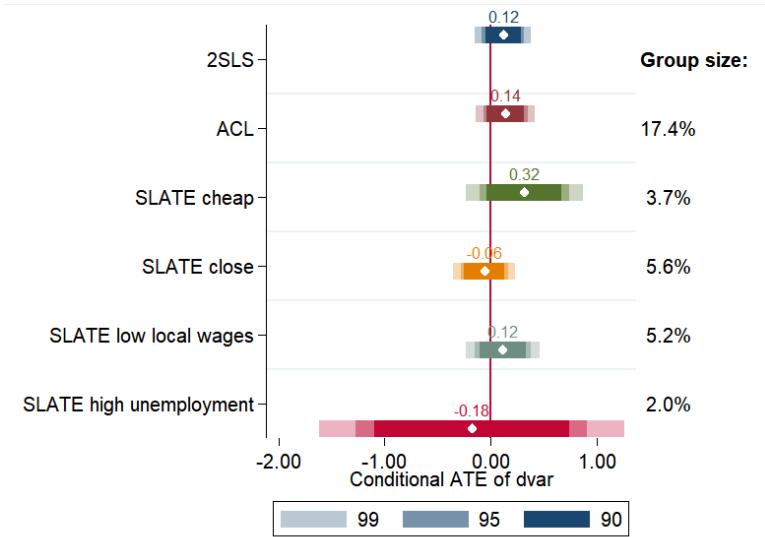


Figure 3: Estimates of various causal parameters identified under VM with all four instruments, alongside fully-saturated 2SLS for comparison. Bars indicate 95% confidence intervals, and “Group Size” refers to the identified quantity $P(C_i = 1)$ for each parameter.

With all four instruments, over 17% of the population are now some type of “complier” and counted in \mathcal{G}^c . However, Table 3 shows that the ACL is not appreciably changed from the case with only two instruments, and we again have that the tuition SLATE is much larger and the proximity SLATE close to zero. The SLATE for low local wages occupies an intermediate value, while the SLATE for high unemployment is estimated to be negative (suggesting that more schooling reduces wages), but with a much larger standard error. The unemployment SLATE is so imprecisely estimated in part because its corresponding complier group is the smallest of the estimands considered: with just 2% of the population.

To compare these results more directly with CHV, recall that the marginal treatment effect function (e.g. Heckman and Vytlacil 2005) is defined as $MTE(u, x) := \mathbb{E}[Y_i(1) - Y_i(0)|U_i =$

$u, X_i = x]$, where U_i is a uniformly distributed heterogeneity parameter that can be thought of as a proclivity against treatment in the selection model

$$D_i(\mathbf{z}, x) = \mathbb{1}(P(\mathbf{z}, x) \geq U_i), \quad (12)$$

with $P(\mathbf{z}, x) := \mathbb{E}[D_i | Z_i = \mathbf{z}, X_i = x]$. Selection model (12)—which Vytlacil (2002) shows is equivalent to Assumption IAM—effectively “flattens” variation in the various instruments into variation in the scalar quantity $P(Z_i, X_i)$. By contrast, VM allows flows into treatment to depend in an essential way on *which* instrument is manipulated, and heterogeneity in selection cannot be summarized by a single “cost” U_i . Meanwhile, this makes the VM estimands reported in Figure 3 directly relevant to hypothetical policies which vary that instrument alone.

The results in Figure 3 can also be compared with estimates reported by Mogstad et al. (2021) that are calculated by 2SLS. While their empirical application focuses on the interpretation of 2SLS under PM (or VM), we have seen that in this particular setting 2SLS tends to yield numerical estimates that are close to the ACL. Similarly, the SLATEs for the proximity and low local wage instruments in Figure 3 align roughly with 2SLS specifications in MTW in which a single instrument is excluded in the second stage. However this similarity will not hold in all contexts: Appendix C provides simulates a data generating process for example in which 2SLS lies outside of the convex hull of treatment effects for any individual in the population.

Finally, observe that in this four instrument setup, there are in principle 167 underlying response groups aside from always- and never-takers, and that together these comprise 17.4% of the population (cf. 7.8% for the four such groups with two instruments). Nevertheless, computing the treatment effect estimates involves regressions with at most 16 terms in addition to the controls, keeping implementation manageable. While the standard errors for the 2SLS estimate are only slightly smaller than for the *ACL*, note that this is sufficient for significance at the 95% level, even in this small sample. This in part reflects the fact 2SLS weighs across the groups to minimize variance rather than pin down a specific target parameter.

7.2 The effect of children on labor supply

In this section I revisit the analysis of Angrist and Evans (1998), who study the effect of family size on parental labor supply. Angrist and Evans consider two types of instruments that induce exogenous variation in family size among families that have at least two children. The

instruments are: i) whether the first two children have the same sex; and ii) whether the second birth was a multiple birth (i.e. twins, triplets, etc.). Since twins account for the overwhelming majority of multiple births, I refer to multiple births as simply “twins”.

Let us first consider instruments, based on the first two children’s sexes (interpreted here as sex assigned at birth, which I’ll refer to as “girl” and “boy”). If the first two children in a family have the same sex, this may cause some parents to have a third child in an effort to have children of both sexes. If some parents’ furthermore have a preference for having at least one boy, they may respond only to this same sex instrument if the first two are girls, and if they have a preference for girls may respond only if the first two are boys. These various sex-preferences can be modeled by introducing two separate binary instruments that operate through child-sex preferences: $Z_{1i} = \mathbb{1}(i\text{'s first two children are both girls})$ and $Z_{2i} = \mathbb{1}(i\text{'s first two children are both boys})$.³¹

Let treatment D_i indicate whether mother i has three or more children in total. With the above instruments, vector monotonicity is a reasonable assumption, saying that $D_i(1, 0) > D_i(0, 0)$ and $D_i(0, 1) > D_i(0, 0)$ for all i —no mother would have a third child only when her first two kids are of the opposite sex. VM places no restrictions on $D_i(1, 1)$, since the point $Z_1 = Z_2 = 1$ can be ruled out of the set \mathcal{Z} of possible instrument values. Nevertheless, the two instruments introduce multiple distinct “complier” populations under VM, since some parents’ may specifically seek a girl, some may seek a boy, and some may seek at least one of each. Note that IAM can only hold for these instruments if all mothers who would have a third child with two boys would also have a third child with two girls, or vice versa. This is a strong restriction, which rules out there being some parents who seek at least one girl, and others who seek at least one boy.

Because the instrument value $(1, 1)$ does not belong to \mathcal{Z} , VM and PM are equivalent for the two instruments Z_1 and Z_2 introduced above. However twinning (or triplets, etc.) introduces a third binary instrument for family size: $Z_{3i} = \mathbb{1}(i\text{'s } 2^{\text{nd}} \text{ birth a multiple birth})$. Note that all mothers with a multiple birth after their first child have at least three children. This implies additional restrictions on the conceivable response groups beyond those arising from VM, leaving five possible response groups as summarized in Table 5.

There are only six columns in Table 5, rather than eight (2^3), because the points $Z_i = (1, 1, 0)$

³¹Note that these definitions differ from those used by MTW in an example of PM, who let Z_1 and Z_2 indicate the sex of the first and second child, respectively.

group name	(0, 0, 0)	(1, 0, 0)	(0, 1, 0)	(0, 0, 1)	(1, 0, 1)	(0, 1, 1)
girl compliers	N	T	N	T	T	T
boy compliers	N	N	T	T	T	T
same-sex compliers	N	T	T	T	T	T
twin compliers	N	N	N	T	T	T
always-takers	T	T	T	T	T	T

Table 5: Response groups under VM for three instruments drawn from Angrist and Evans (1998).

and $Z_i = (1, 1, 1)$ do not appear in \mathcal{Z} .³² Since the support of Z_i is not rectangular (Assumption 3), we will need to appeal to the more general identification result from Appendix A. For this result, the rectangular support condition is replaced by Assumption 3*, which states that there exists a family \mathcal{F} of products of the instruments that are linearly independent and span the space of compliance functions $\mathcal{D}_g(Z_i)$ for $g \in \mathcal{G}^c$ (here \mathcal{G}^c consists of the four groups that are not never-takers). From Table 5, one such spanning family is $\mathcal{F} = \{\{1\}, \{2\}, \{3\}, \{1, 3\}, \{2, 3\}\}$, since we have $\mathcal{D}_{girl}(z) = z_1 + z_3 - z_1 z_3$, $\mathcal{D}_{boy}(z) = z_2 + z_3 - z_2 z_3$, $\mathcal{D}_{twin}(z) = z_3$, and $\mathcal{D}_{samesex}(z) = \mathcal{D}_{girl}(z) + \mathcal{D}_{boy}(z)$. Accordingly, I set $\Gamma_i = (Z_{1i}, Z_{2i}, Z_{3i}, Z_{1i}Z_{3i}, Z_{2i}Z_{3i})'$ when implementing the estimator $\hat{\Delta}_c$.

I use the dataset considered by Angrist and Evans, 1998 drawn from the 1980 U.S. census, creating a sample of 394,840 mothers between the ages of 21 and 35 with multiple children. Table 6 reports the distribution of the instruments and the propensity score function. In contrast to the returns-to-schooling example, I do not condition on any observed covariates (this is done for simplicity, although twin births may not be unconditionally exogenous).

		Z_3 (twins)		Z_2 (boys)	
		0	1	0	1
$(Z_1, Z_2) =$	(0,0)	193,567	1,725	195,292	
	(1,0)	94,618	803	95,421	
	(0,1)	103,275	852	104,127	
		391,460	3,380	Z_1 (girls)	
				0	.336 .418
				1	.436 n/a

Table 6: Cross-tabulation (left) and propensity scores (right) for the three instruments. Propensity scores reported are $\hat{\mathbb{E}}[D_i | Z_i = (z_1, z_2, 0)]$ for $z_1, z_2 \in \{0, 1\}$. For all z_1, z_2 : $\hat{\mathbb{E}}[D_i | Z_i = (z_1, z_2, 1)] = 1$ so these values are not reported. Total $N = 394,840$.

From the left panel of Table 6, we can see that having two boys as one's first two children is more likely than having two girls, whether or not the second birth was a multiple birth. In the right panel, we see that the data is consistent with vector monotonicity as described above. The proportion of always-takers is identified as 33.6%, which given that there are no never-takers implies that the remaining 66.4% of mothers in this population respond to the three instruments

³²In addition reducing the support of the instruments by two points, this also reduces the number of response groups: there are no “reluctant complier” groups for instruments Z_1 and Z_2 .

in some way. The propensity score estimates imply that the total of girl compliers and same-sex compliers is 10% of the population, and that the total of boy compliers and same-sex compliers is 15.7%. This indicates that there are nearly 6 percentage points more boy compliers than girl compliers among U.S. mothers.

I consider four choices of the outcome variable Y_i drawn from Angrist and Evans, 1998: i) the mother's labor income in the year prior to the census (1979 dollars); ii) weeks worked in the year prior to the census; ii) average hours worked per week, and iv) an indicator for whether mother i worked for pay (any of i-iii are positive). Treatment effect estimates are reported in Table 7. Note that in this setting (unlike in the standard setting of Section 5 with full instrument support) the ACL is *not* identified.³³ I report SLATEs for each of the three instruments individually, along with fully-saturated 2SLS for comparison.

	(1) 2SLS	(2) SLATE(girls)	(3) SLATE(boys)	(4) SLATE(twins)
Worked for pay	-0.0889*** (-7.13)	-0.102*** (-3.67)	-0.182*** (-4.94)	-0.106*** (-5.55)
Hours worked for pay	-3.573*** (-7.68)	-3.478** (-3.27)	-7.459*** (-5.29)	-3.740*** (-5.23)
Weeks worked for pay	-3.852*** (-6.99)	-4.122** (-3.29)	-9.486*** (-5.74)	-4.347*** (-5.23)
Labor income (1979 USD)	-512.6*** (-4.07)	-755.1** (-2.67)	-1539.1*** (-4.15)	-617.1** (-3.12)
Size of compliant pop.		0.0700*** (35.93)	0.0515*** (27.34)	0.634*** (578.59)
Observations		394840	394840	394840

t statistics in parentheses

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Table 7: Treatment effect estimates. Each row indicates a different choice of the outcome variable Y_i , while columns (1)-(5) correspond to alternative treatment effect estimators. 2SLS is fully saturated, including interactions between the gender instruments and the twin instruments. Estimates of various Δ_c (columns 2-4) use the estimator $\hat{\Delta}_c$ from Section 6. Size of compliant population reports estimates of $P(C_i = 1)$. $N = 394,840$.

Comparing Columns (2)-(4), we see that the set of compliers that respond to the twins instrument is much larger than those for the boys or girls instrument (63% vs. 5% or 7% of the population); however mothers who respond to the twins instrument have similar treatment effects to those who respond to the boys instrument, across outcomes. By contrast, estimates

³³In this setting, the function $c(g, z) = \mathbb{1}(g \in \mathcal{G}^c)$ does not satisfy Property M like it does when \mathcal{Z} is rectangular: we cannot write $c(g, z)$ as $\mathcal{D}_g(1, 1, 1) - \mathcal{D}_g(0, 0, 0)$ (cf. Proposition 4), since instrument cell (1, 1, 1) is not in \mathcal{Z} .

of the LATE among mothers who respond to the two-boys instrument are much larger in magnitude than the other estimates. This suggests that the mothers most likely to reduce their labor supply—and by much more—are those who seek a girl.

Finally, I note that Assumption 1 along with the condition that all mothers with a multiple second birth take treatment implies an overidentification restriction that

$$\mathbb{E}[Y_i|Z_i = (0, 0, 1)] = \mathbb{E}[Y_i|Z_i = (1, 0, 1)] = \mathbb{E}[Y_i|Z_i = (0, 1, 1)]$$

I test this restriction via an F-test by regressing Y on Z_1 , Z_2 and a constant, restricted to the twins subsample. The model implies that both regression coefficients for the cell indicators should be zero. The p-value for the regression F-statistic is about .05 when Y_i indicates worked for pay, .22 for hours worked for pay, .07 for weeks worked for pay, and .19 for mother's income. This provides some evidence against the validity of the model when the outcome is on the extensive margin (worked for pay). The estimates here should thus be interpreted with caution.³⁴

8 Conclusion

In both observational and experimental settings, it is natural to expect individuals to vary both in their treatment effects and in how they select into treatment. This latter type of heterogeneity is likely to be particularly pronounced when a researcher is using multiple instrumental variables for a single (binary) treatment. This paper has shown that causal inference with heterogeneous treatment effects is possible in such settings under a simple assumption about selection behavior that is often motivated by economic theory: vector monotonicity.

In particular, I have defined and characterized a class of interpretable causal parameters that can be point identified under vector monotonicity with discrete instruments, and proposed an estimator that is similar in construction to the familiar method of two-stage least squares (2SLS). While the convenience of implementing the two estimators scales similarly with the number of instruments, 2SLS is not guaranteed to recover an interpretable causal parameter

³⁴One explanation is that the twin instrument is not actually unconditionally exogenous (note that we have not included any X variables, for simplicity). Another explanation is through number of children: mothers who have gender parity between their first two children as well as twins might not only be more likely to have a third child, but also be more likely to have *four* or more children. This challenges the exclusion restriction in this setting when treatment is considered to be a binary variable for three or more children. A useful extension of the methods in this paper would be to cases with multi-valued, ordered treatment like number of children.

under vector monotonicity (though it may in special cases). By contrast, the estimator I propose is always targets a particular well-defined causal parameter. In applications to the labor market returns to college education and the effect of family size on labor supply, I find that estimates based on VM suggest that underlying groups in the population that exhibit different selection behavior also have very different treatment effects from one another.

References

- ANGRIST, B. J. D. and EVANS, W. N. (1998). "Children and Their Parents' Labor Supply: Evidence from Exogenous Variation in Family Size". *American Economic Review* 88 (3), pp. 450–477.
- ANGRIST, J. D. and PISCHKE, J.-S. (2008). *Mostly Harmless Econometrics*. Princeton, NJ: Princeton University Press.
- CARNEIRO, P., HECKMAN, J. J. and VYTLACIL, E. J. (2011). "Estimating marginal returns to education". *American Economic Review* 101 (6), pp. 2754–2781.
- CHAISEMARTIN, C. de (2017). "Tolerating defiance? Local average treatment effects without monotonicity". *Quantitative Economics* 8 (2), pp. 367–396.
- D'HAULTFŒUILLE, X. and FÉVRIER, P. (2015). "Identification of Nonseparable Triangular Models With Discrete Instruments". *Econometrica* 83 (3), pp. 1199–1210.
- FAN, Y. and PARK, S. S. (2010). "Sharp Bounds on the Distribution of Treatment Effects and Their Statistical Inference". *Econometric Theory* 26 (3), pp. 931–951.
- FENG, J. (2020). "Matching Points : Supplementing Instruments with Covariates in Triangular Models". *Working Paper*. URL: <https://arxiv.org/pdf/1904.01159.pdf>.
- FRÖLICH, M. (2007). "Nonparametric IV estimation of local average treatment effects with covariates". *Journal of Econometrics* 139 (1), pp. 35–75.
- GAUTIER, E. (2020). "Relaxing monotonicity in endogenous selection models and application to surveys". *Working Paper*. URL: <https://arxiv.org/abs/2006.10997>.

- GAUTIER, E. and HODERLEIN, S. (2011). “A triangular treatment effect model with random coefficients in the selection equation”. *Working Paper*. URL: <http://arxiv.org/abs/1109.0362>.
- GUNSLIUS, F. F. (2020). “Identifying multivariate models with binary instruments via cyclically monotone dynamics”. *Working Paper*. URL: <https://drive.google.com/open?id=1Q1Yoqdok9G5hZ8hAqRs6Z-mfnVuiMlng>.
- HECKMAN, J. J. and PINTO, R. (2018). “Unordered Monotonicity”. *Econometrica* 86 (1), pp. 1–35.
- HECKMAN, J. J. and VYTLACIL, E. (2005). “Structural Equations, Treatment Effects, and Econometric Policy Evaluation”. *Econometrica* 73 (3), pp. 669–738.
- HOERL, A. and KENNARD, R. (1970). “Ridge regression : Biased estimation for nonorthogonal problems”. *Technometrics* 42 (1), pp. 80–86.
- IMBENS, B. G. W. and NEWEY, W. K. (2009). “Identification and Estimation of Triangular Simultaneous Equations Models Without Additivity”. *Econometrica* 77 (5), pp. 1481–1512.
- IMBENS, G. W. and ANGRIST, J. D. (1994). “Identification and Estimation of Local Average Treatment Effects”. *Econometrica* 62 (2), pp. 467–475.
- KISIELEWICZ, A. (1988). “A solution of Dedekind’s problem on the number of isotone Boolean functions.” *Journal fur die Reine und Angewandte Mathematik* 386, pp. 139–144.
- KLEITMAN, D. J. and MILNER, E. C. (1973). “On the average size of the sets in a Sperner family”. *Discrete Mathematics* 6 (2), pp. 141–147.
- LEE, S. and SALANIÉ, B. (2018). “Identifying Effects of Multivalued Treatments”. *Econometrica* 86 (6), pp. 1939–1963.
- (2020). “Filtered and Unfiltered Treatment Effects with Targeting Instruments”. *Working Paper*. URL: <https://arxiv.org/abs/2007.10432>.

- LEWBEL, A. and YANG, T. T. (2016). “Identifying the average treatment effect in ordered treatment models without unconfoundedness”. *Journal of Econometrics* 195 (1), pp. 1–22.
- MANSKI, C. F. (1990). “Nonparametric Bounds on Treatment Effects”. *The American Economic Review* 80 (2), pp. 829–823.
- MOGSTAD, M., SANTOS, A. and TORGOVITSKY, A. (2018). “Using Instrumental Variables for Inference About Policy Relevant Treatment Parameters”. *Econometrica* 86 (5), pp. 1589–1619.
- MOGSTAD, M., TORGOVITSKY, A. and WALTERS, C. (2021). “The Causal Interpretation of Two-Stage Least Squares with Multiple Instrumental Variables”. *American Economic Review* 111 (11), pp. 3663–98.
- (2022). “Policy Evaluation with Multiple Instrumental Variables”. *Journal of Econometrics* Forthcoming.
- MOUNTJOY, J. (2022). “Community Colleges and Upward Mobility”. *American Economic Review* 112 (8), pp. 2580–2630.
- SUN, Z. and WÜTHRICH, K. (2022). *Pairwise Valid Instruments*. URL: <https://arxiv.org/abs/2203.08050>.
- TORGOVITSKY, A. (2015). “Identification of Nonseparable Models Using Instruments With Small Support”. *Econometrica* 83 (3), pp. 1185–1197.
- VYTLACIL, E. (2002). “Independence , Monotonicity , and Latent Index Models : An Equivalence Result”. *Econometrica: Notes and Comments* 70 (1), pp. 331–341.
- YIN, J. et al. (2010). “Nonparametric covariance model”. *Statistica Sinica* 20 (1), pp. 469–479.

Appendices

A Identification result without rectangular support

This section provides an extension of Theorem 1 for cases when the support \mathcal{Z} of the instruments is not rectangular (i.e. $\text{supp}(Z_i) \neq \mathcal{Z}_1 \times \mathcal{Z}_2 \times \dots \times \mathcal{Z}_J$), and there may be perfect linear dependencies between the instruments (of the form, for example, that would arise from the mapping from discrete to binary instruments presented in Proposition 2).

A weaker version of Assumption 3 is comprised of the following two conditions. Define $Z_{\emptyset i}$ to be a degenerate random variable that takes the value of one with probability one.

Assumption 3a* (existence of instruments). *There exists a family \mathcal{F} of subsets of the instruments $S \subseteq \{1 \dots J\}$, where $\emptyset \in \mathcal{F}$ and $|\mathcal{F}| > 1$, such that random variables Z_{Si} for all $S \in \mathcal{F}$ are linearly independent, i.e. $P(\sum_{S \in \mathcal{F}} \omega_S \cdot Z_{Si} = 0) < 1$ for all vectors $\omega \in \mathbb{R}^{|\mathcal{F}|}/\mathbf{0}$.*

Assumption 3b* (non-degenerate subsets generate the response groups). *There exists a family \mathcal{F} satisfying Assumption 3a*, such that for any $S \notin \mathcal{F}$, $g(F) \notin \mathcal{G}$ for all Sperner families F that contain S .*

Assumption 3a* is in itself very weak, requiring only that there exists some product of the instruments that has strictly positive variance. Assumption 3b* is much more restrictive: it says that selection functions \mathcal{D}_g for all response groups aside from never-takers can be generated from such linearly independent products over subsets of the instruments.

The construction in Proposition 2 mapping discrete instruments to binary instruments yields a case where Assumption 3* will hold, given rectangular support of the original discrete instruments.

Proposition. *Let each Z_j have M_j ordered points of support $z_1^j < z_2^j \dots < z_{M_j}^j$ and let $\tilde{Z}_m^j = \mathbb{1}(Z_{ji} \geq z_m^j)$. If $P(Z_i = \mathbf{z}) > 0$ for $\mathbf{z} \in (\mathcal{Z}_1 \times \mathcal{Z}_2 \times \dots \times \mathcal{Z}_J)$, then Assumption 3* holds with \mathcal{F} the family of all subsets of $\mathcal{M} := \{\tilde{Z}_m^j\}_{\substack{j \in \{1 \dots J\} \\ m=2 \dots M_j}}$ containing at most one Z_m^j for any given $j \in \{1 \dots J\}$.*

Proof. See Appendix E. □

To illustrate the usefulness of the above proposition, consider a case with a single discrete instrument Z_1 having three levels $z_1 < z_2 < z_3$ and instruments 2 to J each binary. Proposition

2 shows that if $Z_1 \dots Z_J$ satisfies VM then so does the set of $J+1$ instruments $\tilde{Z}_2, \tilde{Z}_3, Z_2, \dots Z_J$ where $\tilde{Z}_2 = \mathbb{1}(Z_{1i} \geq z_2^1)$ and $\tilde{Z}_3 = \mathbb{1}(Z_{1i} \geq z_3^1)$. In this case there are 2^{J-1} “redundant” simple response groups vis-a-vis Assumption 3, since for any $S \subseteq \{2 \dots J\}$: $\tilde{Z}_{2i}\tilde{Z}_{3i}Z_{Si} = \tilde{Z}_{3i}Z_{Si}$.

In this example, the vector Γ_i would contain all non-null subsets of $\{\tilde{Z}_2, \tilde{Z}_3, Z_2, \dots Z_J\}$ that do not contain both of \tilde{Z}_2 and \tilde{Z}_3 . In general, \mathcal{F} can be constructed by considering all subsets of the instruments, and for each subset considering all possible assignments of a value to each instrument, with one fixed value for each instrument omitted from consideration throughout. Provided rectangular support on the original instruments, Assumption 3* then follows with this choice of \mathcal{F} , for which a generalized version of Theorem 1 can be stated:

Theorem 1*. *The results of Theorem 1 hold under Assumption 3* replacing Assumption 3, where now $\Gamma_i := \{Z_{Si}\}_{S \in \mathcal{F}, S \neq \emptyset}$, $\lambda := \{\mathbb{E}[c(g(S), Z_i)]\}_{S \in \mathcal{F}, S \neq \emptyset}$ and again $h(Z_i) = \lambda'\Sigma^{-1}(\Gamma_i - \mathbb{E}[\Gamma_i])$ with $\Sigma := \text{Var}(\Gamma_i)$, for any family \mathcal{F} satisfying Assumption 3*.*

Proof. Identical to that of Theorem 1, except as noted therein. \square

Theorem 1* may also be useful in other cases in which the practitioner has auxiliary knowledge that some of the response groups are not present in the population, or are ruled out on conceptual grounds (for example in the second empirical application of Section 7). Note that, as discussed in that example, a parameter Δ_c that satisfies Property M when \mathcal{Z} is rectangular (such as the ACL) may violate Property M when \mathcal{Z} is not. To verify Property M in a given empirical context, Proposition 4 may be useful.

B Identification and estimation with covariates

This section discusses how one can accommodate covariates that need to be conditioned on for the instruments to be valid. In practice, it is often easier to justify a *conditional* version of Assumption 1:

$$\{(Y_i(1), Y_i(0), G_i) \perp Z_i\} | X_i$$

in which X is a set of observed covariates unaffected by treatment (X may be scalar or a vector). VM is assumed as before, that $D_i(\mathbf{z}) \geq D_i(\mathbf{z}')$ for all i , whenever $\mathbf{z} \geq \mathbf{z}'$ componentwise. With binary instruments, this implies that we may continue to take the “1” value of each instrument to be the direction in which potential treatments are including, regardless of the value of X_i .

If Assumption 3 and Property M each hold conditional on $X_i = x$, then Theorem 1 implies that we can identify $\Delta_c(x) := \mathbb{E}[\Delta_i|C_i = 1, X_i = x]$ for Δ_c satisfying Property M, from the distribution of $(Y_i, Z_i, D_i)|X_i = x$. In particular, the function $h(\mathbf{z})$ from Theorem 1 will now depend on the conditioning value of X_i :

$$h(Z_i, x) = \lambda(x)'Var(\Gamma_i|X_i = x)^{-1}(\Gamma_i - \mathbb{E}[\Gamma_i|X_i = x])$$

for each $x \in \mathbb{X}$, where recall that Γ_i is a vector of products Γ_{Si} of Z_{ji} within subsets of the instruments, where S indexes such subsets. Here we define $\lambda_S(x) = \mathbb{E}[c(g(S), Z_i)|X_i = x]$ —which is identified—for each simple response group $g(S)$. Under these assumptions, we have that $\Delta_c(x) = \mathbb{E}[h(Z_i, x)Y_i|X_i = x]/\mathbb{E}[h(Z_i, x)D_i|X_i = x]$.

If the support of X_i corresponds to a small number of “covariate-cells”, it might be feasible to repeat the entire estimation on fixed-covariate subsamples, to estimate $\Delta_c(x)$ for each x . If the number of groups is large, or if X_i includes continuous variables, estimation of $\Delta_c(x)$ could still in principle be implemented by nonparametric regression of each component of Γ_i on X_i as well as nonparametrically estimating the conditional variance-covariance matrix $Var(\Gamma_i|X_i = x)$ (Yin et al. (2010) describe a kernel-based method for this). The vector $\lambda(x)$ can also be computed via nonparametric regression.

When the object of interest is simply the unconditional version of Δ_c , the conditional quantities become nuisance parameters. Notably, they can be integrated over separately in the numerator and the denominator of the empirical estimand. To see this, write:

$$\begin{aligned}\Delta_c &= \mathbb{E}[\Delta_i|C_i = 1] = \int dF_{X|C}(x|1)\Delta_c(x) \\ &= \int dF_{X|C}(x|1)\frac{\mathbb{E}[h(Z_i, x)Y_i|X_i = x]}{\mathbb{E}[h(Z_i, x)Y_i|X_i = x]} = \int dF_{X|C}(x|1)\frac{\mathbb{E}[h(Z_i, x)Y_i|X_i = x]}{P(C_i = 1|X_i = x)} \\ &= \frac{1}{P(C_i = 1)} \int dF_X(x)\mathbb{E}[h(Z_i, X_i)Y_i|X_i = x] = \frac{\mathbb{E}[h(Z_i, X_i)Y_i]}{\mathbb{E}[h(Z_i, X_i)D_i]}\end{aligned}$$

where we have used Bayes’ rule and that $P(C_i = 1|X_i = x) = \mathbb{E}[h(Z_i, x)D_i|X_i = x]$ (and hence $P(C_i = 1) = \mathbb{E}[h(Z_i, X_i)D_i]$ as well). This provides a VM analog to a similar result that holds under IAM. In that context, Frölich (2007) shows that this fact can deliver \sqrt{n} -consistency of a nonparametric analog of the Wald ratio.

Note that by the conditional version of Corollary 6 we have that:

$$\Delta_c = \frac{\mathbb{E}[\tilde{\lambda}(X_i)' A \{\mathbb{E}[Y_i|Z_i = \mathbf{z}, X_i]\}]}{\mathbb{E}[\tilde{\lambda}(X_i)' A \{\mathbb{E}[D_i|Z_i = \mathbf{z}, X_i]\}]}$$

where A is the $2^J \times 2^J$ matrix defined in Corollary 6, $\tilde{\lambda}(x)$ is a vector over $S \subseteq \{1 \dots J\}$ with components $\lambda_S(x)$ for $S \neq \emptyset$ and 0 for $S = \emptyset$, and $\{\cdot\}$ indicates a vector over $z \in \mathcal{Z}$. If the CEFs of Y and D happen to both be separable between Z and X , i.e $\mathbb{E}[Y_i|Z_i = \mathbf{z}, X_i = x] = y(\mathbf{z}) + w(x)$ and $\mathbb{E}[D_i|Z_i = \mathbf{z}, X_i = x] = d(\mathbf{z}) + v(x)$, then the above expression simplifies:

$$\Delta_c = \frac{\mathbb{E}[\tilde{\lambda}(X_i)' A \{y(\mathbf{z})\} + w(X_i)\tilde{\lambda}(X_i)' A \mathbf{1}]}{\mathbb{E}[\tilde{\lambda}(X_i)' A \{d(\mathbf{z})\} + v(X_i)\tilde{\lambda}(X_i)' A \mathbf{1}]} = \frac{\mathbb{E}[\tilde{\lambda}(X_i)' A \{y(\mathbf{z})\}]}{\mathbb{E}[\tilde{\lambda}(X_i)' A \{d(\mathbf{z})\}]}$$

where $\mathbf{1}$ is a vector of ones and we have used that $\tilde{\lambda}(x)' A \mathbf{1} = 0$ for any x .³⁵ Now, since each $\lambda_S(x)$ is defined as $\mathbb{E}[C_i = 1|G_i = g(S), X_i = x]$, its expectation delivers the unconditional analog: $\lambda_S := \mathbb{E}[C_i = 1|G_i = g(S)] = \mathbb{E}[\lambda_S(X_i)]$. Thus we can write $\Delta_c = \frac{\tilde{\lambda}' A \{y(\mathbf{z})\}}{\tilde{\lambda}' A \{d(\mathbf{z})\}}$, i.e. in this separable case the estimand that identifies Δ_c is essentially unchanged from the baseline case without covariates, aside from the need to control semiparametrically for X_i to obtain the functions $y(\mathbf{z})$ and $d(\mathbf{z})$. Given consistent estimators $\hat{y}(\mathbf{z})$ and $\hat{d}(\mathbf{z})$ of the functions $y(\mathbf{z})$ and $d(\mathbf{z})$, we can estimate Δ_c by

$$\hat{\Delta}_c = \frac{\sum_{\mathbf{z} \in \mathcal{Z}} \left(\sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} \hat{\lambda}_S A_{S,\mathbf{z}} \right) \hat{y}(\mathbf{z})}{\sum_{\mathbf{z} \in \mathcal{Z}} \left(\sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} \hat{\lambda}_S A_{S,\mathbf{z}} \right) \hat{d}(\mathbf{z})} \quad (13)$$

The estimates reported in the returns to schooling application of Section 7 use this result, with $w(x)$ and $v(x)$ taken to each be linear in x . Note that as the vector Γ_i contains a full set of interactions between the binary instruments, both $y(\mathbf{z})$ and $d(\mathbf{z})$ are automatically linear in Γ_i . When the functions $w(x)$ and $v(x)$ are also linear in x , Equation (13) is equal to a simple generalization of the estimator $\hat{\rho}(\hat{\lambda})$ from Section 6:

$$\hat{\Delta}_c = \left((0, \hat{\lambda}') (\Gamma' \mathcal{M}_X \Gamma)^{-1} \Gamma' \mathcal{M}_X D \right)^{-1} (0, \hat{\lambda}') (\Gamma' \mathcal{M}_X \Gamma)^{-1} \Gamma' \mathcal{M}_X Y$$

where \mathcal{M}_X is a projection onto the orthogonal complement of the design matrix of X_i . In this

³⁵This follows from the definition of the entries: $A_{S,\mathbf{z}} = \sum_{f \subseteq \mathbf{z}_0} (-1)^{|f|}$ where \mathbf{z}_0 is the set of components of \mathbf{z} that are equal to zero. For any $S \neq \emptyset$, the identity $\sum_{f \subseteq S} (-1)^{|f|} = 0$ implies that $[A\mathbf{1}]_S = \sum_{z_1 \in S} \sum_{f \subseteq (S - z_1)} (-1)^{|f|} = 0$. The first component of $A\mathbf{1}$, corresponding to $S = \emptyset$, does not contribute since the first component of $\tilde{\lambda}(x)$ is always zero, by construction.

case, the only change required to the estimator $\hat{\Delta}_c$ is to augment the linear projections of Y_i and D_i onto the instruments with X_i as additional linear regressors.

C Regularization and asymptotic distribution

In this section I propose a regularization procedure for the estimator presented in Section 6, in order to improve its performance in small samples when Assumption 3 may be close to failing. I then show asymptotically normality of the regularized estimator and give an expression for the variance, based on a general result from Imbens and Angrist (1994).

C.1 Regularization of the estimator

Recall from Section 6 that the simple plug-in estimator of the ACL in fact only uses data at two points in \mathcal{Z} . This issue can be seen as a “near-collinearity” problem: when there are few observations in the points \bar{Z} and \underline{Z} , the $n \times |\mathcal{F}|$ design matrix Γ will have singular values that are close to zero (to see this, note that $\Gamma'\Gamma = A'^{-1}n \cdot \text{diag}\{\hat{P}(Z_i = \mathbf{z})\}A^{-1}$). This observation suggests that the issue might be mitigated by employing a ridge-type shrinkage estimator (see e.g. Hoerl and Kennard, 1970). Accordingly, let us allow in the estimator a regularization parameters α :

$$\hat{\rho}(\hat{\lambda}, \alpha) = \left((0, \hat{\lambda}')(\Gamma'\Gamma + \alpha I)^{-1}\Gamma'D \right)^{-1} (0, \hat{\lambda}')(\Gamma'\Gamma + \alpha I)^{-1}\Gamma'Y \quad (14)$$

The estimator $\hat{\rho}(\hat{\lambda}, \alpha)$ with a choice of $\alpha > 0$ establishes a lower bound of α on the singular values of the matrix Γ .

In the case of the ACL, Corollary 6 can be leveraged to show that $\alpha > 0$ allows the estimator to make use of the full support of Z_i , rather than just the two points \bar{Z} and \underline{Z} . But ridge regression comes at the expense of bias. Proposition 6 below yields a means of navigating this trade-off to choose α in practice. In particular, I propose choosing α to minimize a feasible estimator of the conditional MSE $\mathbb{E}[(\hat{\rho}(\lambda, \alpha) - \Delta_c)^2 | Z_1 \dots Z_n]$.

Proposition 6. *Under the assumptions of Theorem 1, $\mathbb{E}[(\hat{\rho}(\lambda, \alpha) - \Delta_c)^2 | Z_1 \dots Z_n]$ is, up to*

second order in estimation error and a positive constant of proportionality:

$$\begin{aligned} \tilde{\lambda}'(\Gamma'\Gamma + \alpha I)^{-1} & \left\{ \Gamma'(\Omega_Y + \Delta_c^2\Omega_D - 2\Delta_c\Omega_{YD})\Gamma \right. \\ & \left. + \alpha^2(\beta_Y\beta'_Y + \Delta_c^2\beta_D\beta'_D - 2\Delta_c\beta_Y\beta'_D) \right\} (\Gamma'\Gamma + \alpha I)^{-1}\tilde{\lambda} \end{aligned} \quad (15)$$

where $\tilde{\lambda} := (0, \lambda')'$, $\beta_Y := \mathbb{E}[\Gamma_i\Gamma'_i]^{-1}\mathbb{E}[\Gamma_iY_i]$, $\beta_D := \mathbb{E}[\Gamma_i\Gamma'_i]^{-1}\mathbb{E}[\Gamma_iD_i]$, and $\Omega_{VW} = \mathbb{E}[(V - \beta_V\Gamma)(W - \beta_W\Gamma)'|\Gamma]$ for $V, W \in \{Y, D\}$, and all expectations are assumed to exist.

Furthermore, if $\hat{\alpha}_{mse}$ is chosen as the smallest positive local minimizer of the following estimate of the above:

$$\hat{M}(\alpha) := (0, \hat{\lambda}')(\Gamma'\Gamma + \alpha I)^{-1} \left\{ n\hat{\Pi} + \alpha^2(\hat{\beta}\hat{\beta}') \right\} (\Gamma'\Gamma + \alpha I)^{-1}(0, \hat{\lambda})'$$

with $\hat{\beta}_V := (\Gamma'\Gamma)^{-1}\Gamma'V$ for each $V \in \{Y, D\}$, $\hat{\Pi} := \frac{1}{n} \sum_i (Y_i - \hat{\beta}_Y\Gamma_i - \frac{(0, \hat{\lambda}')\hat{\beta}_Y}{(0, \hat{\lambda}')\hat{\beta}_D}(D_i - \hat{\beta}_D\Gamma_i))^2\Gamma_i\Gamma'_i$ and $\hat{\beta} := \hat{\beta}_Y - \frac{(0, \hat{\lambda}')\hat{\beta}_Y}{(0, \hat{\lambda}')\hat{\beta}_D}\hat{\beta}_D$ then

$$\hat{\alpha}_{mse}/\sqrt{n} \xrightarrow{p} 0$$

provided that $\hat{\lambda} \xrightarrow{p} \lambda$, $(0, \lambda')\Sigma^{-1}(\beta_Y + \Delta_c\beta_D) \neq 0$.

Proof. See Appendix E. □

The proposed data-driven choice $\hat{\alpha}_{mse}$ estimates the unknown quantities in Eq. (15) based on an initial guess of $\alpha = 0$, and then minimizes with respect to α . This can be seen as a “one-step” version of a more general iterative algorithm in which a value α_t would be used to compute the function $\hat{M}(\alpha)$, which is then minimized to find α_{t+1} and so on until convergence. I implement the single-step version in Appendix C.3, and find that it indeed improves estimation error considerably for the simulation DGPs considered.³⁶

³⁶The reason that my proposed rule evaluates $\hat{\alpha}_{mse}$ as a *local* minimizer of $\hat{M}(\alpha)$ rather than a global minimizer, is that the function $\hat{M}(\alpha)$ is always positive but approaches zero as $\alpha \rightarrow \infty$ (this stands in contrast with the standard case of ridge regression in which regularization bias always grows with α , eventually dominating any efficiency gains from increasing it further). In the present case, the vector $\hat{\beta}$ as defined above and $(0, \hat{\lambda}')'$ are orthogonal (in sample as well as in the population limit), and thus the “(squared) bias” term vanishes as $\alpha \rightarrow \infty$, along with the variance of the regularized estimator (this is roughly analogous to ridge regularizing a vector of regression coefficients when their true values are all zero). Nevertheless, the function $\hat{M}(\alpha)$ does have a well-defined local minimum that achieves a lower value than $\hat{M}(0)$ at some strictly positive α (see Appendix E for details), and this local minimum is shown to provide a helpful guide to choosing α in the simulations of Appendix C. Note that the condition $(0, \lambda')\Sigma^{-1}(\beta_Y + \Delta_c\beta_D) \neq 0$ in Proposition 6 rules out a knife-edge case in which the Hessian of $\hat{M}(\alpha)$ is zero when the other arguments of \hat{M} are evaluated at their probability limits.

C.2 Asymptotic distribution

Consistency and asymptotic normality of the estimator $\hat{\rho}(\hat{\lambda}, \alpha)$ is straightforward. In particular, with $\alpha = 0$ the asymptotic variance can be computed as a special case of Theorem 3 in Imbens and Angrist (1994). In the notation of that paper, we can view estimation of $h(\mathbf{z})$ in our setting as a parametric problem $h(\mathbf{z}) = g(\mathbf{z}, \theta)$ where the parameter vector θ is the mean and variance of Γ_i , along with the vector λ :

$$\theta = (\mu_\Gamma, \Sigma, \lambda)' = (\{\mu_{\Gamma,l}\}_l, \{\Sigma_{lm}\}_{l \leq m}, \{\lambda\}_l)' \text{ with } l, m \in \{1 \dots |\mathcal{F}|\}$$

Then $\hat{\rho}(\lambda, \alpha) = \widehat{\text{Cov}}(g(Z_i, \hat{\theta}), Y_i)/\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), D_i)$, where $\hat{\theta}$ solves a set of moment conditions $\sum_{i=1}^N \psi(Z_i, \hat{\theta}) = 0$ given explicitly in the theorem below.

Theorem 3 below shows that the expression from Imbens and Angrist (1994) for the asymptotic variance of simple IV estimators remains valid for Eq. (14), provided that the sequence of α_n across sample sizes n converges in probability to zero sufficiently quickly. By Proposition 6, we obtain this rate for the “one-step” minimizer of the feasible MSE estimate given in Eq. (15).

Theorem 3. *Under the Assumptions of Theorem 1, if $\alpha_n = o_p(\sqrt{n})$ then*

$$\sqrt{n}(\hat{\rho}(\hat{\lambda}, \alpha_n) - \Delta_c) \xrightarrow{d} N(0, V)$$

where $V = \mathbf{e}_1' \Pi^{-1} \Omega(\Pi')^{-1} \mathbf{e}_1$ (i.e. the top-left element of $\Pi^{-1} \Omega(\Pi')^{-1}$) with:

$$\Omega = \begin{pmatrix} -\mathbb{E}[D_i g(Z_i, \theta)] & -\mathbb{E}[g(Z_i, \theta)] & \mathbb{E}[U_i d_\theta g(Z_i, \theta)] \\ -\mathbb{E}[D_i] & -1 & 0 \\ 0 & 0 & \mathbb{E}[d_{\theta'} \psi(Z_i, \theta)] \end{pmatrix}$$

$$\Pi = \begin{pmatrix} \mathbb{E}[g(Z_i, \theta)^2] & \mathbb{E}[g(Z_i, \theta)U_i] & \mathbb{E}[g(Z_i, \theta)\psi(Z_i, \theta)]' \\ \mathbb{E}[g(Z_i, \theta)U_i] & \mathbb{E}[U_i^2] & \mathbb{E}[U_i \psi(Z_i, \theta)]' \\ \mathbb{E}[g(Z_i, \theta)U_i \psi(Z_i, \theta)] & \mathbb{E}[U_i \psi(Z_i, \theta)] & \mathbb{E}[\psi(Z_i, \theta)\psi(Z_i, \theta)'] \end{pmatrix}$$

so long as Ω and Π are finite and Π has full rank, with the definitions:

$$U_i := Y_i - \mathbb{E}[Y_i] - \Delta_c(D_i - \mathbb{E}[D_i])$$

$$\theta = (\mu_\Gamma, \Sigma, \lambda)' = (\{\mu_{\Gamma,l}\}_l, \{\Sigma_{lm}\}_{l \leq m}, \{\lambda\}_l)'$$

$$g(\mathbf{z}, \theta) = \lambda' \Sigma^{-1} (\Gamma(Z_i) - \mu_\Gamma)$$

$$\psi(Z_i, \theta) = ((\Gamma(Z_i) - \mu_\Gamma)', \{(\Gamma_l(Z_i) - \mu_{\Gamma,l})(\Gamma_m(Z_i) - \mu_{\Gamma,m}) - \Sigma_{lm}\}_{l \leq m}, \{c_l(Z_i) - \lambda_l\}_l)'$$

Here $\Gamma(Z_i) = (\Gamma_1(Z_i) \dots \Gamma_{|\mathcal{F}|}(Z_i))'$ where $\Gamma(Z_i)_l = \mathbf{z}_{S_l, i}$ for some arbitrary ordering S_l of the sets in \mathcal{F} , and $c_l(\mathbf{z}) = c(g(S_l), \mathbf{z})$ (and thus $P(C_i = 1 | G_i = g(S_l)) = \mathbb{E}[c_l(Z_i)]$).

Proof. See Appendix E. □

C.3 Simulation study

This section reports a Monte Carlo experiment in which the regularized estimator proposed above is compared against its unregularized version and 2SLS. I proceed in two steps. In a first simulation involving three binary instruments, I demonstrate the practical importance of regularization. A second simulation with two binary instruments highlights the potential dangers of instead using 2SLS.

Three instrument DGP: We first let $J = 3$, and put equal weight $P(G_i = g) = .05$ over each of the 20 response groups. To introduce endogeneity, I let $Y_i(0) = G_i \cdot U_i$ where the G_i are numbered arbitrarily from one to 20 and $U_i \sim Unif[0, 1]$. The treatment effect within each group g is chosen to be constant and equal to g , so that

$$Y_i(1) = Y_i(0) + G_i + V_i$$

with $V_i \sim Unif[0, 1]$. With this setup, $ACL = 10$.

For the joint distribution of the instruments, I consider two alternatives, meant to capture different extremes regarding statistical dependence among the instruments:

1. (Z_{1i}, Z_{2i}, Z_{3i}) generated as uncorrelated coin tosses
2. (1) followed by the following transformation: if $Z_{2i} = 1$ set $Z_{3i} = 0$ with probability 95%

I let the sample size be $n = 1000$, and perform one thousand simulations. Our primary goal

is to compare the estimator $\hat{\rho}(1, 1, \dots, 1, \alpha)$, where α chosen by the feasible approximate MSE-minimizing procedure described in Section 6, to the simple Wald estimator of ACL ($\hat{\mathbb{E}}[Y_i|Z_i = (111)] - \hat{\mathbb{E}}[Y_i|Z_i = (000)] / (\hat{\mathbb{E}}[D_i|Z_i = (111)] - \hat{\mathbb{E}}[D_i|Z_i = (000)])$), which is equal to $\hat{\rho}(1 \dots 1, \alpha = 0)$. I benchmark both estimators against fully saturated 2SLS.

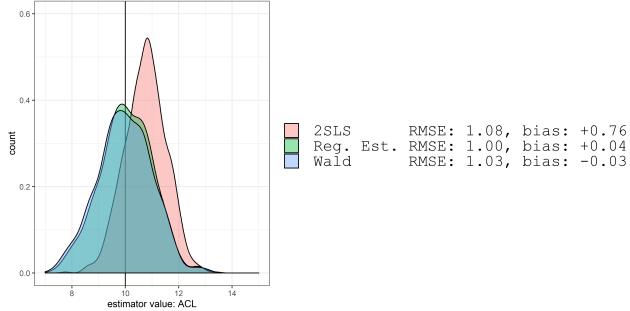


Figure 4: Monte Carlo distributions of estimators, for the first DGP (\mathbf{z} uncorrelated coin tosses) with three binary instruments. “Reg. Est.” indicates $\hat{\rho}(1, \dots, 1, \hat{\alpha}_{mse})$. The vertical line shows the true value of ACL.

Figure 4 shows the results for the first DGP, where the Z_j are independent Bernoulli random variables with mean $1/2$. We see that with the good overlap of the points $\bar{Z} = (1, 1, 1)$ and $\bar{Z} = (0, 0, 0)$ (which are each equal to $1/8$), the Wald estimator performs well. For this DGP, the procedure to choose $\hat{\alpha}_{mse}$, minimizing MSE, results in small values with high probability. Hence the regularized estimator $\hat{\rho}((1, 1, \dots, 1)', \hat{\alpha}_{mse})$ according to Proposition 6 is very close to the Wald estimator (recall that they are identical when $\alpha = 0$). However, the regularized estimator does deliver a slightly smaller RMSE, as expected, at the cost of some bias. Fully saturated 2SLS happens to also perform reasonably well for this DGP.

Figure 5 shows the results for the second DGP, where I modify the joint distribution of (Z_{1i}, Z_{2i}, Z_{3i}) to impose $E(Z_{3i}|Z_{2i} = 1) = 0.05$. In this case, the Wald estimator performs comparatively poorly. We see that regularizing the estimator to use the full sample rather than just the points $\bar{Z} = (1, 1, 1)$ and $\bar{Z} = (0, 0, 0)$ can help considerably.

Two instrument DGP: Note that in both Figures 4 and 5, fully saturated 2SLS (regression on the propensity score) performs well, in the latter case actually outperforming both of the alternative estimators. This is despite the fact that it is not consistent for the ACL , and is in general not even guaranteed to be consistent for Δ_c for any choice of the function $c(g, \mathbf{z})$. To demonstrate that 2SLS can in practice perform very poorly under vector monotonicity, I below report results from an additional simulation in which $J = 2$.

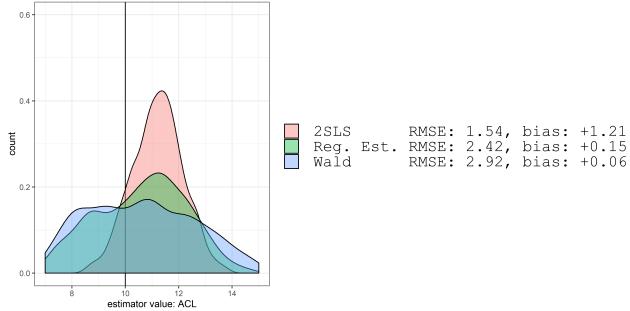


Figure 5: Monte Carlo distributions of estimators, for the first DGP ($P(Z_{3i}|Z_{2i} = 1) = 0.05$) with three binary instruments. “Reg. Est.” indicates $\hat{\rho}(1, \dots, 1, \hat{\alpha}_{mse})$. The vertical line shows the true value of ACL.

For this simulation, the DGP is as follows. Among the six possible response groups under vector monotonicity, I give units a 90% chance of being Z_1 complier and a 10% chance of Z_2 complier. The treatment effect is set to 2 for Z_1 compliers, and -8 for Z_2 compliers, resulting in a *ACL* of unity. I generate negatively correlated binary instruments (with correlation of about -0.1) from a multivariate normal. In particular, with

$$\begin{pmatrix} Z_1^* \\ Z_2^* \end{pmatrix} \sim N \left[\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{bmatrix} 1 & -0.8 \\ -0.8 & 1 \end{bmatrix} \right]$$

I set $Z_{1i} = 1$ when Z_{1i}^* is over its sample median (and $Z_{1i} = 0$ when below), and $Z_{2i} = 1$ when Z_{2i}^* is over its median. I again let the sample size be $n = 1000$, and perform a thousand simulations.

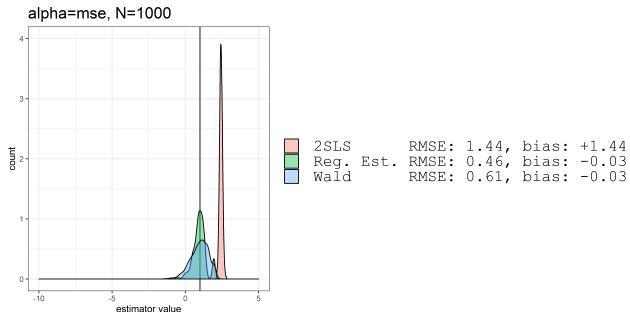


Figure 6: Monte Carlo distributions of estimators, for the two-instrument DGP. “Reg. Est.” indicates $\hat{\rho}(1, \dots, 1, \hat{\alpha}_{mse})$. The vertical line shows the true value of ACL.

Figure 6 shows that in this case, 2SLS is indeed outside of the convex hull of all treatment effects in the population, despite having a small variance. In this setting, my proposed regularized estimator outperforms both of the alternatives by a substantial margin.

D Selection groups in a random coefficients model

The regions (a)-(e) in Figure 1 can be made more concrete by considering a setting of two binary instruments $\mathcal{Z} = \{0, 1\} \times \{0, 1\}$, with a random coefficients threshold crossing model of the form:

$$D_i(z_1, z_2) = \mathbb{1}(\beta_{0i} + \beta_{1i}z_1 + \beta_{2i}z_2 + \beta_{3i}z_1z_2 \geq 0) \quad (16)$$

where $\beta_i = (\beta_{0i}, \beta_{1i}, \beta_{2i}, \beta_{3i})' \perp Z_i$ (Assumption 1). Given the binary treatments, this model is general enough to capture all possible selection functions $D_i(\mathbf{z})$.

Equation (16) could capture a utility maximization model in which individuals trade off an incentive $\beta_{1i}z_1 + \beta_{2i}z_2 + \beta_{3i}z_1z_2$ produced by the instruments against a net cost $-\beta_{0i}$ of treatment. Table 8 gives restrictions on the support of the components of β_i that illustrate each of the points (a)-(e) in Figure 1.

Case	Example of support restriction on β 's	Implied restrictions on selection
(a)	$\beta_1, \beta_2, \beta_3$ homogeneous; $0 \leq \beta_1 \leq \beta_2, \beta_3 = 0$	$D_i(0, 0) \leq D_i(1, 0) \leq D_i(0, 1) \leq D_i(1, 1)$
(b)	$\beta_1, \beta_2, \beta_3$ homogeneous; $-\beta_2 \leq \beta_3 \leq -\beta_1 \leq 0$	$D_i(0, 0) \leq D_i(1, 0) \leq D_i(1, 1) \leq D_i(0, 1)$
(c)	$\beta_{2i} \geq \beta_{1i} \geq 0, -\beta_{2i} \leq \beta_{3i} \leq -\beta_{1i}$ for all i	$D_i(0, 0) \leq D_i(0, 1); D_i(0, 0) \leq D_i(1, 0); D_i(1, 0) \leq D_i(1, 1); D_i(1, 1) \leq D_i(0, 1)$
(d)	$\beta_{3i} = 0, \beta_{1i} \geq 0, \beta_{2i} \geq 0$ for all i $\beta_{2i} < -\beta_{0i} \leq \beta_{1i}$ for some i $\beta_{1i} < -\beta_{0i} \leq \beta_{2i}$ for some i	$D_i(0, 0) \leq D_i(0, 1) \leq D_i(1, 1); D_i(0, 0) \leq D_i(1, 0) \leq D_i(1, 1)$
(e)	a neighborhood of the zero vector in \mathbb{R}^4	none

Table 8: Illustrative examples of each of the cases (a)-(e) in the random coefficients selection model Eq. (16).

Let us now consider the plausibility of the above cases in the returns to schooling example, with “cheap” and “close” the 1 states of Z_1 and Z_2 , respectively. In a utility maximization model β_{0i} might denote the net benefit of attending college when it is far and expensive. If college then became either cheap or close, it is natural to expect this to only increase the net benefit of college, incenting some individuals into enrolling while discouraging none. This motivates the restrictions $\beta_{1i} \geq 0$ and $\beta_{2i} \geq 0$. If we then imagine changing to *(cheap, close)* from either *(expensive, close)* or *(cheap, far)*, it’s reasonable to again expect all students to move weakly towards college, unless there are individuals for whom the interaction coefficient β_{3i} is sufficiently strong and negative.³⁷ The example above for (c) considers for simplicity a

³⁷It is possible to imagine scenarios in which this could happen. As a toy example, suppose there exist students who do not want to live with their parents during college, and feel that they will have to if attending a college near their parents’ home. Accordingly, some such students might go to college only when it is cheap and far. Note that in this case, PM would then require that there be no other individuals in the population that go to

case in which such interactions are totally absent ($\beta_{3i} = 0$ for all i), but this is by no means necessary for VM to hold.

E Proofs

This section provides proofs for the formal results presented in the body of the paper.

The following Lemma gives an alternative characterization of VM, similar in form to the definition of PM, that is useful in establishing Proposition 1.

Lemma. *Let \mathcal{Z} be connected. Then VM holds iff for each $j \in \{1 \dots J\}$ there is an ordering \geq_j on \mathcal{Z}_j such that the following holds for all i : for all $z_j, z'_j \in \mathcal{Z}_j$ and $\mathbf{z}_{-j} \in \mathcal{Z}_{-j}$ such that both $(z_j, \mathbf{z}_{-j}) \in \mathcal{Z}$ and $(z'_j, \mathbf{z}_{-j}) \in \mathcal{Z}$, $D_i(z_j, \mathbf{z}_{-j}) \geq D_i(z'_j, \mathbf{z}_{-j})$ when $z_j \geq_j z'_j$.*

Proof. To simplify notation take each ordering \geq_j to be the ordering on the natural numbers \geq , without loss of generality. The two versions of VM are:

Assumption VM (vector monotonicity). *For $\mathbf{z}, \mathbf{z}' \in \mathcal{Z}$, if $\mathbf{z} \geq \mathbf{z}'$ component-wise, then $D_i(\mathbf{z}) \geq D_i(\mathbf{z}')$ for all i .*

Assumption VM' (alternative characterization). *$D_i(\mathbf{z}_j, \mathbf{z}_{-j}) \geq D_i(\mathbf{z}'_j, \mathbf{z}_{-j})$ for all i when $z_j \geq z'_j$ and both $(\mathbf{z}_j, \mathbf{z}_{-j})$ and $(\mathbf{z}'_j, \mathbf{z}_{-j}) \in \mathcal{Z}$*

The claim is that $VM \iff VM'$.

- $VM \implies VM'$: immediate, since $(\mathbf{z}_j, \mathbf{z}_{-j}) \geq (\mathbf{z}'_j, \mathbf{z}_{-j})$ in a vector sense when $z_j \geq z'_j$
- $VM' \implies VM$: consider $\mathbf{z}, \mathbf{z}' \in \mathcal{Z}$ such that $\mathbf{z} \geq \mathbf{z}'$ in a vector sense, i.e. $z_j \geq z'_j$ for all $j \in \{1 \dots J\}$. Then by VM' and connectedness of \mathcal{Z} , then for some ordering of the instrument labels $1 \dots J$:

$$D_i \begin{pmatrix} z_1 \\ z_2 \\ \vdots \\ z_J \end{pmatrix} \geq D_i \begin{pmatrix} z'_1 \\ z_2 \\ \vdots \\ z_J \end{pmatrix} \quad \forall i, \quad D_i \begin{pmatrix} z'_1 \\ z_2 \\ \vdots \\ z_J \end{pmatrix} \geq D_i \begin{pmatrix} zW'_1 \\ z'_2 \\ \vdots \\ z_J \end{pmatrix} \quad \forall i, \quad etc \dots$$

college only if it is both cheap and close.

and thus:

$$D_i \begin{pmatrix} z_1 \\ z_2 \\ \vdots \\ z_J \end{pmatrix} \geq D_i \begin{pmatrix} z'_1 \\ z_2 \\ \vdots \\ z_J \end{pmatrix} \geq D_i \begin{pmatrix} z'_1 \\ z'_2 \\ \vdots \\ z_J \end{pmatrix} \geq \dots \geq D_i \begin{pmatrix} z'_1 \\ z'_2 \\ \vdots \\ z'_J \end{pmatrix} \text{ for all } i$$

□

E.1 Proof of Proposition 1

Let $P(\mathbf{z}) := \mathbb{E}[D_i | Z_i = \mathbf{z}]$ be the propensity score function. By the law of iterated expectations and Assumption 1:

$$P(\mathbf{z}) = \sum_{g \in \mathcal{G}} P(G_i = g | Z_i = \mathbf{z}) \mathbb{E}[D_i(Z_i) | G_i = g, Z_i = \mathbf{z}] = \sum_{g \in \mathcal{G}} P(G_i = g) \mathcal{D}_g(\mathbf{z})$$

By VM, $\mathcal{D}_g(\mathbf{z})$ is component-wise monotonic for any g in the support of G_i . As a convex combination of component-wise monotonic functions, $P(\mathbf{z})$ will thus also be component-wise monotonic.

In the other direction, note that by PM if $P(\mathbf{z}_j, \mathbf{z}_{-j}) > P(\mathbf{z}'_j, \mathbf{z}_{-j})$, then we must have that $D_i(\mathbf{z}_j, \mathbf{z}_{-j}) \geq D_i(\mathbf{z}'_j, \mathbf{z}_{-j})$ rather than $D_i(\mathbf{z}_j, \mathbf{z}_{-j}) \leq D_i(\mathbf{z}'_j, \mathbf{z}_{-j})$. Thus component-wise monotonicity of $P(\mathbf{z})$ with respect to some collection of orderings $\{\geq_j\}_{j \in \{1 \dots J\}}$ implies $D_i(\mathbf{z}_j, \mathbf{z}_{-j}) \geq D_i(\mathbf{z}'_j, \mathbf{z}_{-j})$ for all choices of $j \in \{1 \dots J\}$, $z_j \geq_j z'_j$, and $\mathbf{z}_{-j} \in \mathcal{Z}_{-j}$ (and all i). This is the equivalent form of VM stated in the Lemma above.

E.2 Proof of Proposition 2

Let $\tilde{\mathcal{Z}}$ be the set of possible values for the new set of instruments $(\tilde{Z}_2, \dots, \tilde{Z}_m, Z_{-1})$, where Z_{-1} is a shorthand for (Z_2, \dots, Z_J) . Note that fixing the value of Z_1 is equivalent to fixing the values of all of $\tilde{Z}_2 \dots \tilde{Z}_M$. Since $P(\tilde{Z}_{mi} = 0 \& \tilde{Z}_{ni} = 1) = 0$ for any $m > n$, we may without loss take $\tilde{\mathcal{Z}}$ to consist only of cases where $\tilde{Z}_2 \dots \tilde{Z}_M$ takes the form $(0, \dots, 0, z, 1, \dots, 1)$ for some $z \in \{0, 1\}$. Let \tilde{Z}_{-m} denote all of the instruments in $\tilde{Z}_2 \dots \tilde{Z}_M$ aside from \tilde{Z}_m .

If \mathcal{Z} is connected, then the $\tilde{\mathcal{Z}}$ given above is also connected. Then, by Proposition E, we simply need to show that for any \mathbf{z}_{-1} and $\tilde{\mathbf{z}}_{-m}$ such that $(0, \tilde{\mathbf{z}}_{-m}, \mathbf{z}_{-1}) \in \mathcal{Z}$ and $(1, \tilde{\mathbf{z}}_{-m}, \mathbf{z}_{-1}) \in \mathcal{Z}$:

$$D_i(1, \tilde{\mathbf{z}}_{-m}; \mathbf{z}_{-1}) \geq D_i(0, \tilde{\mathbf{z}}_{-m}; \mathbf{z}_{-1})$$

where the notation $D_i(a, b; c)$ is understood as $D_i(d, c)$ where d is the value of Z_1 corresponding to \tilde{Z} with value a for \tilde{Z}_m and b for \tilde{Z}_{-m} . For any $\tilde{\mathbf{z}}_{-m}$ satisfying $(0, \tilde{\mathbf{z}}_{-m}, \mathbf{z}_{-1}) \in \mathcal{Z}$ and $(1, \tilde{\mathbf{z}}_{-m}, \mathbf{z}_{-1}) \in \mathcal{Z}$, switching \tilde{Z}_m from zero to ones corresponds to switching instrument Z_1 from value z_{m-1} to value z_m . Since

$$D_i(1, \tilde{\mathbf{z}}_{-m}; \mathbf{z}_{-1}) - D_i(0, \tilde{\mathbf{z}}_{-m}; \mathbf{z}_{-1}) = D_i(z_m, \mathbf{z}_{-1}) - D_i(z_{m-1}, \mathbf{z}_{-1}) \geq 0$$

by vector monotonicity on the original vector $(Z_1 \dots Z_J)$, the result now follows.

E.3 Proof of Proposition 3

For any fixed \mathbf{z} , write the condition $\mathcal{D}_{g(F)}(\mathbf{z}) = 1$ as

$$\{\mathcal{D}_{g(F)}(\mathbf{z}) = 1\} \iff \left\{ \bigcup_{S \in F} \{\mathcal{D}_{g(S)}(\mathbf{z}) = 1\} \right\} \iff \text{not } \left\{ \bigcap_{S \in F} \{\mathcal{D}_{g(S)}(\mathbf{z}) = 0\} \right\}$$

which can be written as

$$\mathcal{D}_g(\mathbf{z}) = 1 - \prod_{S \in F} (1 - \mathcal{D}_{g(S)}(\mathbf{z})) = \sum_{f \subseteq F: f \neq \emptyset} (-1)^{|f|+1} \prod_{S \in f} \mathcal{D}_{g(S)}(\mathbf{z})$$

Let $\mathbf{z}_1(\mathbf{z}) = \{j \in \{1 \dots J\} : z_j = 1\}$ represent a vector of instrument values \mathbf{z} as the subset of instrument indices for which the associated instrument takes the value of one. Then, using that

for a simple response group $\mathcal{D}_{g(S)}(\mathbf{z}) = \mathbb{1}(S \subseteq \mathbf{z}_1(\mathbf{z}))$:

$$\begin{aligned}
\mathcal{D}_g(\mathbf{z}) &= \sum_{f \subseteq F: f \neq \emptyset} (-1)^{|f|+1} \prod_{s \in f} \mathcal{D}_{g(s)}(\mathbf{z}) \\
&= \sum_{f \subseteq F: f \neq \emptyset} (-1)^{|f|+1} \cdot \mathcal{D}_{g((\bigcup_{s \in f} s))}(\mathbf{z}) \\
&= \sum_{f \subseteq F: f \neq \emptyset} (-1)^{|f|+1} \cdot \mathbb{1}\left(\left(\bigcup_{s \in f} s\right) \subseteq \mathbf{z}(\mathbf{z})\right) \\
&= \sum_{\substack{\emptyset \subset f \subseteq F: \\ (\bigcup_{s \in f} s) \subseteq \mathbf{z}(\mathbf{z})}} (-1)^{|f|+1} = \sum_{S' \subseteq \mathbf{z}(\mathbf{z})} \sum_{\substack{\emptyset \subset f \subseteq F: \\ (\bigcup_{s \in f} s) = S'}} (-1)^{|f|+1} \\
&= \sum_{S' \subseteq \{1 \dots J\}} \mathbb{1}(S' \subseteq \mathbf{z}(\mathbf{z})) \sum_{\substack{\emptyset \subset f \subseteq F: \\ (\bigcup_{s \in f} s) = S'}} (-1)^{|f|+1} \\
&= \sum_{S' \subseteq \{1 \dots J\}} \left[\sum_{\substack{\emptyset \subset f \subseteq F: \\ (\bigcup_{s \in f} s) = S'}} (-1)^{|f|+1} \right] \mathcal{D}_{g(S')}(\mathbf{z}) = \sum_{\emptyset \subset S' \subseteq \{1 \dots J\}} \left[\sum_{\substack{f \subseteq F: \\ (\bigcup_{s \in f} s) = S'}} (-1)^{|f|+1} \right] \mathcal{D}_{g(S')}(\mathbf{z})
\end{aligned}$$

Thus, letting $s(F, S') := \left\{ f \subseteq F : (\bigcup_{s \in f} s) = S' \right\}$, we have $\mathcal{D}_{g(F)}(\mathbf{z}) = \sum_{S'} [M_J]_{F, S'} \mathcal{D}_{g(S')}(\mathbf{z})$, where the sum ranges over non-null subsets of the instruments $\emptyset \subset S' \subseteq \{1 \dots J\}$ and $[M_J]_{F, S'} = \sum_{f \in s(F, S')} (-1)^{|f|+1}$.

E.4 Proof of Proposition 4

The if direction is most straightforward. From Proposition 3 we have that for any $\mathbf{z} \in \mathcal{Z}$ and $g \in \mathcal{G}^c$:

$$\mathcal{D}_g(\mathbf{z}) = \sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} [M_J]_{F(g), S} \cdot \mathcal{D}_{g(S)}(\mathbf{z})$$

Thus, for any such $c(g, \mathbf{z})$:

$$\begin{aligned}
c(g, \mathbf{z}) &= \sum_{k=1}^K \sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} [M_J]_{F(g), S} \cdot \mathcal{D}_{g(S)}(h_k(\mathbf{z})) - \sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} [M_J]_{F(g), S} \cdot \mathcal{D}_{g(S)}(l_k(\mathbf{z})) \\
&= \sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} [M_J]_{F(g), S} \cdot \left\{ \sum_{k=1}^K \mathcal{D}_{g(S)}(h_k(\mathbf{z})) - \mathcal{D}_{g(S)}(l_k(\mathbf{z})) \right\} \\
&= \sum_{S \subseteq \{1 \dots J\}, S \neq \emptyset} [M_J]_{F(g), S} \cdot c(g(S), \mathbf{z})
\end{aligned}$$

for any $\mathbf{z} \in \mathcal{Z}$. To finish verifying Property M, we need only observe that $c(a.t., \mathbf{z}) = c(n.t., \mathbf{z}) = 0$ for all \mathbf{z} since $\mathcal{D}_g(h_k(\mathbf{z})) = \mathcal{D}_g(l_k(\mathbf{z}))$ for any h_k, l_k when $g \in \{a.t., n.t.\}$.

Now we turn to the other implication of the Proposition, that any c satisfying Property M has a representation like the above. For shorthand, let $c^{-1}(\mathbf{z})$ indicate the family of $S \subseteq \{1 \dots J\}$ such that $c(g(S), \mathbf{z}) = 1$. The following Lemma establishes that the family $c^{-1}(\mathbf{z})$ and its complement are each closed under unions:

Lemma. *Let c be a function from $\mathcal{G} \times \mathcal{Z}$ to $\{0, 1\}$ satisfies Property M. If $A \in c^{-1}(\mathbf{z})$ and $B \in c^{-1}(\mathbf{z})$, then $A \cup B \in c^{-1}(\mathbf{z})$, and if $A \notin c^{-1}(\mathbf{z})$ and $B \notin c^{-1}(\mathbf{z})$, then $A \cup B \notin c^{-1}(\mathbf{z})$.*

Proof. If the sets A and B are nested, then the result follows trivially. Now suppose neither set contains the other, and consider the Sperner family $A \sqcup B$ constructed of the two sets A and B . By Property M and using Proposition 3:

$$\begin{aligned} c(g(A \sqcup B), \mathbf{z}) &= \sum_{\emptyset \subset S' \subseteq \{1 \dots J\}} \left[\sum_{\substack{f \subseteq \{A, B\}: \\ (\bigcup_{S \in f} S) = S'}} (-1)^{|f|+1} \right] c\left(\bigcup_{S \in f} S, \mathbf{z}\right) \\ &= \sum_{\emptyset \subset f \subseteq \{A, B\}} c\left(\bigcup_{S \in f} S, \mathbf{z}\right) \\ &= c(g(A), \mathbf{z}) + c(g(B), \mathbf{z}) - c(g(A \cup B), \mathbf{z}) \end{aligned}$$

In the first case, if both A and B are in $c^{-1}(\mathbf{z})$, then we must have $c(g(A \cup B), \mathbf{z}) = 1$ to prevent $c(g(A \sqcup B), \mathbf{z})$ from evaluating to 2, which contradicts the assumption that c takes values in $\{0, 1\}$. In the second case, when both $c(g(A), \mathbf{z})$ and $c(g(B), \mathbf{z})$ are zero, we must have $c(g(A \cup B), \mathbf{z}) = 1$ to prevent $c(g(A \sqcup B), \mathbf{z})$ from evaluating to -1. \square

As a consequence of the Lemma, since $c^{-1}(\mathbf{z})$ is a finite set, there exists a member $S_1(\mathbf{z})$ of $c^{-1}(\mathbf{z})$ that satisfies $S_1(\mathbf{z}) = \bigcup_{S \in c^{-1}(\mathbf{z})} S$ (similarly, there exists a $S_0(\mathbf{z}) = \bigcup_{S \notin c^{-1}(\mathbf{z})} S$ with $S_0(\mathbf{z}) \notin c^{-1}(\mathbf{z})$). All members of the family $c^{-1}(\mathbf{z})$ are subsets of $S_1(\mathbf{z})$, and all $S \subseteq \{1 \dots J\}$ that are not in $c^{-1}(\mathbf{z})$ are subsets of $S_0(\mathbf{z})$.

Let \mathbf{z} take some fixed value, and beginning with the set $S_1 = S_1(\mathbf{z})$, define a sequence of sets $\{S_1, S_2, S_3, \dots\}$ as follows:

$$S_{2k} = \bigcup_{\substack{S' \subseteq S_{2k-1}: \\ S' \notin c^{-1}(\mathbf{z})}} S' \quad \text{and} \quad S_{2k+1} = \bigcup_{\substack{S' \subseteq S_{2k}: \\ S' \in c^{-1}(\mathbf{z})}} S'$$

where we take $\bigcup_{S' \in \emptyset} S'$ to evaluate to the empty set. This sequence provides a characterization of the family $c^{-1}(\mathbf{z})$ as follows. For any $\emptyset \subset S \subseteq \{1 \dots J\}$:

$$\begin{aligned}
c(g(S), \mathbf{z}) &= \mathbb{1}(S \in c^{-1}(\mathbf{z})) \\
&= \mathbb{1}(S \subseteq S_1 : S \in c^{-1}(\mathbf{z})) \\
&= \mathbb{1}(S \subseteq S_1) - \mathbb{1}(S \subseteq S_1 : S \notin c^{-1}(\mathbf{z})) \\
&= \mathbb{1}(S \subseteq S_1) - (\mathbb{1}(S \subseteq S_2) - \mathbb{1}(S \subseteq S_2 : S \in c^{-1}(\mathbf{z}))) \\
&= \mathbb{1}(S \subseteq S_1) - \mathbb{1}(S \subseteq S_2) + (\mathbb{1}(S \subseteq S_3) - \mathbb{1}(S \subseteq S_3 : S \notin c^{-1}(\mathbf{z}))) \\
&= \dots \\
&= \sum_{n=1}^N (-1)^{n+1} \cdot \mathbb{1}(S \subseteq S_n) + (-1)^N \cdot \begin{cases} \mathbb{1}(S \subseteq S_N : S \in c^{-1}(\mathbf{z})) & \text{if } N \text{ even} \\ \mathbb{1}(S \subseteq S_N : S \notin c^{-1}(\mathbf{z})) & \text{if } N \text{ odd} \end{cases}
\end{aligned}$$

for any natural number N .

Think of the power set of S_1 as a “first-order” approximation to the family $c^{-1}(\mathbf{z})$. However, in most cases this family is too large, as there will be subsets of S_1 that are not found in $c^{-1}(\mathbf{z})$. Define S_2 to be the union of all such offending sets. The power set of S_2 now provides a possible “overestimate” of the family of offending sets (since they are all in 2^{S_2}) and hence removing all subsets of S_2 as a correction to be applied to 2^{S_1} as an estimate of $c^{-2}(\mathbf{z})$ will overcompensate: we will have removed some sets which are indeed in $c^{-1}(\mathbf{z})$. We thus define S_3 analogously, whose power set provides an approximation to the error in S_2 as an approximation to the error in S_1 , and so on.

Does this process of over-correction eventually terminate, so that the final remainder term is zero? Note that for any n : $S_n \subseteq S_{n-1}$. If $S_n = S_{n-1} \neq \emptyset$, then we have a fixed point S where $\bigcup_{S' \subseteq S : S' \in c^{-1}(\mathbf{z})} S' = \bigcup_{S' \subseteq S : S' \notin c^{-1}(\mathbf{z})} S'$. But by the Lemma, this would imply that S is a member both of $\{S' \subseteq S : S' \in c^{-1}(\mathbf{z})\}$ and of $\{S' \subseteq S : S' \notin c^{-1}(\mathbf{z})\}$, and therefore that both $c(g(S), \mathbf{z}) = 1$ and $c(g(S), \mathbf{z}) = 0$, a contradiction. Thus, $S_n \subset S_{n-1}$, and $|S_n|$ is a decreasing sequence of non-negative integers that is strictly decreasing so long as $|S_n| > 0$. It must thus converge to zero in at most $|S_1|$ iterations, so that $S_n = \emptyset$ for all $n \geq |S_1|$.

Without loss, we can terminate the sequence on an even term, since $\mathbb{1}(S \subseteq \emptyset) = 0$ for any $S \supset \emptyset$. Let $2K$ denote the smallest even number such that $S_n = \emptyset$ for all $n > 2K$, for a fixed \mathbf{z} .

Thus, we have for any $\emptyset \subset S \subseteq \{1 \dots J\}$:

$$c(g(S), \mathbf{z}) = \sum_{n=1}^{2K} (-1)^{n+1} \cdot \mathcal{D}_{g(S)}(S_n) = \sum_{k=1}^K \mathcal{D}_{g(S)}(S_{2k-1}) - \mathcal{D}_{g(S)}(S_{2k})$$

where $2K \leq |S_1| \leq J$, and we have used that $\mathcal{D}_{g(S)}(S') = \mathbb{1}(S \subset S')$ for any S' .

Now recall that we have left the dependence of each of the sets S_n (as well as the integer K) on \mathbf{z} implicit, and have also adopted the notational convention of $\mathcal{D}_g(S)$ as a shorthand for $\mathcal{D}_g(\mathbf{z})$ where \mathbf{z} is a point in \mathcal{Z} that takes a value of one for exactly the instruments in the set S . To obtain the notation of the final result, define for each $k = 1 \dots K$ the point $u_k(\mathbf{z}) \in \mathcal{Z}$ to have a value of one exactly for the elements in S_{2k-1} for that value of \mathbf{z} , and $l_k(\mathbf{z}) \in \mathcal{Z}$ to have a value of one exactly for the elements in S_{2k} for that value of \mathbf{z} . We may thus write, for any $\emptyset \subset S \subseteq \{1 \dots J\}$ and any $\mathbf{z} \in \mathcal{Z}$:

$$c(g(S), \mathbf{z}) = \sum_{k=1}^{K(\mathbf{z})} \mathcal{D}_{g(S)}(u_k(\mathbf{z})) - \mathcal{D}_{g(S)}(l_k(\mathbf{z})) = \sum_{k=1}^K \mathcal{D}_{g(S)}(u_k(\mathbf{z})) - \mathcal{D}_{g(S)}(l_k(\mathbf{z}))$$

where we let K be the maximum of $K(\mathbf{z})$ over the finite set \mathcal{Z} , and we define $u_k(\mathbf{z})$ and $l_k(\mathbf{z})$ to each be a vector of zeros whenever $k > K(\mathbf{z})$. For each \mathbf{z} , the relations $u_k(\mathbf{z}) \geq l_k(\mathbf{z})$ and $l_k(\mathbf{z}) \geq u_{k+1}(\mathbf{z})$ component-wise now follow from $S_n \subseteq S_{n+1}$.

Now we may apply Property M to construct $c(g, \mathbf{z})$ for any of the non-simple response groups as well. Recall that Property M says that $c(g(F), \mathbf{z}) = \sum_{\emptyset \subset S \subseteq \{1 \dots J\}} [M_J]_{F,S} \cdot c(g(S), \mathbf{z})$ for all \mathbf{z} , for any Sperner family F . Thus:

$$\begin{aligned} c(g(F), \mathbf{z}) &= \sum_{\emptyset \subset S \subseteq \{1 \dots J\}} [M_J]_{F,S} \cdot \sum_{k=1}^K \{\mathcal{D}_{g(S)}(u_k(\mathbf{z})) - \mathcal{D}_{g(S)}(l_k(\mathbf{z}))\} \\ &= \sum_{k=1}^K \left\{ \sum_{\emptyset \subset S \subseteq \{1 \dots J\}} [M_J]_{F,S} \cdot \mathcal{D}_{g(S)}(u_k(\mathbf{z})) \right\} - \left\{ \sum_{\emptyset \subset S \subseteq \{1 \dots J\}} [M_J]_{F,S} \cdot \mathcal{D}_{g(S)}(l_k(\mathbf{z})) \right\} \\ &= \sum_{k=1}^K \mathcal{D}_{g(F)}(u_k(\mathbf{z})) - \mathcal{D}_{g(F)}(l_k(\mathbf{z})) \end{aligned}$$

Finally, note that $\mathcal{D}_g(u_k(\mathbf{z})) = \mathcal{D}_g(l_k(\mathbf{z}))$ for any $g \in \{a.t., n.t.\}$ so the following expression holds for all $g \in \mathcal{G}$:

$$c(g, \mathbf{z}) = \sum_{k=1}^K \mathcal{D}_g(u_k(\mathbf{z})) - \mathcal{D}_g(l_k(\mathbf{z}))$$

E.5 Proof of Lemma 1

Any indicator $\mathbb{1}(Z_i = \mathbf{z})$ for a value $z \in \{0, 1\}^J$ can be expanded out as a polynomial in the instrument indicators as $\mathbb{1}(Z_i = \mathbf{z}) = \prod_{j \in \mathbf{z}_1} Z_{ji} \prod_{j \in \mathbf{z}_0} (1 - Z_{ji}) = \sum_{f \subseteq \mathbf{z}_0} (-1)^{|f|} Z_{(\mathbf{z}_1 \cup f), i}$, where $(\mathbf{z}_1, \mathbf{z}_0)$ is a partition of the indices $j \in \{1 \dots J\}$ that take a value of zero or one in \mathbf{z} , respectively. With $J = 2$ for example,

$$((1 - Z_{1i})(1 - Z_{2i}), Z_{1i}(1 - Z_{2i}), Z_{2i}(1 - Z_{1i}), Z_{1i}Z_{2i}) = (1, Z_{1i}, Z_{2i}, Z_{1i}Z_{2i})A = (1, \Gamma'_i)A$$

where $A = \begin{pmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ 1 & -1 & -1 & 1 \end{pmatrix}$. Denote the random vector of such indicators \mathfrak{Z}_i . Then $(1, \Gamma'_i)A = \mathfrak{Z}'_i$, with the matrix of coefficients $A_{S, \mathbf{z}} = \sum_{\substack{f \subseteq \mathbf{z}_0 \\ (\mathbf{z}_1 \cup f) = S}} (-1)^{|f|}$. The matrix A so defined must be invertible, because any product of the instruments Z_{Si} for $S \subseteq \{1 \dots J\}$ can similarly be expressed as a linear combination of the components of \mathfrak{Z}_i , where we define $Z_{\emptyset i} = 1$. Specifically, $Z_{Si} = \sum_{z \in \mathcal{Z}} \mathbb{1}(\forall j \in S, z_j = 1) \mathbb{1}(Z_i = \mathbf{z})$.

Consider the matrix

$$\Sigma^* := \mathbb{E}[(1, \Gamma'_i)'(1, \Gamma'_i)] = A'^{-1} \mathbb{E}[\mathfrak{Z}_i \mathfrak{Z}'_i] A^{-1} = A'^{-1} \text{diag}\{P(Z_i = \mathbf{z})\} A^{-1}$$

where $\mathbb{E}[\mathfrak{Z}_i \mathfrak{Z}'_i]$ is diagonal since the events that Z_i take on two different values are exclusive. Since A^{-1} exists, the rank of Σ^* must be equal to the rank of $\text{diag}\{P(Z_i = \mathbf{z})\}$, which is in turn equal to the cardinality of \mathcal{Z} . Assumption 3 thus holds if and only if Σ^* has full rank of 2^J . Note that although A^{-1} diagonalizes the matrix Σ^* , it does not provide its eigen-decomposition, as $A^{-1} \neq A'$ (A is not orthogonal).

Now we prove that Σ^* has full rank whenever Σ has full rank, and vice versa. Note that $\Sigma = \text{Var}(\Gamma_i)$ has full rank if and only if $\omega' \mathbb{E}[(\Gamma_i - E\Gamma_i)(\Gamma_i - E\Gamma_i)']\omega = \mathbb{E}[\omega'(\Gamma_i - E\Gamma_i)(\Gamma_i - E\Gamma_i)\omega] > 0$, i.e. $P(\omega'(\Gamma_i - E\Gamma_i) = 0) < 1$ for any $\omega \in \mathbb{R}^{2^J-1}/\mathbf{0}$. Similarly Σ^* has full rank if $P((\omega_0, \omega)'((1, \Gamma_i) = 0) < 1$ for any $\omega_0 \in \mathbb{R}, \omega \in \mathbb{R}^{2^J-1}$ where (ω_0, ω) is not the zero vector in \mathbb{R}^{2^J} . But if for some ω , $\omega'(\Gamma_i - E\Gamma_i) = 0$ w.p.1., then we also have $(\omega_0, \omega)'(1, \Gamma_i) = 0$ w.p.1. by choosing $\omega_0 = -\omega' \mathbb{E}[\Gamma_i]$. In the other direction, note that $(\omega_0, \omega)'(1, \Gamma_i) = 0$ w.p.1. implies that $\omega' \Gamma_i = -\omega_0$ and hence $\omega'(\Gamma_i - E\Gamma_i) = -\omega_0 - \omega' E\Gamma_i = -\omega_0 - \mathbb{E}[\omega' \Gamma_i] = -\omega_0 + \omega_0 = 0$.

E.6 Proof of Theorem 1

We first note that any measurable function $f(Y)$ preserves Assumption 1, that is

$$(f(Y_i(1)), f(Y_i(0)), G_i) \perp Z_i$$

and Assumptions 2-3 are unaffected by such a transformation to the outcome variable. Thus, we continue without loss with Y_i , $Y_i(1)$ and $Y_i(0)$ possibly redefined as $f(Y_i)$, $f(Y_i(1))$ and $f(Y_i(0))$ respectively.

Note that the function $h(\cdot)$ given in Theorem 1 has the property that $\mathbb{E}[h(Z_i)] = 0$, for any distribution of the instruments. Consider the quantity $\mathbb{E}[Y_i D_i h(Z_i)]$ for a function h having this property. By the law of iterated expectations, and the independence assumption:

$$\begin{aligned} \mathbb{E}[Y_i D_i h(Z_i)] &= \sum_g P(G_i = g) \mathbb{E}[Y_i D_i h(Z_i) | G_i = g] \\ &= \sum_g P(G_i = g) \mathbb{E}[Y_i(1) \mathcal{D}_g(Z_i) h(Z_i) | G_i = g] \\ &= \sum_g P(G_i = g) \mathbb{E}[Y_i(1) | G_i = g] \mathbb{E}[\mathcal{D}_g(Z_i) h(Z_i)] \end{aligned} \quad (17)$$

where $\mathcal{D}_g(\mathbf{z})$ denotes the selection function for response group g . Similarly,

$$\begin{aligned} \mathbb{E}[Y_i(1 - D_i) h(Z_i)] &= \sum_g P(G_i = g) \mathbb{E}[Y_i(0)(1 - D_i) h(Z_i) | G_i = g] \\ &= \sum_g P(G_i = g) \{ \mathbb{E}[Y_i(0) | G_i = g] \mathbb{E}[h(Z_i)] \\ &\quad - \mathbb{E}[Y_i(0) | G_i = g] \mathbb{E}[\mathcal{D}_g(Z_i) h(Z_i)] \} \\ &= \sum_g -P(G_i = g) \mathbb{E}[Y_i(0) | G_i = g] \mathbb{E}[\mathcal{D}_g(Z_i) h(Z_i)] \end{aligned} \quad (18)$$

where we have used that $Z_i \perp (Y_i(0), Z_i)$ and $\mathbb{E}[h(Z_i)] = 0$.

Combining these two results:

$$\mathbb{E}[Y_i h(Z_i)] = \mathbb{E}[Y_i D_i h(Z_i)] + \mathbb{E}[Y_i(1 - D_i) h(Z_i)] = \sum_g P(G_i = g) \mathbb{E}[\mathcal{D}_g(Z_i) h(Z_i)] \Delta_g \quad (19)$$

where $\Delta_g := \mathbb{E}[Y_i(1) - Y_i(0)|G_i = g]$. By the law of iterated expectations, we also have that

$$\mathbb{E}[D_i h(Z_i)] = \sum_g P(G_i = g) \mathbb{E}[\mathcal{D}_g(Z_i) h(Z_i)] \quad (20)$$

Note that in all of Equations (17), (18) and (19), the weighing over various groups g is governed by the quantity $\mathbb{E}[\mathcal{D}_g(Z_i) h(Z_i)]$. It can be seen that never takers and always takers receive no weight, since $\mathbb{E}[\mathcal{D}_{n.t}(Z_i) h(Z_i)] = \mathbb{E}[0] = 0$ and since $\mathbb{E}[\mathcal{D}_{a.t}(Z_i) h(Z_i)] = \mathbb{E}[h(Z_i)] = 0$.

Let \mathcal{F} denote the set of non-empty subsets of the instrument indices: $\mathcal{F} := \{S \subseteq \{1, 2, \dots, J\}, S \neq \emptyset\}$, and recall that these correspond each to a simple response group $g(S)$, where $\mathcal{D}_{g(S)}(Z_i) = z_{Si}$. I first show that for any $\lambda \in \mathbb{R}^{|\mathcal{F}|}$, Assumption 3 allows us to define an $h(Z_i)$ such that $\mathbb{E}[\mathcal{D}_{g(S)}(Z_i) h(Z_i)] = \mathbb{E}[z_{Si} h(Z_i)] = \lambda_S$. Note that since $\mathbb{E}[h(Z_i)] = 0$, this is the same as tuning each covariance $Cov(Z_{Si}, h(Z_i))$ to λ_S (c.f. Proposition 6). In particular, consider the choice $h(Z_i) = (\Gamma_i - \mathbb{E}[\Gamma_i])'\Sigma^{-1}\lambda$, where recall that Γ_i is a vector of Z_{Si} for each $S \in \mathcal{F}$.

$$\begin{aligned} (\mathbb{E}[h(Z_i)], \mathbb{E}[h(Z_i), \Gamma_{i1}], \mathbb{E}[h(Z_i), \Gamma_{i2}], \dots, \mathbb{E}[h(Z_i), \Gamma_{ik}])' &= \mathbb{E}[(\Gamma_i - \mathbb{E}[\Gamma_i])h(Z_i)] \\ &= \mathbb{E}[(\Gamma_i - \mathbb{E}[\Gamma_i])(\Gamma_i - \mathbb{E}[\Gamma_i])']\Sigma^{-1}\lambda \\ &= \Sigma\Sigma^{-1}\lambda = \lambda \end{aligned}$$

We can understand the algebra of this result as follows. Let $V = \text{span}(\{Z_{Si} - \mathbb{E}[Z_{Si}]\}_{S \in \mathcal{F}})$. V is a subspace of the vector space \mathcal{V} of random variables on \mathcal{Z} , with the zero vector being a degenerate random variable equal to zero. Since the matrix Σ is positive semidefinite by construction, Assumption 3 is equivalent to the statement that for all $\omega \in \mathbb{R}^{|\mathcal{F}|}/\mathbf{0}$, $\omega'\mathbb{E}[(\Gamma_i - \mathbb{E}[\Gamma_i])(\Gamma_i - \mathbb{E}[\Gamma_i])']\omega = \mathbb{E}[|\omega'(\Gamma_i - \mathbb{E}[\Gamma_i])|^2] > 0$: i.e. $P(\sum_{S \in \mathcal{F}} \omega_S(z_{Si} - \mathbb{E}[z_{Si}])) = 0 < 1$ for all $\omega \in \mathbb{R}^{|\mathcal{F}|}/\mathbf{0}$. In other words, the random variables $(z_{Si} - \mathbb{E}[z_{Si}])$ for $S \in \mathcal{F}$ are linearly independent, and hence form a basis of V . Since V is finite dimensional, there exists an orthonormal basis of random vectors of the same cardinality, $|\mathcal{F}|$, where orthonormality is defined with respect to the expectation inner product: $\langle A, B \rangle := \mathbb{E}[A_i B_i]$. It is this orthogonalized version of the Z_{Si} that affords the ability to separately tune each of the $\mathbb{E}[h(Z_i)Z_{Si}]$ to the desired value λ_S , without disrupting the others.

Note that under Assumption 1:

$$\Delta_c = \sum_{g \in \mathcal{G}} \left\{ \frac{P(G_i = g)P(C_i = 1|G_i = g)}{P(C_i = 1)} \right\} \cdot \Delta_g = \frac{\sum_{g \in \mathcal{G}} P(G_i = g)P(C_i = 1|G_i = g) \cdot \Delta_g}{\sum_{g \in \mathcal{G}} P(G_i = g)P(C_i = 1|G_i = g)}$$

Comparing with Equations (19) and (20), the equality $\Delta_c = \mathbb{E}[h(Z_i)Y_i]/\mathbb{E}[D_i h(Z_i)]$ follows (provided that $P(C_i = 1) > 0$) if the coefficients match. That is: $\mathbb{E}[\mathcal{D}_g(Z_i)h(Z_i)] = P(C_i = 1|G_i = g)$, for all $g \in \mathcal{G}^c$. By the above, this is guaranteed under Property M if we choose $\lambda_S = P(C_i = 1|G_i = g(S)) = \mathbb{E}[c(g(S), Z_i)]$, since the quantity $\mathbb{E}[\mathcal{D}_g(Z_i)h(Z_i)]$ appearing in Eq. (19) is linear in $\mathcal{D}_g(Z_i)$. The same logic follows for causal parameters of the form $\mathbb{E}[Y_i(d)|C_i = 1]$ for $d \in \{0, 1\}$, using Equations (17) and (18) and

$$\begin{aligned} \mathbb{E}[Y_i(d)|C_i = 1] &= \sum_{g \in \mathcal{G}} P(G_i = g|C_i = 1) \mathbb{E}[Y_i(d)|G_i = g, c(g, Z_i) = 1] \\ &= P(C_i = 1)^{-1} \sum_{g \in \mathcal{G}} P(G_i = g) P(C_i = 1|G_i = g) \mathbb{E}[Y_i(d)|G_i = g] \end{aligned}$$

by independence. Note that the quantity λ_S for each S can be computed from the observed distribution of Z .

To replace Assumption 3 with Assumption 3* from Appendix A, simply replace \mathcal{F} as defined here with a maximal \mathcal{F} from Assumption 3a*.

Finally, to see that $P(C_i = 1) = \mathbb{E}[h(Z_i)D_i]$ we can simply use independence and Property M:

$$\begin{aligned} \mathbb{E}[h(Z_i)D_i] &= \sum_g P(G_i = g) \mathbb{E}[h(Z_i)\mathcal{D}_g(Z_i)] \\ &= \sum_g P(G_i = g) \mathbb{E} \left[h(Z_i) \left\{ \sum_S [M_J]_{F(g), S} \mathcal{D}_{g(s)}(Z_i) \right\} \right] \\ &= \sum_g P(G_i = g) \sum_S [M_J]_{F(g), S} \mathbb{E}[h(Z_i)\mathcal{D}_{g(s)}(Z_i)] \\ &= \sum_g P(G_i = g) \sum_S [M_J]_{F(g), S} P(C_i = 1|\mathcal{D}_{g(s)}(Z_i) = 1) \\ &= \sum_g P(G_i = g) P(C_i = 1|G_i = g) = P(C_i = 1) \end{aligned}$$

E.7 An Equivalence Result

The proofs of Theorem 2 and Proposition 5 make use of the following equivalence result:

Proposition 7. *Let the support \mathcal{Z} of the instruments be finite. Fix a function $c(g, \mathbf{z})$. Let \mathcal{P}_{DZ} denote the joint distribution of D_i and Z_i . Then the following are equivalent:*

1. Δ_c is (point) identified by \mathcal{P}_{DZ} and a finite set of IV-like estimands $\beta_s := \mathbb{E}[s(D_i, Z_i)Y_i]$, where each function $s(d, \mathbf{z})$ is known or identified from \mathcal{P}_{DZ}
2. $\Delta_c = \beta_s$ for a single such s
3. Δ_c is identified from the set of CEFs $\{\mathbb{E}[Y_i|D_i = d, Z_i = \mathbf{z}]\}$ for $d \in \{0, 1\}$, $\mathbf{z} \in \mathcal{Z}$ along with the joint distribution \mathcal{P}_{DZ}

Proof. Let us use the notation that \mathcal{S} denotes a set of measurable functions $s(d, \mathbf{z})$ defining IV-like estimands $\{\beta_s\}_{s \in \mathcal{S}}$. We can show each of the following implications:

- **2 → 1** Immediate, since **2** is a special case of **1** with \mathcal{S} a singleton
- **3 → 1** Let $\mathcal{S} = \{s(d, \mathbf{z}) = \mathbb{1}(D_i = d)\mathbb{1}(Z_i = \mathbf{z})\}_{d \in \{0, 1\}, z \in \mathcal{Z}}$. Then each β_s is equal to $P(D_i = d, Z_i = \mathbf{z})\mathbb{E}[Y_i|D_i = d, Z_i = \mathbf{z}]$ for some d, \mathbf{z} . The coefficient is known from \mathcal{P}_{DZ} , thus **3** is a case of **1**
- **1 → 3** Any β_s can be written: $\beta_s = \sum_{d, \mathbf{z}} P(D_i = d, Z_i = \mathbf{z})s(d, \mathbf{z})\mathbb{E}[Y_i|D_i = d, Z_i = \mathbf{z}]$, and is thus pinned down by the CEFs $\mathbb{E}[Y_i|D_i = d, Z_i = \mathbf{z}]$, the joint distribution \mathcal{P}_{DZ} , and the known function s .
- **3 → 2** Given independence:

$$\begin{aligned} \mathbb{E}[Y_i \mathbb{1}(D_i = d)|Z_i = \mathbf{z}] &= \mathbb{E}[Y_i(d) \mathbb{1}(D_i = d)|Z_i = \mathbf{z}] \\ &= \sum_g P(G_i = g|Z_i = \mathbf{z})\mathbb{E}[Y_i(d) \mathbb{1}(D_i = d)|G_i = g, Z_i = \mathbf{z}] \\ &= \sum_{g: \mathcal{D}_g(\mathbf{z})=d} P(G_i = g)\mathbb{E}[Y_i(d)|G_i = g] \end{aligned}$$

If we stack the unknown quantities $P(G_i = g)\mathbb{E}[Y_i(d)|G_i = g]$ for all $g \in \mathcal{G}, d \in \{0, 1\}$ into a vector x , and the identified quantities $\mathbb{E}[Y_i \mathbb{1}(D_i = d)|Z_i = \mathbf{z}]$ for all $d \in \{0, 1\}, \mathbf{z} \in \mathcal{Z}$ into a vector b , then we have a system of linear equations $Ax = b$, where A is a fixed matrix of entries of the form $[A]_{d\mathbf{z}, g} = \mathbb{1}(\mathcal{D}_g(\mathbf{z}) = d)$ (note the matrix A here is not the

same as the matrix A defined in Corollary 6).

Similarly, as we have seen Δ_c can be written as a linear combination of the components of the vector \mathbf{z} . Specifically, from Equation(5):

$$\Delta_c = \sum_g \frac{\mathbb{E}[c(g, Z_i)]}{\mathbb{E}[c(G_i, Z_i)]} P(G_i = g) \{ \mathbb{E}[Y_i(1)|G_i = g] - \mathbb{E}[Y_i(0)|G_i = g] \}$$

We can now write $\Delta_c = \theta'_c x$, where θ_c is the vector of coefficients $\pm \frac{\mathbb{E}[c(g, Z_i)]}{\mathbb{E}[c(G_i, Z_i)]}$ from the above equation.

The set of vectors x compatible with the set of identifying restrictions $Ax = b$ can be written as $\{A^\dagger b + (I - A^\dagger A)w\}$ for all arbitrary vectors $w \in \mathbb{R}^{2|\mathcal{G}|}$, where A^\dagger is the Moore-Penrose pseudo-inverse of A . The corresponding set of values for Δ_c is $\{\theta'_c A^\dagger b + \theta'_c (I - A^\dagger A)w\}$. For this set to be a singleton for all w , we must either have $A^\dagger A = I$ (i.e. A has full column rank, which is only possible in the $J = 1$ case), or the vector θ_c must lie in the row space of the matrix A , so that in either case $\theta'_c (I - A^\dagger A)$ is equal to the zero vector. If the set were not a singleton, then Δ_c would not be identified absent additional restrictions, since an infinite collection of values of Δ_c would be compatible with the full set of restrictions $Ax = b$ placed by the CEF's $\mathbb{E}[Y_i|D_i = d, Z_i = \mathbf{z}]$. Thus, by **3.**, we have that $\Delta_c = \theta'_c A^\dagger b$. This then implies **2.**, if we take $s(d, \mathbf{z}) = \frac{P(D_i=d|Z_i=\mathbf{z})}{P(D_i=d, Z_i=\mathbf{z})} \cdot [\theta'_c A^\dagger]_{(d, \mathbf{z})}$, where $[\theta'_c A^\dagger]_{(d, \mathbf{z})}$ is the component of the vector $\theta'_c A^\dagger$ corresponding to the pair (d, \mathbf{z}) . Note that A^\dagger is a known matrix (without looking at the data), and θ_c is a known function of the marginal distribution of Z_i , up to the factor $\mathbb{E}[c(G_i, Z_i)]$, for a fixed function c .

It only remains to be shown that $\mathbb{E}[c(G_i, Z_i)]$ is also identified under assumption of **1**. For Δ_c to be pinned down for all joint distributions of $(G_i, Y_i(1), Y_i(0))$, it must be pinned down in the special case where in which $Y_i(d) = d$ with probability one. In this case each $\mathbb{E}[Y_i|D_i = d, Z_i = \mathbf{z}] = d$, and $\Delta_c = 1$. Thus, using our result above we have that $\mathbb{E}[c(G_i, Z_i)] = \mathbb{E}[\tilde{s}(d, \mathbf{z})D_i]$, where $\tilde{s}(d, \mathbf{z}) := \frac{P(D_i=d|Z_i=\mathbf{z})}{P(D_i=d, Z_i=\mathbf{z})} \cdot [\tilde{\theta}'_c A^\dagger]_{(d, \mathbf{z})}$, where $\tilde{\theta}_c := \mathbb{E}[c(G_i, Z_i)]\theta_c = \pm E(c(g, Z_i))$.

□

E.8 Proof of Theorem 2

Note that if θ_c^{fd} is identified for all d and f , so must $\Delta_c = \theta_c^{y1} - \theta_c^{y0}$ where $f(y) = y$. If Δ_c is identified from a finite set of IV-like estimands and \mathcal{P}_{DZ} , then it can also be written as a single one: $\Delta_c = \beta_s$ with $s(d, \mathbf{z})$ an identified functional of \mathcal{P}_{DZ} (see Appendix Proposition 7). Now, using that $Y_i = Y_i(0) + D_i\Delta_i$ where $\Delta_i := Y_i(1) - Y_i(0)$,

$$\begin{aligned}\Delta_c &= \beta_s = \{\mathbb{E}[s(D_i, Z_i)Y_i(0)] + \mathbb{E}[s(D_i, Z_i)D_i\Delta_i]\} \\ &= \sum_g P(G_i = g) \{\mathbb{E}[s(\mathcal{D}_g(Z_i), Z_i)Y_i(0)|G_i = g] + \mathbb{E}[s(\mathcal{D}_g(Z_i), Z_i)\mathcal{D}_g(Z_i)\Delta_i|G_i = g]\} \\ &= \sum_g P(G_i = g) (\underbrace{\mathbb{E}[s(\mathcal{D}_g(Z_i), Z_i)]}) \mathbb{E}[Y_i(0)|G_i = g] \\ &\quad + \sum_g P(G_i = g) (\mathbb{E}[s(\mathcal{D}_g(Z_i), Z_i)\mathcal{D}_g(Z_i)]) \mathbb{E}[\Delta_i|G_i = g] \\ &= \sum_g P(G_i = g) (\mathbb{E}[s(1, Z_i)\mathcal{D}_g(Z_i)]) \Delta_g\end{aligned}$$

where I've used independence in the second equality. In the final equality, I've used that $s(\mathcal{D}_g(Z_i), Z_i)\mathcal{D}_g(Z_i) = s(1, Z_i)\mathcal{D}_g(Z_i)$ for all i . The crossed out term must be equal to zero, because we've assumed $\beta_s = \Delta_c$ holds for *every* joint distribution of response groups and potential outcomes compatible with \mathcal{P}_{DZ} . Consider two such distributions, identical except that the second distribution describes a case in which for all units $Y_i(0) \rightarrow Y_i(0) + \delta$ and $Y_i(1) \rightarrow Y_i(1) + \delta$. Then the Δ_i and hence Δ_c are unchanged, but if the crossed out term were not zero, β_s would change as $\mathbb{E}[Y_i(0)|G_i = g] \rightarrow \mathbb{E}[Y_i(0)|G_i = g] + \delta$.

Recall that from Equation (5) that Δ_c can also be written as a weighted average of group-specific average treatment effects $\Delta_g = \mathbb{E}[Y_i(1) - Y_i(0)|G_i = g]$ as:

$$\Delta_c = \frac{1}{P(C_i = 1)} \sum_g P(G_i = g) \mathbb{E}[c(g, Z_i)] \cdot \Delta_g$$

Since $\beta_s = \Delta_c$ holds for any vector of $\{\Delta_g\}$ across all of the g for which $P(G_i = g) > 0$ is compatible with \mathcal{P}_{DZ} , we can match coefficients within this group to establish that $\mathbb{E}[c(g, Z_i)] =$

$P(C_i = 1)\mathbb{E}[s(1, Z_i)\mathcal{D}_g(Z_i)]$. This set of weights satisfies Property M, since for any $g \in \mathcal{G}^c$:

$$\begin{aligned}\mathbb{E}[c(g, Z_i)] &= P(C_i = 1)\mathbb{E}[s(1, Z_i) \sum_S [M_J]_{F(g), S} \mathcal{D}_{g(S)}(Z_i)] \\ &= \sum_S [M_J]_{F(g), S} (P(C_i = 1)\mathbb{E}[s(1, Z_i)\mathcal{D}_{g(S)}(Z_i)]) \\ &= \sum_S [M_J]_{F(g), S} \cdot \mathbb{E}[c(Z_i, g(S))]\end{aligned}$$

If this holds for any distribution of Z_i satisfying Assumption 3, then we must have $c(g, \mathbf{z}) = \sum_S [M_J]_{F(g), S} \cdot c(g(S), \mathbf{z})$ for all $\mathbf{z} \in \mathcal{Z}$. To see this, consider a sequence of distributions for Z_i that converges point-wise to a degenerate distribution at any single point \mathbf{z} , but satisfies Assumption 3 for each term in the sequence. Applying the dominated convergence theorem to $\mathbb{E}[c(g, Z_i)] - \sum_S [M_J]_{F(g), S} \cdot \mathbb{E}[c(g(S), Z_i)] = 0$ along this sequence, we have that $c(g, \mathbf{z}) = \sum_S [M_J]_{F(g), S} \cdot c(g(S), \mathbf{z})$. A similar argument establishes that $c(a.t., \mathbf{z}) = c(n.t., \mathbf{z}) = 0$ for all $\mathbf{z} \in \mathcal{Z}$ given that $\mathbb{E}[c(g, Z_i)] = P(C_i = 1)\mathbb{E}[s(1, Z_i)\mathcal{D}_g(Z_i)]$ and $\mathbb{E}[s(1, Z_i)] = 0$.

E.9 Proof of Proposition 5

Step 1: documenting the response groups. I first show that with two binary instruments, if PM holds but not VM or IAM, then \mathcal{G} consists of 7 response groups, given in Table 10.

To see this, let us begin by supposing PM holds but VM does not. For VM to be violated, there must be a “flip” in which value of one of the instruments—say Z_2 —is the “pro-treatment” state, depending on the value of the other instrument. In other words, for some choice of which instrument is called Z_2 , and some choice of labeling for the values of each instrument, we have that:

$$D_i(0, 1) \geq D_i(0, 0) \text{ for all } i, \text{ and } D_i(1, 1) \leq D_i(1, 0) \text{ for all } i$$

with both

$$D_i(0, 1) > D_i(0, 0) \text{ for some } i, \text{ and } D_i(1, 1) < D_i(1, 0) \text{ for some } i$$

(this is without loss of generality, given the choice to arbitrarily assign the labels 0, 1).

Now consider the set of possible response groups that satisfy PM but not VM, denoted as \mathcal{G}^{PM-VM} . By PM, any response group $g \in \mathcal{G}^{PM-VM}$ must be either a complier, always-taker, or never-taker with respect to Z_2 , when $Z_1 = 0$. Similarly, any response group g must then be either a “defier”, always-taker, or never-taker with respect to Z_2 when $Z_1 = 1$. The set of

pairs (g_0, g_1) , where $g_0 \in \{c, a, n\}$ and $g_1 \in \{d, a, n\}$ exhaust the possible response groups, since knowing g_0 and g_1 pins down $D_i(\mathbf{z})$ for all four values of \mathbf{z} . This generates an exhaustive set of 9 possible response groups, shown in Table 9.

name	$\mathbf{z}_1 = \mathbf{0}, \mathbf{z}_2 = \mathbf{0}$	$\mathbf{z}_1 = \mathbf{0}, \mathbf{z}_2 = \mathbf{1}$	$\mathbf{z}_1 = \mathbf{1}, \mathbf{z}_2 = \mathbf{0}$	$\mathbf{z}_1 = \mathbf{1}, \mathbf{z}_2 = \mathbf{1}$
odd compliers (c, d)	N	T	T	N
eager compliers (c, a)	N	T	T	T
1-only takers (c, n)	N	T	N	N
reluctant defiers (a, d)	T	T	T	N
always takers (a, a)	T	T	T	T
Z_1 defiers (a, n)	T	T	N	N
2-only takers (n, d)	N	N	T	N
Z_1 compliers (n, a)	N	N	T	T
never takers (n, n)	N	N	N	N

Table 9: Rows are possible response groups in the set \mathcal{G}^{PM-VM} . T and N indicate treatment, or non-treatment, respectively. Not all of these groups can coexist in the population without violating PM.

However, all nine of the response groups cannot coexist at the same time. For example, if both odd compliers and Z_1 defiers both exist in the population, there will be two-way flows when varying Z_1 when Z_2 is fixed at zero. We can thus distinguish two cases, depending on whether there are “odd compliers” in the population. If there are odd compliers, then there can be neither Z_1 compliers nor Z_1 defiers in the population, leaving the seven groups listed in Table 10.³⁸

group name	$\mathbf{z}_1 = \mathbf{0}, \mathbf{z}_2 = \mathbf{0}$	$\mathbf{z}_1 = \mathbf{0}, \mathbf{z}_2 = \mathbf{1}$	$\mathbf{z}_1 = \mathbf{1}, \mathbf{z}_2 = \mathbf{0}$	$\mathbf{z}_1 = \mathbf{1}, \mathbf{z}_2 = \mathbf{1}$
odd compliers	N	T	T	N
eager compliers	N	T	T	T
reluctant defiers	T	T	T	N
1-only takers	N	T	N	N
2-only takers	N	N	T	N
always takers	T	T	T	T
never takers	N	N	N	N

Table 10: Case 1, when “odd compliers” exist.

If, on the other hand, there are no odd compliers, then there can be either Z_1 compliers, or Z_1 defiers, but not both. This creates a second type of case. Supposing that there are Z_1 compliers, there can be no Z_1 defiers, 1-only takers, or reluctant defiers. This leaves five possible groups, depicted in Table 11.

If there are Z_1 defiers instead of Z_1 compliers, this yields the same as Table 11, up to a

³⁸The case described in Table 10 applies to the family-size instruments considered by MTW as an example of PM. In particular, if we reverse the 0/1 labels for Z_2 , we recover the inequalities $D_i(0, 1) \leq D_i(1, 1)$, $D_i(1, 0) \leq D_i(1, 1)$, $D_i(0, 1) \leq D_i(0, 0)$, $D_i(1, 0) \leq D_i(0, 0)$ from MTW.

group name	$\mathbf{z}_1 = \mathbf{0}, \mathbf{z}_2 = \mathbf{0}$	$\mathbf{z}_1 = \mathbf{0}, \mathbf{z}_2 = \mathbf{1}$	$\mathbf{z}_1 = \mathbf{1}, \mathbf{z}_2 = \mathbf{0}$	$\mathbf{z}_1 = \mathbf{1}, \mathbf{z}_2 = \mathbf{1}$
always takers	T	T	T	T
eager compliers	N	T	T	T
Z_1 compliers	N	N	T	T
2-only takers	N	N	T	N
never takers	N	N	N	N

Table 11: Case 2, when there are no odd compliers but there are Z_1 compliers.

relabeling of “0” and “1” for Z_1 : in addition to Z_1 defiers, there can be reluctant defiers, 1-only units, always takers and never takers. A final case occurs if there are neither Z_1 compliers, Z_2 compliers, nor odd compliers, which is a special case of Table 10.

Tables 10 and 11 have very different implications for identification. In the case of Table 11, Assumption IAM holds, and group-specific average treatment effects Δ_g are thus identified for all groups aside from always takers and never takers. For Proposition 5, we must therefore be the setting described by Table 10.

Step 2: non-identification of the ACL. In the absence of any further restrictions on selection behavior, any of the 7 groups in Table 10 are possibly present in the population. By Appendix Proposition 7, point identification from a finite set of IV-like estimands is equivalent to point identification of the moments $\mathbb{E}[Y_i D_i | Z_i = \mathbf{z}]$ and $\mathbb{E}[Y_i (1 - D_i) | Z_i = \mathbf{z}]$ for all $\mathbf{z} \in \mathcal{Z}$, along with the observable joint distribution of (D_i, Z_i) .

Using Table 10 and Assumption 1, the eight such observable moments can be written in matrix

form as

$$\begin{pmatrix} \mathbb{E}[Y_i D_i | Z_i = (0,0)] \\ \mathbb{E}[Y_i D_i | Z_i = (0,1)] \\ \mathbb{E}[Y_i D_i | Z_i = (1,0)] \\ \mathbb{E}[Y_i D_i | Z_i = (1,1)] \\ \mathbb{E}[Y_i(1 - D_i) | Z_i = (0,0)] \\ \mathbb{E}[Y_i(1 - D_i) | Z_i = (0,1)] \\ \mathbb{E}[Y_i(1 - D_i) | Z_i = (1,0)] \\ \mathbb{E}[Y_i(1 - D_i) | Z_i = (1,1)] \end{pmatrix} = \left[\begin{array}{cccc|cccc|cccc} 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 1 & 1 & 0 \\ \end{array} \right] \cdot \begin{pmatrix} p_{odd} \cdot \mathbb{E}[Y_i(1) | G_i = odd] \\ p_{eager} \cdot \mathbb{E}[Y_i(1) | G_i = eager] \\ p_{reluct.} \cdot \mathbb{E}[Y_i(1) | G_i = reluct.] \\ p_1 \cdot \mathbb{E}[Y_i(1) | G_i = 1only] \\ p_2 \cdot \mathbb{E}[Y_i(1) | G_i = 2only] \\ p_a \cdot \mathbb{E}[Y_i(1) | G_i = a.t.] \\ p_n \cdot \mathbb{E}[Y_i(1) | G_i = n.t.] \\ \hline p_{odd} \cdot \mathbb{E}[Y_i(0) | G_i = odd] \\ p_{eager} \cdot \mathbb{E}[Y_i(0) | G_i = eager] \\ p_{reluct.} \cdot \mathbb{E}[Y_i(0) | G_i = reluct.] \\ p_1 \cdot \mathbb{E}[Y_i(0) | G_i = 1only] \\ p_2 \cdot \mathbb{E}[Y_i(0) | G_i = 2only] \\ p_a \cdot \mathbb{E}[Y_i(0) | G_i = a.t.] \\ p_n \cdot \mathbb{E}[Y_i(0) | G_i = n.t.] \end{pmatrix},$$

If this equation is written as $b = Ax$, where b is the 8×1 vector of identified quantities, and x the 14×1 unknown vector of potential outcome moments (note the matrix A here is not the same as the matrix A defined in Corollary 6), then the ACL can be written as

$$ACL = \frac{1}{1 - p_a - p_n} \cdot \underbrace{\left(1 \ 1 \ 1 \ 1 \ 1 \ 0 \ 0 \ -1 \ -1 \ -1 \ -1 \ -1 \ 0 \ 0 \right)}_{:=\lambda}' x \quad (21)$$

ACL is identified only if the vector λ is in the row space of matrix A , which follows from the proof of **3 → 2** in Proposition 7. It can be readily verified that

$$A'(AA')^{-1}A\lambda \approx \left(1.5 \ .82 \ .82 \ .73 \ .73 \ .18 \ 0 \ -1.5 \ -.73 \ -.73 \ -.82 \ -.82 \ 0 \ -.18 \right)$$

where $A'(AA')^{-1}A$ is the orthogonal projector into the row space of A (which has full row rank).

Since the RHS of the above is not equal to λ (given explicitly in Eq. 21), λ is not in the row space of A .

E.10 Proof of the Corollary to Theorem 1

The proof of Lemma 1 shows that $(1, \Gamma'_i)A = \mathbf{z}'_i$, where \mathbf{z}_i is a vector of indicators $\mathbb{1}(Z_i = \mathbf{z})$ for each of the values $\mathbf{z} \in \mathcal{Z}$. We can thus write $h(Z_i)$ from Theorem 1 as

$$\begin{aligned} h(Z_i) &= \lambda' \Sigma^{-1} (\Gamma_i - \mathbb{E}[\Gamma_i]) = (0, \lambda') \mathbb{E}[(1, \Gamma'_i)'(1, \Gamma'_i)]^{-1} (1, \Gamma'_i)' \\ &= (0, \lambda') \mathbb{E}[A'^{-1} A'(1, \Gamma'_i)'(1, \Gamma'_i) A A^{-1}]^{-1} A'^{-1} \mathbf{z}_i \\ &= (0, \lambda') A \mathbb{E}[\mathbf{z}_i \mathbf{z}'_i] \mathbf{z}_i \end{aligned}$$

This is useful because $\mathbb{E}[\mathbf{z}_i \mathbf{z}'_i]$ is diagonal, since the events that Z_i take on two different values are exclusive: $\mathbb{E}[\mathbf{z}_i \mathbf{z}'_i] = \text{diag}\{P(Z_i = \mathbf{z})\}_{\mathbf{z} \in \mathcal{Z}}$.

Now, for $V \in \{Y, D\}$, $\mathbb{E}[h(Z_i)V_i] = (0, \lambda') A \text{diag}\{P(Z_i = \mathbf{z})\}_{\mathbf{z} \in \mathcal{Z}}^{-1} \{\mathbb{E}[\mathbb{1}(Z_i = \mathbf{z})V_i]\}_{\mathbf{z} \in \mathcal{Z}} = (0, \lambda') A \{\mathbb{E}[V_i | Z_i = \mathbf{z}]\}_{\mathbf{z} \in \mathcal{Z}}$. Thus $(0, \lambda') A$ describes the coefficients in an expansion of $\mathbb{E}[h(Z_i)V_i]$ into CEFs of V_i across the support of Z_i .

E.11 Proof of the Appendix A Proposition

Introduce the notation that \sqcup indicates inclusion of a new set among a family of sets (while \cup continues to indicate taking the union of elements across sets).

For any $S \subseteq \mathcal{M}$ that contains both Z_m^j and $Z_{m'}^j$ for some j and $m < m'$, $g(F \sqcup S)$ and $g(F \sqcup S/\{Z_m^j\})$ generate the same selection behavior for any Sperner family F on all of \mathcal{Z} (this can be seen by mapping the implied selection behavior back to the original discrete instrument Z_j). Thus, we can take \mathcal{G} to exclude such S without loss of generality.

Now, consider the family \mathcal{F} of all $S \subset \mathcal{M}$ that contain at most one Z_m^j for any given j . By the above, this choice of \mathcal{F} satisfies Assumption 3b*. Suppose it did not satisfy Assumption 3a*. Then, there would need to exist a non-zero vector ω such that $P(\sum_{S \in \mathcal{F}} \omega_S Z_{Si} = 0) = 1$ with $Z_{Si} := \prod_{(j,m) \in S} \tilde{Z}_m^j$. This would imply non-invertibility of $\Sigma^* := \mathbb{E}[(1, \Gamma_i)(1, \Gamma_i)']$, where $\Gamma_i := \{Z_{Si}\}_{S \in \mathcal{F}, S \neq \emptyset}$ by the same argument as in the proof of Lemma 1 (Γ_i and a vector of indicators for all $\mathbf{z} \in \mathcal{Z}$ are each related by an invertible linear map), which in turn contradicts the assumption of full support. Note that invertibility of Σ^* is again equivalent to invertibility of $\text{Var}(\Gamma_i)$ as before.

E.12 Proof of Appendix C Proposition 6

Write the parameter of interest Δ_c as θ_Y/θ_D , where for $V \in \{Y, D\}$, $\theta_V = \tilde{\lambda}'\beta_V$ with $\beta_V := \mathbb{E}[\Gamma_i\Gamma_i']^{-1}\mathbb{E}[\Gamma_i'V_i]$ and $\tilde{\lambda} = (0, \lambda')'$. Denote the estimator $\hat{\rho}(\hat{\lambda}, \alpha)$ as $\hat{\Delta}_c$ for shorthand. It takes the form $\hat{\Delta}_c = \hat{\theta}_Y/\hat{\theta}_D$, where $\hat{\theta}_V := (0, \hat{\lambda}')(\Gamma'\Gamma + K)^{-1}\Gamma'V$, and $K = \alpha I$. I keep the notation in terms of K as the first part of the argument below will go through with any diagonal matrix of positive entries, allowing a different regularization parameter corresponding to each singular vector of $\Gamma'\Gamma$. Write each $\hat{\theta}_V := (0, \hat{\lambda}')\hat{\beta}_V^*$ where $\hat{\beta}_V^*$ is the ridge-regression estimate of β_V , and let $\hat{\beta}_V = (\Gamma'\Gamma)^{-1}\Gamma'V$ be the unregularized regression coefficient estimator.

Consider the conditional MSE $M = \mathbb{E}[(\hat{\Delta}_c - \Delta_c)^2 | \Gamma]$. It can be rearranged as:

$$\begin{aligned} M &= \mathbb{E} \left[\left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} - \frac{\theta_Y}{\theta_D} \right)^2 \middle| \Gamma \right] = \frac{1}{\theta_D^2} \mathbb{E} \left[\left((\hat{\theta}_Y - \theta_Y) - \hat{\Delta}_c(\hat{\theta}_D - \theta_D) \right)^2 \middle| \Gamma \right] \\ &= \frac{1}{\theta_D^2} \mathbb{E} \left[(\hat{\theta}_Y - \theta_Y)^2 + \hat{\Delta}_c^2(\hat{\theta}_D - \theta_D)^2 - 2\hat{\Delta}_c(\hat{\theta}_Y - \theta_Y)(\hat{\theta}_D - \theta_D) \middle| \Gamma \right] \end{aligned} \quad (22)$$

For any $V, W \in \{Y, D\}$, and $m \geq 1$:

$$\begin{aligned} \mathbb{E} \left[(\hat{\Delta}_c)^m (\hat{\theta}_V - \theta_V)(\hat{\theta}_W - \theta_W) \middle| \Gamma \right] &= \mathbb{E} \left[(\hat{\Delta}_c)^m (0, \hat{\lambda})'(\hat{\beta}_V^* - \beta_V)(\hat{\beta}_W^* - \beta_W)'(0, \hat{\lambda})' \middle| \Gamma \right] \\ &= (\Delta_c)^m \tilde{\lambda}' \mathbb{E} \left[(\hat{\beta}_V^* - \beta_V)(\hat{\beta}_W^* - \beta_W)' \middle| \Gamma \right] \tilde{\lambda} + R_n^m \end{aligned}$$

where the first term in the above is viewed as an approximation that ignores terms that are of third or higher order in estimation errors. The asymptotic rate at which the approximation error captured by the R_n^m converges to zero is considered explicitly at the end of this section.

Let $Z = (\Gamma'\Gamma + K)^{-1}\Gamma'\Gamma$ and notice that $\hat{\beta}_V^* = \mathbf{z}\hat{\beta}_V$. Using that $\mathbb{E}[\hat{\beta}_V|\Gamma] = \beta_V$ (as Γ includes all products of the instruments the CEF must be linear) for $V \in \{Y, D\}$:

$$\begin{aligned} \mathbb{E} \left[(\hat{\beta}_V^* - \beta_V)(\hat{\beta}_W^* - \beta_W)' \middle| \Gamma \right] &= \mathbf{z} \mathbb{E} \left[(\hat{\beta}_V - \beta_V)(\hat{\beta}_W - \beta_W)' \middle| \Gamma \right] \mathbf{z}' + (\mathbf{z} - I)\beta_V\beta_W'(\mathbf{z} - I)' \\ &= (\Gamma'\Gamma + K)^{-1}(\Gamma'\Omega_{VW}\Gamma + K\beta_V\beta_W'K)(\Gamma'\Gamma + K)^{-1} \end{aligned}$$

where we define the $n \times 1$ vector $U_V = V - \Gamma\beta_V$ and $\Omega_{VW} = \mathbb{E}[U_V U_W' | \Gamma]$. Thus, total conditional

MSE is, by Equation (22):

$$M \approx \frac{1}{\theta_D^2} \tilde{\lambda}' (\Gamma' \Gamma + K)^{-1} \left\{ \Gamma' (\Omega_Y + \Delta_c^2 \Omega_D - 2\Delta_c \Omega_{YD}) \Gamma + K (\beta_Y \beta'_Y + \Delta_c^2 \beta_D \beta'_D - 2\Delta_c \beta_Y \beta'_D) K \right\} (\Gamma' \Gamma + K)^{-1} \tilde{\lambda}$$

This development follows and generalizes that of Hoerl and Kennard (1970), who consider MSE optimal regularization via ridge regression for estimating a single regression vector, under homoscedasticity. Our case targets the ratio $\hat{\theta}_Y/\hat{\theta}_D$ rather than a vector of regression coefficients, and also allows for heteroscedasticity.

We now prove that $\alpha/\sqrt{n} \xrightarrow{P} 0$ if α is chosen to minimize the following “single-step” estimator of the MSE (ignoring the positive factor of θ_D^{-2} that does not depend on K):

$$\hat{M} := \tilde{\lambda}' (\Gamma' \Gamma + K)^{-1} \left\{ \Gamma' \left(\hat{\Omega}_Y + \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right)^2 \hat{\Omega}_D - 2 \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right) \hat{\Omega}_{YD} \right) \Gamma + K \left(\hat{\beta}_Y \hat{\beta}'_Y + \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right)^2 \hat{\beta}_D \hat{\beta}'_D - 2 \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right) \hat{\beta}_Y \hat{\beta}'_D \right) K \right\} (\Gamma' \Gamma + K)^{-1} \tilde{\lambda}$$

where $\left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right)$ is the un-regularized estimator of Δ_c . The problem can be re parameterized as a choice of $b := \alpha/n$, where

$$\begin{aligned} \hat{M}(b) &:= \tilde{\lambda}' \left(\frac{\Gamma' \Gamma}{n} + bI \right)^{-1} \left\{ \frac{1}{n} \frac{\Gamma' \left(\hat{\Omega}_Y + \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right)^2 \hat{\Omega}_D - 2 \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right) \hat{\Omega}_{YD} \right) \Gamma}{n} + \right. \\ &\quad \left. b^2 \left(\hat{\beta}_Y - \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right) \hat{\beta}_D \right) \left(\hat{\beta}_Y - \left(\frac{\hat{\theta}_Y}{\hat{\theta}_D} \right) \hat{\beta}_D \right)' \right\} \left(\frac{\Gamma' \Gamma}{n} + bI \right)^{-1} \tilde{\lambda} \\ &:= m(b, \hat{\Pi}, \hat{\beta}, \hat{\Sigma}, \hat{\lambda}) \end{aligned}$$

where $\hat{\Pi} := \frac{1}{n} \sum_i (\hat{U}_{Yi} - \hat{\theta}_Y/\hat{\theta}_D \hat{U}_{Di})^2 \Gamma_i \Gamma'_i$, $\hat{\beta} := (\hat{\beta}_Y - \hat{\theta}_Y/\hat{\theta}_D \hat{\beta}_D)$, and $\hat{\Sigma}^* := \frac{1}{n} \sum_i \Gamma_i \Gamma'_i$. Note that $\hat{\beta} \xrightarrow{P} \beta := \beta_Y - \Delta_c \beta_D$, $\hat{\Sigma}^* \xrightarrow{P} \Sigma^* := \mathbb{E}[(1, \Gamma'_i)'(1, \Gamma'_i)]$, $\sqrt{n} (\hat{\Pi} - \Pi) \xrightarrow{d} N(0, V)$ for some V provided that the variance of $(\hat{U}_{Yi} - \hat{\theta}_Y/\hat{\theta}_D \hat{U}_{Di})^2 \Gamma_i \Gamma'_i$ exists, where $\Pi := \mathbb{E}[(\hat{U}_{Yi} - \hat{\theta}_Y/\hat{\theta}_D \hat{U}_{Di})^2 \Gamma_i \Gamma'_i]$.

The function m is

$$m(b, \Pi/n, \beta, \Sigma^*, \lambda) = (0, \lambda') (\Sigma^* + bI)^{-1} \{ \Pi/n + b^2 \beta \beta' \} (\Sigma^* + bI)^{-1} (0, \lambda')'$$

We wish to show that $\sqrt{nb} = \alpha/\sqrt{n} \xrightarrow{P} 0$, when b is chosen as the smallest positive minimizer of $m(\cdot, \hat{\Pi}/n, \hat{\beta}, \hat{\Sigma}, \hat{\lambda})$. The strategy will be to show that $nb \xrightarrow{P} X$ where X is a finite degenerate random variable. Since Π and $\beta\beta'$ are positive definite, it is clear that $m(b, \Pi/n, \beta, \Sigma^*, \lambda)$ is weakly positive for any choice of b . Further, $m(b, \Pi/n, \beta, \Sigma^*, \lambda)$ is typically strictly positive at $b = 0$, and it can also be seen that $\lim_{b \rightarrow \infty} m(b, \Pi/n, \beta, \Sigma^*, \lambda) = 0$ (see Section C.1 for discussion). However, m is generally not monotonically decreasing in between, as we shall see below.

Observe that $b = 0$ minimizes $m(b, \mathbf{0}, \beta, \Sigma^*, \lambda)$ with respect to b regardless of the values of β, Σ^*, λ , where $\mathbf{0}$ is a $k \times k$ matrix of zeros (the dimension of Π), since $m(\cdot)$ is always positive and when its second argument vanishes can be made equal to zero by choosing $b = 0$. Furthermore, $b = 0$ is a local minimizer when $\Pi/n = \mathbf{0}$, since m_b vanishes when evaluated at $(0, \mathbf{0}, \beta, \Sigma^*, \lambda)$ —see below, while the second derivative of m with respect to b , evaluated at $(0, \mathbf{0}, \beta, \Sigma^*, \lambda)$, is equal to

$$(0, \lambda') \Sigma^{*-1} \beta \beta' \Sigma^{*-1} \lambda = ((0, \lambda') \Sigma^{*-1} \beta)^2$$

up to a strictly positive constant. We have assumed that the quantity in parenthesis is non-zero. By the implicit function theorem, there then exists a unique function $g(\Pi/n; \beta, \Sigma^*, \lambda)$ such that $g(\mathbf{0}; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) = 0$ and $m_b(g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}), \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) = 0$, in a neighborhood \mathcal{N} of the probability limits $(\mathbf{0}, \beta, \Sigma^*, \lambda)$ of $(\hat{\Pi}/n, \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})$, and this function is continuously differentiable with respect to all parameters, (including, in particular, the elements of Π). Since the second derivative of m is strictly positive at $(0, \mathbf{0}, \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})$ and continuous with respect to all arguments, \mathcal{N} can furthermore be chosen such that the critical point at $(g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}), \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})$ is always a local minimum within \mathcal{N} .

Since for any realization of $\hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}$:

$$m_b(0, \mathbf{0}, \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) = 2\tilde{\lambda}'(\hat{\Sigma}^* + bI)^{-1} \left\{ bI - b^2(\hat{\Sigma}^* + bI)^{-1} \right\} \hat{\beta}\hat{\beta}'(\hat{\Sigma}^* + bI)^{-1} \tilde{\lambda} \Big|_{b=0} = 0$$

we see that m has a critical point at $b = 0$ for values $(\mathbf{0}, \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})$ of the other arguments. By uniqueness of the function $g(\Pi/n; \beta, \Sigma^*, \lambda)$, this implies then that $g(\mathbf{0}, \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) = 0$. By the

mean value theorem, we can write

$$\begin{aligned} g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) &= g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) - g(\mathbf{0}, \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) \\ &= \frac{\partial}{\partial x} g(vec(cn^{-1}\hat{\Pi}); \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) \cdot \frac{vec(\hat{\Pi})}{n} \end{aligned}$$

for some $c \in [0, 1]$, where $vec(\Pi)$ denotes the vectorization x of the matrix Π , and we let $\frac{\partial}{\partial x} g(x; \beta, \Sigma^*, \lambda)$ denote a gradient of g with respect to that vector (recall that existence of the derivative is a consequence of the implicit function theorem). By continuity of $\frac{\partial}{\partial x} g(x; \beta, \Sigma^*, \lambda)$ and the continuous mapping theorem then,

$$n \cdot g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda}) \xrightarrow{p} \frac{\partial}{\partial x} g(\mathbf{0}, \beta, \Sigma^*, \lambda) vec(\Pi) \quad (23)$$

which is a finite scalar.

To complete the proof, we now simply note that with probability approaching unity, $(\hat{\Pi}/n, \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})$ is within the neighborhood \mathcal{N} , and thus if b is chosen as the smallest positive local minimizer of $m(b, \hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})$ we have that $b = g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})$. We have now established the result, since for any $B > 0$:

$$\begin{aligned} P(|\alpha/\sqrt{n}| > B) &\leq P(|\alpha/\sqrt{n}| > B \text{ and } b = g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})) + P(b \neq g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})) \\ &= P(|n \cdot g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})| > \sqrt{n}B) + P(b \neq g(\hat{\Pi}/n; \hat{\beta}, \hat{\Sigma}^*, \tilde{\lambda})) \\ &\xrightarrow{n} 0 + 0 \end{aligned}$$

Finally, I consider the error involved in the approximation made to Equation (22). Write this as:

$$\begin{aligned} R_n &:= R_n^m + R_n^m = \\ &= \frac{1}{\theta_D^2} \tilde{\lambda}' (\Gamma' \Gamma + K)^{-1} \left\{ (\hat{\Delta}_c^2 - \Delta_c^2) (\Gamma' \Omega_D \Gamma + K \beta_D \beta'_D K) \right. \\ &\quad \left. - 2(\hat{\Delta}_c - \Delta_c) (\Gamma' \Omega_{YD} \Gamma + K \beta_Y \beta'_D K) \right\} (\Gamma' \Gamma + K)^{-1} \tilde{\lambda} \\ &= \frac{1}{\theta_D^2 \cdot n^{3/2}} \cdot \tilde{\lambda}' \left(\frac{\Gamma' \Gamma}{n} + \frac{K}{n} \right)^{-1} \left\{ \sqrt{n}(\hat{\Delta}_c^2 - \Delta_c^2) \left(\frac{\Gamma' \Omega_D \Gamma}{n} + \frac{K}{\sqrt{n}} \beta_D \beta'_D \frac{K}{\sqrt{n}} \right) \right. \\ &\quad \left. - 2\sqrt{n}(\hat{\Delta}_c - \Delta_c) \left(\frac{\Gamma' \Omega_{YD} \Gamma}{n} + \frac{K}{\sqrt{n}} \beta_Y \beta'_D \frac{K}{\sqrt{n}} \right) \right\} \left(\frac{\Gamma' \Gamma}{n} + \frac{K}{n} \right)^{-1} \tilde{\lambda} \end{aligned}$$

Provided that $\alpha/\sqrt{n} \xrightarrow{p} 0$ as above, we will show in Theorem 3 that $\hat{\Delta}_c$ is \sqrt{n} -consistent for Δ_c . In this case, the approximation error term is $O_p(n^{-3/2})$.

E.13 Proof of Appendix C Theorem 3

When $\alpha_n = 0$, the result follows from Theorem 3 of Imbens and Angrist (1994). To see that $o_p(\sqrt{n})$ regularization has no asymptotic effect, note that

$$\begin{aligned}(0, \hat{\lambda}')'(\Gamma'\Gamma + \alpha I)^{-1}\Gamma'Y &= (0, \hat{\lambda}')'(\Gamma'\Gamma + \alpha I)^{-1}(\Gamma'\Gamma + \alpha I - \alpha I)(\Gamma'\Gamma)^{-1}\Gamma'Y \\ &= (0, \hat{\lambda}')'(\Gamma'\Gamma)^{-1}\Gamma'Y - \alpha(0, \hat{\lambda}')'(\Gamma'\Gamma + \alpha I)^{-1}(\Gamma'\Gamma)^{-1}\Gamma'Y\end{aligned}$$

and similarly for D , thus:

$$\begin{aligned}\rho(\hat{\lambda}, \alpha) &= \frac{(0, \hat{\lambda}')'(\Gamma'\Gamma)^{-1}\Gamma'Y - \alpha(0, \hat{\lambda}')'(\Gamma'\Gamma + \alpha I)^{-1}(\Gamma'\Gamma)^{-1}\Gamma'Y}{(0, \hat{\lambda}')'(\Gamma'\Gamma)^{-1}\Gamma'D - \alpha(0, \hat{\lambda}')'(\Gamma'\Gamma + \alpha I)^{-1}(\Gamma'\Gamma)^{-1}\Gamma'D} \\ &= \frac{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), Y_i) - \frac{\alpha}{n}(0, \hat{\lambda}')'(\frac{1}{n}\Gamma'\Gamma + \frac{\alpha}{n}I)^{-1}(\frac{1}{n}\Gamma'\Gamma)^{-1}\frac{1}{n}\Gamma'Y}{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), D_i) - \frac{\alpha}{n}(0, \hat{\lambda}')'(\frac{1}{n}\Gamma'\Gamma + \frac{\alpha}{n}I)^{-1}(\frac{1}{n}\Gamma'\Gamma)^{-1}\frac{1}{n}\Gamma'D} \\ &= \frac{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), Y_i)}{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), D_i)} + \frac{\alpha}{n} \cdot \frac{(0, \hat{\lambda}')'(\frac{1}{n}\Gamma'\Gamma + \frac{\alpha}{n}I)^{-1}(\frac{1}{n}\Gamma'\Gamma)^{-1} \left\{ \frac{1}{n}\Gamma'D \cdot \frac{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), Y_i)}{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), D_i)} - \frac{1}{n}\Gamma'Y \right\}}{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), D_i) - \frac{\alpha}{n}(0, \hat{\lambda}')'(\frac{1}{n}\Gamma'\Gamma + \frac{\alpha}{n}I)^{-1}(\frac{1}{n}\Gamma'\Gamma)^{-1}\frac{1}{n}\Gamma'D}\end{aligned}$$

and thus the asymptotic distribution of $\sqrt{n}(\hat{\rho}(\hat{\lambda}, 0) - \Delta_c)$ is the same as that of $\sqrt{n} \left(\frac{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), Y_i)}{\widehat{\text{Cov}}(g(Z_i, \hat{\theta}), D_i)} - \Delta_c \right)$, provided that $\alpha_n/\sqrt{n} \xrightarrow{p} 0$ (in which case the second term above is $o_p(n^{-1/2})$).

**Chapter 3: The Career Impact of First Jobs: Evidence and Labor Market
Design Lessons from Randomized Choice Sets**

The Career Impact of First Jobs: Evidence and Labor Market Design Lessons from Randomized Choice Sets*

Ashna Arora
University of Chicago
& Statistics Norway

Leonard Goff
Columbia University

Jonas Hjort
Columbia University
& CEPR & NBER

May 19, 2021

PRELIMINARY AND INCOMPLETE

Abstract

Does a worker's first job affect her long-run career? If so, can policy improve upon a "free" labor market by altering initial matches with employers? We begin to study the impact of market design on the performance of entry-level labor markets, by comparing 20 years when Norway assigned doctors to their first job—residencies—through a Random Serial Dictatorship (RSD), with the post-2013 era when the RSD mechanism was replaced with decentralized job-finding. We first estimate the consequences of different employers for the earnings and various long-run outcomes for male and female workers. We do so by exploiting random, individual level variation in workers' initial choice set over first-job employers generated by the RSD, using an instrumental variables framework. We then decompose preferences over first jobs into a component that is due to earnings first-job effects and another that is due to the "amenity value" workers of a given type associate with employers of a given type. This allows us to show how realized first job effects, amenity values, and overall worker welfare differ in a decentralized labor market compared to the RSD system, by observing how worker-employer matches changed after 2013.

*ashnaarora@uchicago.edu, leonard.goff@columbia.edu, hjort@columbia.edu. We thank Josh Angrist, Chris Conlon, Francois Gerard, Adam Kapor, José Luis Montiel Olea, Ben Olken, Paul Oyer, Parag Pathak, Jonah Rockoff, Till von Wachter, and seminar audiences at several universities and conferences for helpful suggestions. We are especially grateful to Andreas Fagereng, with whom we started this project. The project received funding from the Research Council of Norway (#256678).

1 Introduction

An individual's first job may have important consequences for her career trajectory. This view—common among researchers—appears widely held also among those entering the labor market. New job-seekers may therefore put weight on the expected impact of different jobs on their trajectories when choosing a first job to pursue. Policy, on the other hand—whether centralized mechanisms allocating doctors, teachers, and other groups of workers serving the public to first jobs¹, or the rules and regulations that influence initial worker-job matches in the decentralized labor market—is typically designed without accounting for expected “first job effects” (FJEs).

If FJEs are non-negligible in magnitude and heterogeneous across types of workers, then FJE-responsive policy design could in principle be used to increase welfare. However, even in such a scenario, whether alternative policies *actually* affect initial worker-job matches—and hence realized FJEs—differentially is an empirical question. Unusual types of data and variation are necessary for researchers to be able to identify FJEs and graduates' and policymakers' actual and ideal response. To estimate the causal, long-term effect of an individual's first job, random variation in her match, holding all else constant, is needed. To estimate how individuals' distribution of FJEs across jobs influence their job search choices, causal estimates of the long-term effect of each type of job for each type of individual—and knowledge of individuals' choice set when entering the labor market—are needed.

In this paper we take advantage of Norway's 1997-2013 allocation of doctors' first job—their residency—through a Random Serial Dictatorship (RSD) mechanism², and the replacement of the RSD with decentralized job-finding in 2013, to overcome these challenges. We first estimate the consequences for earnings, place of residence, and specialization in the long-term of each type of job characteristic separately for men and women. We do this by exploiting RSD-generated random, individual level variation in new doctors' choice sets over first jobs. In the last part of the paper, we use a 2013 policy change—which replaced the RSD system with a regular, decentralized job market for new doctors—to assess how total worker welfare, initial worker-job matches, and the associated realized FJEs, differ in a market system relative to RSD.

The unique suitability of doctors' residencies in Norway for studying FJEs is due to the combination of choice sets over jobs being assigned randomly, and the unusually high quality of the registry data available on the universe of Norwegian workers. While our quantitative results may not generalize to other occupations, it is worth noting that (i) the differences between the possible pathways a doctor's career can take share many features with those in other occupations³, and

¹The following is an incomplete list of countries that use centralized mechanisms to assign workers in some (in some of the countries, almost all) public service occupations to first jobs: Australia, Bangladesh, Bhutan, Botswana, Canada, Denmark, France, Ghana, India, Iran, Ireland, Israel, Italy, Japan, Malaysia, Malta, Nepal, Norway, Pakistan, Philipines, Saudi Arabia, Senegal, Singapore, South Africa, South Korea, Taiwan, Tanzania, Uganda, U.K., USA.

²A Random Serial Dictatorship mechanism starts with a lottery. The person who draws number 1 then chooses her preferred object freely among all available options. After that, the person who draws number 2 chooses among the remaining objects, and so on.

³For example, doctors' jobs are located in many different parts of a country; there is considerable dispersion in employer size and "quality" (which is highly correlated with doctors' earned income); and there are ample opportunities for

(ii) the literature generally finds that highly skilled workers are least affected by temporary career shocks (Oreopoulos *et al.*, 2012; von Wachter & Bender, 2006). Most likely our results thus represent a lower bound on FJEs and the associated responses in other occupations.

This paper contributes to the literature on how temporary shocks to a worker's employment status affects her career trajectory. Existing studies have convincingly and carefully documented the consequences of job displacement (see, among many others, Bender *et al.*, 2009; Sullivan & von Wachter, 2009; von Wachter & Bender, 2006; ?); exposure to high unemployment rates later in life (Coile *et al.*, 2012); regional labor demand composition (Arellano-Bover, 2020); and, most closely related to this paper, graduating in a weak labor market (Genda *et al.*, 2010; Heisz *et al.*, 2012; Kahn, 2010; Oyer, 2006, 2008). These influential studies have shown how cohort and group-level labor market shocks affect individual workers' trajectories.

In addition to taking advantage of an explicit randomization for identification, this paper to our knowledge provides among the first causal evidence on the career consequences of *individual* level shocks to a graduate's first job.⁴ The distinction is essential because cohort level studies may not be informative about the career consequences of individual level labor market shocks, which are ubiquitous. When the cohort or group an individual belongs to is hit, for example, by a recession or a mass layoff, then the individual's peers are also affected. Peers' exposure to the shock could adversely affect the trajectory of the individual in question (if for example she now faces more competition for current jobs) or benefit her (if for example she's now competing against less employable other workers for future jobs).⁵

To leverage the RSD to estimate first job effects, we develop an instrumental variables (IV) approach that explicitly allows for the effect of beginning one's career at a given employer to vary by individual. This setting extends beyond standard results in the treatment effects literature, as a worker's first job "treatment" is the identity of their first-job employer. We thus contribute to a recent literature that extends IV research designs to unordered, multivalued treatment variables (Heckman & Pinto, 2018; Lee & Salanié, 2020; ?). Furthermore, relative to recent work that has used other centralized assignment mechanisms for causal inference (Abdulkadiroğlu *et al.*, 2017; Kirkeboen *et al.*, 2017), we do not observe workers preference orderings over outcomes of the matching mechanism. We show that first job effects can still be partially identified, by making use of the observation that variation in mean outcomes across choice sets is informative about heterogeneity in the counterfactual outcomes for a given employer across workers of different unobserved preference types. This builds upon a recent linear programming approach to IV analysis (Kamat, 2020; Mogstad *et al.*, 2018), as well as a recent literature on causal inference with a collection of distinct

doctors to undertake horizontal specialization (choice of medical field) as well as vertical specialization (e.g. becoming a specialist as opposed to a General Practitioner).

⁴To our knowledge, the only existing causal evidence on the long-term effects of individual level shocks to first jobs comes from Angrist (1990)'s seminal study of the Vietnam draft. He shows that being drafted lowered earnings by 15 percent long after the veterans' service ended (see also Angrist, 1995; Angrist & Chen, 2011). Staiger (2020) use job openings at the employer of young worker's parents to generate individual-level variation in early job opportunities.

⁵Ruhm (2000) shows that mortality tends to improve during recessions, while Sullivan & von Wachter (2009) show that *own* job displacements increase mortality for U.S. workers.

instrumental variables ([Goff, 2020](#); [Mogstad *et al.*, 2020](#)).

With estimates of FJEs in hand, our final contribution is to assess the impact of the 2013 reform to a decentralized labor market on the welfare of workers. To do so, we decompose preferences over employers into a component that is due to first job effects and another that is due to the “amenity value” workers of a given type associate with employers of a given type.⁶ This leverages lottery draws as a reduced form measure of preferences, coupled with a high-level assumption on the distribution of preferences in the lottery. We show how realized first job effects, amenity values, and overall worker welfare differ, for each group and in total, in a decentralized labor market compared to the RSD system, by examining how w_{ib} by examining how worker-employer matches changed after 2013.

The paper is organized as follows. In Section 2 we discuss background on the setting and institutional setup. In Section 3 we present the datasets used in our empirical analysis. In Section 4 we lay out our instrumental variables strategy to estimating first job effects, and present results in Section 5. Section 6 then applies these results to investigate worker preferences and evaluate the reform to a decentralized labor market from those workers’ perspective.

2 Setting

2.1 The Random Serial Dictatorship mechanism for Norwegian doctors

The “turnus” (roster) system that was used to match medical graduates with residency positions in Norway from 1954 to 2013 was a Random Serial Dictatorship (RSD) mechanism. Theorists have shown that, among other important properties, the RSD is incentive-compatible, inducing participants to reveal their true preferences ([Abdulkadiroğlu & Sönmez, 1998](#)).

Equitable access to primary healthcare across regions was the main motivation behind the use of a lottery system in Norway. Like other countries, Norway had had trouble filling doctor vacancies in rural areas, and the RSD mechanism was expected to distribute the best doctors more equally across space.⁷ In addition, the mechanism appealed to policymakers because it was perceived to be fair to the participating medical graduates.⁸

First, graduating students would enter a lottery, either in February or in August, and be assigned a random draw number. Next, the student with the lowest draw number would choose freely between all available positions. Then, the student with the second lowest draw number would choose from the remaining residency positions. This would continue until the student with

⁶This paper is of course not the first to recognize that workers care about non-income job characteristics and may choose occupations and employers partly based on those preferences. Recently, for example, [Sorkin \(2016\)](#) showed evidence of compensating differentials revealed in workers’ job-to-job transitions in the U.S.

⁷Such considerations have been studied in the context of the US National Medical Residency Matching Program, see e.g. (e.g. ??).

⁸The government also wished to incentivize doctors to work in rural locations in other ways. For instance, doctors who agreed to intern at hospitals in the largely rural counties of Sogn og Fjordane and Finnmark could skip the lottery entirely.

the highest draw number remained, who would take whichever spot was available.⁹

Three categories of new doctors received special treatment: couples, who were allowed to draw a shared lottery draw number and to choose residencies simultaneously; doctors with children; and doctors with maternity or health issues. The latter two categories were allowed to choose between positions deemed especially suitable for them before the lottery took place. Since these three types of doctors were not subject to randomization via the lottery, we exclude them from our analysis.

In the late 2000s, the system began to concern the government, because of the growth of the number of applicants and the rise in proportion of students from foreign universities.¹⁰ The number of medical graduates participating in the lottery would routinely exceed the number of training positions available. As a result, it became increasingly difficult for the government to guarantee a six-month maximum waiting time to obtain a residency. In 2013, the Norwegian Health Minister replaced the lottery system with direct qualification after six years of medical school. Medical graduates now apply to residencies directly, as in a regular labor market, and hospital trusts are responsible for selection and recruitment.

2.2 Doctors in Norway 1993-2017

This section profiles doctors that worked in Norway during the study period; we go through the data we use in detail in Section 3. Medical students in Norway begin their studies in the Fall or Spring semester, and usually take ten semesters to graduate. Starting in the 1950s, the Norwegian government mandated an eighteen month residency period, after which medical school graduates could become fully licensed physicians and practice independently. The first twelve months were to be spent at a hospital, while the remaining six months were to be spent as a General Practitioner (i.e. one who works in Primary Care) within the same county.¹¹

Table C.1 summarizes a range of socioeconomic information on doctors including age, proportion born abroad, proportion that studied abroad, family size, field of specialization and income and assets. The last two columns summarize this information separately for men and women. Women comprise over 40 percent of doctors, and tend to be over-represented in fields like gynecology and psychiatry. Male doctors are older and tend to be over-represented in fields like surgery and internal medicine. A fifth of all doctors were born abroad, of which an overwhelming proportion are citizens of Denmark and Sweden. Finally, recorded income, asset and debt holdings are higher on average for male doctors.

There are 30 basic medical specialties, and specialization is usually in the form of training on the job.¹² The average length of time required to complete a specialty is five years, but it can take

⁹If the number of students exceeded the number of residency positions, the unassigned students would get priority in the next lottery.

¹⁰Norway was compelled by its participation in the EU common labor market system to accept any European medical graduate who could pass a Norwegian language test into the system.

¹¹The last six months could be spent at an institution that was disjoint from the hospital of the first twelve months.

¹²The Norwegian Medical Association evaluates whether the candidate has met the requirements to become a spe-

longer with large variations between the specialties. Figure C.4 indicates that there appear to be substantial returns to specializing—job retention rates are higher for specialists, and the salary bump from specialization increases with age.

2.3 Hospitals in Norway 1993-2017

The employer-employee database contains information on all registered employers that employ doctors in Norway.¹³ Figure C.5 depicts the steady growth in both the number of hospitals and average hospital size (number of doctors employed) since 1995.

Hospitals vary across multiple dimensions. Table C.1 summarizes information on salaries, geographical remoteness, number of doctors and other medical staff, proportion of specialists, as well as the presence of fifteen distinct specialist fields. Most hospitals in Norway are located in urban municipalities; on average, municipalities with hospitals have only 10 per cent of their population living in rural areas. This is noteworthy because the average municipality in Norway had 49 percent of its population living in rural areas.

3 Data

We combine information on lottery outcomes with Norwegian administrative data from 1993 to 2017. We obtained information on lottery draw numbers for all lottery participants who were assigned a residency position during 1993-2013 from the Norwegian Registration Authority for Health Personnel (SAFH). This information was linked with the employer-employee registry to match medical graduates to their residency hospitals, as well as employer information in the years following the residency. This data was then linked to administrative registers provided by Statistics Norway, a rich longitudinal database that includes information on medical graduates' socioeconomic information (sex, age, marital status, educational attainment, specialization, income, and gross wealth), geographical identifiers and year-end asset holdings and liabilities (such as real estate, stock holdings, etc) for each year. These data have several valuable attributes. There is no attrition from the sample, and most components of income and wealth are third-party reported without any top or bottom coding.

The final dataset tracks the career path of each graduate, starting with her lottery number and choice of residency hospital. After excluding people belonging to special lottery categories and hospitals with missing information, we end up with a sample of about 9000 individuals and 55 hospitals, which participate in 34 lotteries from 1996 to 2012.¹⁴ Figure C.6 displays the number of individuals and hospitals that participated in each lottery. Figure C.7 splits participants by gender and by birth location. It is evident that there is an increase in the proportion of women and foreign students over time. Most foreign doctors are citizens of the European Economic Area (EEA).

cialist. Specialist titles are formally awarded by the Health Directorate.

¹³We define hospitals as employers that hire at least 10 doctors. These account for around 80 per cent of doctor employment.

¹⁴Data is missing for the lottery in January 1998.

We observe employment outcomes for all doctors up until the year 2017. This allows us to track doctors who graduated in the earliest lotteries (during the 1990s) for over fifteen years, while participants in the last few lotteries (in the 2010s) can only be tracked for a few years. Data from the 2013-2017 period is used to analyze the last cohorts in the lottery system, as well as to observe hospital matches after the 2013 reform. Figure C.8 displays the distribution of the number of times doctors are observed in the years following their residencies. Our sample consists of roughly 4500 individuals five years after their residency, but less than a quarter of these are observed 10 years down the line.

We construct the choice set of hospitals faced by each lottery participant using her lottery number and the residency hospitals chosen in that lottery. We know that if a hospital h was chosen by someone with a higher (worse) lottery number than individual i , i must have been given the option of choosing h as well (since hospitals cannot reject applicants). Assuming that no residency spots were left unfilled,¹⁵ we can thus impute the choice sets C_i that were offered to each lottery participant. Most medical graduates have a sizable number of residency options to choose from, as displayed in Figure C.8.

4 Identification of first job effects

In this section we describe how we use these RSD-generated choice sets to generate instruments for a doctor’s residency hospital. To abstract somewhat from our specific context, we will typically refer to “workers” choosing a first-job “employer”, rather than “doctors” choosing a residency “hospital”.

The basic intuition behind our approach is that randomization of the RSD lottery ensures that each worker’s choice set is independent of her observable and unobservable characteristics. At the same time, this choice set affects that worker’s realized first-job, which is constrained to be within the randomly assigned choice set. This allows us to use the RSD lottery to construct instruments for first job effects (FJEs) that are both exogenous and relevant.

We work in three steps. First, we show in subsection 4.2 that FJEs are identified in the RSD without any substantive assumptions on workers’ selection behavior—but that this strategy requires conditions on the support of the data that do not hold in our context. We then show in subsection 4.3 that FJEs can be identified under weaker conditions on the support of the data—but that this strategy requires strong assumptions on workers’ selection patterns. This culminates in subsection 4.4, which simultaneously relaxes both assumptions at the expense of yielding bounds rather than a point identification result.

¹⁵This assumption is reasonable, in part because excess demand was one of the reasons for replacing the turnus system after 2013.

4.0 Notation

We begin by establishing some notation that will be used throughout this section. Let h denote employers, \mathcal{H} the set of all employers, and i workers in population \mathcal{I} . Let $Y_i(h, c)$ be the potential outcome (e.g. earnings four years later) of worker i if their first job is at employer h , and their choice set from the lottery was c . We assume that choice sets do not effect outcomes except through a doctor's first job employer, that is: $Y_i(h, c) = Y_i(h)$ for all $i \in \mathcal{I}$, $h \in \mathcal{H}$, and $c \subseteq \mathcal{H}$. Since choice sets play the role of instruments, this constitutes the standard IV exclusion restriction in our context.

Let $C_i \subseteq \mathcal{H}$ be a worker's realized choice set, and H_i their realized choice from C_i . Let n be the number of workers and $J = |\mathcal{H}|$ the number of employers. Let $L_i \in \{1 \dots 34\}$ be an identifier for the lottery among the 34 between 1996 and 2013 in which worker i was allocated their first job. Let $\mathcal{R}_i \in \{1, 2, \dots\}$ denote worker i 's random place-in-line in their lottery, and let $R_i = F_{\mathcal{R}|L=L_i}(\mathcal{R}_i)$ be this lottery draw normalized to the unit interval within each lottery.

We will assume that each worker has a complete preference relation over hospitals, and is indifferent between no two hospitals. We denote by $h \succ_i h'$ if i prefers h to h' , and let \succ_i alone denote i 's entire preference relation over \mathcal{H} . For any choice set $c \subseteq \mathcal{H}$ denote the most-preferred $h \in c$ according to \succ as $H_\succ(c)$, and write $H_i(c) = H_{\succ_i}(c)$ for shorthand. Note that θ_i is isomorphic to the vector $\{H_i(c)\}_{c \subseteq \mathcal{H}}$. Similarly, let \mathcal{Y}_i be defined in isomorphism with the vector $\{Y_i(h)\}_{h \in \mathcal{H}}$.

We will often treat \succ_i and C_i as random variables, although a realization of each is a set and a relation on a set, respectively. However, in light of the isomorphism mentioned above, we can view each preference relation \succ_i as element of \mathbb{Z}^{2^J} , with integers indicating the index of $H_i(c)$ for each $c \subseteq \mathcal{H}$. Similarly, we can also view any choice set C_i as an element of $\{0, 1\}^{\times J}$, with each component indicating the presence or absence of a hospital $h \in \mathcal{H}$.

4.1 Exogeneity of choice sets

The randomization of each lottery ensures that \mathcal{R}_i is independent of potential outcomes $Y_i(h, c)$, conditional on lottery L_i (we will sometimes use the term "cohort" interchangeably). However, with a finite number of workers participating in the lottery, it does not immediately follow that a worker i 's probability distribution over possible choice sets C_i is perfectly independent of her characteristics.¹⁶ Rather, any two workers having *the same preferences* \succ will receive a C_i drawn from the same probability distribution within a cohort (Abdulkadiroğlu *et al.* 2017). Thus choice sets are independent of potential outcomes, conditional on preferences (and lottery). Since preferences are unobserved in our data, we cannot directly control for them.

Let us partition the population of workers into a set of groups $g \in \mathcal{G}$ on the basis of observable demographic variables G_i . These groups will play a central role in our analysis, allowing us to

¹⁶ To see this, consider a small economy in which there are two workers and two employers, each with one spot available. Worker 1 prefers employer A to B, and Worker 2 prefers B to A. Worker A then has a 50% chance of having $C_i = \{A, B\}$ (if she goes first in the lottery), and a 50% chance of having $C_i = \{A\}$ (if B chooses first). By contrast, Worker B has a 50% chance of having $C_i = \{A, B\}$ and a 50% chance of having $C_i = \{B\}$. The probability distribution over choice sets facing worker i depends on the preferences of each worker *except for i*.

examine observable heterogeneity in first-job effects. We will assume that after conditioning on instance of the lottery and a value of g , choice-sets are as good as randomly assigned:

Assumption 1 ((independence of choice sets)).

$$\{(\mathcal{Y}_i, \succ_i) \perp C_i\} \mid (G_i, L_i)$$

One instance in which Assumption 1 would follow directly from randomization of the lottery is if preferences were perfectly homogeneous within each value of g , e.g. all Norwegian men have the same ranking over employers.¹⁷ In this case, Assumption 1 echoes Proposition 1 of [Abdulkadirloğlu et al. \(2017\)](#) and first-job effects could be assessed without the use of instrumental variables. Indeed, such preference homogeneity along with Assumption would imply that conditional on G_i , a worker's actual choice of first-job H_i is exogenous, since then \succ_i has a degenerate distribution conditional on G_i and $H_i = H_{\succ_i}(C_i)$.

However, we will not maintain the strong assumption that available observable proxies are sufficient to control for unobserved preferences \succ_i . Rather, we offer a second justification for Assumption 1. When the number of workers is “large” in comparison with the number of employers, independence will hold approximately even when there is heterogeneity in preferences within groups. As described formally in Appendix B, the actual set of n workers is viewed as a sample from an underlying continuum of workers, with each employer accounting for a fixed proportion of the available jobs. In this “continuum economy”, choice sets are random unconditionally. The IV estimators we use are then consistent along an asymptotic sequence in which $n \rightarrow \infty$ with the share of jobs belonging to each employer fixed. In Appendix B we provide evidence that this asymptotic approximation is a good one in our context, by simulating the lottery many times with a number of workers and employers chosen to match our dataset, and plausible heterogeneity in preferences.

4.2 Choice sets as instruments

As a parameter of interest, we are interested in the quantity

$$\mu_{gh} := \mathbb{E}[Y_i(h)|G_i = g]$$

for an individual employer h and demographic group g . The parameter μ_{gh} is the average counterfactual outcome that would occur for a worker in group g if their first job were at employer h , and $\mu_{gh'} - \mu_{gh}$ is the average effect of “moving” workers in g from employer h to employer h' . Note that μ_{gh} generally differs from the observable $E[Y_i|H_i = h, G_i = g] = E[Y_i(h)|H_i = h, G_i = g]$, which unlike μ_{gh} further conditions on the worker's endogenous choice of employer h . Given that randomization of choice sets holds only when conditioning on lottery $L_i = \ell$, we also define a

¹⁷We can formalize randomization of the lottery as the statement that $\{\mathcal{R}_i \perp (\mathcal{Y}_i, \succ_i, G_i)\}|L_i$.

lottery-specific analog of μ_{gh} that will be useful as an intermediate quantity:

$$\mu_{gh\ell} := \mathbb{E}[Y_i(h)|G_i = g, L_i = \ell]$$

With Assumption 1 in hand, we seek to use features of a worker's choice set C_i to construct instruments for the causal effect μ_{gh} of her first job. Assuming that a worker will always choose some employer from their choice-set, we can in principle identify first job effects μ_{gh} without placing any further restrictions on selection. The following Proposition is a consequence of Theorem 1 in (Goff, 2020); however it admits of a very simple proof that we present here.

Proposition 1 ((impractical identification)). *Make Assumption 1. If for a given h, ℓ : $P(C_i = \{h\}|G_i = g, L_i = \ell) > 0$, then*

$$\mu_{gh\ell} = \mathbb{E}[Y_i|C_i = \{h\}, G_i = g, L_i = \ell]$$

provided that $\forall i, H_i(\{h\}) \neq \emptyset$.

Proof. Given that $H_i \in C_i$ and $H_i = H_i(C_i) \neq \emptyset$, we must have $H_i = h$ whenever $C_i = \{h\}$, so: $\mathbb{E}[Y_i|C_i = \{h\}, G_i = g] = \mathbb{E}[Y_i(h)|C_i = \{h\}, G_i = g] = \mu_{gh}$, where the last equality follows from Assumption 1. \square

Proposition 1 shows that if there is some probability that each singleton $\{h\}$ emerges as a worker's choice set, then the μ_{gh} are identified. This requires us to assume that the worker will choose each h over no employer (all hospitals are preferred to the relevant outside option), but requires no other assumptions on workers' choices. For instance, it does not require us to assume that worker's choose rationally, according to well-defined preferences.

In practice however, the event $C_i = \{h\}$ is unlikely to occur for popular employers h , and even for an unpopular h it will only occur at most for the last few workers in a given lottery. As a result, Proposition 1 is not directly useful in estimation. In the next section, we thus consider a more practical route to identification under a restriction on heterogeneity, before returning to the general case in a partial identification framework in Section 4.4.

4.3 Identification with a restriction on treatment effect heterogeneity

The discussion of the last section has shown that first job effects μ_{hg} are identified even with complete heterogeneity of treatment effects, provided that there is a positive probability that some workers will face a choice-set containing only the single employer h . In practice, this result is not immediately useful given our moderately sized sample.

To make estimation tractable, we first impose the following restriction on treatment effect heterogeneity (cf. Kolesár 2013):

Assumption 2 ((limited selection on gains)). *For any h_0, h_1, h, c, g , the quantity:*

$$\mathbb{E}[Y_i(h_1) - Y_i(h_0)|H_i = h, C_i = c, G_i = g, L_i = \ell]$$

depends only on h_1 , h_0 , and g .

Assumption 2 states that for any pair of employers h_0 and h_1 , the contrast $Y_i(h_1) - Y_i(h_0)$ is not correlated with actual employer choice H_i within a group and lottery. This rules out selection on unobserved heterogeneity in gains *within* group and choice set—what Heckman *et al.* (2006) call *essential heterogeneity*. Assumption 2 also requires that treatment effects are not correlated with lottery/cohort, however this can be relaxed. However, Assumption 2 is strictly weaker than assuming treatment effects are homogenous within each group. It still allows sorting on *levels*, that is that workers choosing $H_i = h$ have a different average value of $Y_i(h)$ than those who do not.

To see this, it is illustrative to write potential outcomes in the two-way-fixed-effects form:

$$Y_i(h) = \alpha_i + \beta_{G_i h} + u_{ih} \quad (1)$$

where $\alpha_i := Y_i(h_0)$ with h_0 a fixed reference employer, $\beta_{gh} := \mathbb{E}[Y_i(h) - Y_i(h_0)|G_i = g]$ and $u_{ih} := \{Y_i(h) - Y_i(h_0)\} - \mathbb{E}[Y_i(h) - Y_i(h_0)|G_i]$. Assumption 2 implies that idiosyncratic gains u_{ih} are (conditionally) mean independent of first-job choice: $\mathbb{E}[u_{ih}|H_i, G_i = g, L_i = \ell] = 0$, but not that H_i is in any way uncorrelated with the “worker-effects” α_i .¹⁸

To operationalize the use of choice sets of instruments, it will be convenient to represent a choice set as a vector of indicators Z_{hi} for the presence of each employer h in C_i , where $Z_{hi} = \mathbb{1}(h \in C_i)$. A realization of C_i is equivalent to a realization of the full vector $\mathbf{Z}_i := (Z_{1i}, Z_{2i} \dots Z_{Ji})'$, for some arbitrary ordering of the employers. Similarly, let \mathbf{D}_i be a vector of $D_{hi} := \mathbb{1}(H_i = h)$ across all employers h . Again, the random vector \mathbf{D}_i encodes exactly the same information as H_i . For any group g and lottery ℓ , let $\Sigma_{gl} = \mathbb{E}[\mathbf{Z}_i \mathbf{D}'_i | G_i = g, L_i = \ell]$.

Assumption 3 ((relevance)). Σ_{gl} has full rank for each $g \in \mathcal{G}$ and $\ell \in \mathcal{L}$.

Assumption 3 imposes the standard IV relevance condition that the J instruments have independent predictive power for the J treatments $h \in \mathcal{H}$, and that this holds within each (g, ℓ) cell.

For any group g , collect the μ_{gh} over all the employers into a vector $\boldsymbol{\mu}_g$. Proposition 2 shows that the assumptions given are sufficient to identify this full set of first job effects for each group:

Proposition 2 ((identification of FJEs)). *Make Assumptions 1, 2 and 3. Then for each $g \in \mathcal{G}$ and any ℓ :*

$$\boldsymbol{\mu}_g = \Sigma_{gl}^{-1} \mathbb{E}[\mathbf{Z}_i Y_i | G_i = g, L_i = \ell]$$

Proof. See Appendix D. □

Note that based on Proposition 2, the vector $\boldsymbol{\mu}_g$ is in fact over-identified: in principle it can be estimated using the data from any single cohort ℓ . Indeed, the implication of Assumption 2 that

¹⁸Note that $\mathbb{E}[u_{ih}|H_i, g, \ell] = 0$ is not sufficient for identification given observations of (Y, H) alone, on account of the unobserved α_i . Unlike “AKM” settings in which the α_i can be differenced out by workers moving between firms (cf. Abowd *et al.* 1999), a worker by definition has only one actual first-job H_i , yielding the single cross-section of observed earnings: $Y_i = \alpha_i + \beta_{G_i H_i} + u_i$ where $u_i := u_{iH_i}$.

treatment effects $Y_i(h) - Y_i(h_0)$ are mean independent of L_i could be relaxed to identify FJEs that vary by cohort. However, in practice, it is desirable to pool across lotteries given our limited sample size. The proof of Proposition 2 shows that μ_g can be estimated from a sample that pools across L_i but conditions on $G_i = g$ by two stage least squares (2SLS) with the inclusion of cohort fixed effects, which pick up $\mathbb{E}[Y_i(h_0)|G_i = g, L_i = \ell]$ across the lotteries ℓ .

4.4 Partial identification with essential heterogeneity

While the results of the last section yield a straightforward route to identification of FJEs based on random choice set variation, the required assumption that workers do not sort into first jobs on the basis of their idiosyncratic FJE's is restrictive. For instance, when the outcome variable is earnings, it is incompatible with a Roy-type selection model in which there are worker \times employer match effects and workers choose in part on the basis of earnings. Or, if the outcome of interest is mobility after residency, we must believe that workers are not more likely to move away from their residency locations if they ended up in a location that they preferred less during the lottery.

To accommodate violations of Assumption 2—*essential heterogeneity*—we develop a partial identification approach based upon the observation that the instruments provide a system of moment conditions that are linear in the FJEs μ_{gh} . This builds upon an existing literature that maps IV identification into a linear programming problem (Mogstad *et al.* 2018; Kamat 2020).

Let $D_{hi}(c)$ be an indicator for whether worker i would choose first-job employer h given choice set c , i.e. $D_{hi}(c) = \mathbb{1}(H_i(c) = h)$, and recall the notation of D_{hi} as an indicator for i actually choosing h , i.e. $D_{hi} = D_{hi}(C_i)$. Note that $D_{hi}(c)$ only depends on i through i 's preference relation \succ_i , and we may instead index the function D by \succ rather than i . For ease of notation let $X_i = (G_i, L_i)$ be a vector composed of demographic group g and lottery ℓ . For any h and z , consider the observable quantity $\mathbb{E}[Y_i D_{hi}|C_i = c, X_i = x]$.

By the law of iterated expectations and Assumption 1:

$$\begin{aligned}\mathbb{E}[Y_i D_{hi}|C_i = c, X_i = x] &= \sum_{\succ} P_{\succ|x} \cdot \mathbb{E}[Y_i D_{hi}|\succ_i=\succ, C_i = c, X_i = x] \\ &= \sum_{\succ} P_{\succ|x} \cdot \mathbb{E}[Y_i(h) D_{h\succ}(c)|\succ_i=\succ, C_i = c, X_i = x] \\ &= \sum_{\succ} D_{h\succ}(c) \cdot \{P_{\succ|x} \cdot \mu_{\succ h|x}\}\end{aligned}\tag{2}$$

where $P_{\succ|x} := P(\succ_i=\succ | X_i = x)$ and $\mu_{\succ h|x} := \mathbb{E}[Y_i(h)|\succ_i=\succ, X_i = x]$ is the average counterfactual outcome that would occur for a worker with preference relation \succ if their first job were at employer h . The above expression reveals that the observable $\mathbb{E}[Y_i D_{hi}|C_i = c, X_i = x]$ identifies a linear combination of the $\mu_{\succ h|x}$ over all preferences \succ under which h is the best choice for the fixed choice set c . Note that a linear combination of the $P_{\succ|x}$ alone with the same weights is also

identified by removing Y_i from Eq. (2):

$$\mathbb{E}[D_{hi}|C_i = c, X_i = x] = \sum_{\succ} D_{h\succ}(c) \cdot P_{\succ|x} \quad (3)$$

As an example, consider an instance of the lottery in which there three choice sets occur: $\{1, 2, 3\}$, $\{1, 2\}$ and $\{1\}$, and we are interested in the FJE of $h = 2$. The coefficients in the system of linear equations (2) or (3) can be summarized by Table 1. In this example $J := |\mathcal{H}| = 3$, and

		\succ_i					
		$1 \succ 2 \succ 3$	$1 \succ 3 \succ 2$	$2 \succ 1 \succ 3$	$2 \succ 3 \succ 1$	$3 \succ 1 \succ 2$	$3 \succ 2 \succ 1$
C_i	$\{1, 2, 3\}$	0	0	1	1	0	0
	$\{1, 2\}$	0	0	1	1	0	1
	$\{1\}$	0	0	0	0	0	0
B_{\succ}^h	{1}	{1, 3}	\emptyset	\emptyset	{1, 3}	{3}	

Table 1: Example response matrix $D_{hi}(c)$ with three employers and one lottery.

the 6 columns of Table 1 represent the $J!$ distinct preference orderings over employers. The first three rows represent the three choice sets observed in the lottery, where first employer 3 becomes unavailable, then employer 2, and for workers with the lowest lottery numbers only employer 1 remains. The entries $\{0, 1\}$ indicate the value of $D_{h\succ}(c)$ for each row and column pair, forming what Heckman & Pinto (2018) call the *response matrix* of the model.

The last row of Table 1 summarizes the set of employers that are preferred to h according to the preference relation \succ for that column, which we denote as $B_{\succ}^h := \{h' \in \mathcal{H} : h' \succ h\}$. Note that any two columns sharing a value of B_{\succ}^h have an identical response $D_{h\succ}(c)$ for all choice sets c . This is because the event that a worker with preferences \succ chooses h only depends on whether h is their most-preferred employer in c , and not what the relative ordering is between the other available employers c/h . In particular, i will choose h if and only if $h \in c$ and their “better-than- h ” set $B_i^h := B_{\succ_i}^h$ does not intersect the choice set c , i.e. $D_{hi}(c) = \mathbb{1}(h \in c \text{ and } B_i^h \cap c = \emptyset)$.

We can thus coarsen the columns of Table 1 to combine all preferences \succ that share a value of B_{\succ}^h , by rewriting Equation (2) as:

$$\mathbb{E}[Y_i D_{hi}|C_i = c, X_i = x] = \sum_{B \subseteq \mathcal{H}/h} \mathbb{1}(h \in c \text{ and } B \cap c = \emptyset) \cdot \{P_{Bh|x} \cdot \mu_{Bh|x}\} \quad (4)$$

where $P_{Bh|x} := P(B_i^h = B | X_i = x) = \sum_{\succ: B_{\succ}^h = B} P_{\succ|x}$ and $\mu_{Bh|x} = \sum_{\succ: B_{\succ}^h = B} P_{\succ|x} \cdot \mu_{\succ|x}$. Equation (3) can be similarly rewritten as a summation over the B^h rather than \succ .

The value of moving from a summation over preferences to a summation over better-than- h sets B^h is that rather than a system of $J!$ unknowns $P_{\succ|x} \cdot \mu_{\succ|h|x}$ for each x we now have a system in 2^{J-1} unknowns $P_{B|x} \cdot \mu_{Bh|x}$ for each x . In the example of Tables 1 and 2 this only reduces the number of columns from $3! = 6$ to $2^2 = 4$; however the gain quickly becomes dramatic for $J > 3$.

Note that there is still redundancy in the columns of Table 2: workers who prefer only employer

		B_i^h			
		\emptyset	{1}	{3}	{1, 3}
C_i	{1, 2, 3}	1	0	0	0
	{1, 2}	1	0	1	0
	{1}	0	0	0	0
$j_i^{h\ell}$		1	0	2	0

Table 2: Response matrix from Table 1 written in terms of better-than- h sets B_i^h ; $h = 2$.

1 to employer 2 have the same response as workers who prefer both employers 1 and 3 to employer 2, for all choice sets $c \in \text{supp}\{C_i\}$. This is because the choice sets $\text{supp}\{C_i\}$ have a nesting property arising from the sequential nature of the RSD: once an employer's position has been filled, it never re-enters the choice sets of doctors choosing later in the lottery. This leads to a close connection with Heckman & Pinto (2018), who generalize the notion of "monotonicity" from Angrist & Imbens (1994) to treatments that take on multiple unordered values.

To appreciate this connection, we introduce some further notation. For a given instance ℓ of the lottery, label the choice sets as $C_{1\ell} \supset C_{2\ell} \supset \dots \supset C_{J_\ell,\ell}$, where J_ℓ is the number of employers in lottery ℓ . Let J_ℓ^h be the last set along this sequence that contains employer h . Selection behavior $D_{hi}(c)$ towards employer h within a single lottery ℓ can now be characterized by just $J_\ell^h + 1$ distinct groups. The reason is that if $D_{hi}(C_{j\ell}) = 1$ then it must be that $D_{hi}(C_{j'\ell}) = 1$ for all $j \leq j' \leq J_\ell^h$ (if h is chosen from a larger set, it must be chosen from any smaller subset that still contains h). Thus we can infer $D_{hi}(c)$ for all $c \in \text{supp}\{C_i | L_i = \ell\}$ from the lowest value of j such that $D_{hi}(C_{j\ell}) = 1$: call this $J_i^{h\ell}$. If alternatively $D_{hi}(C_{J_\ell^h,\ell}) = 0$, i.e. i does not choose h even in the smallest choice set in which it appears, then we can call i an " h -never-taker" in lottery ℓ , and write $j_i^{h\ell} = 0$. The final row in Table 2 lists the value of $j_i^{h\ell} = 0$ corresponding to each B_i^h . Note that while there are $2^2 = 4$ better-than- h sets in this example, there are just $J_\ell^h + 1 = 3$ distinct values of $j_i^{h\ell}$.

The analysis from the preceding paragraph reveals that within each lottery ℓ , selection behavior satisfies what Heckman & Pinto (2018) call *unordered monotonicity*: for each $h \in \mathcal{H}$, there exists an ordering on the points in $c \in \text{supp}\{C_i | L_i = \ell\}$ such that $D_{hi}(c)$ is weakly increasing along that order.¹⁹ Heckman & Pinto (2018) provide point identification results under unordered monotonicity with discrete instruments: in particular their results imply that

$$\mathbb{E}[Y_i(h) | i \text{ is not an } h\text{-never-taker}, X_i = x] \quad (5)$$

is point identified.²⁰ In particular, it is equal to $\frac{E[Y_i D_{hi}|C_i=C_{j_\ell^h,\ell}, X_i=x] - E[Y_i D_{hi}|C_i=C_{1,\ell}, X_i=x]}{E[D_{hi}|C_i=C_{j_\ell^h,\ell}, X_i=x] - E[D_{hi}|C_i=C_{1,\ell}, X_i=x]}$, where ℓ

¹⁹When viewed across all possible choice-sets $c \subseteq \mathcal{H}$, $D_{hi}(c)$ is increasing according to a *partial* order on the c , which depends on h . In particular $D_{hi}(c) \geq D_{hi}(c')$ whenever $c/h \subseteq c'/h$ and $h \in c$ if $h \in c'$. This generalizes the notions of "partial" and "vector" monotonicity analyzed by (Goff, 2020; Mogstad *et al.*, 2020) to the unordered-treatment case.

²⁰Lee & Salanié (2020) also consider identification under unordered monotonicity with discrete instruments. They introduce the concept of particular instrument values *targeting* particular treatments h in such a setting. In their language, we can say that our choice sets $C_{1\ell}$ to $C_{j_\ell^h,\ell}$ strictly target employer h (interpreting the event $h \notin C_i$ as endowing h with

is the lottery indicator appearing in $X_i = x$. Indeed, it is also easily shown that the mean of $Y_i(h)$ can be further disaggregated within each of the j_ℓ^h individual complier groups (that occur with positive probability), by considering adjacent values of j in the above ratio.

However parameter (5) is by itself not sufficient to identify μ_{gh} , as it does not capture $Y_i(h)$ among workers who would never choose h from any choice set present in their lottery. The proportion of such h -never-takers within a given (ℓ, h) pair can be quite large, leading to wide bounds even if we assume $\mathbb{E}[Y_i(h)|i \text{ is an } h\text{-never-taker}, X_i = x]$ belongs to some bounded set $[Y^L, Y^U]$.

By contrast, the approach of Table 2 based on better-than- h sets keeps the $j_i^{th} = 0$ group disaggregated into those workers with different better-than- h sets— $\{1\}$ and $\{1, 3\}$ in the example—in a way that is consistent across different lotteries ℓ (unlike the unordered monotonicity approach, in which the meaning of the j_i^{th} groups depend on lottery). This allows us to combine the identifying information of Equation (4) across multiple lotteries, that have different sequences of choice sets. In the following example, for instance, employer 4 offers jobs only in lottery $\ell = 2$, and the order in which vacancies are filled differs in the two years:

		C_i								B_i^h							
										\emptyset	{1}	{3}	{4}	{1, 3}	{1, 4}	{3, 4}	{1, 3, 4}
$\ell = 1$	$\{1, 2, 3\}$	1	0	0	1	0	0	0	0								
	$\{1, 2\}$	1	0	1	1	0	0	1	0								
	$\{1\}$	0	0	0	0	0	0	0	0								
	$j_i^{h,1}$	1	0	2	1	0	2	1	0								
$\ell = 2$	$\{1, 2, 3, 4\}$	1	0	0	0	0	0	0	0								
	$\{2, 3, 4\}$	1	1	0	0	0	0	0	0								
	$\{2, 3\}$	1	1	0	0	0	1	0	0								
	$\{2\}$	1	1	1	0	1	1	1	1								
	$j_i^{h,2}$	1	2	4	0	4	3	4	4								

Table 3: Response matrix across two lotteries, with $\mathcal{H} = \{1, 2, 3, 4\}$ and $h = 2$.

To make efficient use of our data spanning multiple lotteries, we assume that both first job effects and group sizes are stable across lotteries:

Assumption 4 ((stability)). Neither $\mu_{Bh|x}$ nor $P_{Bh|x}$ depends on the ℓ component of x (for all $h \in \mathcal{H}$, $B \subseteq \mathcal{H}$ and $g \in \mathcal{G}$)

Note that Assumption 4 does not nest Assumption 2, which only requires FJE differences to be stable over lotteries.²¹ Assumption 4 implies that we can write FJEs and group sizes for a given

so low a utility to worker i that they would never choose it).

²¹Recall that we include lottery fixed effects in the strategy from Section 4.3 to absorb the dependence of $\mathbb{E}[Y_i(h_0)|G_i = g, L_i = \ell]$ on ℓ). On the other hand, Assumption 4 allows selection on gains, so neither Assumption 4 nor Assumption 2 nest each other, but emphasize different relaxations of identifying assumptions.

better-than- h set introduced earlier ($\mu_{gh|x}$ and $P_{Bh|x}$) simply as $\mu_{B|g}$ and $P_{B|g}$, depending only on demographic group g and not on lottery ℓ .

With Assumption 4 in mind, we now turn to the formal identification analysis. We also make explicit the assumption that workers' choices are made rationally:

Assumption 5 ((rationality)). *Each worker i has a preference relation \succ_i over \mathcal{H} without indifferences, such that $H_i(c) = \{h \in \mathcal{H} : h \succ_i h', \forall h' \in \mathcal{H}, h' \neq h\}$*

Note that Assumption 5 ensures that the better-than- h sets B_i^h are always well-defined. We return to a discussion of Assumption 5 below.

Recall that our parameters of interest are the $\mu_{gh} = \mathbb{E}[Y_i(h)|G_i = g]$ for each h and g , and note that under Assumption 4 this can be written as

$$\mu_{gh} = \sum_{B \subseteq \mathcal{H}/h} Q_{Bh|g}$$

where we define $Q_{Bh|g} := P_{Bh|g} \cdot \mu_{Bh|g}$. Note that $Q_{Bh|g}$ is precisely the quantity appearing in the summand of each of the linear restrictions (4), with the known coefficients $\mathbb{1}(h \in c \text{ and } B \cap c = \emptyset)$. Thus our identification problem involves a linear optimization problem over the $Q_{Bh|g}$, involving a set of linear constraints on them. Proposition 3 below makes this precise.

Before stating the result, we introduce one final assumption: that the $\mu_{Bh|g}$ are uniformly bounded by known constants Y^L and Y^U :

Assumption 6 ((boundedness)). $Y^L \leq \mu_{Bh|g} \leq Y^U$ (for all $h \in \mathcal{H}, g \in \mathcal{G}, B \subseteq \mathcal{H}/h$)

Assumption 6 is useful because the data will give no direct information about $\mu_{Bh|g}$ for a given B if $P(H_i = h|B_i^h = B, G_i = g) = 0$.²² Thus, depending on the support of the choice sets C_i , even partial identification of μ_{gh} may require such auxiliary assumptions, with boundedness being the simplest example. Note that a simple sufficient condition for Assumption 6 is that $Y^L \leq Y_i(h) \leq Y^U$ for all doctors i and employers h , that is the bounds hold individually rather than on average.

Since $Q_{Bh|g} = P_{Bh|g} \cdot \mu_{Bh|g}$, Assumption 6 implies that:

$$Y^L P_{Bh|g} - Q_{Bh|g} \leq 0 \text{ and } Y^U P_{Bh|g} - Q_{Bh|g} \geq 0 \text{ for each } B \subseteq \mathcal{H}/h \quad (6)$$

Fix a g and h . Letting \mathbf{Q} be a vector of the $Q_{Bh|g}$ across all $B \subseteq \mathcal{H}/h$, and similarly \mathbf{P} for the $P_{Bh|g}$, we denote by \mathcal{M} the set of all (\mathbf{Q}, \mathbf{P}) pairs such that (6) holds and that:

$$\sum_{B \subseteq \mathcal{H}/h} P_{Bh|g} = 1 \text{ and } P_{Bh|g} \geq 0 \text{ for each } B \subseteq \mathcal{H}/h \quad (7)$$

We may now characterize the identified set of μ_{gh} as an optimization problem over \mathcal{M} :

²²In particular, this is most likely to happen for $B = \mathcal{H}/h$, for which a doctor will only choose h if their choice set consists only of h . $C_i = \{h\}$ does in fact occur in the example of Table 3, and thus every column of the table has at least one entry of 1.

Proposition 3. Under Assumptions 1, 4, 5 and 6, $\mu_{gh} \in [\theta_{gh}^L, \theta_{gh}^U]$, where

$$\theta_{gh}^L := \min_{(\mathbf{Q}, \mathbf{P}) \in \mathcal{M}} \sum_{B \subseteq \mathcal{H}/h} Q_{Bh|g} \quad \text{and} \quad \theta_{gh}^U := \max_{(\mathbf{Q}, \mathbf{P}) \in \mathcal{M}} \sum_{B \subseteq \mathcal{H}/h} Q_{Bh|g}$$

subject to the following restrictions:

$$\sum_{B \subseteq \mathcal{H}/h} \mathbb{1}(h \in c \text{ and } B \cap c = \emptyset) \cdot Q_{Bh|g} = \mathbb{E}[Y_i D_{hi} | C_i = c, G_i = g, L_i = \ell] \quad (8)$$

$$\sum_{B \subseteq \mathcal{H}/h} \mathbb{1}(h \in c \text{ and } B \cap c = \emptyset) \cdot P_{Bh|g} = \mathbb{E}[D_{hi} | C_i = c, G_i = g, L_i = \ell] \quad (9)$$

for each $\ell \in \mathcal{L}$ and $c \in \text{supp}\{C_i | L_i = \ell\}$.

Comparison with Mogstad et al. (2018) and Kamat (2020):

Proposition 3 is closely related to recent results in Mogstad et al. (2018) and Kamat (2020), who also express IV estimands of solutions to a linear programming problem. Mogstad et al. (2018) develops an approach to identifying treatment effect parameters that depend on marginal counterfactual means of the form $\mathbb{E}[Y_i(d) | U_i = u]$, in the classic LATE setting in which a binary treatment is driven by a separable threshold crossing model: $D_i(z) = \mathbb{1}(p(z) \geq U_i)$. In their case, the latent groups correspond to the values of U_i , which is uniformly distributed on the unit interval. This avoids the need for a distinction between the vectors \mathbf{Q} and \mathbf{P} , which must be both optimized over in our setting, since the analog of $P_{Bh|g}$ for the is simply a uniform measure on $[0, 1]$.

Similar to us, Kamat (2020) also optimizes over latent group probabilities jointly with outcomes. However, they make use of a discrete outcome variable to optimize directly over the full joint distribution of potential outcomes and potential treatments. This keeps the constraints and objective functions linear in parameters without a need to introduce the final set of inequality constraints (6). They also consider a richer class of parameters of interest, taking the form $\sum_{B \subseteq \mathcal{H}/h} w_B \cdot Q_{Bh|g} / \sum_{B \subseteq \mathcal{H}/h} w_B \cdot P_{Bh|g}$. This introduces a non-linear objective function, which they handle by introducing an additional variable and re-paramaterizing the problem.

Remark: Note that point identification obtains in Proposition 3 if $\sum_{\mathcal{H}/h} Q_{Bh|g} = (1, 1, \dots, 1)'\mathbf{Q}$ can take on just a single value subject to restrictions (8), (9) and $(\mathbf{Q}, \mathbf{P}) \in \mathcal{M}$. This holds for example whenever the vector $(1, 1, \dots, 1)'$ lies in the column space of the response matrix depicted in Tables 2 and 3, describing the coefficients appearing in expansion (8). This yields an alternative way to understand the point identification result of Proposition 1; whenever $P(C_i = \{h\} | G_i = g) > 0$, the rows corresponding to this choice set in the response matrix are composed of all ones.

Estimation and Inference:

Given our finite sample of data $\{(Y_i, H_i, C_i, G_i, L_i)\}_{i=1\dots n}$, the endpoints of the identified set

$\Theta_{gh} = [\theta_{gh}^L, \theta_{gh}^U]$ could be consistently estimated by solving the linear program of Proposition 3 upon replacing the expectations appearing in (8) and (9) with their finite sample analogs \mathbb{E}_n , e.g. $\mathbb{E}_n[Y_i] = \frac{1}{n} \sum_{i=1}^n Y_i$. Note however that the resulting interval estimate $\hat{\Theta}_{gh} = [\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U]$ may be empty even if the model is correctly specified. For example, in the data we sometimes observe cases where $\mathbb{E}_n[D_{hi}|C_i = c, G_i = g, L_i = \ell] > \mathbb{E}_n[D_{hi}|C_i = c', G_i = g, L_i = \ell]$ where c' is a strict subset of c (and both of which contain h). On its face, this appears to be evidence against the joint hypothesis of choice-set independence (Assumption 1) and utility maximization on the part of workers. However, it is also not at all unlikely when the events $(C_i = c \cap L_i = \ell)$ and $(C_i = c' \cap L_i = \ell)$ have only a few observations each, as is often the case in our data.

For the above reason, we solve a relaxation of the linear programs in Proposition 3 to obtain point estimates for θ_{gh}^L and θ_{gh}^U , which also forms the basis for our approach to constructing confidence intervals for the underlying parameter μ_{gh} . Firstly, we pool data across lotteries ℓ , replacing moments like $\mathbb{E}_n[D_{hi}|C_i = c, G_i = g, L_i = \ell]$ by $\mathbb{E}_n[D_{hi}|C_i = c, G_i = g]$, and similarly for the moments of $Y_i D_{hi}$. This is justified under Assumption 4, which implies by (4) and the law of iterated expectations that:²³

$$\mathbb{E}[Y_i D_{hi}|C_i = c, G_i = g] = \sum_{B \subseteq \mathcal{H}/h} \mathbb{1}(h \in c \text{ and } B \cap c = \emptyset) \cdot \{P_{Bh|g} \cdot \mu_{Bh|g}\} \quad (10)$$

An analogous expression holds for Eq. (3) with the lottery conditioning removed.

Secondly, we do not require the moment conditions to be satisfied exactly in sample. Define the quantities:

$$s_{hcg}^Y = \left(\sum_{B \subseteq \mathcal{H}/h} \mathbb{1}(h \in c \text{ and } B \cap c = \emptyset) \cdot Q_{Bh|g} - \mathbb{E}_n[Y_i D_{hi}|C_i = c, G_i = g] \right) \quad (11)$$

$$s_{hcg}^D = \left(\sum_{B \subseteq \mathcal{H}/h} \mathbb{1}(h \in c \text{ and } B \cap c = \emptyset) \cdot P_{Bh|g} - \mathbb{E}_n[D_{hi}|C_i = c, G_i = g] \right) \quad (12)$$

which measure the deviation of a given (\mathbf{Q}, \mathbf{P}) pair from the identifying restrictions (8) and (9). Let \mathbf{s} be a vector of all of the s_{hcg}^V over $V \in \{Y, D\}$, c and g . Similarly to Mogstad *et al.* (2018), we consider the smallest deviation \mathbf{s} attainable, in an L^1 norm-sense. In particular, let

$$T_n := \min_{\substack{(\mathbf{Q}, \mathbf{P}) \in \mathcal{M} \\ \mathbf{s}}} \left(\sum_{\ell \in \mathcal{L}} \sum_{B \subseteq \mathcal{H}/h} \sum_{V \in \{Y, D\}} a_{hcg}^V \cdot |s_{hcg}^V| \right)$$

subject to (11) and (12) for each $c \in \text{supp}\{C_i|L_i = \ell\}$ and $\ell \in \mathcal{L}$, and where we introduce a set of positive scaling coefficients a . We set the scaling coefficients as $a_{hcg}^D = \sqrt{n_{cg}/Var_n(D_{hi}|G_i = g)}$ and $a_{hcg}^Y = \sqrt{n_{cg}/Var_n(Y_i D_{hi}|G_i = g)}$, where n_{cg} is the number of observations for which $C_i = c$

²³Note that (10) is the same expression we would arrive at after assuming choice-sets exogeneity without conditioning on lottery. However Assumption 4 coupled with Assumption 1 as stated is still slightly weaker, as it e.g. only requires that the response-group specific conditional means of $Y_i(h)$ —rather than their full distributions—do not depend on L_i .

and $G_i = g$.²⁴

Given T_n , now estimate θ_{gh}^L as

$$\hat{\theta}_{gh}^L := \min_{\substack{(\mathbf{Q}, \mathbf{P}) \in \mathcal{M} \\ \mathbf{s}}} \sum_{B \subseteq \mathcal{H}/h} Q_{Bh|g} \quad \text{s.t.} \quad \left(\sum_{\ell \in \mathcal{L}} \sum_{B \subseteq \mathcal{H}/h} \sum_{V \in \{Y, D\}} a_{hcg}^V \cdot |s_{hcg}^V| \right) \leq T_n + \kappa_n \quad (13)$$

Equation (13) finds the smallest value of μ_{gh} among the (\mathbf{Q}, \mathbf{P}) that are “closest” to satisfying the identifying restrictions (8) and (9) in finite sample. The tuning parameter κ_n broadens the notion of “closest” such (\mathbf{Q}, \mathbf{P}) , and must converge to zero with n for consistency. We report estimates with $\kappa_n = 0$ and $\kappa_n = T_n/10$. The estimate $\hat{\theta}_{gh}^U$ is defined analogously to Eq. (13) but with the min operator replaced by a max.

Although the absolute value function is not linear, the objective function defining T_n can be reformulated by adding an additional variable for each component of \mathbf{s} and reparameterizing the problem slightly. In particular, one can replace each instance of s by $p - n$, and add constraints $a \geq 0, b \geq 0$ to the problem. The simplex algorithm (standard for solving linear programs) will then ensure that these correspond to positive and negative parts of s : $p = \max(s, 0)$ and $n = \max(-s, 0)$, with respect to which the absolute value of s is the linear function $|s| = p + n$. We use this strategy to compute $\hat{\theta}_{gh}^L$ using the mixed integer linear programming package `lpSolveAPI` in R, for each g, h first computing T_n and then evaluating Eq. (13).

Let us now turn to building confidence intervals for the parameter μ_{gh} . To test the null hypothesis that $H_0 : \theta = \theta_0$ against the alternative $H_1 : \theta \neq \theta_0$ for a generic value $\theta_0 \in \mathbb{R}$ we use a test statistic that augments T_n to enforce the null hypothesis, i.e.

$$T_n(\theta_0) := \min_{\substack{(\mathbf{Q}, \mathbf{P}) \in \mathcal{M} \\ \mathbf{s}}} \left(\sum_{\ell \in \mathcal{L}} \sum_{B \subseteq \mathcal{H}/h} \sum_{V \in \{Y, D\}} a_{hcg}^V \cdot |s_{hcg}^V| \right) \quad \text{s.t.} \quad \sum_{B \subseteq \mathcal{H}/h} Q_{Bh|g} = \theta_0 \quad (14)$$

and restrictions (11) and (12) for each $c \in \text{supp}\{C_i | L_i = \ell\}$ and $\ell \in \mathcal{L}$. The statistic $T_n(\theta_0)$ can be interpreted as measuring the minimum weighted deviation from constraints (8) and (9) that is necessary for a (\mathbf{Q}, \mathbf{P}) pair to deliver that value of θ_0 .

We construct confidence intervals by comparing $T_n(\theta_0)$ to a critical value and collecting those values for which we fail to reject, i.e. $C_n = \{\theta \in \mathbb{R} : T_n(\theta_0) \leq \hat{c}\}$, where \hat{c} is a critical value estimated from the data. In particular, construct a collection of $\{T_{bn}^*(\theta_0)\}_{b=1 \dots B}$ by non-parametric bootstrap, and compute \hat{c} as the $1 - \alpha$ quantile of the $T_{bn}^*(\theta_0)$ for $\alpha = 0.05$. Each $T_{bn}^*(\theta_0)$ replaces the moments in (11) and (12) with bootstrap analogues \mathbb{E}_b^* and re-centers with respect to the “full-

²⁴The goal of this choice is to normalize the sampling variance of each s_{hcg}^V to unity. However, we cannot divide by the conditional sample variance specific to each choice set, because D_{hi} has no variation within some choice sets in which h appears (since no employees in fact choose h). Thus, $\text{Var}_n(D_{hi}|C_i = c, G_i = g)$ and $\text{Var}_n(Y_i D_{hi}|C_i = c, G_i = g)$ would be zero.

sample” estimates, for example:

$$s_{bhcg\ell}^{*D} = \sum_{B \subseteq \mathcal{H}/h} \mathbb{1}(h \in c \text{ and } B \cap c = \emptyset) \cdot (P_{Bh|g} - P_{BH|g}^0) \\ - (\mathbb{E}_l^*[D_{hi}|C_i = c, G_i = g, L_i = \ell] - \mathbb{E}_n[D_{hi}|C_i = c, G_i = g, L_i = \ell])$$

where $P_{BH|g}^0$ denotes the optimizer from (14), and $s_{bhcg\ell}^{*Y}$ is defined analogously. We repeat this entire exercise over a grid of θ_0 between the values Y^L and Y^U from Assumption 6, and report the maximum and minimum value along that grid for which we fail to reject.

This approach to inference is close to that of Kamat (2020), who uses a quadratic objective function of the form $\mathbf{s}'\Omega^{-1}\mathbf{s} = \|\Omega^{-1/2}\mathbf{s}\|_2$ rather than $\|\Omega^{-1/2}\mathbf{s}\|_1$, where in our case $\Omega^{-1/2}$ is a diagonal matrix defined by the a coefficients. We use the L^1 norm in order to keep all quantities computable by linear-programming algorithms, which are faster to solve than quadratic programs (computational limitations loom large for us, as we discuss in the next section). Kamat (2020) also uses subsampling rather than bootstrap to compute the critical values \hat{c} , on the basis of results from Kalouptsidi *et al.* (2020). We choose bootstrap to avoid the need to choose the subset size, which represents an additional tuning parameter. Our setting is also related to methods that treat inference under partial identification in moment equality/inequality models (Andrews & Soares, 2010; Chernozhukov *et al.*, 2007, 2013), as well as specification tests for random utility models (Kitamura & Stoye, 2018; Smeulders *et al.*, 2021).

Implementation:

In implementing the above methods in our data, we face very real constraints on computational tractability (as well as statistical power). Workers choose among 55 employers in our final sample, with a typical lottery including most of these employers. With $|\mathcal{H}| = 55$, the vectors \mathbf{Q} and \mathbf{P} would each contain about 2×10^{-16} entries, which is clearly infeasible from a computational standpoint.²⁵

For this reason, we group the employers (hospitals) into a manageable *categories*, and ignore the distinction between hospitals within a category. We define the categories on the basis of employers’ overall desirability, as evidenced by the average lottery number R_i among workers who choose it in the RSD, across the study period (recall that R_i is normalized to the unit interval within each lottery). We find that $|\mathcal{H}| = 10$ is about the largest linear program the software will support, yielding roughly 8000 parameters to be optimized over. However for ease of interpretation, we for now use just four categories. Categories 1-3 are defined by terciles of the distribution of \bar{r}_h across hospitals, where $\bar{r}_h := \mathbb{E}_n[R_i|H_i = h]$. To have a well-defined reference group, we for Category 4 use the hospitals in the remote counties of Finnmark and Sogn og Fjordane, the most remote regions of northern and western Norway.²⁶ Table 4 reports some observable characteristics of the

²⁵Assuming one byte for each entry, simply storing the constraint matrix for the linear program requires about 2 gigabytes for $|\mathcal{H}| = 20$, 2 terabytes for $|\mathcal{H}| = 30$, 2 petabytes for $|\mathcal{H}| = 40$, and so on.

²⁶First jobs at employers in these regions are very unappealing to most graduates. Because of this, the government in some years introduced special incentives for residents in Finnmark and Sogn og Fjordane. Despite this, the large majority of candidates that chose employers in these regions drew very high (poor) lottery numbers and therefore had

categories: for example the lower category numbers tend to be larger and more urban.

Grouping the employers together in this way embodies a substantive assumption: to make use of the methods of this section we must be willing to assume that workers are indifferent between hospitals within a single category.²⁷ A failure of this assumption could explain deviations from the model of the type we previously attributed to sampling variation: e.g. workers being more likely to choose a Category 2 hospital when $C_i = \{1, 2, 3\}$ than when $C_i = \{2, 3\}$.²⁸ One could explore alternative data-driven ways to define the categories: for example to minimize a statistic like T_n over such choices. An alternative approach may be retain the full set of $|\mathcal{H}| = 55$ employers but find some efficient way select among the 2^{54} columns of the response matrix depicted in Table 2.²⁹ Indeed, only 1,802 distinct choice sets C_i ever occur across the study period, a tiny fraction of those that are conceptually possible. As a result, the column rank of the response matrix can at most be 1,802, which is within our maximum manageable number of columns from a computational standpoint.

5 Results

This section presents estimates of first job effects obtained by the methods presented in Section 4. Throughout, we let the “employers” h correspond to the four Categories of hospitals defined above. The first four columns of Table 4 report some characteristics of these groups. For simplicity, we also focus throughout this section on just two demographic groups g of workers: male and female.

5.1 Results of the point-identification approach with limited selection on gains

Recall that Proposition 2 shows that FJEs are point identified under an assumption of no selection on unobserved gains. The final column of Table 4 presents estimates based on this result, where the outcome Y_i is chosen to represent earned income four years after a doctor’s residency (five years after graduation). The gap of four years chosen to maximize the number of workers that can be included, but later drafts will consider FJEs across various time horizons. To increase power, we first report FJEs that furthermore do not condition on gender, i.e. unconditional counterfactual means $\mathbb{E}[Y_i(h)]$ rather than μ_{gh} .

Differences in $\mathbb{E}[Y_i(h)]$ across residency hospitals h reveal that the location of one’s residency affects their earnings *after* they’ve moved on to their post-residency position. We estimate that, relative to Category 4 employers, a first job in the most-in-demand employer category raises annual

unappealing choice sets. Hospitals in these regions are excluded from the definition of Categories 1-3.

²⁷This distinguishes our categories from the related notion of a *filtered* treatment introduced by Lee & Salanié (2020). In the latter case, agents select according to their preferences over a fine set of treatment states, and the researcher observes only coarsened categories of that choice.

²⁸Since the order at which hospitals are removed from the available set via the RSD may change year to year, this could happen if doctors tend to prefer the best Category 3 hospitals to the worst Category 2 hospitals, and the best Category 2 hospitals tend to be gone from the RSD before the last hospitals in Category 1 are.

²⁹A related idea is pursued in Smeulders *et al.* (2021).

earnings five years post-graduation by about \$28,000, in 2020. The corresponding estimates for categories 3 and 2 are \$38,000 and (an insignificant) \$16,000.

Category	Avg. Draw	# Hospitals	Avg. Emp.	Proportion Urban	Earnings FJE diff.
1	0.17	17	1634	0.82	28.21* (14.64)
2	0.44	17	1457	0.65	16.03 (13.00)
3	0.73	16	453	0.50	37.70** (18.99)
4	0.89	5	502	0.20	-

Table 4: FJEs measure the impact of a first job in each category on earnings 5 years post-graduation, in thousands of 2020 USD. First stage F-statistics for Category 1, 2, and 3 are 322.65, 327.44, and 102.40. $N = 9,049$. Robust standard errors in parentheses. *** $p < 0.01$ ** $p < 0.05$ * $p < 0.10$.

However, statistical significance of the differences $\mu_{gh} - \mu_{g4}$ disappears when we estimate FJEs $\mu_{gh} = \mathbb{E}[Y_i(h)|G_i = g]$ separately by gender. However the point estimates reveal an annual earnings gap five years out of at least \$20,000 between men and women—about 13%—across first employers. Appendix Figure C.9 plots the μ_{gh} against average realized earnings $\mathbb{E}[Y_i|H_i = h, G_i = g]$, as a rough indication of the extent of the endogeneity that the IV approach is correcting for. See also Table 8 of Section 6 for the point estimates.

5.2 Results of the general partial identification approach

We now turn to the partial identification approach from Section 4.4 that relaxes the assumption of no selection on gains within gender g . Continuing with the earnings outcome variable, Table 5 reports estimates of the identified set $[\theta_{gh}^L, \theta_{gh}^U]$ and confidence intervals for μ_{gh} . We set Y^L at 10 thousand dollars in 2020 USD, and Y^U at \$300,000 (about 2% of the sample is outside of this range in either direction). For all of our outcome variables, the 95% confidence interval C_n takes a grid of 20 values of θ_0 across the range $[Y^L, Y^U]$, and uses 200 bootstrap replications.

As Table 5 shows, the $\kappa_n = 0$ point estimates of θ_{gh}^L and θ_{gh}^U in all cases suggest point identification, i.e. $\hat{\theta}_{gh}^L = \hat{\theta}_{gh}^U$. However, this point identification is “spurious”, as there can be a unique (\mathbf{Q}, \mathbf{P}) that minimizes the sample statistic T_n even when there is no unique (\mathbf{Q}, \mathbf{P}) setting it to zero in the population. We thus treat the slightly “nudged” estimates with $\kappa_n = T_n/10$ as preferred. Recall that these seek the largest and smallest values of μ_{gh} compatible with a T_n within 10% of its minimum value. Overall, the results suggest that Category 1 hospitals cause doctors’ earnings to be highest, especially for women. Note that confidence intervals are missing for women in Category 4 and for men in Category 2. In these cases, the null hypothesis not accepted for any θ_0 between Y^L and Y^U . This suggests that the model is rejected in these cases, and warrants further investigation. In the other cases, the confidence intervals are also quite wide, so we focus subsequent attention on the point estimates with $\kappa_n = T_n/10$ in this section.

In Table 6, the outcome variable Y_i is whether a doctor ever specializes during their career

g	h (category)	[$\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U$]				C_n	
		$\kappa_n = 0$		$\kappa_n = 10\%$			
Women	1	152.49	152.49	129.71	157.91	61.58	190.53
	2	90.93	90.93	89.06	96.31	61.58	138.95
	3	133.31	133.31	130.71	134.12	35.79	164.74
	4	146.47	146.47	133.35	158.06	-	-
Men	1	187.64	187.64	151.65	192.78	61.58	216.32
	2	201.10	201.10	85.10	216.76	-	-
	3	148.61	148.61	132.98	151.38	61.58	216.32
	4	116.71	116.71	110.79	119.26	35.79	190.53

Table 5: First job effects $\mu_{gh} = \mathbb{E}[Y_i(h)|G_i = g]$ for earned income four years post residency, in thousands of 2020 USD. Average earnings across the sample are about \$150,000. Table reports estimates of the identified set $[\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U]$ and 95% confidence intervals for μ_{gh} . (“-” indicates that $C_n = \emptyset$).

(Appendix Table C.3 reports results for the *number* of specializations). This outcome variable is bounded by definition, where $Y^L = 0$ and $Y^U = 1$. The results suggest that working at a Category 4 hospital causes the greatest rates of specialization.

g	h (category)	[$\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U$]				C_n	
		$\kappa_n = 0$		$\kappa_n = 10\%$			
Women	1	0.50	0.50	0.49	0.64	0.21	1.00
	2	0.50	0.50	0.50	0.56	0.26	1.00
	3	0.63	0.63	0.62	0.64	0.37	1.00
	4	0.71	0.73	0.71	0.86	0.11	1.00
Men	1	0.58	0.58	0.50	0.68	0.21	1.00
	2	0.50	0.50	0.49	0.66	0.21	1.00
	3	0.83	0.83	0.83	0.87	0.26	1.00
	4	1.00	1.00	1.00	1.00	0.16	1.00

Table 6: First job effects $\mu_{gh} = \mathbb{E}[Y_i(h)|G_i = g]$ for whether doctor ever specializes during their career. Overall, about 55% of doctors specialize. Table reports estimates of the identified set $[\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U]$ and 95% confidence intervals for μ_{gh} .

Table 7 takes the outcome of interest to be whether the doctor ever moves municipalities, starting with their residency year. Looking at the preferred estimates ($\kappa_n = T_n/10$ column), μ_{gh} is increasing in Category number for both men and women. This is as expected, since lower Category numbers correspond to hospitals that tend to be more sought-after, and thus doctors are more likely to want to stay there longer term.

g	h (category)	$[\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U]$				\mathcal{C}_n (95% CI)	
		$\kappa_n = 0$		$\kappa_n = 10\%$			
		$\hat{\theta}_{gh}^L$	$\hat{\theta}_{gh}^U$	$\hat{\theta}_{gh}^L$	$\hat{\theta}_{gh}^U$		
Women	1	0.07	0.07	0.07	0.13	0.05	0.74
	2	0.17	0.17	0.17	0.20	0.05	1.00
	3	0.14	0.14	0.14	0.16	0.05	1.00
	4	0.16	0.16	0.16	0.25	0.05	1.00
Men	1	0.33	0.33	0.33	0.38	0.05	1.00
	2	0.33	0.33	0.33	0.44	0.05	1.00
	3	0.46	0.51	0.36	0.58	0.05	1.00
	4	0.50	0.50	0.50	0.55	0.05	1.00

Table 7: First job effects $\mu_{gh} = \mathbb{E}[Y_i(h)|G_i = g]$ for whether doctor ever changes municipalities after residency. About 13% of doctors in fact move, across the sample. Table reports estimates of the identified set $[\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U]$ and 95% confidence intervals for μ_{gh} .

6 Using the FJEs to assess the consequences of decentralization

The last section has presented estimates of first job effects on earnings—among other outcomes variables—based on data from the era in which residencies were allocated by the random serial dictatorship mechanism. We now combine these estimates with data from after the replacement of the RSD with a decentralized labor market in 2013, to understand this reform from the perspective of workers.

Recall that in the RSD era, choice sets were allocated to workers independently of their preferences and potential outcomes. In the market era by contrast, the opportunities available to a worker are likely to be highly correlated with her unobserved ability and preferences. Thus, the reform may have affected average outcomes within each demographic group, by changing the distribution of choice sets the workers in that group face and hence their actual employer matches. Given that we observe the distribution of demographic group \times employer matches in both periods, we can calculate the implied welfare changes across groups given a suitable measure of welfare.

We do this by assuming a particular aggregate relationship in the RSD period between a worker’s indirect utility at her chosen employer and her lottery number. Given our FJE estimates from the last section, we further decompose this utility into an earnings component and an “amenity” component, allowing us to track changes in both across the reform. To keep the analysis simple, we use FJE estimates based on the point-identification approach from Section 4.3, again focusing on the four employer categories described in Section 4.4.

6.1 Estimating first-job amenity values

The first step of our approach is to define an average “amenity” value for each employer category h . To this end, we take preferences of workers defined over the employer categories h to have the

form:

$$U_i(h) = \mu_{G_i h} + A_{G_i h} + \eta_{hi} \quad (15)$$

where μ_{gh} is the first job effect of category h for group g , A_{gh} captures the average “amenity” value of employer category h , and $\mathbb{E}[\eta_{hi}|G_i = g] = 0$ for each h and g .³⁰ The term $\mu_{gh} + A_{gh}$ represents a systematic component of utility for employer h among members of group g , while η_{hi} captures variation in utility arising from individual heterogeneity in preferences. This specification allows “typical” preferences to differ flexibly between genders through the $A_{G_i h}$, and higher moments of η_{ih} beyond the mean may also depend upon G_i (e.g. if men or women have greater variability in preferences).

The form of Equation (15) embodies three substantive assumptions. The first is that preferences can be defined at the level of employer categories rather than individual employers, which we require for reasons of statistical power in estimation of μ_g . The second is quasi-linearity in these first-job-effects, which allows us to separate amenities additively from FJEs. Finally, we take workers to anticipate the mean earnings within their group at a given employer, rather than knowing what their exact outcome will be, so that $\mu_{G_i h} = \mathbb{E}[Y_i(h)|G_i]$ appears in utility rather than $Y_i(h)$ itself. This is consistent with Assumption 2, while η_{ih} can still be correlated with $Y_i(h)$ (thus creating endogeneity in FJEs). Both A_{gh} and the distribution of η are taken to be static over the years in which the RSD system was in place.

While the quasi-linearity assumption pins down a unique scale for utility (such that it is measured in dollars), we are also free to fix a location normalization for each i . For an arbitrary fixed employer category h_0 , we may define $U_i(h_0) = 0$ for all i . This yields the following interpretation for amenities at any other employer: A_{gh} is the average amount in excess of their expected earnings μ_{gh} at h that workers in group g would be willing to pay to move from h_0 to h . If group g tends to prefer h to h_0 and would be willing to give up *part* of their earnings to stay at h , then $A_{gh} \in [-\mu_{gh}, 0]$. In practice, we choose h_0 to represent Category 4, the hospitals in Finnmark and Sogn og Fjordane.

Let $v_{gh} := \mu_{gh} + A_{gh}$ denote the total systematic component of utility. Define $r_{gh} := \mathbb{E}[R_i|H_i = h, G_i = g]$, where recall that R_i is worker i 's random lottery number draw, normalized to the unit interval within each lottery. We make the following assumption:

Assumption 7. $r_{gh} = \alpha_g - \beta \cdot v_{gh}$ for some $\beta > 0$ and α_g .

Assumption 7 formalizes, in a specific way, the intuition that the average lottery number among workers choosing employer h is a proxy for their aggregate preference v_{gh} for that employer. If two employers share a value of r_{gh} (for some g), but differ in their FJEs μ_{gh} , then the difference in amenities at the two employers must offset this difference. This intuition supports assuming that $r_{gh} = \phi_g(v_{gh}, \cdot)$ for some decreasing, possibly non-linear function ϕ_g that itself depends on all of

³⁰Equation 15 can be obtained from a general additive-in-FJEs form: $U_i(h) = \mu_{G_i h} + \epsilon_{ih}$ with some generic ϵ_{ih} if we define $A_{gh} = \mathbb{E}[\epsilon_{hi}|G_i = g]$ and $\eta_{hi} := \epsilon_{ih} - A_{G_i h}$. Note that Equation (15) also nests the canonical conditional logit model (McFadden, 1974), when $\eta_{ih} = \lambda \cdot (u_{ih} - \mathbb{E}[u_{ih}])$ with u_{ih} distributed across h as independent extreme value random variables for all i , and λ a scale parameter.

the other v_{hg} , the distribution of (η_{ih}, G_i) , and the number of slots available for each h in each run of the lottery. But even parametric assumptions on the η_{hi} (such as the logit model) do not appear to readily imply reduced-form expressions for ϕ_g . Assumption 7 reflects the simplest functional form assumption that can reasonably fit the data.³¹

Our goal is to use Assumption 7 along with the estimated FJE's and observable r_{gh} to pin down the α_g and β , and hence the amenities A_{gh} . Given the four employer categories, two demographic groups, and utility normalization that implies $A_{gh_0} = -\mu_{gh_0}$ for each g , Assumption 7 involves 9 unknowns (six A_{gh} , two α_g , and β), from 8 equations. Thus, one more restriction is needed for identification. Figure 1 plots the estimated μ_{gh} against the r_{gh} , which we use to motivate an eighth restriction.

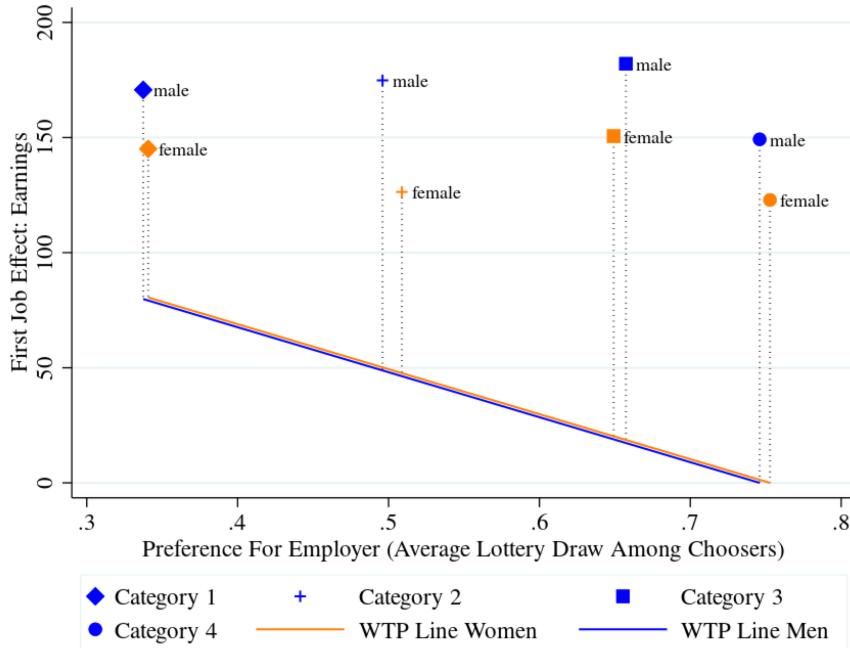


Figure 1: Employer amenity values, based on earnings first job effects (thousands 2020 USD) and the average lottery number at which each employer is chosen. The line for each demographic group g yields the systematic component of utility $\mu_{gh} + A_{gh}$, indicating worker's average willingness to pay to move from Category 4, as a function of average lottery number. Amenities A_{gh} are weakly negative for all categories (with magnitudes depicted by vertical dotted lines), reflecting that workers would give up only part of their income to stay at their employer rather than Catgeory 4.

Observe that Assumption 7 implies that for any employer category $h \neq h_0$, we can net out the α_g and β parameters to write:

$$\frac{\Delta r_{fh}}{\Delta r_{mh}} = \frac{\Delta A_{fh} + \Delta \mu_{fh}}{\Delta A_{mh} + \Delta \mu_{mh}} \quad (16)$$

³¹In particular, the group-specific intercept α_g allows us to reconcile the data with reasonable values of β ; although women have significantly lower FJEs for all h , they tend to have similar r_{gh} to men. This is natural: in the limit of a constant earnings gap $\mu_{mh} = \mu_{fh} + \delta$ and no differences in amenities across genders, we would expect that $r_{fh} = r_{mh}$, which requires $\alpha_m = \alpha_f + \beta\delta$. We note however that linearity in Assumption 7 can only hold as an approximation for some range of v_{gh} , since R_i only has support on the unit interval. In practice, our r_{gh} range between 0.34 and 0.75.

where for any quantity X , ΔX_{hg} denotes the difference between employer categories h and h_0 : $\Delta X_{gh} := X_{gh} - X_{gh_0}$. Comparing categories 1 and 4 in Figure 1, we observe that both Δr_{g1} and $\Delta \mu_{g1}$ are nearly identical across genders $g \in \{f, m\}$. By Equation (16), this suggests that $\Delta A_{m1} \approx \Delta A_{f1}$, regardless of the values of β and α_g .³² We thus assume that $\Delta A_{m1} = \Delta A_{f1}$ exactly as a reasonable ninth equation, allowing us to point identify all parameters. Intuitively, this restriction says that men and women exhibit the same willingness to pay—in excess of the earnings difference—for the mostly large, urban employers in Category 1, compared with the smaller, rural employers in Category 4. Given α_g and β we can extract each amenity value A_{gh} , as described in the caption of Figure 1.

6.2 The effects of decentralization

Estimates of amenities A_{gh} now allow us to construct the systematic component of utility at each employer $\mu_{gh} + A_{gh}$ for each group h (in dollar terms), in turn allowing us to approximate average welfare given the new distribution of workers over employers in the post-reform period. And given that we know μ_{gh} and A_{gh} separately, we can decompose this change into changes in earnings and changes in the amenity value of realized employer matches.

Specifically, let $P_{gh} := P(H_i = h | G_i = g)$ be the match probability for group g at employer h , in the pre-reform period, and let \tilde{P}_{gh} be the corresponding probability in the post-reform period. The change in total welfare for group g can be calculated as

$$\sum_h (P_{gh} - \tilde{P}_{gh})(\mu_{gh} + A_{gh})$$

and this change can further be decomposed as a change in earnings $\sum_h (P_{gh} - \tilde{P}_{gh}) \cdot \mu_{gh}$ and an amenity value $\sum_h (P_{gh} - \tilde{P}_{gh}) \cdot A_{gh}$. In addition to assuming first job effects μ_{gh} and average amenity values A_{gh} are stable over time, these calculations consider welfare as captured by these systematic components of utility only. Table 8 reports the results.

The second column of Table 8 reports the earnings FJEs μ_{gh} (also plotted in Figure 1) while the third column reports amenitie values A_{gh} . Amenities fall in the range $[-\mu_{gh}, 0]$, indicating that workers would give up some fraction of the earnings at their chosen employer to remain there instead of moving to Category 4. The A_{gh} are generally increasing (decreasing in magnitude) in category popularity while earnings FJEs exhibit a flatter trend. Workers' combined surplus falls at nearly identical rates between men and women as a function of average lottery draw.

Overall, both men and women lose with regards to earnings FJEs, while gaining—to a greater extent—in employer amenities, with the post-reform distribution of workers over employers. The net effect of the decentralized labor market on worker welfare is positive but not large, representing about 4.71% of the pre-reform average of v_{gh} . Men gain more than women in employer

³²In principle, ΔA_{m1} and ΔA_{f1} could still be arbitrarily far apart, but the values of the parameters in Assumption 7 required to sustain such differences then imply unreasonable values of ΔA_{g1} . We calculated the ΔA_{m1} implied by Equation (16) as a function of ΔA_{f1} for $\Delta A_{f1} \in [\$10,000, \$200,000]$. The maximum relative difference $(\Delta A_{m1} - \Delta A_{f1})/\Delta A_{f1}$ is about 2.5% (which occurs for the smallest ΔA_{f1} in that range).

Employer Category	FJEs By Gender	Amenity Values	Distribution of Workers (%)			
			Pre-Reform (RSD) Women	Pre-Reform (RSD) Men	Post-Reform Women	Post-Reform Men
1	W: 145.07 M: 170.71	W: -68.55 M: -94.87	33.89	33.43	36.89	39.22
2	W: 126.38 M: 174.83	W: -81.11 M: -128.38	37.85	35.38	37.62	34.05
3	W: 150.61 M: 182.04	W: -131.41 M: -165.59	22.49	24.50	19.31	20.94
4	W: 122.89 M: 149.22	W: -122.89 M: -149.22	5.77	6.70	6.18	5.79
Average Predicted Earnings (5 Years Out)			137.96	173.50	137.74	173.24
Average Post-Reform Difference			-0.22	-0.26		
Average Predicted Amenity Values			-90.57	-127.69	-88.77	-124.24
Average Post-Reform Difference			1.80	3.45		
Total Change in Welfare (Per Worker)					1.58	3.19
Workers	9,049	9,049	4,855	4,194	1,781	1,122

Table 8: FJEs measure earnings five year post graduation (four years post residency). Earnings and amenity values in thousands of 2020 USD. Pre/post-reform total welfare_g = $\sum_h prob(h|g)(\mu_{gh} + A_{gh})$ where $prob(h|g)$ = columns 6-7/8-9.

amenities. We conclude that, in the setting we study, first jobs affect workers' long-run career trajectories; they do so differentially for men and women; and "market design" policy can affect the aggregate realized effects of workers' first jobs.

References

- ABDULKADIROĞLU, ATILA, & SÖNMEZ, TAYFUN. 1998. Random Serial Dictatorship and the Core from Random Endowments in House Allocation Problems. *Econometrica*, **66**(3), 689–702.
- ABDULKADIROĞLU, ATILA, ANGRIST, JOSHUA D., NARITA, YUSUKE, & PATHAK, PARAG A. 2017. Research Design Meets Market Design: Using Centralized Assignment For Impact Evaluation. *Econometrica*, **85**(5), 1373–1432.
- ABOWD, JOHN M., KRAMARZ, FRANCIS, & MARGOLIS, DAVID N. 1999. High Wage Workers and High Wage Firms. *Econometrica*, **67**(2), 251–333.
- ANDREWS, DONALD W. K., & SOARES, GUSTAVO. 2010. Inference for Parameters Defined by Moment Inequalities Using Generalized Moment Selection. *Econometrica*, **78**(1), 119–157.
- ANGRIST, JOSHUA D. 1990. Lifetime Earnings and the Vietnam Era Draft Lottery: Evidence from Social Security Administrative Records. *American Economic Review*, **80**(3), 313–336.
- ANGRIST, JOSHUA D. 1995 (July). *Estimating the Labor Market Impact of Voluntary Military Service Using Social Security Data on Military Applicants*. Working Paper 5192. National Bureau of Economic Research.
- ANGRIST, JOSHUA D., & CHEN, STACEY H. 2011. Schooling and the Vietnam-Era GI Bill: Evidence from the Draft Lottery. *American Economic Journal: Applied Economics*, **3**(2), 96–118.
- ANGRIST, JOSHUA D., & IMBENS, GUIDO W. 1994. Identification and Estimation of Local Average Treatment Effects. *Econometrica*, **62**(2), 467–475.
- ARELLANO-BOVER, JAIME. 2020. *Career Consequences of Firm Heterogeneity for Young Workers: First Job and Firm Size*. Working paper.
- BENDER, STEFAN, SCHMIEDER, JOHANNES, & VON WACHTER, TILL. 2009. *The Long-Term Impact of Job Displacement in Germany During the 1982 Recession on Earnings, Income, and Employment*. mimeo Columbia University.
- CHERNOZHUKOV, VICTOR, HONG, HAN, & TAMER, ELIE. 2007. Estimation and Confidence Regions for Parameter Sets in Econometric Models1. *Econometrica*, **75**(5), 1243–1284.
- CHERNOZHUKOV, VICTOR, LEE, SOKBAE, & ROSEN, ADAM M. 2013. Intersection Bounds: Estimation and Inference. *Econometrica*, **81**(2), 667–737.
- COILE, COURTNEY C., LEVINE, PHILLIP B., & MCKNIGHT, ROBIN. 2012 (September). *Recessions, Older Workers, and Longevity: How Long Are Recessions Good For Your Health?* Working Paper 18361. National Bureau of Economic Research.
- GENDA, YUJI, KONDO, AYAKO, & OHTA, SOUICHI. 2010. Long-Term Effects of a Recession at Labor Market Entry in Japan and the United States. *Journal of Human Resources*, **45**(1), 157–196.
- GOFF, LEONARD. 2020. A Vector Monotonicity Assumption for Multiple Instruments. *arXiv*.
- HECKMAN, JAMES J., & PINTO, RODRIGO. 2018. Unordered Monotonicity. *Econometrica*, **86**(1), 1–35.

- HECKMAN, JAMES J., URZUA, SERGIO, & VYTLACIL, EDWARD. 2006. Understanding Instrumental Variables in Models with Essential Heterogeneity. *The Review of Economics and Statistics*, **88**(3), 389–432.
- HEISZ, ANDREW, OREPOULOS, PHIL, & VON WACHTER, TILL MARCO. 2012. Short- and Long-Term Career Effects of Graduating in a Recession. *American Economic Journal: Applied Economics*, **4**(1), 1–29.
- KAHN, LISA B. 2010. The Long-Term Labor Market Consequences of Graduating from College in a Bad Economy. *Labour Economics*, **17**(April), 303–316.
- KALOUPTSIDI, MYRTO, KITAMURA, YUICHI, LIMA, LUCAS, & SOUZA-RODRIGUES, EDUARDO A. 2020 (February). *Partial Identification and Inference for Dynamic Models and Counterfactuals*. Working Paper 26761. National Bureau of Economic Research.
- KAMAT, VISHAL. 2020. *Identification of Program Access Effects with an Application to Head Start*.
- KIRKEBOEN, LARS, LEUVEN, EDWIN, & MOGSTAD, MAGNE. 2017. Field of Study, Earnings, and Self-selection. *Quarterly Journal of Economics*.
- KITAMURA, YUICHI, & STOYE, JÖRG. 2018. Nonparametric Analysis of Random Utility Models. *Econometrica*, **86**(6), 1883–1909.
- KOLESÁR, MICHAL. 2013. *Estimation in an Instrumental Variables Model With Treatment Effect Heterogeneity*. Working paper.
- LEE, SOKBAE, & SALANIÉ, BERNARD. 2020. Filtered and Unfiltered Treatment Effects with Targeting Instruments. *arXiv*.
- MCFADDEN, DANIEL. 1974. *Conditional Logit Analysis of Qualitative Choice Behavior*. New York: Academic Press. Pages 105–142.
- MOGSTAD, MAGNE, SANTOS, ANDRES, & TORGOVITSKY, ALEXANDER. 2018. Using Instrumental Variables for Inference About Policy Relevant Treatment Parameters. *Econometrica*, **86**(5), 1589–1619.
- MOGSTAD, MAGNE, TORGOVITSKY, ALEXANDER, & WALTERS, CHRISTOPHER R. 2020. Policy Evaluation with Multiple Instrumental Variables.
- OREPOULOS, PHILIP, VON WACHTER, TILL, & HEISZ, ANDREW. 2012. The Short- and Long-Term Career Effects of Graduating in a Recession. *American Economic Journal: Applied Economics*, **4**(1), 1–29.
- OYER, PAUL. 2006. Initial Labor Market Conditions and Long-Term Outcomes for Economists. *Journal of Economic Perspectives*, **20**(3), 143–160.
- OYER, PAUL. 2008. The Making of an Investment Banker: Macroeconomic Shocks, Career Choice, and Lifetime Income. *Journal of Finance*.
- RUHM, CHRISTOPHER J. 2000. Are Recessions Good for Your Health? *The Quarterly Journal of Economics*, **115**(2), 617–650.

- SMEULDERS, BART, CHERCHYE, LAURENS, & DE ROCK, BRAM. 2021. Nonparametric Analysis of Random Utility Models: Computational Tools for Statistical Testing. *Econometrica*, **89**(1), 437–455.
- SORKIN, ISAAC. 2016. *Ranking Firms Using Revealed Preference*. mimeo Stanford University.
- STAIGER, MATTHEW. 2020. *The Intergenerational Transmission of Employers and the Earnings of Young Workers*. Working paper.
- SULLIVAN, DANIEL, & VON WACHTER, TILL. 2009. Average Earnings and Long-Term Mortality: Evidence from Administrative Data. *American Economic Review*, **99**(2), 133–38.
- SULLIVAN, DANIEL, & VON WACHTER, TILL. 2009. Job Displacement and Mortality: An Analysis using Administrative Data. *Quarterly Journal of Economics*.
- VON WACHTER, TILL, & BENDER, STEFAN. 2006. In the Right Place at the Wrong Time: The Role of Firms and Luck in Young Workers' Careers. *The American Economic Review*, **96**(5), 1679–1705.

A Formal treatment of the RSD mechanism

In this appendix we provide a formal definition of the Random Serial Dictatorship Mechanism in order to motivate our instrumental variables analysis. We do so by adopting the notion of a continuum economy from [Abdulkadiroğlu et al. \(2017\)](#) that provides an approximation in large samples. Throughout this section we focus on a single instance of the lottery, and suppress the lottery indicator L_i .

A.1 RSD mechanism

Recall the notation from Section 4 that we let i index individual doctors within some population \mathcal{I} . We now consider two cases, one in which \mathcal{I} is a finite set of n doctors (referred to as the *finite economy*), and another in which \mathcal{I} is considered to be the unit interval (referred to as the *continuum economy*). In either case, we assume that there is a fixed set \mathcal{H} of hospitals, and that doctors are indifferent between (residency) jobs within a hospital. We assume each doctor i has a well-defined preference ordering \succ_i over hospitals $h \in \mathcal{H}$. Each hospital each can accommodate proportion q_h of the doctors in that year. That is, in the finite case hospital h has $Q_h = nq_h$ positions available, and in the continuum case hospital h can accommodate proportion q_h of the unit measure of doctors. In our actual data, we have ~ 250 doctors per year across ~ 60 hospitals, so a typical q_h can be thought of as being in the vicinity of $1/60$.

RSD begins by allocating each doctor a lottery number R_i , taken to lie in the unit interval $[0, 1]$ (see Section 4). While in our setting the RSD mechanism is decentralized (doctors indicate a selection from their choice set in real time before the procedure moves to the next doctor), it delivers the same outcome as one in which all doctors submit their full preference ordering \succeq_i and allocations are made centrally, under the assumption that doctors choose according to well-defined and stable such preferences. In this framework, RSD is a special case of the deferred acceptance (DA) mechanism, in which hospitals have no priorities over doctors beyond lottery number. [Abdulkadiroğlu et al. \(2017\)](#), characterize DA in terms of a set of cutoffs $\tau = \{\tau_h\}_{h \in \mathcal{H}}$ such that hospital h is available to i iff $R_i \leq \tau_h$ (and thus i would be centrally assigned to h if they prefer h to any other h' such that $R_i \leq \tau_{h'}$).

Let the *type* θ_i of doctor i be the tuple of their demographic group, potential outcomes and preferences: $\theta_i = (G_i, \mathcal{Y}_i, \succ_i)$, and let Θ be the possible values of θ_i . In a finite economy, the cutoffs τ arising from RSD are determined by the set of pairs $\{(\theta_i, R_i)\}_{i=1\dots n}$. Given a fixed realization of the lottery, we may equivalently represent this set by a discrete uniform distribution over the pairs (θ_i, R_i) in the economy. Denote this probability distribution as F_n . Taking \mathcal{I} to be the underlying sample space, we write $F_n(\mathcal{I}_0) = |\mathcal{I}_0|/n$ for any set $\mathcal{I}_0 \subseteq \mathcal{I}$ of individuals. In the continuum economy, we again work in terms of a distribution over pairs (θ_i, R_i) , denoted as F_0 . To construct F_0 begin with an underlying “population” distribution F_0^θ over types, and take the product measure with a uniform $U[0, 1]$ measure for the lottery draws. This allows for a unified probability space both over individuals and over lottery draws, while maintaining independence between θ_i and R_i .

As described in [Abdulkadiroğlu et al. \(2017\)](#), the cutoffs can be expressed as $\tau_h = \lim_{t \rightarrow \infty} \tau_h^t$ where we imagine a set of “rounds” t in which initially all hospitals are available: $\tau_h^0 = 1$ for all h , and in subsequent rounds the thresholds are lowered for hospitals that were “over-subscribed” given last rounds’ thresholds. With F equal to either F_n or F_0 , we can write this as:

$$\tau_h^{t+1} = \begin{cases} 1 & \text{if } F(Q_h(\tau^t)) < q_h \\ \max\{t \in [0, 1] : F(Q_h(\tau^t) \cap \{i : R_i \leq t\}) \leq q_h\} & \text{if } F(Q_h(\tau^t)) \geq q_h \end{cases} \quad (\text{A.1})$$

where

$$Q_h(\tau) = \{i : R_i \leq \tau_h \text{ and } h \succeq_i h' \text{ for all } h' \text{ s.t. } R_i \leq \tau_{h'}\}$$

is the set of doctors who prefer hospital h from their choice set.³³ Intuitively, Equation (A.1) reduces the threshold for each oversubscribed hospital to the largest value t such that it is no longer overcapacity, ignoring indirect effects of this change from space in other hospitals being made available by the doctors who will now newly choose h . We can write the final choice set C_i for doctor i as a function of their lottery number and the final vector of cutoffs: $C_i = \{h \in \mathcal{H} : R_i \leq \tau_h\}$.

A.2 Asymptotics

We will motivate choices of instruments and treatment based on observations about the continuum economy, which can be expected to provide an accurate approximation to finite economies of sufficient size n . Recall that in either case, the thresholds characterizing the outcome of RSD are determined by the function F (either F_n or F_0 , depending on the case).

Fix a continuum economy with joint distribution F_0 over types and lottery numbers, and vector of hospital capacities q . Formally, we will view a finite economy as a random sample $\{\theta_i, R_i\}_{i=1\dots n}$ of n individuals from this fixed continuum economy. Let F_n be the empirical distribution over (θ_i, R_i) , noting that this coincides with the definition of F_n given in the last section. Asymptotic arguments will consider a sequence of such $\{F_n\}$ with increasing sample size. By the Glivenko-Cantelli theorem, $F_n \rightarrow F_0$ almost surely. This will provide a basis for consistent estimation of “population” quantities defined with respect to the continuum economy.

For a given sample, the econometrician observes a realized outcome Y_i , doctor i 's choice of job H_i and rank R_i according to the lottery. Their choice set can be imputed as

$$C_i = \{h \in \mathcal{H} \text{ such that } |\{j : H_j = h \text{ and } R_j < R_i\}| < Q_h\}$$

where Q_h is the actual number of job openings at hospital h (the continuum economy is defined to have $q_h = Q_h/n$ for the actual sample size n , in our case ~ 250). Let $R_i = R_i/n$ be lottery numbers normalized to the unit interval. The econometrician can compute empirical cutoffs as the largest lottery number such that hospital h is available:

$$\hat{\tau}_h = \max\{R_i : h \in C_i\}_{i=1\dots n}$$

Let $\hat{\tau}$ be the vector of cutoffs for a given sample of size n . Lemma 3 of [Abdulkadiroğlu et al. \(2017\)](#) shows that $\hat{\tau} \xrightarrow{a.s.} \tau$, where τ is the set of cutoffs arising from the continuum economy F . This property will prove useful in the following section.

A.3 Independence of choice sets

Recall that Assumption 1 from Section 4 does generally not hold exactly in a finite economy (c.f. footnote 16). However, we can justify it by appealing to the continuum economy. One way to think about the example given in footnote 16 is that in a finite economy, the probability distribution over C_i depends on $\mathcal{P}_i = \{\succeq_j\}_{j \in I, j \neq i}$, the set of preferences of all individuals in the economy that are not i . Each of these other individuals has the opportunity to choose before i for some realizations of the lottery, thus having an impact on the choices remaining for i . As n gets large, it is reasonable

³³ [Abdulkadiroğlu et al. \(2017\)](#) define the function F over sets taking the form $\mathcal{I}_0 = I(\Theta_0, r_0) := \{i \in I : \theta_i \in \Theta_0, R_i \leq r_0\}$ for any subset $\Theta_0 \subset \Theta$. $F(\mathcal{I}_0)$ is again defined as $|\mathcal{I}_0|/n$ in the finite case, and as $P(\theta_i \in \Theta_0) \cdot r_0$ in the continuum case. However, as the set $Q_h(\tau)$ does not take this form, we do not pursue this definition here.

to expect the magnitude of this effect to attenuate, as \mathcal{P}_i and $\mathcal{P}_{i'}$ become nearly the same “overall” for any $i \neq i'$. We can now formalize this notion, based on the asymptotic sequence of economies introduced in the last section.

Let us now consider a binary “instrument” Z_i that indicates a particular value of i ’s choice set: $Z_i = \mathbb{1}(C_i = c)$ for some fixed $c \subseteq \mathcal{H}$. For any economy (whether finite or continuum) with cutoff vector τ , we can write $Z_i = f(R_i, \tau)$ where

$$f(R_i, \tau) := \mathbb{1}(\forall h \in c : R_i \leq \tau_h \text{ and } \forall h \notin c : R_i > \tau_h)$$

For each n , let τ_n be the cutoffs according to F_n . Let $Z_i^n = f(R_i, \tau_n)$ denote the instrument defined with respect to the finite economy’s cutoffs τ_n , and let $Z_i = f(R_i, \tau)$ represent the “population” analog defined with respect to the continuum limiting cutoffs τ .

Proposition 4 (independence in the continuum economy). $Z_i \perp \theta_i$.

Proof. Immediate, since with τ fixed, Z_i is a measurable function of R_i , and $R_i \perp \theta_i$. \square

Since the above holds for any c , and the events $C_i = c$ and $C_i = c'$ are exclusive for $c \neq c'$, Proposition 4 implies that $C_i \perp \theta_i$, with C_i interpreted with respect to the continuum economy. Thus Assumption 1 as stated in Section 4, since $\theta_i = (G_i, \succ_i, \mathcal{Y}_i)$ and thus Proposition 4 implies that $\{\mathcal{Y}_i, \succ_i\} \perp C_i | G_i$ (recall that lottery L_i is conditioned on implicitly in this section).

Intuitively, Proposition 4 makes use of the notion that with a continuum of doctors and a continuum of positions available at each hospital, any two doctors A and B share the same function that maps lottery numbers to choice sets. With any single doctor a measure zero set from a continuum of doctors, \mathcal{P}_i defined above does not differ between doctors. Note that we do not have an analog of Proposition 4 for the finite economy, since the finite economy cutoffs τ_n depend on F_n , itself a random quantity. As F_n is not independent of θ_i for any fixed i in the realized sample, we cannot expect $Z_i^n = f(R_i, \tau_n)$ to be exactly, except in special cases.

Furthermore, even Z_i^n itself is not directly observed. In a finite sample, we can only compute the estimate $\hat{Z}_i^n = f(R_i, \hat{\tau})$ which uses the empirical cutoffs $\hat{\tau}$ observed in the sample. Nevertheless, \hat{Z}_i^n becomes close to the unobserved ideal instrument Z_i for large n , as the empirical cutoffs τ_n approach their continuum analogs τ . To establish consistency of standard IV estimators, it is sufficient to apply the following Lemma to each of the random variables $V_i \in \{Y_i, D_{hi}\}$:

Proposition 5. For any V_i such that $E[V_i^2] < \infty$, $\frac{1}{n} \sum_{i=1}^n V_i \hat{Z}_i^n \xrightarrow{p} P(C_i = c) E[V_i | C_i = c]$

Proof. See Appendix D. \square

Note that Proposition 5 also implies that $\frac{1}{n} \sum_{i=1}^n V_i (1 - f(R_i, \hat{Z}_i^n)) \xrightarrow{p} P(C_i \neq c) E[V_i | C_i \neq c]$, since $\frac{1}{n} \sum_{i=1}^n V_i (1 - \hat{Z}_i^n) = \frac{1}{n} \sum_{i=1}^n V_i - \frac{1}{n} \sum_{i=1}^n V_i \hat{Z}_i^n$ and $\frac{1}{n} \sum_{i=1}^n V_i \xrightarrow{p} E[V_i]$ by the weak law of large numbers. Then apply the law of iterated expectations over Z_i .

B Simulation evidence on the random choice-set approximation

In this section we present simulation evidence that the asymptotic approximation in Section is a reasonable one in our context.

The simulation DGP constructs an environment with 235 doctors and 60 hospitals, roughly matching a typical year from our data. Hospitals have between 2 and 6 spots available, with a distribution reported in Figure B.1, again intended to match the empirical setting. The total number of spots is 258, allowing each doctor to be placed.

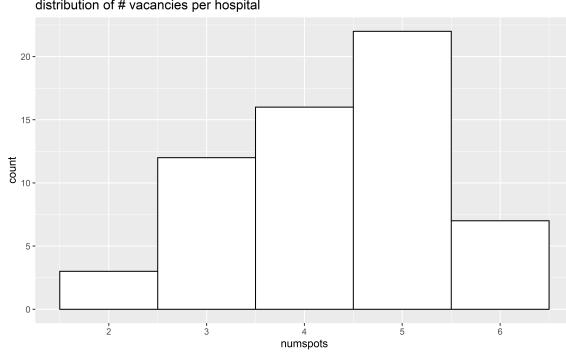


Figure B.1: Simulation distribution of number of spots Q_h in hospital h , across the 60 hospitals.

Preferences over hospitals for the 235 doctors are constructed by first introducing two types of hospitals to create a structured basis for preference heterogeneity. 80% of hospitals (47 in total) are “urban”, and the remaining 20% “rural”. There are two broad types of doctors, those who typically prefer urban hospitals, and those who typically prefer rural ones. 90% of the doctors (199 in total) are urbanites. Within each doctor type, we introduce a “typical” ordering over hospitals, which places all urban hospitals ahead of any rural hospitals, or vice versa. This is meant to reflect a standard ranking over which hospitals are a good place to live/work. Doctors are indifferent between spots in the same hospital.

For 75% of doctors of each type, we start with the archetype ordering for that type and perturb it by performing a series of random swaps of adjacent hospitals in the ordering. Swaps occur with an increasing probability further down the list, reflecting the notion that there is the most agreement among the most desired hospitals, and more heterogeneity among less desired options.³⁴ Since swaps may permute urban with non-urban hospitals, this procedure softens slightly the constraint that all urbanite doctors prefer all urban hospitals to all rural hospitals. The remaining 25% of doctors within each type receive a completely random preference ordering within hospital type, then ordered lexicographically across urban/rural.

The simulation then proceeds by running the RSD lottery 500,000 times, and allocating doctors to hospitals based upon their preferences. Since there are enough jobs for all of the doctors, and none prefer an outside option (by construction), all doctors receive a position. Figures B.2 and B.3 compare the distribution across simulation runs of features of the choice sets facing two doctors: “Doctor 1” and “Doctor 2”. Doctor 1 is a single randomly chosen urbanite doctor, and Doctor 2 is a single randomly chosen non-urbanite doctor. If choice sets are unconditionally random, then the doctors should face the same probability distribution over any function of their choice set. In both cases, statistical tests reject this null-hypothesis. However, the figures reveal that the differences are quite minor, almost imperceptible without close inspection. This provides evidence that the asymptotic approximation of choice-set independence is likely to be quite reasonable in our context.

Figure B.2 compares the distributions facing the two doctors over lottery draws of the proportion of hospitals in their choice set that are urban. If these distributions were substantially different from one another, it would cast doubt on using something like the proportion urban of one’s choice set as an instrument for choosing an urban hospital. In particular, it would suggest that this in-

³⁴Specifically, a vector of 30 “swap positions” are introduced through the preference list, with a CDF increasing as the square of number in the list. For each swap position j , a random draw determines whether the hospital in that position is swapped with the one below, the one above, or no change is made.

strument is not independent of preferences, since the only difference between Doctors 1 and 2 in the simulation are their preferences (recall that Doctor 1 prefers urban hospitals and Doctor 2 rural ones). A two sample Kolmogorov-Smirnov test strongly rejects the null that the two distributions are identical, which is not surprising given the large number of simulation draws. Statistics of the distribution also differ: for example the average proportion urban for Doctor 1 is 68.0% vs. 67.3% for Doctor 2. Nevertheless, the differences are quite small in practical terms, as evident in the histograms. Across all of the 235 Doctors, the minimum value of the average proportion urban is 67.3%, and its maximum is 68.0%.

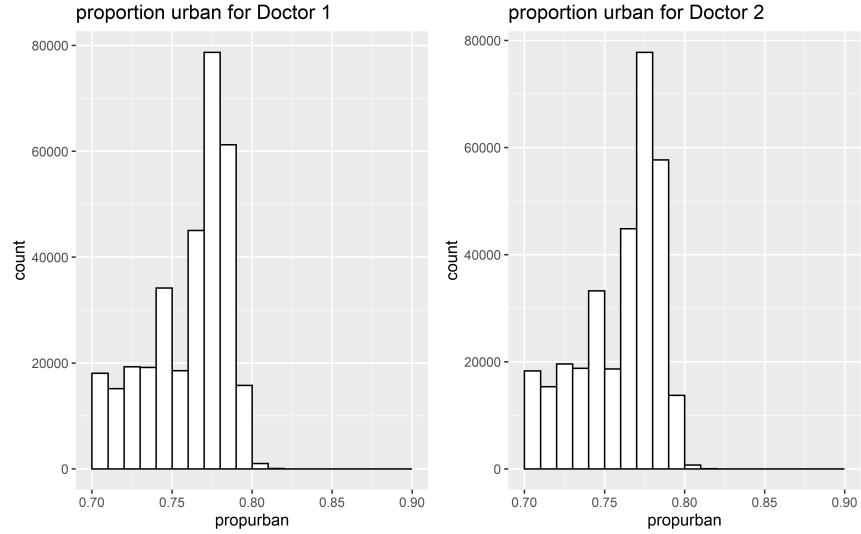


Figure B.2: Simulation distribution over the proportion of one's choice set that is urban hospitals, between Doctors 1 and 2)

As a benchmark, Figure B.3 compares the distribution over number of distinct hospitals present in each of the two doctors' choice sets. Given that there is no statistical relationship between hospital size Q_h and whether h is urban or rural, we might expect $|C_i|$ to be independent of θ_i in this case, even in a finite sample. A chi-squares test also strongly rejects this null hypothesis, however the distributions appear nearly identical in practical terms.

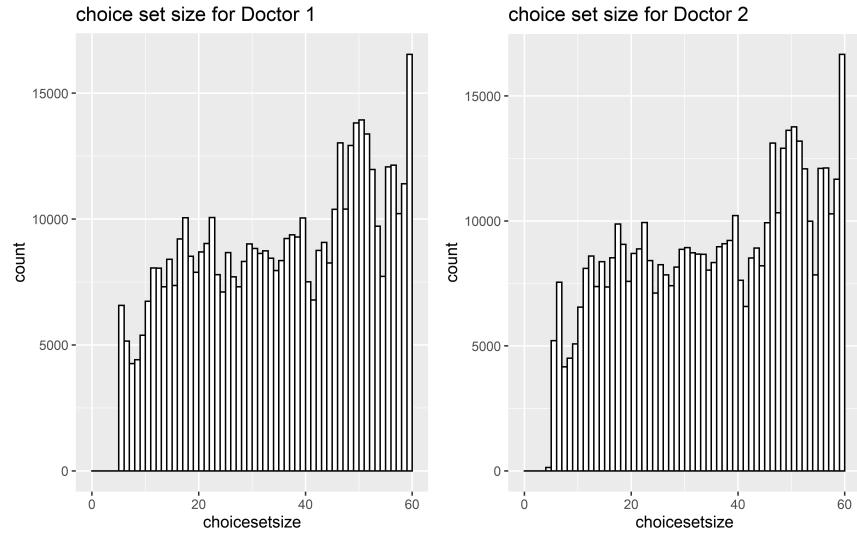


Figure B.3: Comparison of the simulation distribution over number of hospitals in one's choice set, between Doctors 1 and 2)

C Additional tables and figures

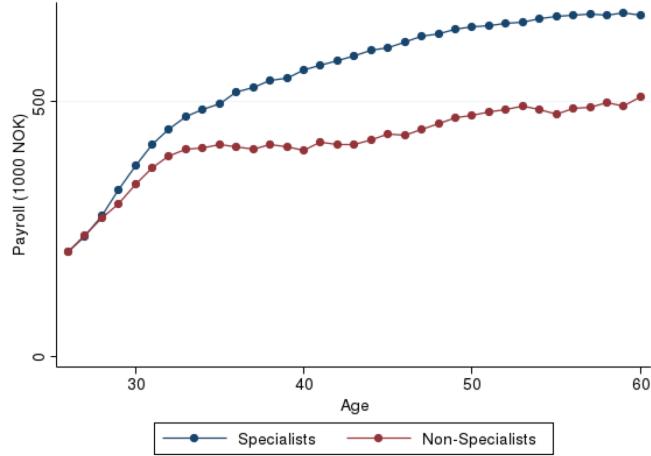


Figure C.4: Average income by age among specialists and non-specialists.

Table C.1: Summary Statistics

Hospital Characteristics	Good Hospitals		Other Hospitals	
	Mean	SD	Mean	SD
Number of doctors	186.92	198.88	47.53	47.13
Mean doctor experience (years)	12.07	2.28	10.88	3.36
Number of specialists	90.46	104.63	78.21	88.24
Proportion of doctors who are specialists	0.54	0.10	0.51	0.15
Mean doctor income	161.47	30.05	163.55	34.76
Average doctor age	43.75	2.04	44.80	4.41
Proportion of male doctors	0.67	0.10	0.72	0.15
Proportion of foreign doctors	0.27	0.14	0.38	0.21
<hr/>				
Observations	703		698	
<hr/>				
Doctor Characteristics	Female		Male	
	Mean	SD	Mean	SD
Age	39.19	9.80	45.73	11.40
Cohabit	0.63	0.48	0.75	0.43
Number of Children	1.05	1.14	0.97	1.18
Born Abroad	0.23	0.42	0.20	0.40
After-Tax Income	71.51	49.83	90.91	135.37
Real Estate	30.68	46.74	48.49	55.24
Debt	133.98	163.10	209.00	255.96
<hr/>				
Specialization				
Specialist	0.34	0.47	0.56	0.50
General Practice	0.09	0.28	0.12	0.32
Internal Medicine	0.13	0.34	0.18	0.39
Surgery	0.05	0.22	0.18	0.39
<hr/>				
Observations	80,025		134,458	

Notes: This table presents summary statistics using annual data on hospitals and doctors in Norway during 1995-2011. Hospitals (and doctors) that are observed twice will count as two separate observations, since observable characteristics may change over time. Rural location is defined as the proportion of population in the municipality that lives in rural areas. Doctor income is in thousands in 2011 USD.

Table C.2: Randomization Via Lottery

Individual Characteristic	ω_C	t-stat	R ²	N
Male	0.003	0.561	0.082	9828
Age	-0.001	-1.370	0.082	9828
Rural Residence at Age 15	-0.011	-0.846	0.085	8267
Born Abroad	-0.005	-0.655	0.082	9828
Study Abroad	-0.0002	-0.021	0.095	2871

Notes: This table presents evidence that the lottery number was not influenced by doctor characteristics. Each row presents estimates from a separate regression, where the dependent variable is the lottery draw number normalized to lie between 0 and 1, and the independent variable is an individual (doctor) characteristic. Regressions include lottery fixed effects to allow for demographic changes in the participant pool over time. The number of observations is much lower for the last row because data on study location is only available for the last few years of the sample.

g	h	$[\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U]$				\mathcal{C}_n	
		$\kappa_n = 0$		$\kappa_n = 10\%$		(95% CI)	
		(category)					
Women	1	0.50	0.50	0.49	0.60	0.26	3.95
	2	1.00	1.00	0.99	1.10	0.26	4.74
	3	0.63	0.71	0.62	0.78	0.53	4.74
	4	0.86	0.88	0.85	1.12	0.26	5.00
Men	1	0.79	0.79	0.74	0.91	0.26	4.74
	2	0.50	0.50	0.49	0.68	0.53	3.68
	3	1.00	1.03	1.00	1.13	0.53	3.95
	4	1.50	1.50	1.50	1.56	0.26	5.00

Table C.3: First job effects $\mu_{gh} = \mathbb{E}[Y_i(h)|G_i = g]$ for career number of specializations. Table reports estimates of the identified set $[\hat{\theta}_{gh}^L, \hat{\theta}_{gh}^U]$ and 95% confidence intervals for μ_{gh} .

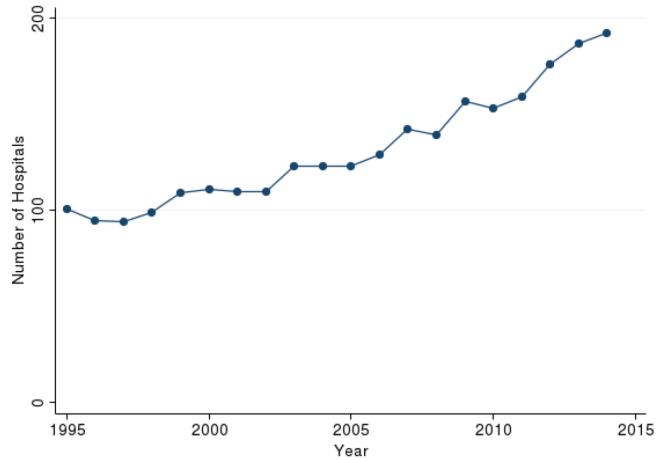


Figure C.5: Number of hospitals by year.

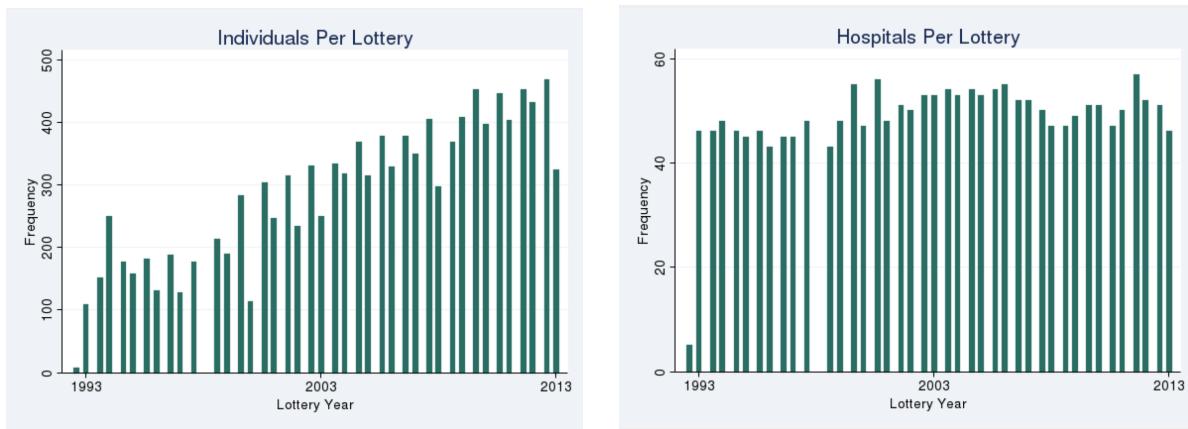


Figure C.6: Number of individuals and hospitals by lottery.

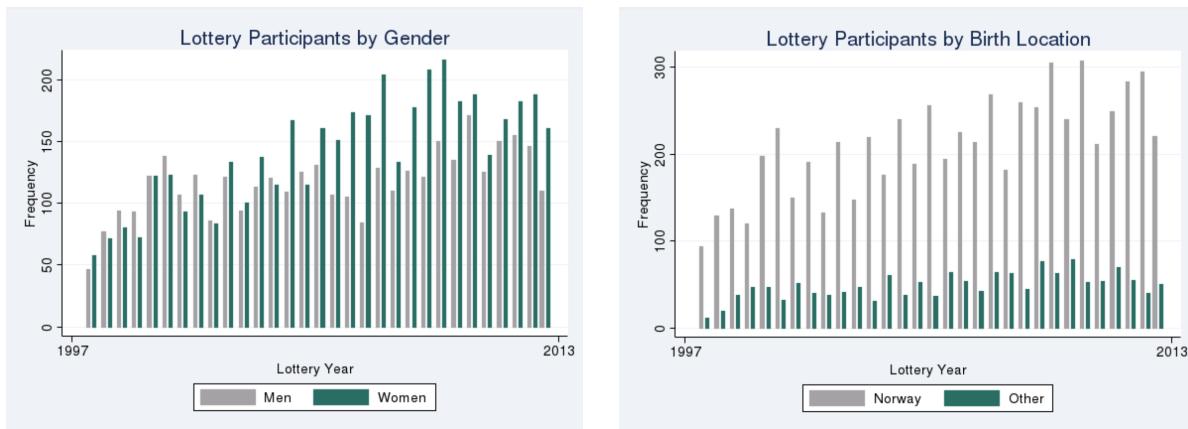


Figure C.7: Number of individuals and hospitals by lottery.

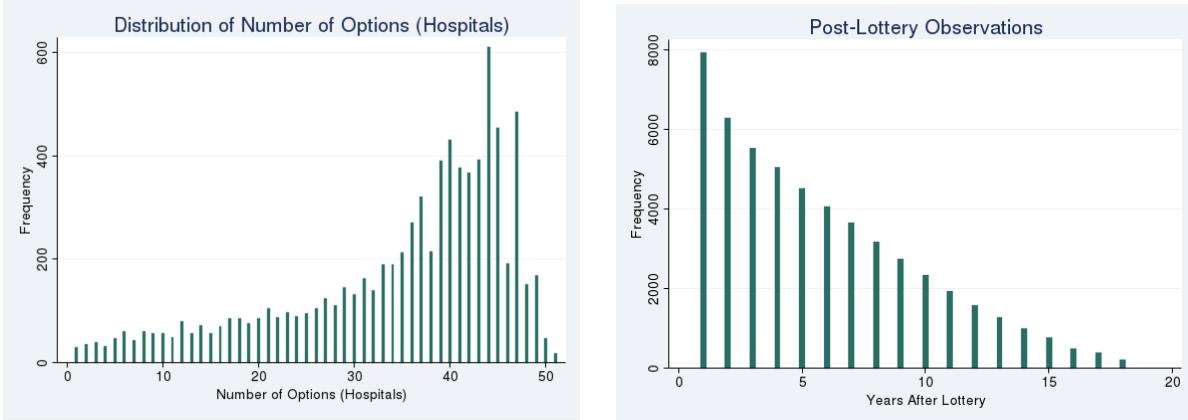


Figure C.8: Distribution of the number of choices for residency hospital $|C_i|$ (left), and the number of observations post-residence for a given doctor (right).

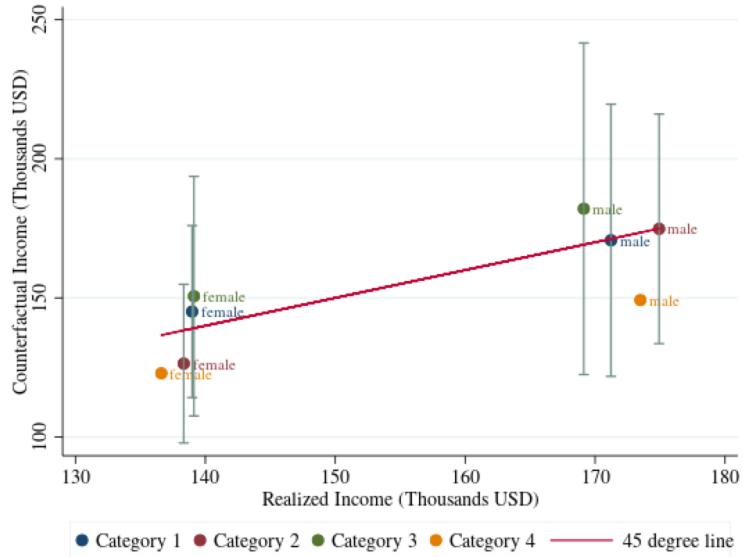


Figure C.9: Earnings FJE's μ_{gh} vs. average realized earnings $\mathbb{E}[Y_i|H_i = h, G_i = g]$, with 45 degree line in red. Brackets represent 95% intervals on the *difference* between μ_{gh} and μ_{g4} for Category 4.

D Proofs

D.1 Proof of Proposition 2

This proof follows the logic of Kolesár (2013) from the binary treatment case. Assumption 1 implies that

$$(Y_i(h) \perp C_i) | G_i, L_i \quad (\text{D.2})$$

where L_i is the lottery/cohort of doctor i . Assumption 2 says that

$$\mathbb{E}[Y_i(h) - Y_i(h_0)|H_i = h', C_i = c, L_i = \ell, G_i = g] = \beta_{hg} \quad (\text{D.3})$$

where β_{hg} is a number that does not depend on h' , h_0 , or x (or c). Fix an arbitrary choice of h_0 , which will serve as a comparison hospital throughout intermediate steps of the proof.

By the law of iterated expectations over H_i and C_i , β_{hg} must be equal to $\mu_{hg} - \mu_{h_0g}$. Now substituting $h' = h$ into Equation (D.3), we have:

$$\mathbb{E}[Y_i(h) - Y_i(h_0)|H_i = h, C_i = c, L_i = \ell, G_i = g] = \beta_{hg} \text{ for all } h, c.$$

Collect the β_{hg} across h for a fixed g into a vector β_g . Then we can rewrite this as:

$$\mathbb{E}[Y_i - Y_i(h_0) - \beta'_g \mathbf{D}_i | \mathbf{D}_i, \mathbf{Z}_i = z, L_i = \ell, G_i = g] = 0$$

for any z , which implies by the law of iterated expectations over \mathbf{D}_i that

$$\mathbb{E}[Y_i - Y_i(h_0) - \beta'_g \mathbf{D}_i | \mathbf{Z}_i = z, L_i = \ell, G_i = g] \quad (\text{D.4})$$

Consider first the case in which there is a single cohort, and we can thus ignore the conditioning on L_i . Then Assumption 1 implies that $\mathbb{E}[Y_i(h_0)|\mathbf{Z}_i = z, G_i = g] = \mu_{h_0g}$. and we can thus write:

$$\mathbb{E}[Y_i - \mu_{h_0g} - \beta'_g \mathbf{D}_i | \mathbf{Z}_i = z, L_i = \ell, G_i = g] = \mathbb{E}[Y_i - \mu'_g \mathbf{D}_i | \mathbf{Z}_i = z, L_i = \ell, G_i = g] = 0$$

This implies in particular that $\mathbb{E}[\mathbf{Z}_i Y_i | G_i = g] = \mathbb{E}[\mathbf{Z}_i \mathbf{D}'_i | G_i = g] \boldsymbol{\mu}_g$. $\mathbb{E}[\mathbf{Z}_i \mathbf{D}'_i | G_i = g]$ is invertible by Assumption 3, yielding the result as stated in Proposition 2.

In actual estimation, we pool over cohorts with cohort fixed effects. To see that this is valid under Equations D.2 and D.3, let \mathbf{L}_i be a vector of indicators for each value of L_i . Note that $\mathbb{E}[Y_i(h_0)|L_i = \ell, G_i = g]$ must be linear in \mathbf{L}_i for each g . Let

$$\mathbb{E}[Y_i(h_0)|\mathbf{Z}_i = z, L_i = \ell, G_i = g] = \delta'_g \mathbf{L}_i$$

Thus, picking up from Equation D.4:

$$\mathbb{E}[Y_i - \delta'_g \mathbf{L}_i - \beta'_g \mathbf{D}_i | \mathbf{Z}_i = z, \mathbf{L}_i, G_i = g] = 0$$

which gives us moment conditions to identify β_g and δ_g for each g . The FJE's are now recoverable as:

$$\mu_{hg} = \mathbb{E}[Y_i(h)|G_i = g] = \mathbb{E}[Y_i(h_0)|G_i = g] + \beta_{gh} = \mathbb{E}[\mathbf{L}_i'] \boldsymbol{\delta}_g + \beta_{gh}$$

D.2 Proof of Lemma 5

Recall that $Z_i = \mathbb{1}(C_i = c) = f(R_i, \tau)$ and $\hat{Z}_i^n = f(R_i, \hat{\tau})$. We show that

$$plim \left(\frac{1}{n} \sum_{i=1}^n V_i f(R_i, \hat{\tau}) - \frac{1}{n} \sum_{i=1}^n V_i Z_i \right) = 0.$$

This suffices to prove the Lemma since $plim \left(\frac{1}{n} \sum_{i=1}^n V_i Z_i \right) = P(C_i = c) E[V_i | C_i = c]$ by the weak law of large numbers.

For a given value r , the function $f(r, \hat{\tau})$ may be discontinuous at values of $\hat{\tau}$ such that $\hat{\tau}_h = r$ for some $h \in S$. However, the continuous mapping theorem nevertheless implies that $f(r, \hat{\tau}) \xrightarrow{a.s.} f(r, \tau)$ pointwise for all r such that $r \neq \tau_h$ for all h . To simplify notation, let $G = \{r \in [0, 1] : r \neq \tau_h \text{ for all } h\}$. By the Severini-Egorov theorem, for any $\tilde{\epsilon}' > 0$, there exists a set $J \subset [0, 1]$ of Lebesgue measure smaller than $\tilde{\epsilon}'$ such that for any $\tilde{\delta} > 0$: $P \left(\sup_{r \in G \setminus J} |f(r, \hat{\tau}) - f(r, \tau)| > \tilde{\delta} \right) \xrightarrow{n} 0$.

Fix any $\delta > 0$ and $\epsilon > 0$, and for now, consider an arbitrary set J . Expanding over the two cases:

$$\begin{aligned} P \left(\left| \frac{1}{n} \sum_{i=1}^n V_i f(R_i, \hat{\tau}) - \frac{1}{n} \sum_{i=1}^n V_i Z_i \right| \geq \delta \right) &= P \left(\left| \frac{1}{n} \sum_{i=1}^n V_i (f(R_i, \hat{\tau})) - Z_i \right| \geq \delta \right) \\ &\leq P \left(\left| \frac{1}{n} \sum_{i: R_i \in G \setminus J} V_i (f(R_i, \hat{\tau})) - Z_i \right| + \frac{1}{n} \sum_{i: R_i \notin G \setminus E} V_i (f(R_i, \hat{\tau})) - Z_i \right| \geq \delta \right) \\ &\leq P \left(\left| \frac{1}{n} \sum_{i: R_i \in G \setminus J} V_i (f(R_i, \hat{\tau})) - Z_i \right| \geq \delta/2 \right) + P \left(\left| \frac{1}{n} \sum_{i: R_i \notin G \setminus J} V_i (f(R_i, \hat{\tau})) - Z_i \right| \geq \delta/2 \right) \end{aligned}$$

Considering the first term, and applying the Markov and Cauchy-Schwarz inequalities:

$$\begin{aligned} P \left(\left| \frac{1}{n} \sum_{i: R_i \in G \setminus J} V_i (f(R_i, \hat{\tau})) - Z_i \right| \geq \delta/2 \right) &\leq P \left(\left| \frac{1}{n} \sum_{i: R_i \in G \setminus K} |V_i (f(R_i, \hat{\tau})) - f(R_i, \tau))| \geq \delta/2 \right) \\ &\leq \frac{2}{\delta} E [\mathbb{1}(R_i \in G \setminus J) \cdot |V_i (f(R_i, \hat{\tau})) - f(R_i, \tau))|] \\ &\leq \frac{2}{\delta} E [|V_i| \cdot \mathbb{1}(R_i \in G \setminus J) \cdot |(f(R_i, \hat{\tau})) - f(R_i, \tau))|] \\ &\leq \frac{2}{\delta} \sqrt{E[V_i^2]} \cdot \sqrt{E [\mathbb{1}(R_i \in G \setminus J) \cdot |f(R_i, \hat{\tau}) - f(R_i, \tau)|^2]} \\ &\leq \frac{2}{\delta} \sqrt{E[V_i^2]} \cdot \sqrt{P(R_i \in G \setminus J) \cdot E [(f(R_i, \hat{\tau})) - f(R_i, \tau))^2 | R_i \in G \setminus J]} \\ &\leq \frac{2}{\delta} \sqrt{E[V_i^2]} \cdot \sqrt{E [(f(R_i, \hat{\tau})) - f(R_i, \tau))^2 | R_i \in G \setminus J]} \\ &\leq \frac{2}{\delta} \sqrt{E[V_i^2]} \cdot \sqrt{E \left[\left\{ \sup_{r \in G \setminus J} |f(r, \hat{\tau}) - f(r, \tau)| \right\}^2 \right]} \end{aligned}$$

For any $\tilde{\epsilon} > 0$, there exists an N_1 such that for all $n \geq N_1$, $P \left(\sup_{r \in G \setminus J} |f(r, \hat{\tau}) - f(r, \tau)| > \tilde{\delta} \right) < \tilde{\epsilon}$. Given that $|f(r, \hat{\tau}) - f(r, \tau)| \leq 1$ for any r and $\hat{\tau}$, it then follows that for $n \geq N_1$:

$$E \left[\left\{ \sup_{r \in G \setminus J} |f(r, \hat{\tau}) - f(r, \tau)| \right\}^2 \right] \leq \tilde{\delta}^2 (1 - \tilde{\epsilon}) + \tilde{\epsilon}$$

By choosing $\tilde{\epsilon}$ and $\tilde{\delta}$ such that $\tilde{\delta}^2 (1 - \tilde{\epsilon}) + \tilde{\epsilon} \leq \frac{\epsilon^2 \delta^2}{36E[V_i^2]}$, we will have

$$P \left(\left| \frac{1}{n} \sum_{i: R_i \in G \setminus J} V_i (f(R_i, \hat{\tau})) - Z_i \right| \geq \delta/2 \right) < \epsilon/3$$

In particular, we can choose $\tilde{\epsilon} = 1/2 \cdot \min \left\{ \frac{\epsilon^2 \delta^2}{36E[V_i^2]}, 1 \right\}$ and $\tilde{\delta} = \frac{\tilde{\epsilon}}{1 - \tilde{\epsilon}}$.

Now we turn to the second term.

$$\begin{aligned}
P \left(\left| \frac{1}{n} \sum_{i: R_i \notin G \setminus J} V_i(f(R_i, \hat{\tau})) - Z_i \right| \geq \delta/2 \right) &\leq P \left(\frac{1}{n} \sum_{i: R_i \notin G \setminus J} |V_i(f(R_i, \hat{\tau})) - Z_i| \geq \delta/2 \right) \\
&\leq P \left(\frac{1}{n} \sum_i \mathbb{1}(R_i \notin G \setminus J) \cdot |V_i(f(R_i, \hat{\tau})) - Z_i| \geq \delta/2 \right) \\
&\leq \frac{2}{\delta} \sqrt{E[V_i^2]} \cdot \sqrt{P(R_i \notin G \setminus J)} \\
&\leq \frac{2}{\delta} \sqrt{E[V_i^2]} \cdot \sqrt{P(R_i \notin G) + P(R_i \in J)} \\
&\leq \frac{2}{\delta} \sqrt{E[V_i^2]} \cdot \left(\sqrt{P(R_i \notin G)} + \sqrt{P(R_i \in J)} \right)
\end{aligned}$$

by similar steps as above. Firstly, we choose J such that $P(R_i \in J) = \tilde{\epsilon}' = \frac{\delta\epsilon^2}{36E[V_i^2]}$. Secondly, note that since $R_i \xrightarrow{d} U[0, 1]$ and G is a finite set of points in $[0, 1]$, $P(R_i \in G) \xrightarrow{n} 0$: that is, given any $\tilde{\epsilon}'' > 0$ there exists a N_2 such that $P(R_i \in G) > 1 - \tilde{\epsilon}''$ for all $n \geq N_2$. In particular, choose $\tilde{\epsilon}'' = \frac{\delta\epsilon^2}{36E[V_i^2]}$.

All together, for $n \geq \max\{N_1, N_2\}$ we have that

$$P \left(\left| \frac{1}{n} \sum_{i=1}^n V_i(f(R_i, \hat{\tau})) - Z_i \right| \geq \delta \right) \leq \epsilon/3 + \epsilon/3 + \epsilon/3 = \epsilon$$

and thus $plim (\frac{1}{n} \sum_{i=1}^n V_i f(R_i, \hat{\tau}) - \frac{1}{n} \sum_{i=1}^n V_i Z_i) = 0$.